


EECS565 Intro to Computer and Information Security

# Intro to Database Security

Bo Luo  
bluo@ku.edu



# Database Security



# DBMS

- Collection of
  - interrelated data and
  - set of programs to access the data
- Convenient and efficient processing of data
- Database Application Software



# Database Security

- **Protect Sensitive Data from**
  - Unauthorized disclosure
  - Unauthorized modification
  - Denial of service attacks
- **Security Controls**
  - Security Policy
  - Access control models
  - Integrity protection
  - Privacy problems
  - Fault tolerance and recovery
  - Auditing and intrusion detection



# Data Confidentiality

- Access control
  - which data users can access
- Information flow control
  - what users can do with the accessed data
- Inference Attacks
- Data Mining



# Access Control

- Ensures that all direct accesses to object are authorized
- Protects against accidental and malicious threats by regulating the read, write and execution of data and programs
- Requires:
  - Proper user identification
  - Information specifying the access rights is protected from modification



# Access Control

- Access control components:
  - Access control policy
    - specifies the authorized accesses of a system
  - Access control mechanism
    - implements and enforces the policy



# Access Control Models

- How to describe the policies?
  - “Who can access what?”
  - Subject: active entity that requests access to an object
    - - e.g., user or program
  - Object: passive entity accessed by a subject
    - - e.g., record, relation, file
  - Access right (privileges): how a subject is allowed to access an object
    - - e.g., subject  $s$  can read object  $o$





# Access Control Models

- Mandatory Access Control (MAC)
- Discretionary Access Control (DAC)
- Role-Based Access Control (RBAC)



# Mandatory Access Control- MAC

- Motivated by government in late 1980's/early 1990's
- Utilize security classifications
  - TS: Top Secret, S: Secret, C: Classified, U: Unclassified
  - $TS > S > C > U$
- Each subject and object are classified into one of the security classifications (TS, S, etc.)
- Bell-LaPadulla properties (restrictions on data access)
  - simple property: No READ UP
  - star (\*) property: No WRITE DOWN (write at own level)



# MLS

- Multilevel relational (MLS) schema
  - classification attribute  $C$
  - tuple classification  $TC$
  - $R(A_1, C_1, A_2, C_2, \dots, A_n, C_n, TC)$  *Jajodia-Sandhu*



# MLS Relation Example

<u>Vessel</u>	Objective	Destination
Micra	Shipping	Moon
Vision	Spying	Saturn
Avenger	Spying	Mars
Logos	Shipping	Venus



# MLS Relation Example

<u>Vessel</u>	Objective	Destination
Micra	Shipping	Moon
Vision	Spying	Saturn
Avenger	Spying	Mars
Logos	Shipping	Venus

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	TC
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S



# MLS

- Read from MLS tables
  - Access to each record is determined by TC
  - Level U sees first 2 tuples
  - Level C sees first 3 tuples
  - Level S sees all tuples

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	TC
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S



# MLS Insert

- What if a U user wants to insert a tuple with vessel = Avenger?
- If reject the insert – what will happen?
  - Covert channel
- If insert another Avenger, what about the primary key? Will have 2 Avengers
  - PK + Classification

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	TC
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S



# MLS Insert

- What if a U user wants to insert a tuple with vessel = Avenger?
- If reject the insert – what will happen?
  - Covert channel
- If insert another Avenger, what about the primary key? Will have 2 Avengers
  - New primary key: PK + Classification

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	TC
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S





# MLS Insert

- What if a U user wants to insert a tuple with vessel = Avenger?
- If reject the insert – what will happen?
  - Covert channel
- If **insert** another Avenger, what about the primary key? Will have 2 Avengers
  - **New primary key: PK + Classification**

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	<u>TC</u>
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S
Avenger	U	Shipping	U	Mars	U	U



# MLS Update

- What if the S level wants to update one of the tuples at the U level?
  - U cannot see the update
  - Replicate the tuple

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	<u>TC</u>
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S
Avenger	U	Shipping	U	Mars	U	U



# MLS Update

- What if the S level wants to update one of the tuples at the U level?
  - U cannot see the update
  - Replicate the tuple
    - Replicate (Vision, ....., U)
    - Update it to (Vision, ....., S)

<u>Vessel</u>	C-V	Objective	C-O	Destination	C-D	<u>TC</u>
Micra	U	Shipping	U	Moon	U	U
Vision	U	Spying	U	Saturn	U	U
Avenger	C	Spying	C	Mars	C	C
Logos	S	Shipping	S	Venus	S	S
Avenger	U	Shipping	U	Mars	U	U
Vision	U	Spying	U	Venus	S	S



# MLS

- Wrap up
  - Simple
  - Easy to manage: just assign a security classification to each user
  - The security classifications are used in government and military
  - MLS is widely adopted in government and military.



## Discretionary Access Control (DAC)

- All commercial DB systems adopt DAC
- Current discretionary authorization models for relational DBMS are based on the System R authorization model
  - P. P. Griffiths and B. W. Wade, “An Authorization Mechanism for a Relational Database System,” *ACM Transactions on Database Systems (TODS)* Volume 1 Issue 3, Sept. 1976.
- It is based on ownership administration with administration delegation



# Discretionary Access Control (DAC)

- For each subject access right to the objects are defined
  - (subject, object, +/- access mode)
  - E.g. (Black, Employee-relation, read)
- Based on granting and revoking privileges
- Assign privileges
  - to account level (subject)
    - independent of the relations
    - create schema, create table, create view
  - on relation level (object)
    - on a particular base relation or view



# Authorization ID's

- A user is referred to by *authorization ID*, typically their login name.
- There is an authorization ID PUBLIC.
  - Granting a privilege to PUBLIC makes it available to any authorization ID.
- The objects on which privileges exist include stored tables and views.
- Other privileges are the right to create objects of a type, e.g., triggers.



# Privileges

- A file system identifies certain privileges on the objects (files) it manages.
  - Typically read, write, execute.
- SQL identifies a more detailed set of privileges on objects (relations) than the typical file system.
- Nine privileges in all, some of which can be restricted to one column of one relation.





# Privileges

- Some important privileges on a relation:
  1. **SELECT** = right to query the relation.
    - ♦ May apply to only one attribute.
  2. **INSERT** = right to insert tuples.
  3. **DELETE** = right to delete tuples.
  4. **UPDATE** = right to update tuples.
    - ♦ May apply to only one attribute.



# Granting Privileges

- You have all possible privileges on the objects, such as relations, that you create.
- You may grant privileges to other users (authorization ID's), including PUBLIC.
- You may also grant privileges WITH GRANT OPTION, which lets the grantee also grant this privilege.



# The GRANT Statement

- To grant privileges, say:  
    GRANT <list of privileges>  
    ON <relation or other object>  
    TO <list of authorization ID's>;
- If you want the recipient(s) to be able to  
    pass the privilege(s) to others add:  
    WITH GRANT OPTION



## Example: GRANT

- Suppose you are the owner of Sells. You may say:

```
GRANT SELECT, UPDATE (price)  
ON Sells  
TO sally;
```

- Now Sally has the right to issue any query on Sells and can update the price component only.



## Example: Grant Option

- Suppose we also grant:  
GRANT UPDATE ON Sells TO sally  
WITH GRANT OPTION;
- Now, Sally not only can update any attribute of Sells, but can grant to others the privilege UPDATE ON Sells.
  - Also, she can grant more specific privileges like UPDATE (price) ON Sells.



# Revoking Privileges

REVOKE <list of privileges>

ON <relation or other object>

FROM <list of authorization ID's>;

- Your grant of these privileges can no longer be used by these users to justify their use of the privilege.
  - But they may still have the privilege because they obtained it independently from elsewhere.



# REVOKE Options

- We must append to the REVOKE statement either:
  1. **CASCADE**. Now, any grants made by a revokee are also not in force, no matter how far the privilege was passed.
  2. **RESTRICT**. If the privilege has been passed to others, the REVOKE fails as a warning that something else must be done to “chase the privilege down.”



# RBAC

- Role-based access control (RBAC)
- Sandhu, R., Coyne, Feinstein, Youman:  
“Role-Based Access Control Models,”  
*IEEE Computer*, 29 (2): 38–47.
  - Semantic construct
  - System administrator creates roles according to job functions





# RBAC

- Role
  - Specific task competency
  - duty assignments
  - Embody authority and responsibility
- Grant permissions to users in these roles
  - Roles & permissions
  - Users & roles



# Motivation

- Roles define individuals and extent of resource access
- Combination of users and permissions can change
  - E.g. user membership in roles
- Permissions associated with roles stable
- Administration of roles rather than permissions
- Role permission predefined
  - Easier to add/remove users membership than create new roles/permissions
- Roles part of SQL3
- Supported by many software products
  - Roles used in Windows NT, XP (system admin)



# RBAC basics

- Access control in RBAC exists in:
  - Role-permission (stable)
  - User-role (dynamic)
  - Role-role relationships (stable)
- RBAC supports principles:
  - Least privilege
  - Separation of duties- mutually exclusive roles
  - Data abstraction- abstract permissions (not just R/W)
- Limitations
  - RBAC cannot enforce way principles applied – system admin could configure to violate



# Constraints

- Mutually exclusive roles
  - User at most 1 role in ME set
  - Combinations of roles and permissions can be prohibited
- Cardinality
  - Maximum number of members in a role
  - Minimum cardinality difficult to implement
- Prerequisite role
  - User assigned to role B, only if assigned to A
  - Permission p assigned to role only if role has permission q



# DAC, MAC vs. RBAC

- DAC vs. MAC emerged from defense security research
- RBAC independent of access control
- RBAC can be used to implement DAC, MAC



# Database Encryption

- Application level encryption
- Database encryption
- Protect keys



# Application Level Encryption

- Data protected in database & storage
- Data protected in use and transit
- Programming needed in the applications
- Power of database limited -indexing, searching, stored procedures
- All access must go through application
- Key management



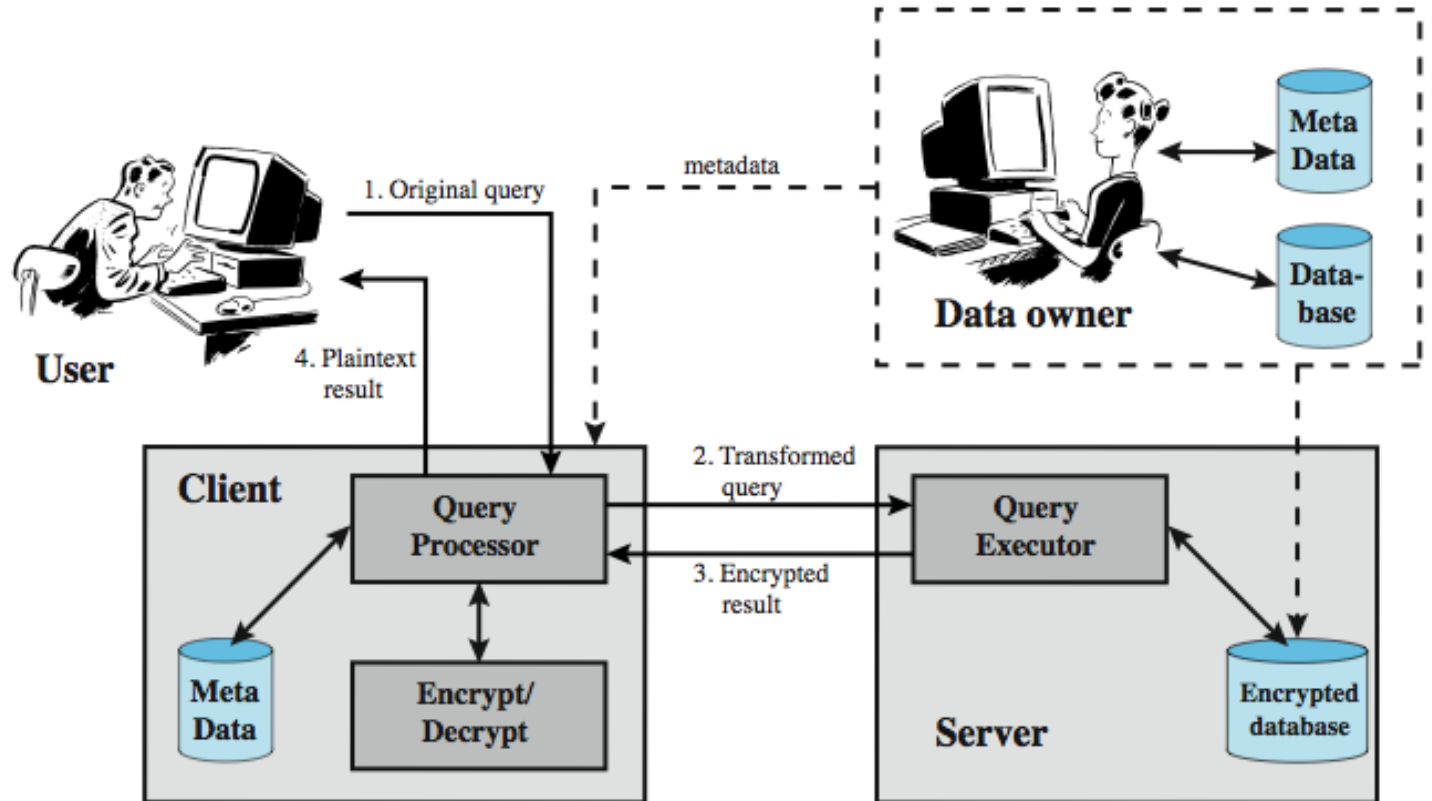
# Database Encryption

- Protects data as it is written to and read from a database
- Secures data in the file-system used by database
- Enables field level encryption
- Transparent to applications
- Watch data in transit





# Database Encryption



# Limit the amount of encrypt/decrypts

- Three Important Techniques
  - Column level encryption
  - Search without decrypting data
  - Conduct data operations without decrypting data



# Sharing Non-Sensitive Data

- A database may contain both sensitive (e.g. salary) and non-sensitive (e.g. name) data.
- Security Goal: To disclose only non-sensitive data.
  - In some cases: aggregated sensitive data (e.g. average salary)
- Precision Goal: To protect all sensitive data while disclosing as much non-sensitive data as possible.
- Ideal combination is to maintain perfect security with maximum precision.

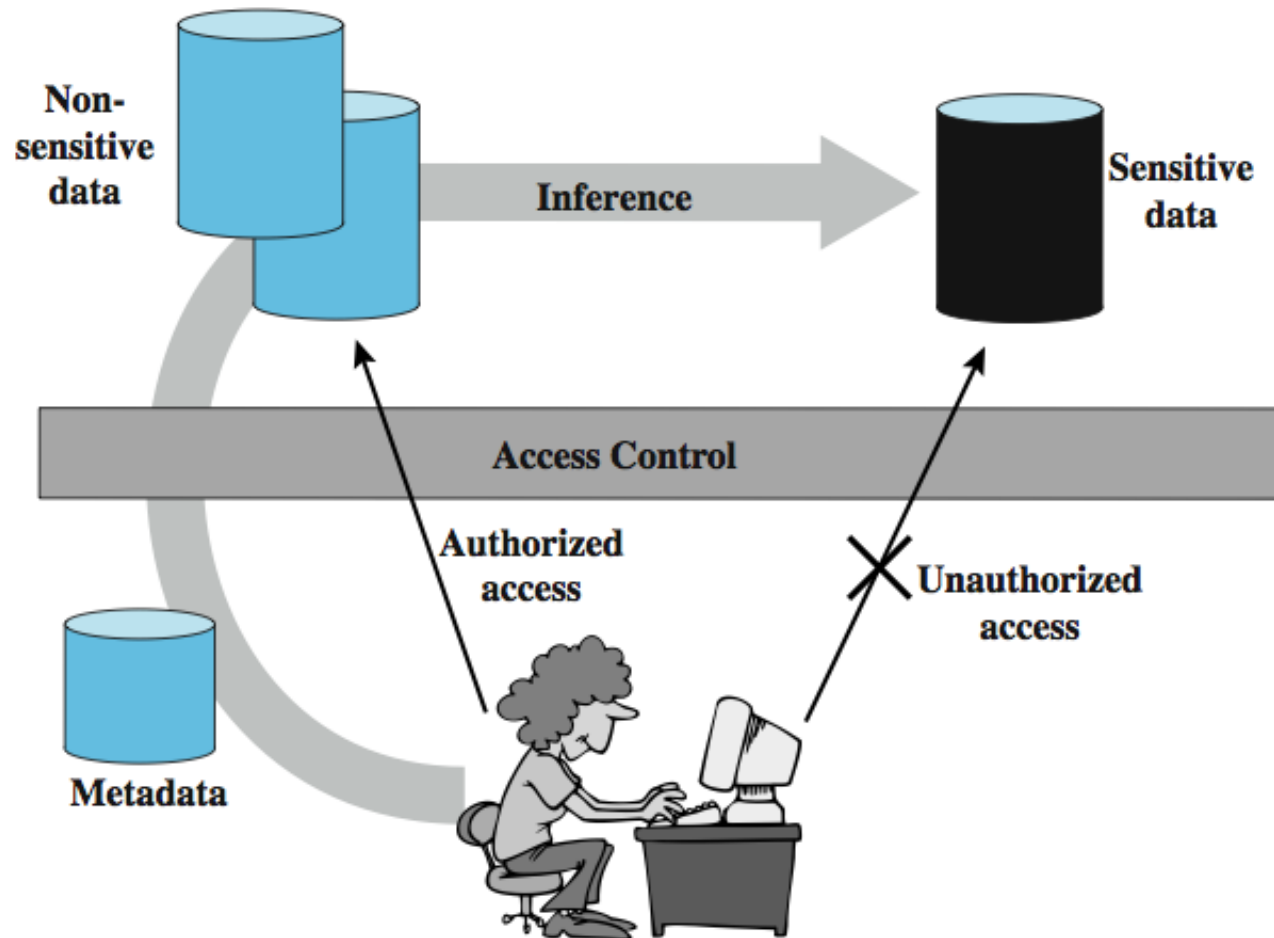


# Sensitive Data Disclosure

- Exact data: exact value of a sensitive data item
- Bounds: knowing the lower and upper bounds values of a sensitive data item
- Negative Result: queries may be made to determine a negative result from which sensitive data may be disclosed (e.g. Alice's salary is not lower than 150K/year).
- Existence: the existence of data is itself a sensitive piece of data
- Probable value: it may be possible to determine the probability that a certain element has a certain value.



# The Inference Problem



# The Inference Problem

- Inference problem is to infer or derive sensitive data from non-sensitive data.
- Sample table with 3 sensitive fields: Aid, Fines and Drugs when disclosed for a given individual

Name	Sex	Race	Aid	Fines	Drugs	Dorm
Adams	M	C	5000	45.	1	Holmes
Bailey	M	B	0	0.	0	Grey
Chin	F	A	3000	20.	0	West
Dewitt	M	B	1000	35.	3	Grey
Earhart	F	C	2000	95.	1	Holmes
Fein	F	C	1000	15.	0	West
Groff	M	C	4000	0.	3	West
Hill	F	B	5000	10.	2	Holmes
Koch	F	C	0	0.	1	West
Liu	F	A	0	10.	2	Grey
Majors	M	C	2000	0.	2	Grey



## Direct Attack

```
SELECT Name  
FROM Sample  
WHERE Sex='M' AND Drugs=1;
```

```
SELECT Name  
FROM Sample  
WHERE (Sex='M' AND Drugs=1)  
      OR (Sex<>'M' AND Sex<>'F')  
      OR Dorm='Ayres';
```

**To confuse  
the DB**



# Indirect Attack

- Statistical databases (queries)
- provides data of a statistical nature
  - e.g. counts, averages
- two types:
  - pure statistical database
  - ordinary database with statistical access
    - some users have normal access, others statistical
- access control objective to allow statistical use without revealing individual entries
- security problem is one of inference
- **Indirect attack - infer a result based on one or more statistical results.**





# Indirect Attack

- **Sum:** a reported sum may be used to infer a value.

```
SELECT Dorm, Sex, SUM(Aid) As Sum_Aid
FROM Sample
Group BY Dorm, Sex;
```

Dorm	Sex	Sum_Aid
Holmes	M	5000
Holmes	F	7000
Grey	M	3000
Grey	F	0
West	M	4000
West	F	4000

Name	Sex	Race	Aid	Fines	Drugs	Dorm
Adams	M	C	5000	45.	1	Holmes
Bailey	M	B	0	0.	0	Grey
Chin	F	A	3000	20.	0	West
Dewitt	M	B	1000	35.	3	Grey
Earhart	F	C	2000	95.	1	Holmes
Fein	F	C	1000	15.	0	West
Groff	M	C	4000	0.	3	West
Hill	F	B	5000	10.	2	Holmes
Koch	F	C	0	0.	1	West
Liu	F	A	0	10.	2	Grey
Majors	M	C	2000	0.	2	Grey



# Indirect Attack

- **Sum:** a reported sum may be used to infer a value.

```
SELECT Dorm, Sex, COUNT(*) As Tot_Students  
FROM Sample  
Group BY Dorm, Sex;
```


Dorm	Sex	Tot_Students
Holmes	M	1
Holmes	F	2
Grey	M	3
Grey	F	1
West	M	1
West	F	3

Name	Sex	Race	Aid	Fines	Drugs	Dorm
Adams	M	C	5000	45.	1	Holmes
Bailey	M	B	0	0.	0	Grey
Chin	F	A	3000	20.	0	West
Dewitt	M	B	1000	35.	3	Grey
Earhart	F	C	2000	95.	1	Holmes
Fein	F	C	1000	15.	0	West
Groff	M	C	4000	0.	3	West
Hill	F	B	5000	10.	2	Holmes
Koch	F	C	0	0.	1	West
Liu	F	A	0	10.	2	Grey
Majors	M	C	2000	0.	2	Grey



# Tracker Attacks

- divide queries into parts
  - $C = C1 \text{ AND } C2$
  - $\text{count}(C) = \text{count}(C1) - \text{count}(C1 \text{ AND } \sim C2)$
- combination is called a tracker
- each part acceptable query size
- overlap is desired result



# Tracker Attacks

- Tracker attacks: Generate the desired data by using additional queries that generate small results.

```
SELECT COUNT(*)  
FROM Sample  
WHERE Sex = 'F' AND Race = 'C' AND Dorm  
= 'Holmes';
```

- Use rules of logic and algebra to rewrite query:

```
count (a AND b AND c)  
= count(a) - count(a AND NOT (b AND c))  
= count(a) - count(a AND (NOT b OR NOT c))
```



# Tracker Attacks

```
SELECT COUNT(*)
```

```
FROM Sample
```

```
WHERE Sex = 'F';
```

```
SELECT COUNT(*)
```

```
FROM Sample
```

```
WHERE Sex = 'F' AND ((Race <> 'C') OR  
    (Dorm <> 'Holmes'));
```



# Controls for Inference Attacks

- Query controls
  - Limit overlap between new and previous queries
- Partitioning
  - Cluster records into exclusive groups and only allow queries on entire groups
- Item controls
  - Suppression: query is rejected without sensitive data provided.
  - Concealing: the answer provided is close to but not exactly the actual answer.



# Controls for Inference Attacks

- No perfect solution
- Three paths to follow:
  - Suppress obvious sensitive information (easily).
  - Track what the user knows (costly).
    - They are used to limit queries accepted and data provided.
  - Disguise the data (problem with precision).
    - It's applied only to the released data.

