

Introduction to OS Security

Bo Luo
bluo@ku.edu



Introduction

- OS: still software
 - All software security vulnerabilities still apply
- OS must protect users from each other
 - memory protection 存储保护
 - file protection
 - general control and access to objects
 - user authentication



Introduction

- The fundamental tradeoff of OS security
 - operating systems tradeoff between:
 - *Sharing*
 - *Protection*
- sharing is desirable
- protection is difficult



Introduction

- Early History
 - no OS
 - programs entered directly in binary through switches
 - user's program only one on system
每个软件用完一次，才能开下一个
 - user responsible for:
 - loading dependent libraries, other tools
 - scheduling time to use computer
 - OS security?



Introduction

- Later
 - machines very expensive
 - people less expensive
 - maximize use of machine
 - allow many users



Introduction

- OS protection – separation.
- Physical separation, e.g. 1 user/printer
- Temporal separation 时间分离
- Logical separation: user thinks own machine
- Cryptographic separation
- Combinations of these



Levels of Protection

- no protection
- isolation
- share all or nothing
- share via access limitation
- share by capabilities
- limit use of an object



Levels of Protection

- No protection
 - *e.g.* early versions of windows
 - some embedded environments
 - designed for one user
 - no need for isolation, access control, etc.




Uniprogramming w/o memory protection

- Simplest.
 - Each application runs within a hardwired range of physical memory addresses
- One application runs at a time
 - Application can use the same physical addresses every time, across reboots



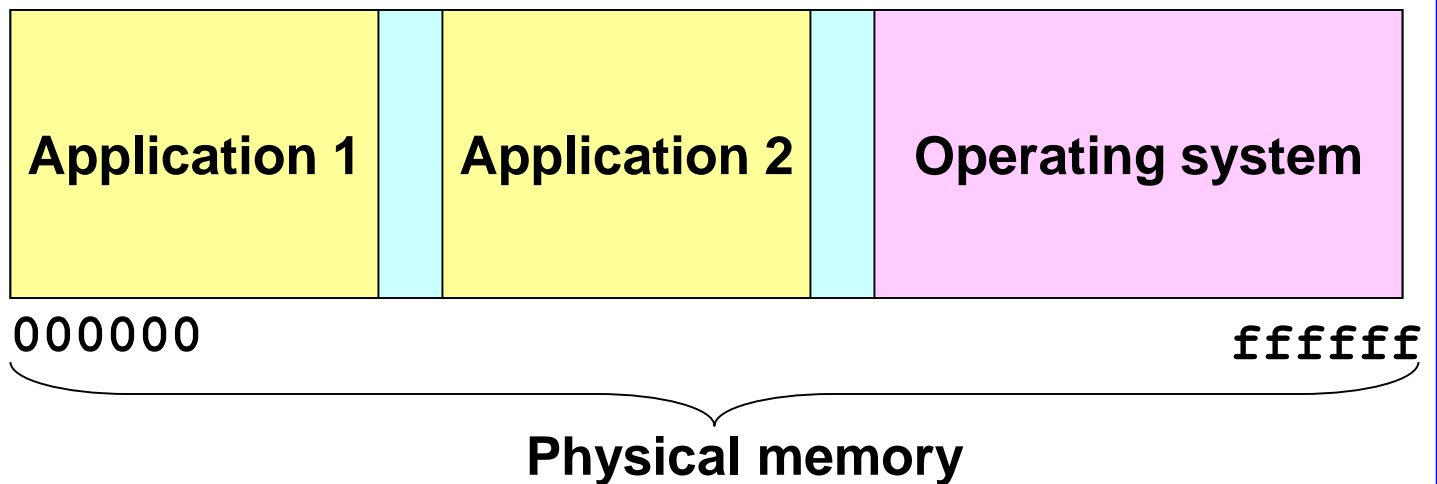
Uniprogramming w/o memory protection

- Applications typically use the lower memory addresses
- An OS uses the higher memory addresses
- An application can address any physical memory location



Levels of Protection

- Isolation
 - processes unaware of other processes
 - each process: own address space, files, *etc.*
 - OS provides confinement
 - Virtual machines



Levels of Protection

- Share all or nothing
 - owner of object declares it: share 所有人
 - *Public*: available to all users
 - *Private*: not available



Levels of Protection

- Share via access limitation
 - Resource/files are shared
 - Who can access what?
 - Access control lists
 - Access Control Matrices
 - Capabilities



Levels of Protection

- Limit use of an object 权限限制更小
 - Sophisticated, fine-grained access control
 - Examples:
 - can view a file, but can't print
 - given aggregate info from database, but not individual records




Memory and Access Protection

- In a multi-user multi-task environment, what's in the memory?
 - OS, processes from different users
 - Data (keys!) in plaintext!
- Memory management
 - Fences
 - Relocation
 - Base/Bounds Registers
 - Tagged Architecture
 - Segmentation
 - Paging
 - Combined Paging with Segmentation



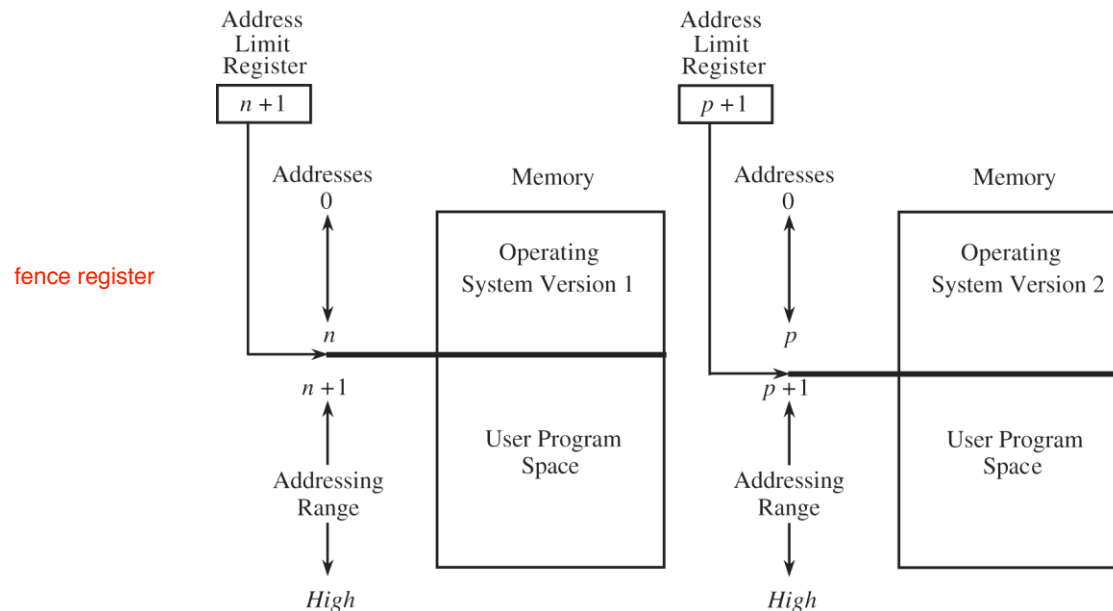
Memory and Access Protection

- Fences: protect OS from user program
 - confine users to one side of a boundary
 - predefined memory address: user code on one side, OS on the other



Memory and Access Protection

- Fences: protect OS from user program
 - Problem?
 - fixed boundary too restrictive
 - doesn't protect users from each other
 - moveable fence:
 - store fence location in fence register



Memory and Access Protection

- Relocation address 不连续性
 - programs written to run starting at address 0
 - can be run at any address
 - addresses in source are symbolic:
 - e.g., numStudents
 - compiler binds these to relocatable addresses.
 - e.g. 20 bytes from beginning of module func
 - then linker or loader binds to absolute addresses
 - e.g. 20114
 - logical addresses mapped to physical by MMU
 - program never sees real addresses



Memory and Access Protection

- We skip the details of segmentation and paging.
- They have been covered in your OS class.



Memory and Access Protection

- Wrap-up
 - Each process
 - has its own address space
 - thinks it's the only process on machine
 - MMU provides translation between process's address space and physical space
 - process cannot generate address not in its own space



Memory and Access Protection

- Relocation
 - Offer protection if the translation tables cannot be altered by applications
 - An application can only touch its address space under the user mode
 - Hardware requires the CPU to be in the kernel mode to modify the address translation tables



Memory and Access Protection

- Switching from Kernel mode to User mode
 - To run a user program, the kernel:
 - Creates a process and initialize the address space
 - Loads the program into the memory
 - Initializes translation tables
 - Sets the hardware pointer to the translation table
 - Sets the CPU to user mode
 - Jumps to the entry point of the program



Memory and Access Protection

- Switching from User mode to Kernel mode
- Voluntary
 - System calls: a user process asks the OS to do something on the process's behalf
- Involuntary
 - Hardware interrupts (e.g., I/O)
 - Program exceptions (e.g., segmentation fault)



Memory and Access Protection

- For all cases, hardware atomically performs the following steps
 - Sets the CPU to kernel mode
 - Saves the current program counter
 - Jumps to the handler in the kernel
 - The handler saves old register values



Memory and Access Protection

- Context switching between processes
 - Need to save and restore pointers to translation tables
- To resume process execution
 - Kernel reloads old register values
 - Sets CPU to user mode
 - Jumps to the old program counter



Memory Attacks

- Motivation
 - You can encrypt HD, USB drives, network traffic, etc.
 - You cannot encrypt memory!
 - A lot of sensitive information: keys, passwords, etc.
- Two categories of memory attacks
 - Software attacks
 - Hardware attacks



Memory Attacks

- Software attacks
 - System bugs: allow a process to read any address
 - Swap and dump: memory contents are written to hard drives: swap, core dump, hibernation, crash reports, etc
 - Attackers: trigger a core dump and examine the dump file, looking for keys.
 - It has been reported that core dumps of FTP servers and email servers contained passwords.
 - Hypervisors: suspend the current state of a VM to a check point file.
 - Uncleared buffers

交换和储存



Memory Attacks

- Physical attacks
 - Bypass OS, bypass CPU
 - Directly read from RAM
 - Cold boot attacks
 - The remanence effect of RAM: the contents in RAM fade away gradually after power off in several minutes or hours (low temperature)
 - Attacker: reboot the computer with OS from USB drive; or move the RAM chips to another machine
 - Read from RAM – reduce the temperature so that data stays longer
 - Completely bypass access control, encryption, task isolation, authentication, etc
 - DMA: direct memory access



Control of Access to General Objects

- control of any kind of object
- examples:
 - memory
 - secondary storage
 - hardware devices
 - some data structure
 - instructions
 - passwords and user-authentication mechanism
 - the protection mechanism itself



Control of Access to General Objects

- Goals in protecting objects
 - check every access
 - user permitted doesn't mean always permitted
 - enforce least privilege
 - grant access to minimum set of objects required to complete a task
 - verify acceptable usage
 - stack: push(), pop(), ...
 - Shouldn't be able to do anything else to stack

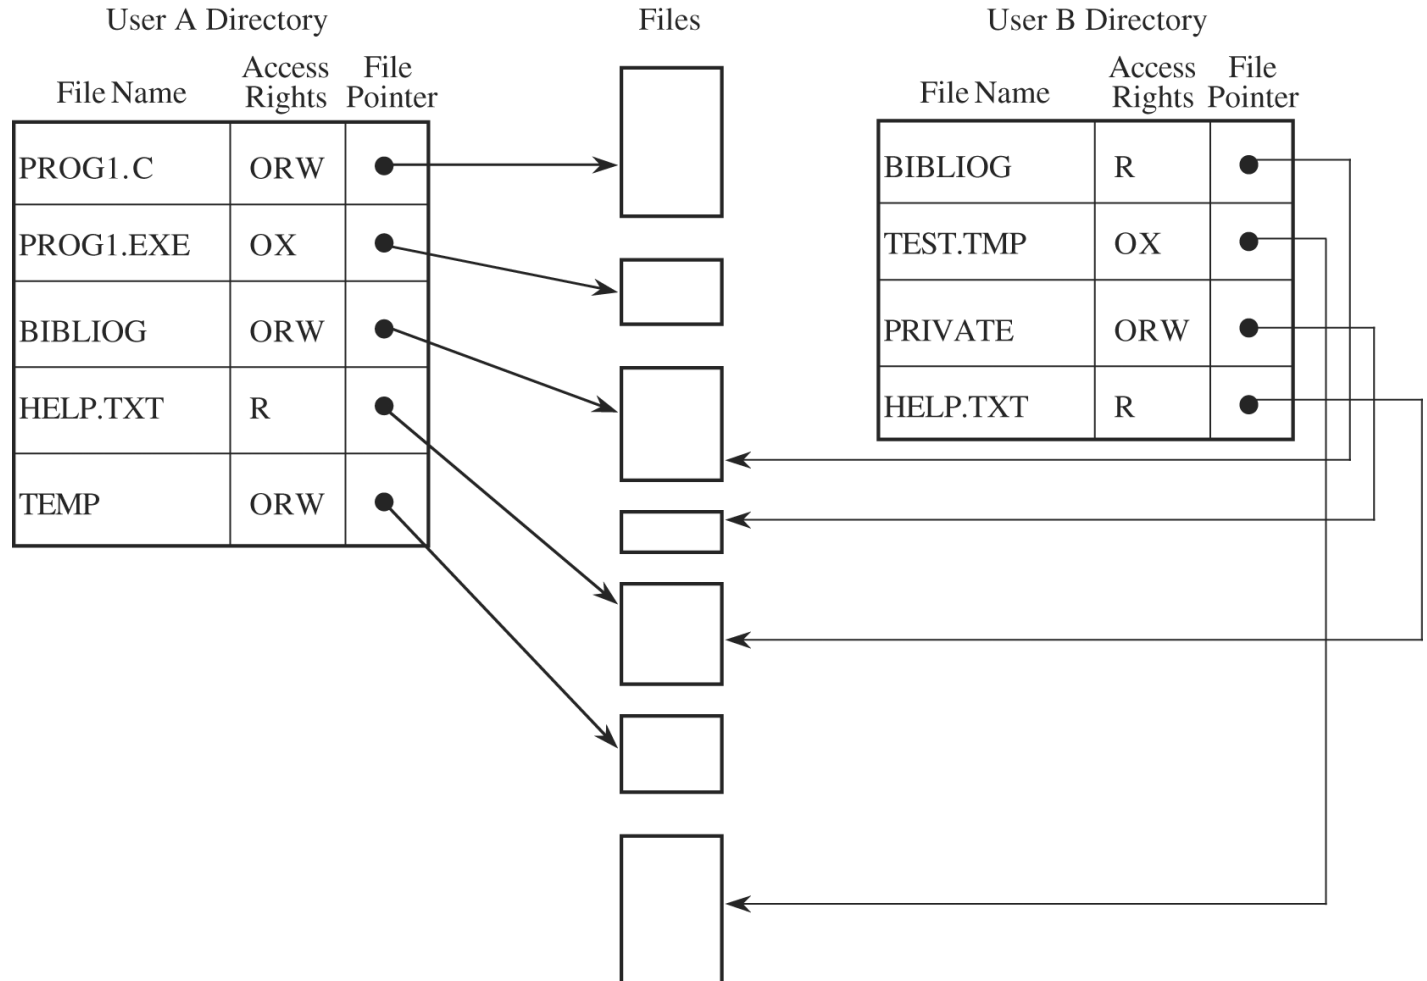


Control of Access to General Objects

- Directory
 - not directory as in FS directory
 - each user has a list (directory) of objects the user owns or has access to
 - no user should be able to write to the directory
 - for each file, directory contains list of permissions, e.g. R, W, X, and owner



Control of Access to General Objects



Control of Access to General Objects

- Directory
 - simple *but*
 - lists can get very long
 - what about shared libraries, programs?
 - same item in many lists
 - revoking permissions?
 - have to go through everyone's lists

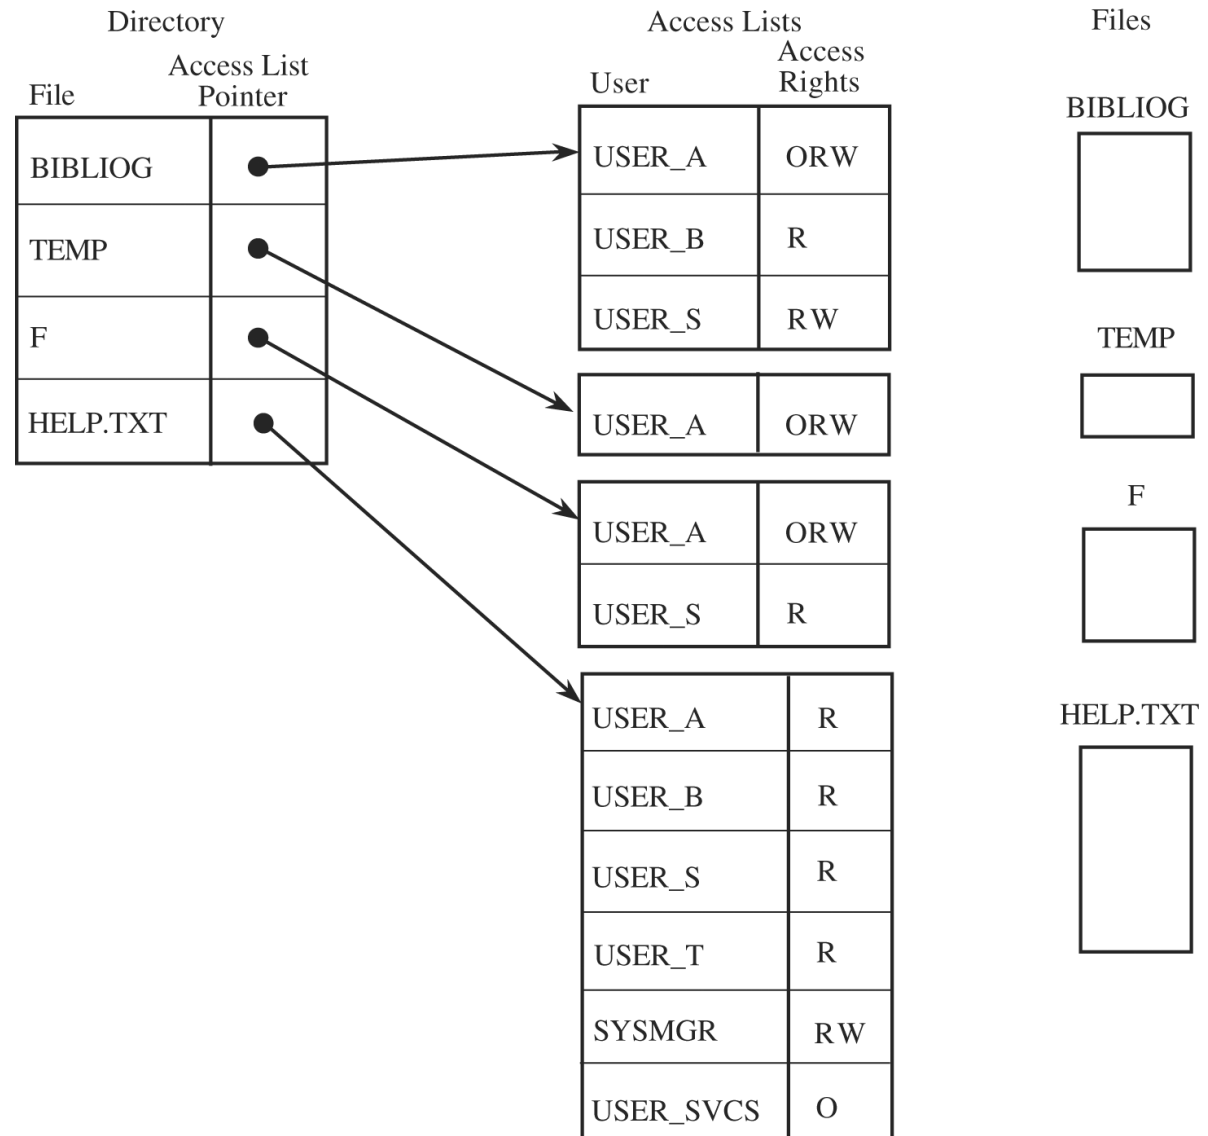


Control of Access to General Objects

- Access control lists
 - maintain a list per object, not user
 - use wildcards (*) to grant permission to a group
 - e.g. administrator-*



Control of Access to General Objects



Control of Access to General Objects

- Access control matrix
 - row for each user
 - column for each protected object
 - simple lookups
 - but probably lots of empty spaces

	BIBLIOG	TEMP	F	HELP.TXT	C_COMP	LINKER	SYS_CLOCK	PRINTER
USER A	ORW	ORW	ORW	R	X	X	R	W
USER B	R	-	-	R	X	X	R	W
USER S	RW	-	R	R	X	X	R	W
USER T	-	-	-	R	X	X	R	W
SYS_MGR	-	-	-	RW	OX	OX	ORW	O
USER_SVCS	-	-	-	O	X	X	R	W

