

# EECS 645

# Computer Architecture

## Lecture 9 Memory Hierarchy Design – Part 1

Chenyun Pan

Department of Electrical Engineering &

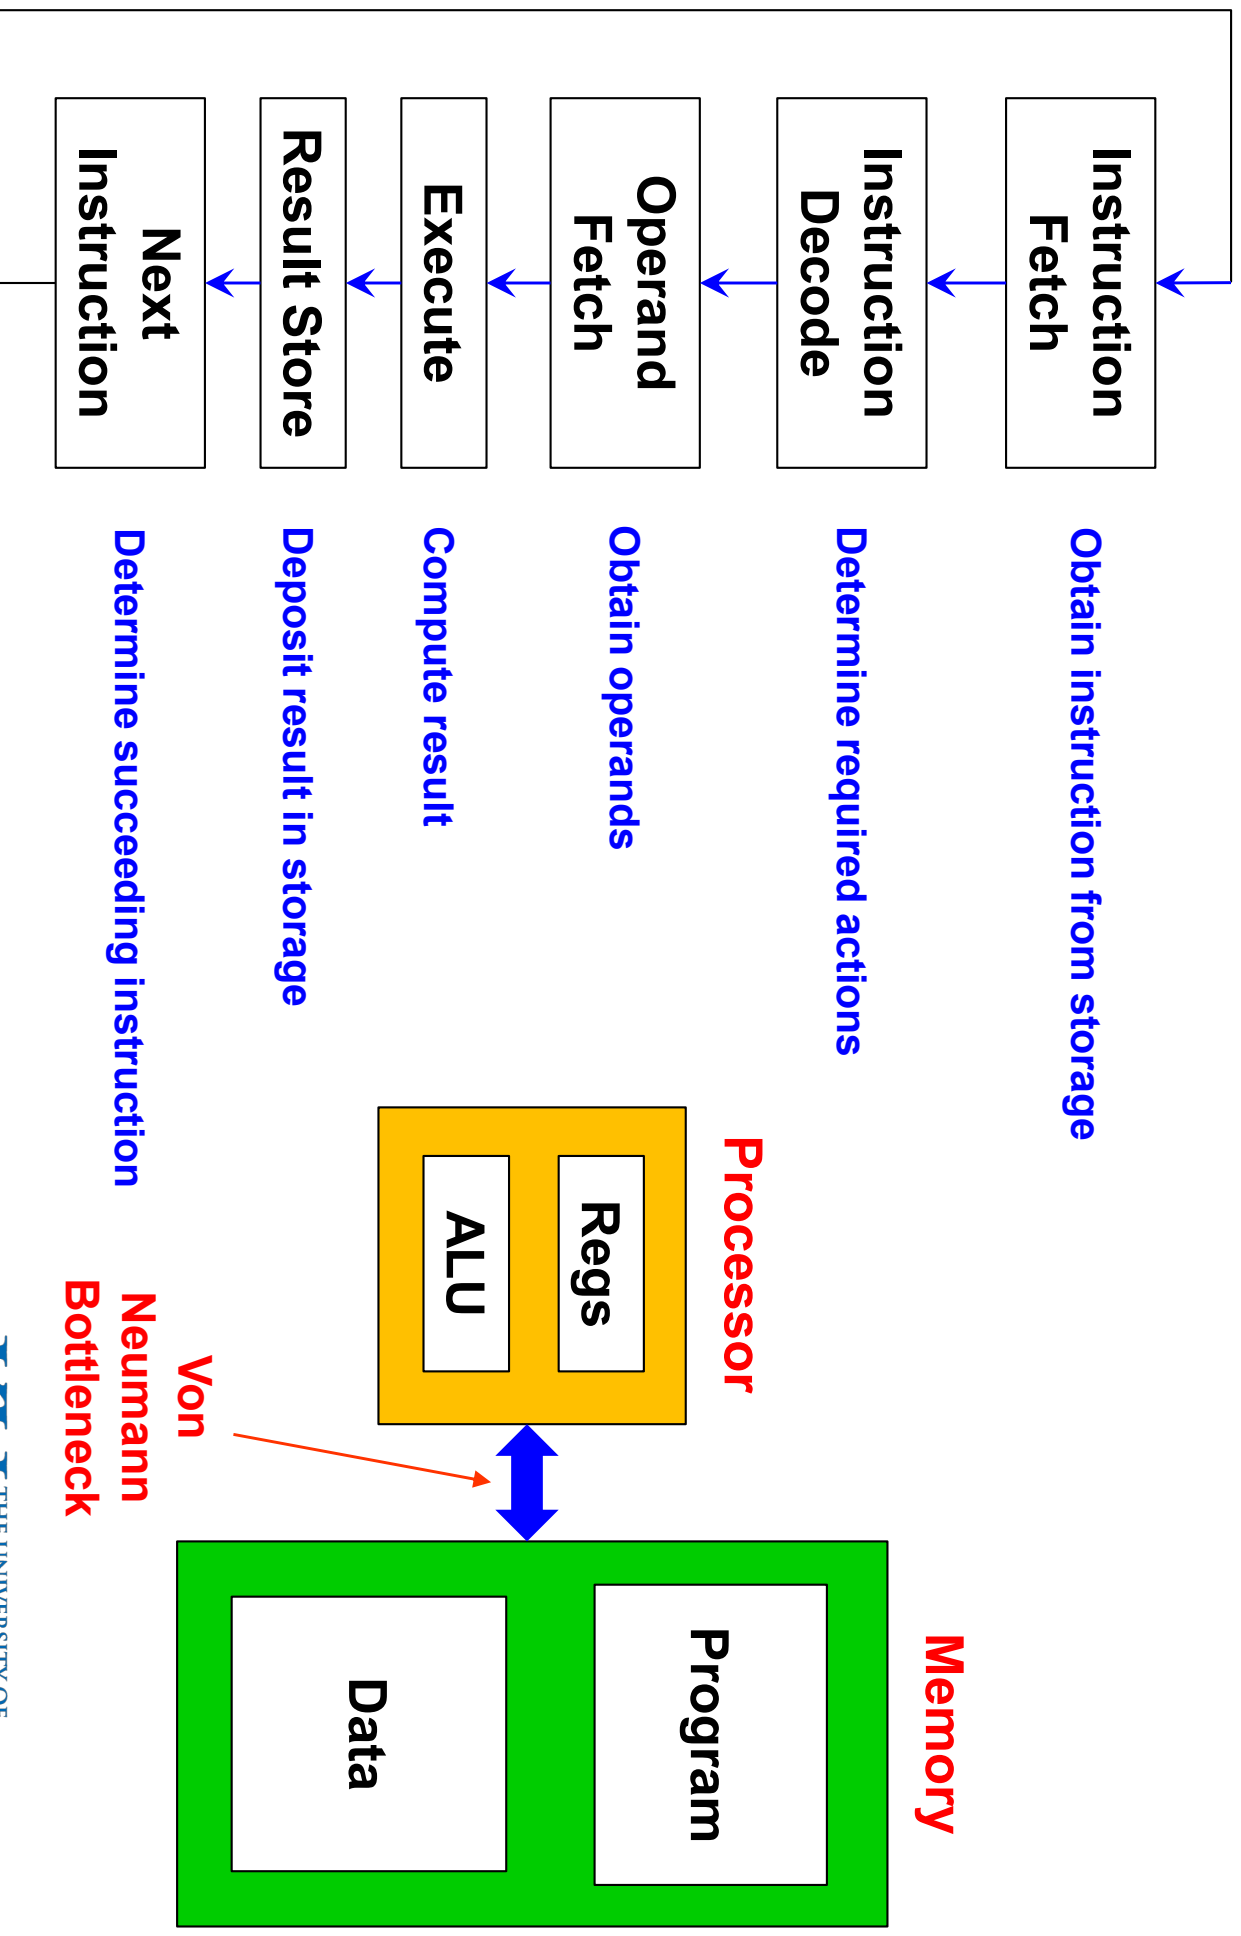
Computer Science

University of Kansas



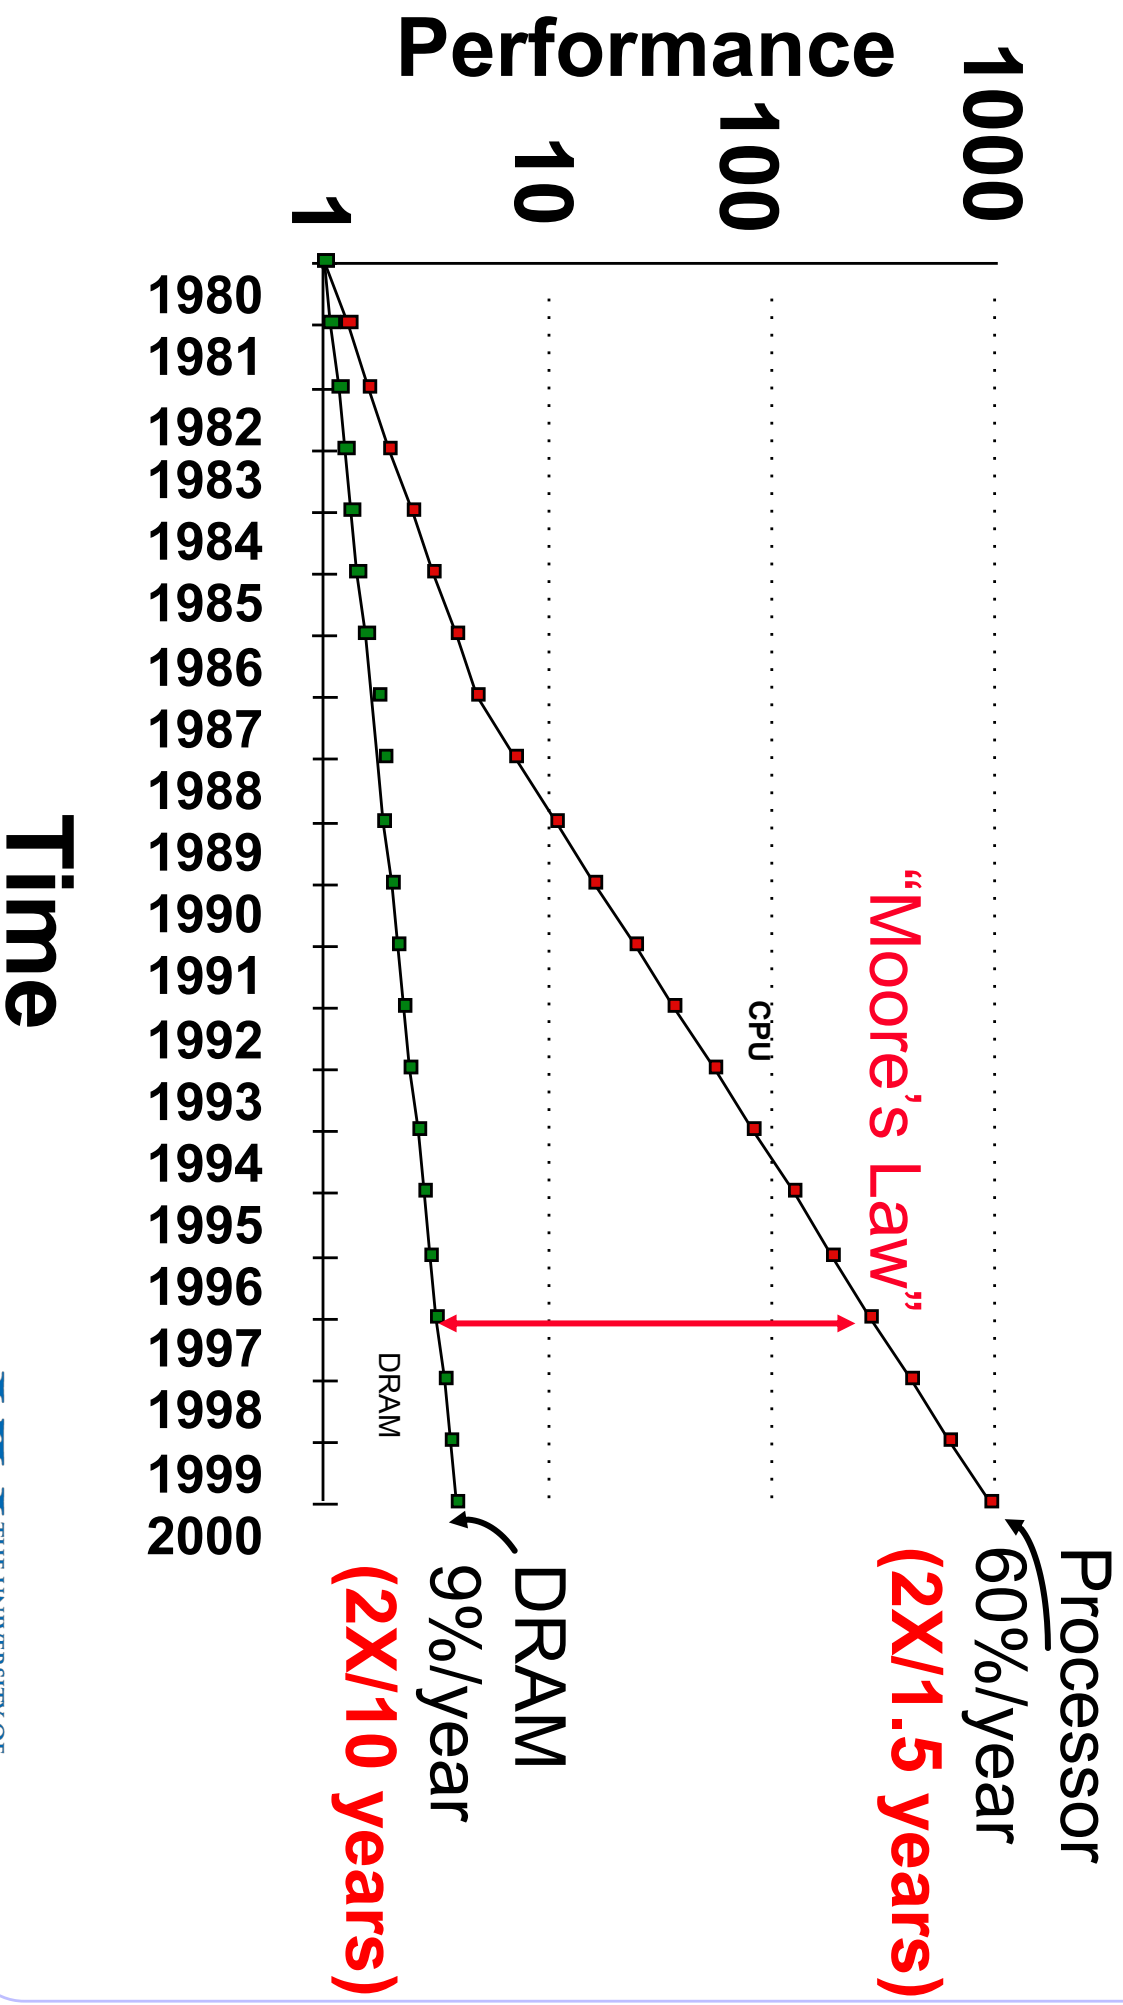
Slide Courtesy of Dr. Hsien-Hsin Sean Lee

# Execution Style of von Neumann Machines



# Why Care About Memory Hierarchy?

Processor-DRAM Performance Gap grows 50% / year



# Energy Comparison

Energy table for 45 nm CMOS process [1]

Operation	Energy [pJ]	Relative Cost
32 bit int ADD	0.1	1
32 bit float ADD	0.9	9
32 bit int MULT	3.1	31
32 bit float MULT	3.7	37
<b>32 bit DRAM</b>	<b>640</b>	<b>6400</b>

DRAM access uses 100x – 1000x more energy than arithmetic operations

[1] M. Horowitz. Energy table for 45nm process, Stanford VLSI wiki.

# An Unbalanced System




Source: Bob Colwell keynote ISCA'29 2002

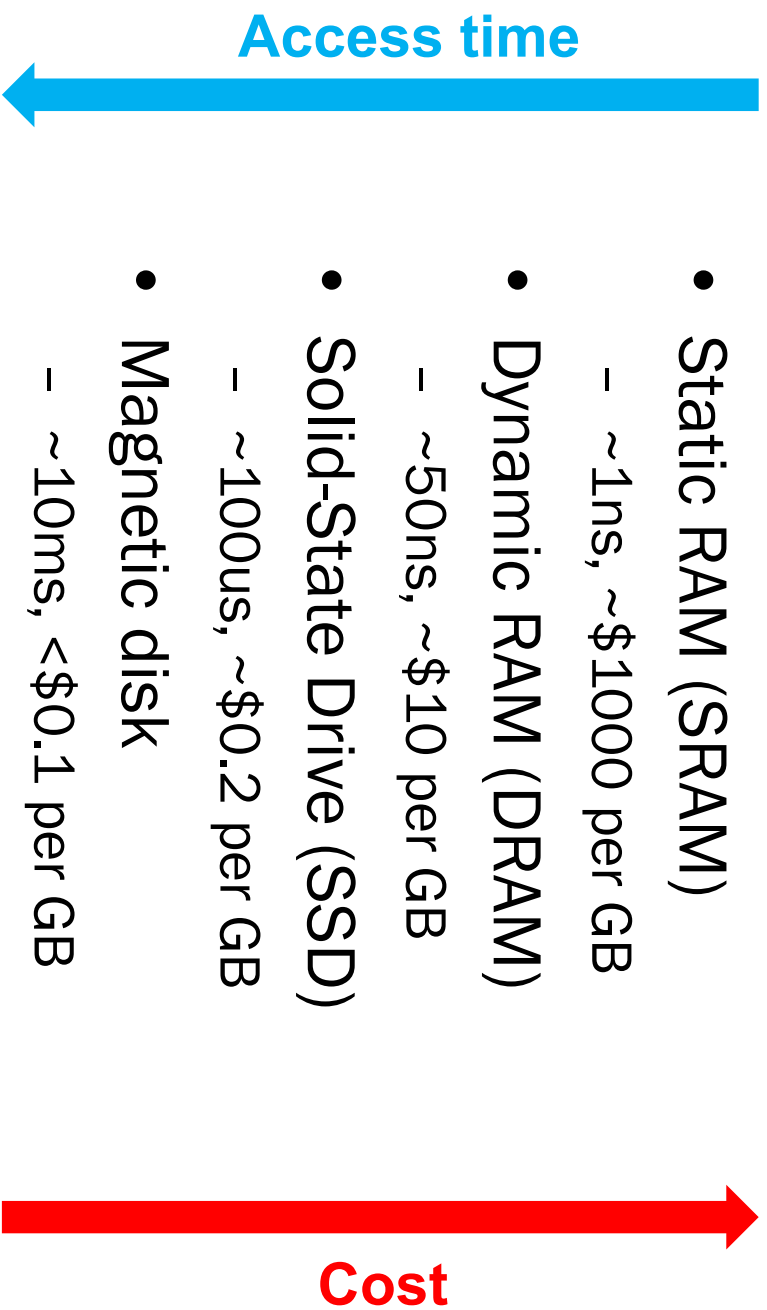
# Memory Issues

- Latency
  - Time to move through the longest circuit path (from the start of request to the response)
- Bandwidth
  - Number of bits transported at one time
- Capacity
  - Size of memory
- Energy
  - Cost of accessing memory (to read and write)

# Model of Memory Hierarchy



# Memory Technology



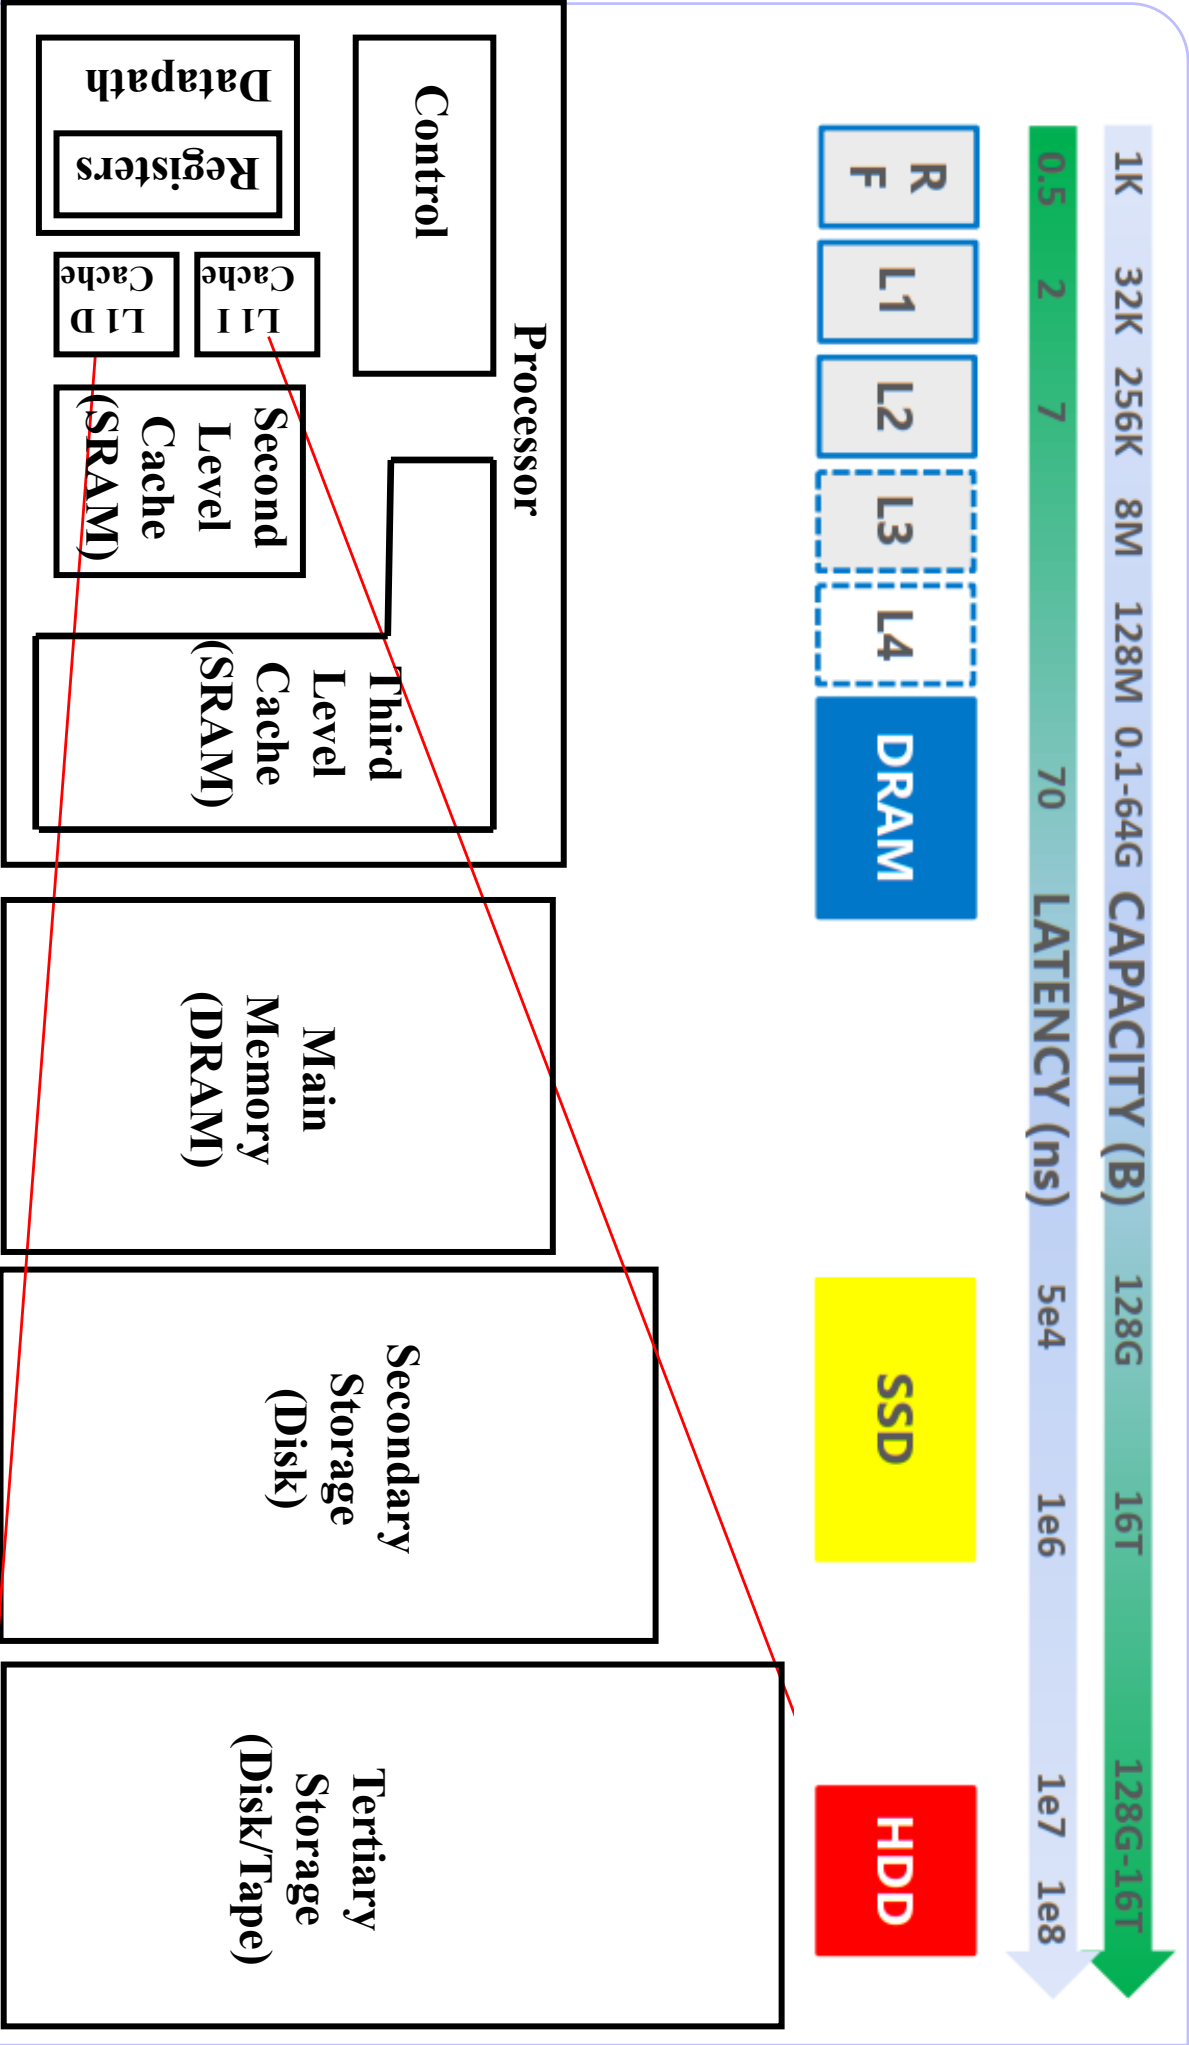
- Static RAM (SRAM)
  - ~1ns, ~\$1000 per GB
- Dynamic RAM (DRAM)
  - ~50ns, ~\$10 per GB
- Solid-State Drive (SSD)
  - ~100us, ~\$0.2 per GB
- Magnetic disk
  - ~10ms, <\$0.1 per GB

## Ideal memory


- Access time of SRAM
- Capacity and cost/GB of disk



# Modern Memory Hierarchy



# Topics covered



- Why do caches work
  - Principle of **program locality**
- Cache hierarchy
  - Average memory access time ( $T_{\text{access}}$ )
- Types of caches
  - Direct mapped
  - Set-associative
  - Fully associative
- Cache policies
  - Write back vs. write through
  - Write allocate vs. No write allocate

# Principle of Locality

- Programs access a relatively small portion of address space at any instant of time.
- Two Types of Locality:
  - Temporal Locality (Locality in Time): If an address is referenced, it tends to be referenced again
    - e.g., loops, reuse
  - Spatial Locality (Locality in Space): If an address is referenced, neighboring addresses tend to be referenced
    - e.g., array access

**Locality is a program property that is exploited in machine design.**

# Example of Locality

```

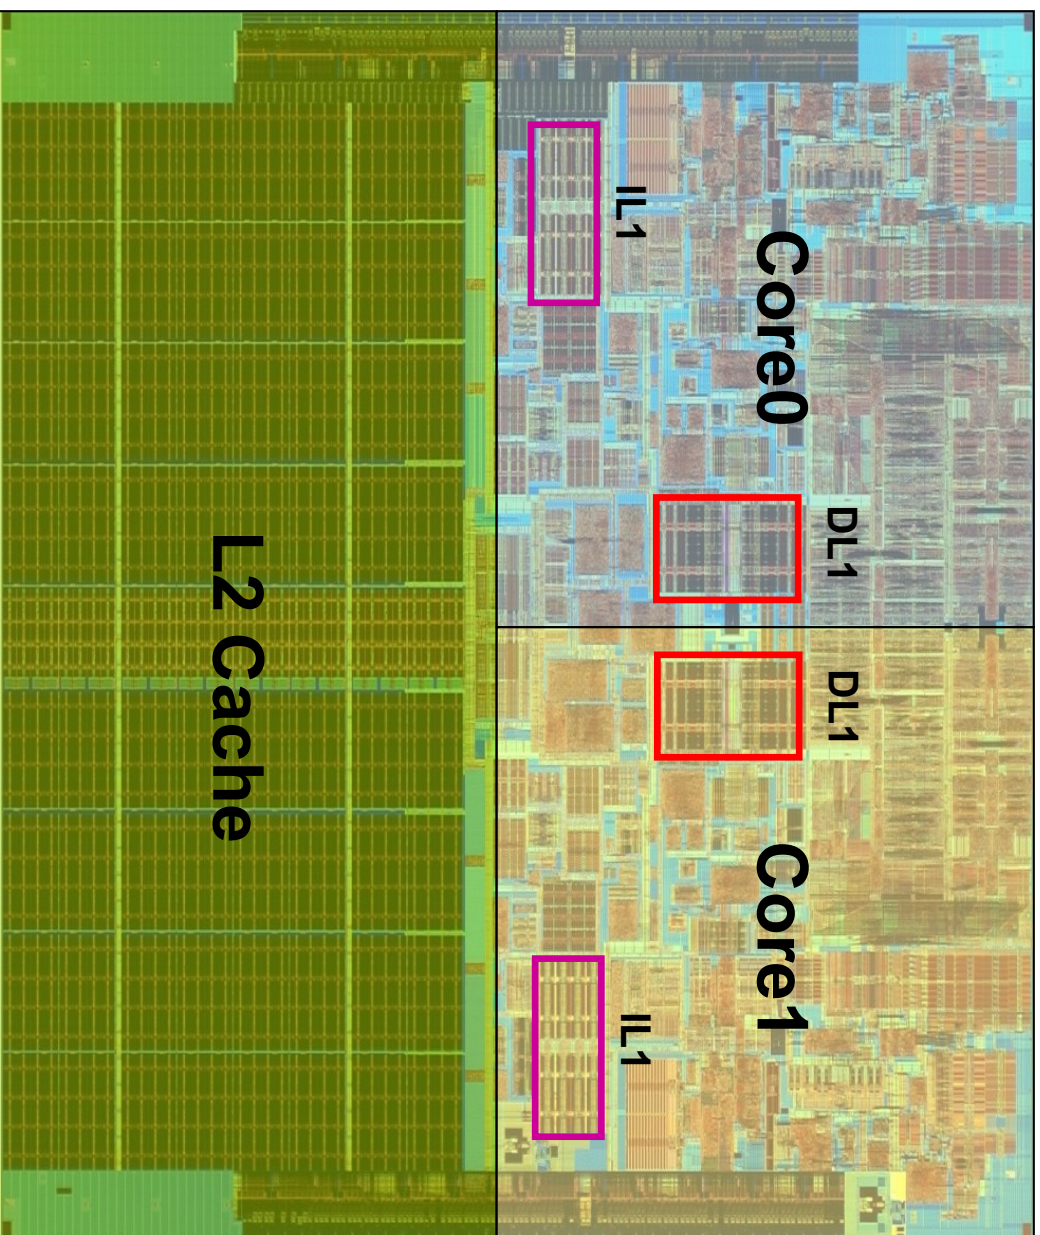
int A[100], B[100], C[100], D;
for (i=0; i<100; i++) {
    C[i] = A[i] * B[i] + D;
}
    
```

				D	C[99]	C[98]	C[97]	C[96]
C[7]	C[6]	C[5]	C[4]	C[3]	C[2]	C[1]	C[0]	
B[11]	B[10]	B[9]	B[8]	B[7]	B[6]	B[5]	B[4]	
B[3]	B[2]	B[1]	B[0]	A[99]	A[98]	A[97]	A[96]	
A[7]	A[6]	A[5]	A[4]	A[3]	A[2]	A[1]	A[0]	


 A Cache Line/Block (One fetch)

# Example: Intel Core2 Duo

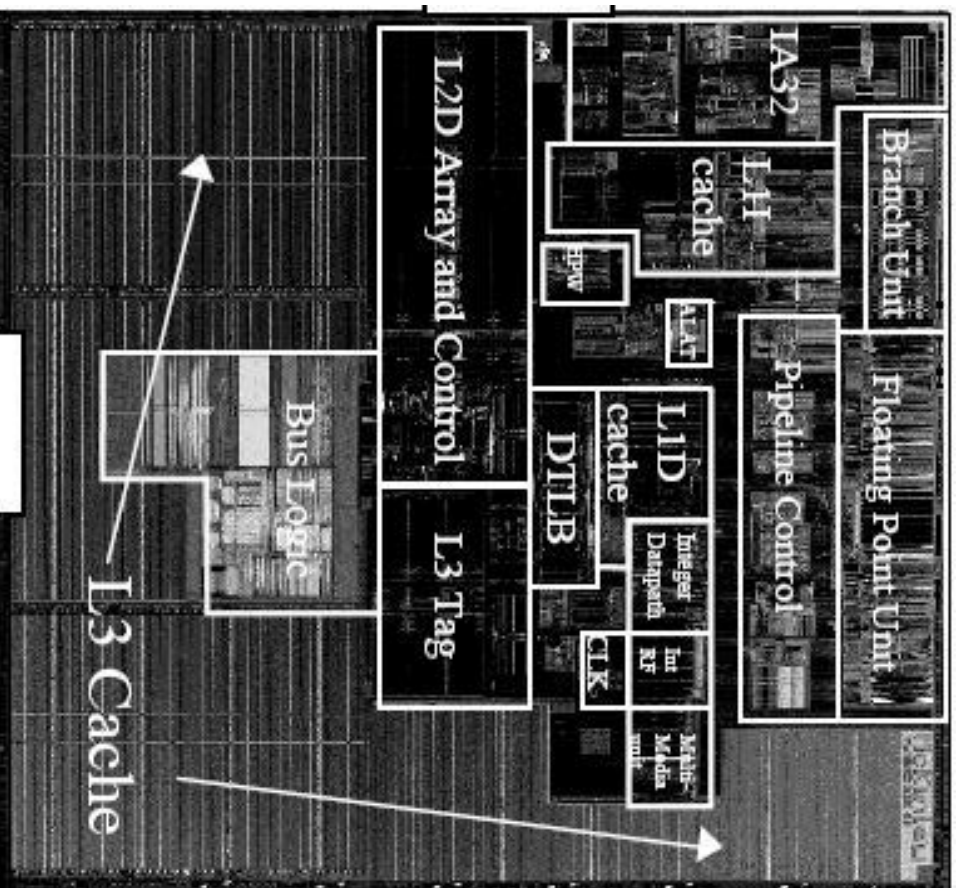
L1	32 KB, 8-Way, 64 Byte/Line, LRU, WB 3 Cycle Latency
L2	4.0 MB, 16-Way, 64 Byte/Line, LRU, WB 14 Cycle Latency



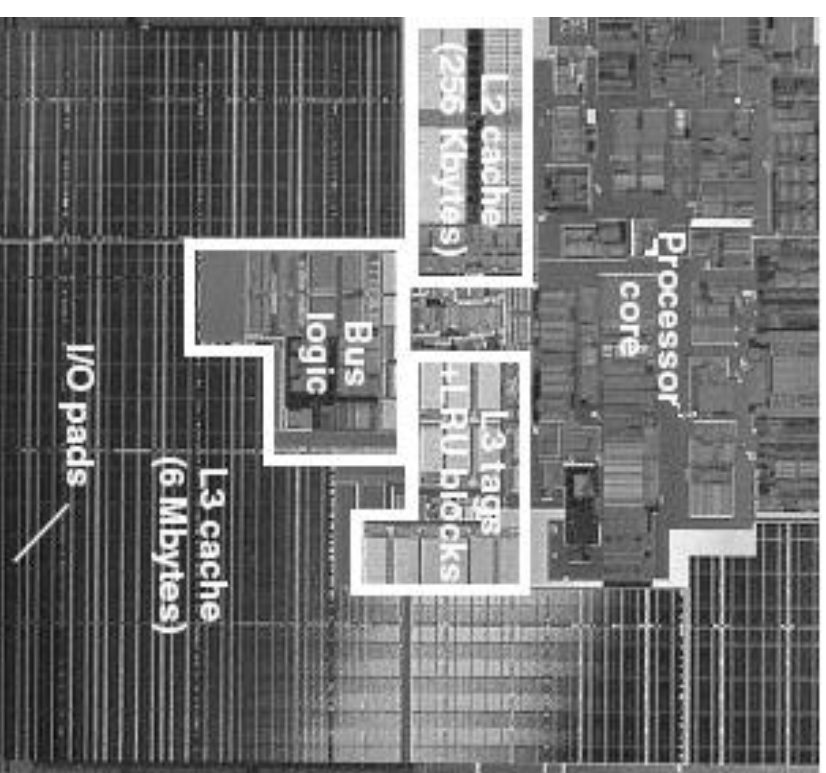
Source: <http://www.sandpile.org>



# Example : Intel Itanium 2




**3MB**  
**Version**  
**180nm**  
**421 mm<sup>2</sup>**



**6MB**  
**Version**  
**130nm**  
**374 mm<sup>2</sup>**

# Intel Nehalem




24MB L3



# Example : STI Cell Processor


Local Storage



**SPE = 21M transistors (14M array; 7M logic)**




# Topics covered



- Why do caches work
  - Principle of **program locality**
- Cache hierarchy
  - Average memory access time ( $T_{\text{access}}$ )
- Types of caches
  - Direct mapped
  - Set-associative
  - Fully associative
- Cache policies
  - Write back vs. write through
  - Write allocate vs. No write allocate

# Cache Terminology

- **Cache Line/Block**: amount of data being transferred at a time
- **Hit**: data appears in some line/block
  - **Hit Rate**: the fraction of memory accesses found in the level
  - **Hit Time**: Time to access the level (including the time to check if hit)
- **Miss**: data needs to be retrieved from a line/block in the lower level memory (e.g. DRAM)
  - **Miss Rate** =  $1 - (\text{Hit Rate})$
  - **Miss Penalty**: Time to replace a **line/block** in the upper level cache + Time to deliver the line/block to the processor
- **Hit Time** should be much shorter than **Miss Penalty**



# Memory Hierarchy Performance



How to estimate the Average Memory Access Time?

$$T_{\text{access}} = \text{Hit Time} + \text{Miss rate} * \text{Miss Penalty} \\ = T_{\text{hit}}(\text{L1}) + \text{Miss}_{\text{L1}} * T_{\text{mem}}$$

- Example:
  - Cache Hit = 1 cycle
  - Miss rate = 10%
  - Miss penalty = 300 cycles
  - Access time =  $1 + 0.1 * 300 = 31$  cycles
  - **9.7X Speed Up!**
- How to further improve it?

# Reducing Penalty: Multi-Level Cache

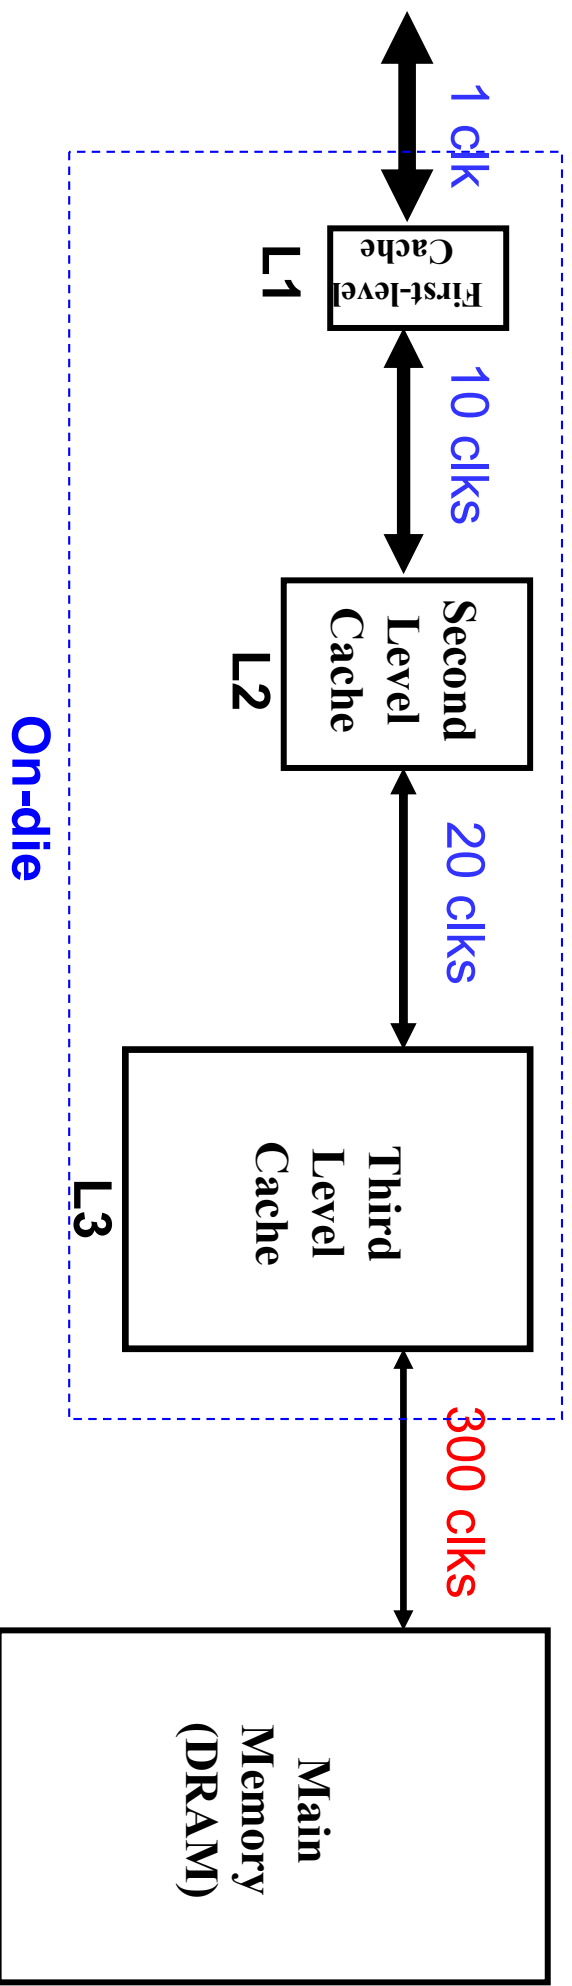


What is the Average Memory Access Time ?

$$T_{\text{access}} = T_{L1} + \text{Miss}_{L1} \cdot (T_{L2} + \text{Miss}_{L2} \cdot T_{\text{mem}})$$

A **larger** cache has a **lower** miss rate but a **longer** hit time.

# Reducing Penalty: Multi-Level Cache



Average Memory Access Time

$$T_{\text{access}} = T_{L1} + \text{Miss}_{L1} \cdot [T_{L2} + \text{Miss}_{L2} \cdot (T_{L3} + \text{Miss}_{L3} \cdot T_{\text{mem}})]$$

# $T_{\text{access}}$ Example

- Example:
  - Miss rate L1=10%,  $T_{\text{hit}}(\text{L1}) = 1$  cycle
  - Miss rate L2=5%,  $T_{\text{hit}}(\text{L2}) = 10$  cycles
  - Miss rate L3=1%,  $T_{\text{hit}}(\text{L3}) = 20$  cycles
  - $T(\text{memory}) = 300$  cycles
- $T_{\text{access}} = ?$ 
  - 2.115 (compare to 31 with only L1 cache)

**14.7x speed-up!**

# Topics covered

- Why do caches work
  - Principle of **program locality**
- Cache hierarchy
  - Average memory access time ( $T_{\text{access}}$ )
- Types of caches
  - Direct mapped
  - Set-associative
  - Fully associative
- Cache policies
  - Write back vs. write through
  - Write allocate vs. No write allocate

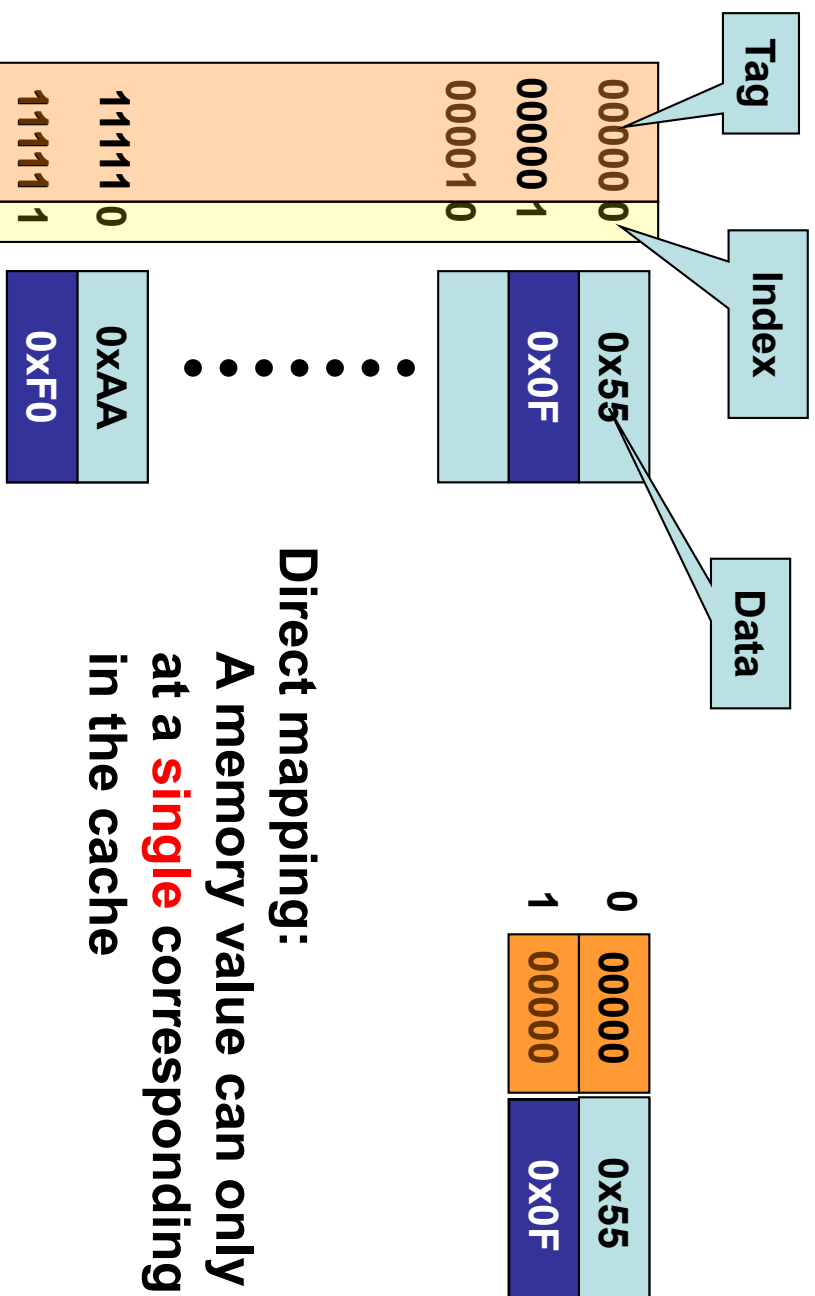


# Types of Caches

Type of cache	Mapping of data from memory to cache	Complexity of searching the cache
Direct mapped (DM)	A memory value can be placed at <b>a single corresponding location</b> in the cache	<b>Fast</b> indexing mechanism
Set-associative (SA)	A memory value can be placed in <b>different locations of a set</b> in the cache	Slightly more involved search mechanism
Fully-associative (FA)	A memory value can be placed in <b>any location</b> in the cache	<b>Extensive</b> hardware resources required to search (CAM)



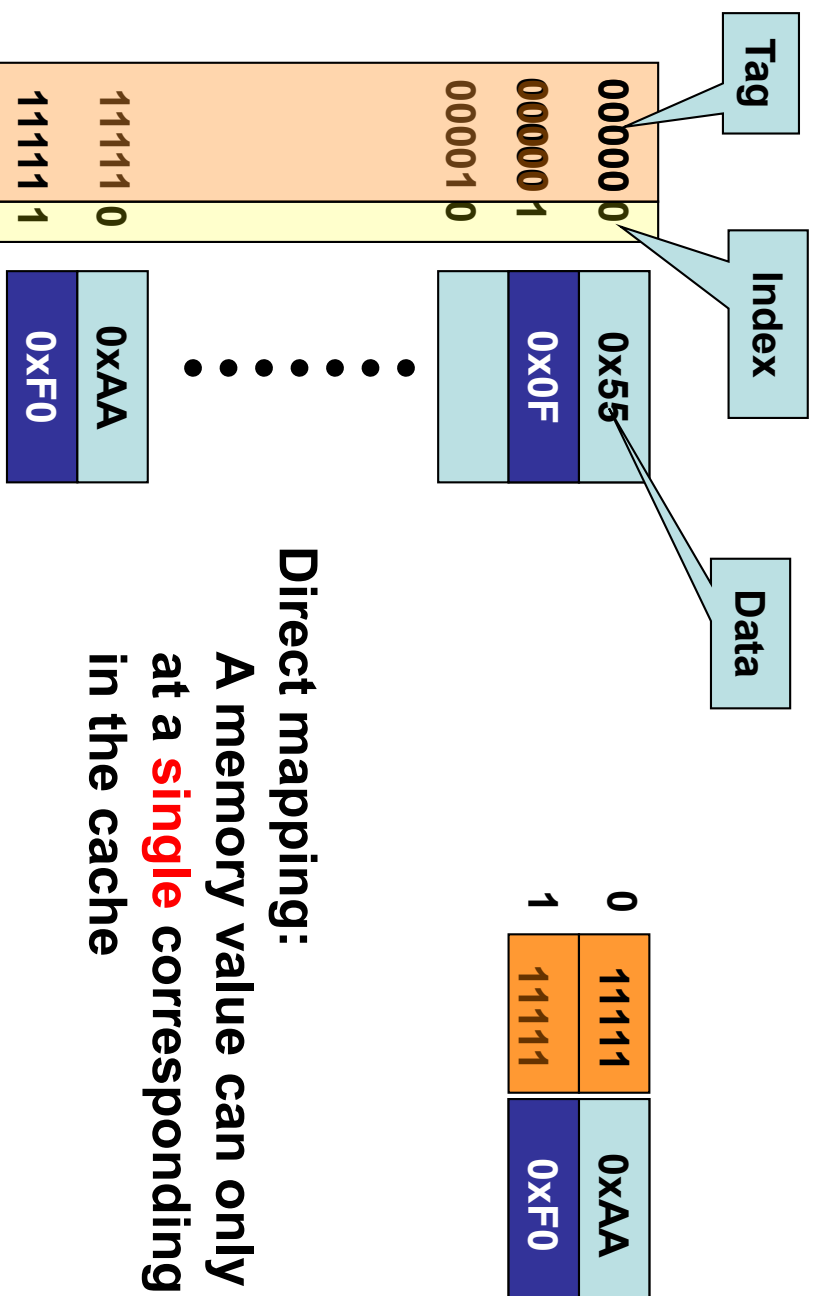
# Direct Mapping



**Direct mapping:**

A memory value can only be placed at a **single** corresponding location in the cache


# Direct Mapping



**Direct mapping:**


A memory value can only be placed at a **single** corresponding location in the cache

# Address Subdivision



Any drawback of direct map cache?

# Set Associative Mapping (2-Way)




	Way 0		Way 1	
0	0000 0	0x55	1111 0	0xAA
1	0000 1	0x0F	1111 1	0xF0

## Set-associative mapping:

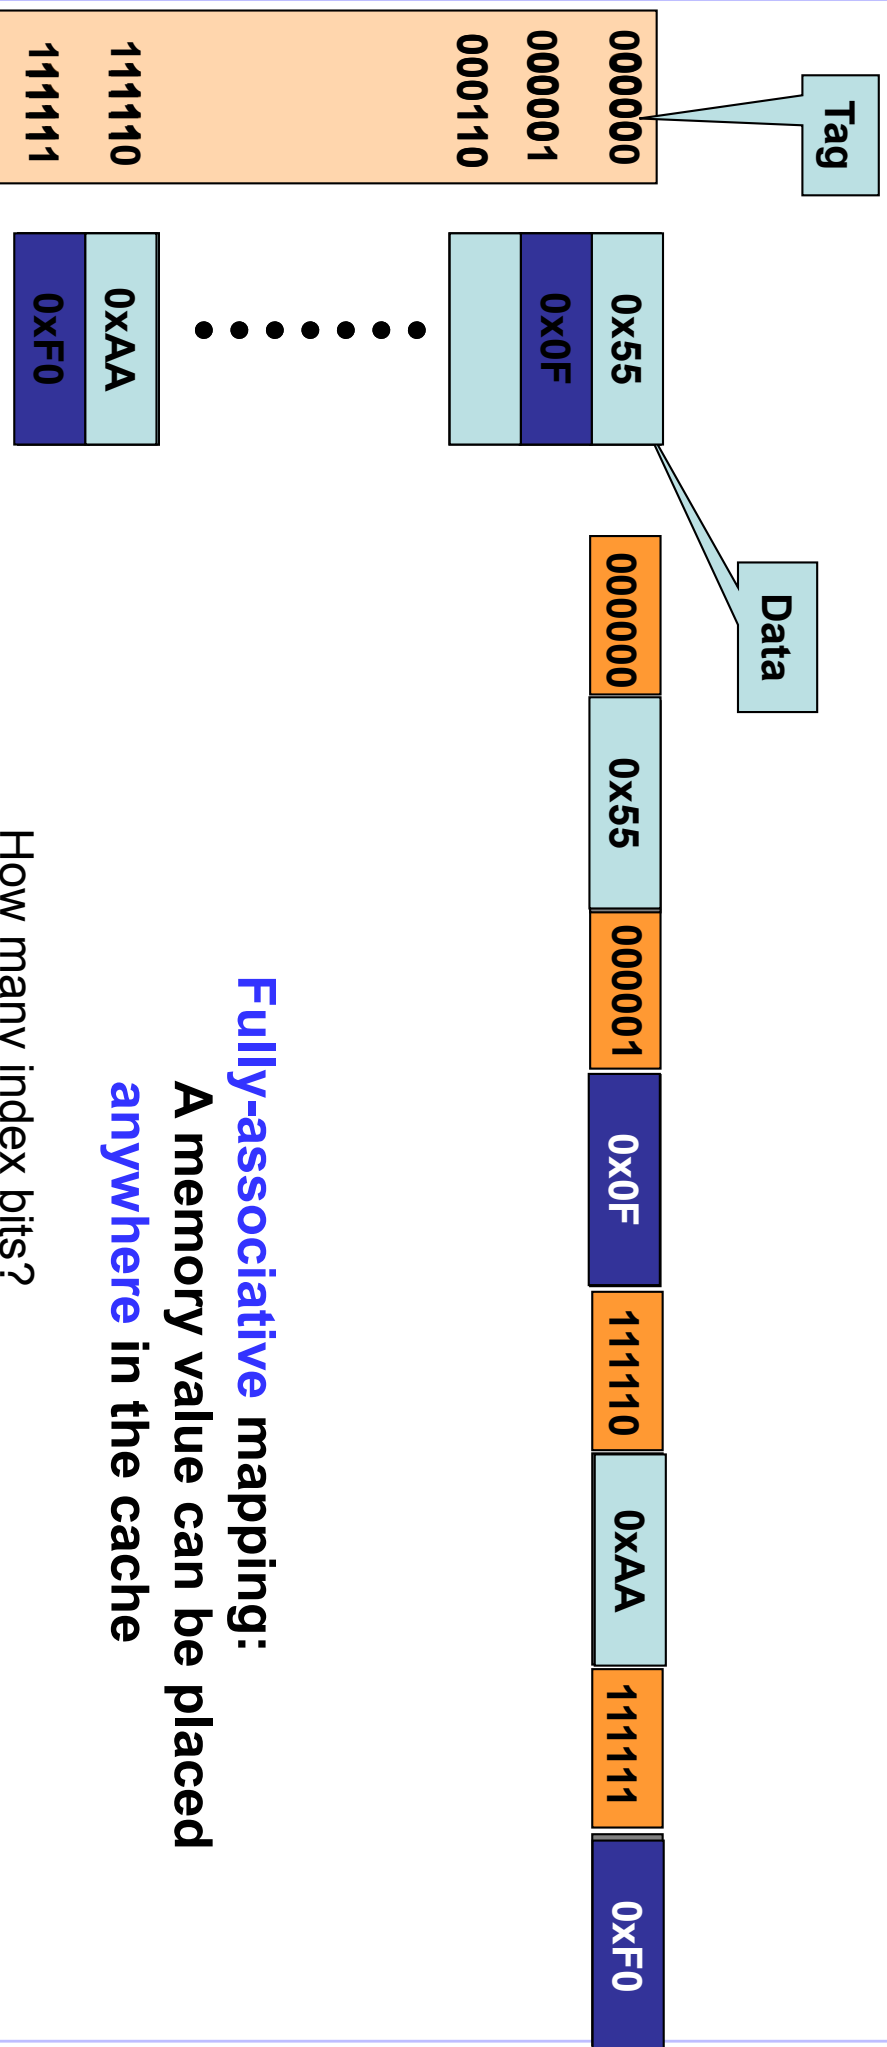
A memory value can be placed in **multiple** location of a set in the cache

Any drawback of set-associative map cache?

# 4-Way Set Associative Cache Organization



# Fully Associative Mapping



## Fully-associative mapping:

A memory value can be placed **anywhere** in the cache

How many index bits?

- None

What are the advantage and disadvantage of fully-associative map cache?

# Cache Replacement Policy

- Random
  - Replace a randomly chosen line
- FIFO
  - Replace the oldest line
- LRU (Least Recently Used)
  - Replace the least recently used line

# How Much Associativity

- Increased associativity **decreases** miss rate
- Simulation of a system with 64KB D-cache, 32-byte blocks, SPEC2000
  - 1-way: 10.3%
  - 2-way: 8.6%
  - 4-way: 8.3%
  - 8-way: 8.1%




# Four Cs (Cache Miss Terms)

- Compulsory Misses:
  - cold start misses (Caches do not have valid data at the start of the program)
- Capacity Misses:
  - Increase cache size
- Conflict Misses:
  - Increase cache size and/or associativity.
  - Associative caches reduce conflict misses
- Coherence Misses:
  - In multiprocessor systems (next lecture)

# Example: 1KB DM Cache, 32-byte Lines

- The lowest M bits are the Offset (Line Size =  $2^M$ )
- Index =  $\log_2$  (# of sets)



# Example of Caches

- Given a 2MB, direct-mapped physical caches, line size=64bytes
- Support 32-bit physical address
- Tag size?
  - # of lines =  $2\text{MB}/64\text{B} = 2^{21}/2^6 = 2^{15}$
  - Direct-mapped  $\Rightarrow$  # of sets = # of lines =  $2^{15} \Rightarrow$  Index = 15 bits
  - Tag = 32 - 15 - 6 = 11 bits
- Now change it to 16-way, Tag size?
  - # of lines =  $2\text{MB}/64\text{B} = 2^{21}/2^6 = 2^{15}$
  - 16-way per set  $\Rightarrow$  # of sets = # of lines/16 =  $2^{11} \Rightarrow$  Index = 11 bits
  - Tag = 32 - 11 - 6 = 15 bits
- How about if it's fully associative, Tag size?
  - # of sets = 1  $\Rightarrow$  No Index
  - Tag = 32 - 6 = 26 bits

# Topics covered

- Why do caches work
  - Principle of **program locality**
- Cache hierarchy
  - Average memory access time ( $T_{\text{access}}$ )




- Types of caches
  - Direct mapped
  - Set-associative
  - Fully associative
- Cache policies
  - Write back vs. write through
  - Write allocate vs. No write allocate


# Cache Write Policy

- Write through - The value is written to **both** the cache line and to the lower-level memory (e.g. main memory).
- Write back - The value is written **only** to the cache line. The modified cache line is written to main memory **only when** it has to be replaced.
  - Cache line is **clean**: holds the **same** value as memory
  - Cache line is **dirty**: holds a **different** value than memory


# Write-through Policy



# Write-through Policy




# Write-through Policy






# Write-through Policy




The value is written to **both** the cache line and to the lower-level memory (e.g. main memory).

# Write Buffer




- Processor writes data into the cache and the **write buffer**
- Memory controller writes contents of the buffer to memory
- Write buffer is a **FIFO** structure:
  - Typically 4 to 8 entries
  - Desirable: Occurrence of Writes << DRAM write cycles


# Writeback Policy



# Writeback Policy




# Writeback Policy




The value is written **only** to the cache line. The modified cache line is written to main memory **only when** it has to be replaced.

# Writeback Policy



The value is written **only** to the cache line. The modified cache line is written to main memory **only when** it has to be replaced.

# Writeback Policy



The modified cache line is written to main memory **only** **when** it has to be replaced.

# What happens if Write Miss?

- Write allocate

- After write miss, the data will be loaded to cache first before write
- Write misses first act like read misses

- No write allocate

- Write misses do **not** interfere cache
- Line is **only** modified in the **lower level** memory



# Topics covered

- Why do caches work
  - Principle of **program locality**
- Cache hierarchy
  - Average memory access time ( $T_{\text{access}}$ )
- Types of caches
  - Direct mapped
  - Set-associative
  - Fully associative
- Cache policies
  - Write back vs. write through
  - Write allocate vs. No write allocate

# Supplementary Reading Materials

## Computer Organization and Design: The Hardware/Software Interface Fifth Edition

Chapter 5:  
Large and Fast: Exploiting Memory  
Hierarchy

Section 5.1 - 5.4

