

# EECS 645

## Computer Architecture

### Lecture 10 Multiprocessor and Memory Coherence

Chenyun Pan

Department of Electrical Engineering &

Computer Science

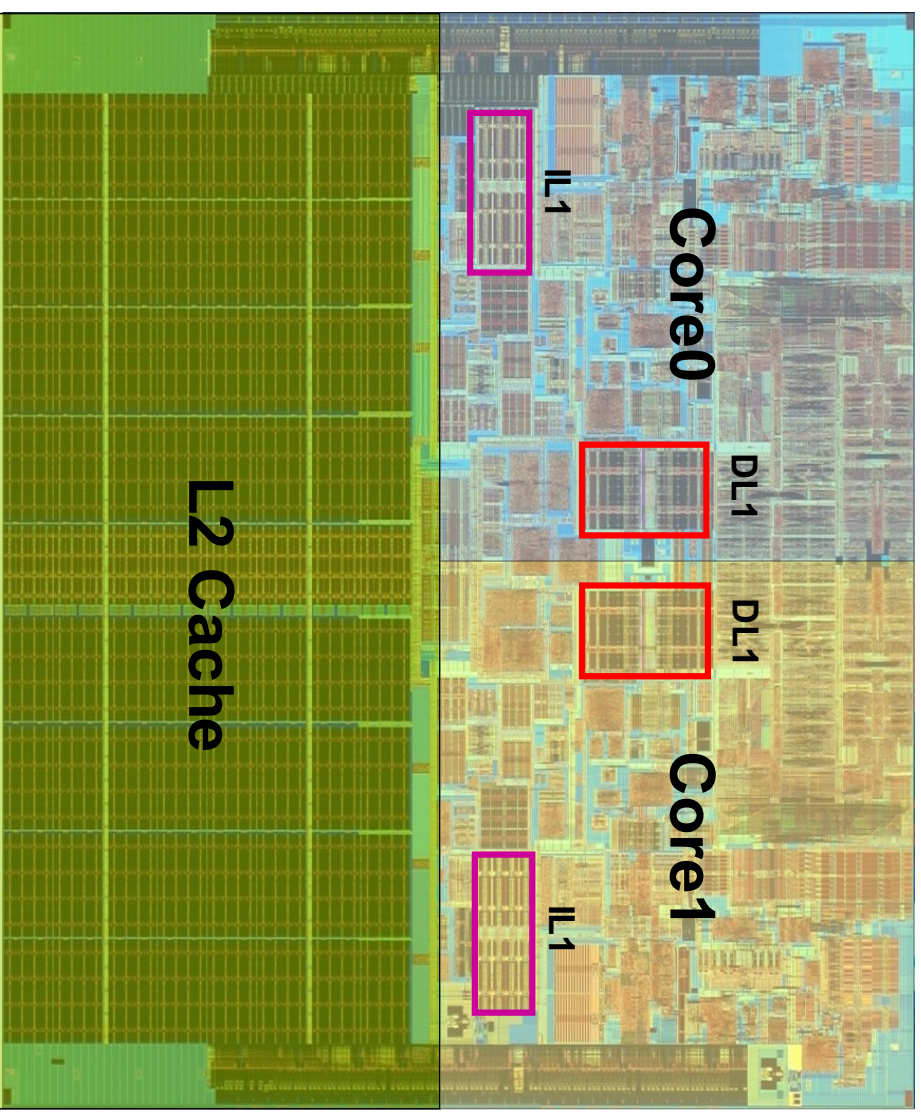
University of Kansas



Slide Courtesy of Dr. Hsien-Hsin Sean Lee


# Intel Core2 Duo

|    |   |
|----|---|
| L1 | 32 KB, 8-Way, 64 Byte/Line, LRU, WB<br>3 Cycle Latency    |
| L2 | 4.0 MB, 16-Way, 64 Byte/Line, LRU, WB<br>14 Cycle Latency |



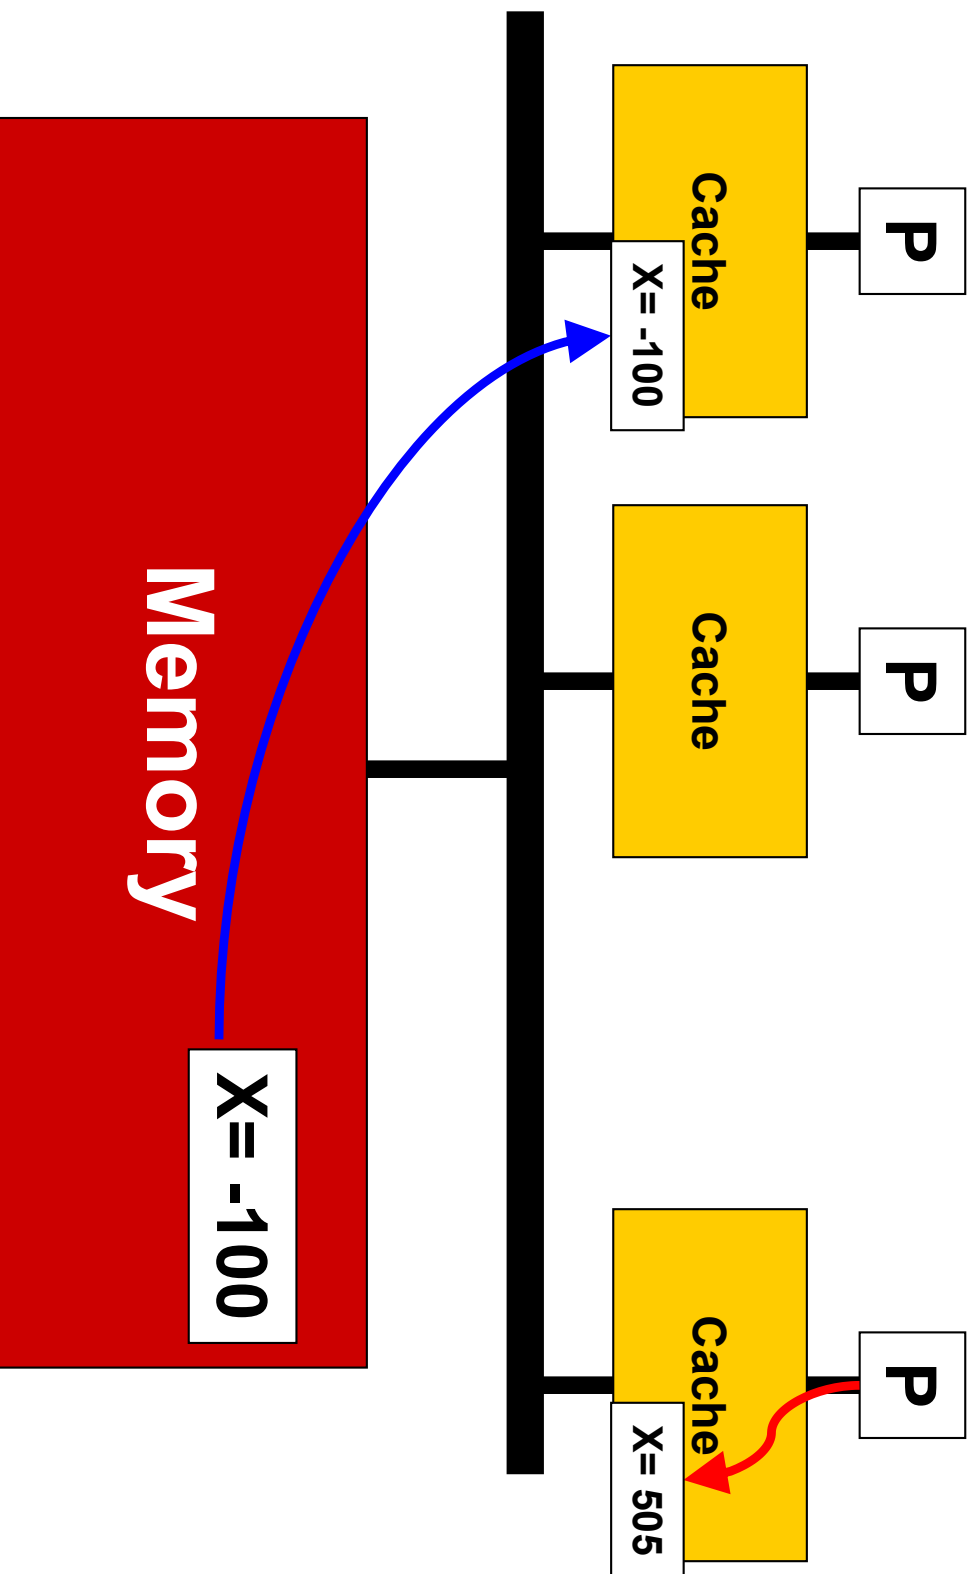
What happens if two cores are running the same program and modifying the same variable?

# Example (Writeback Cache)



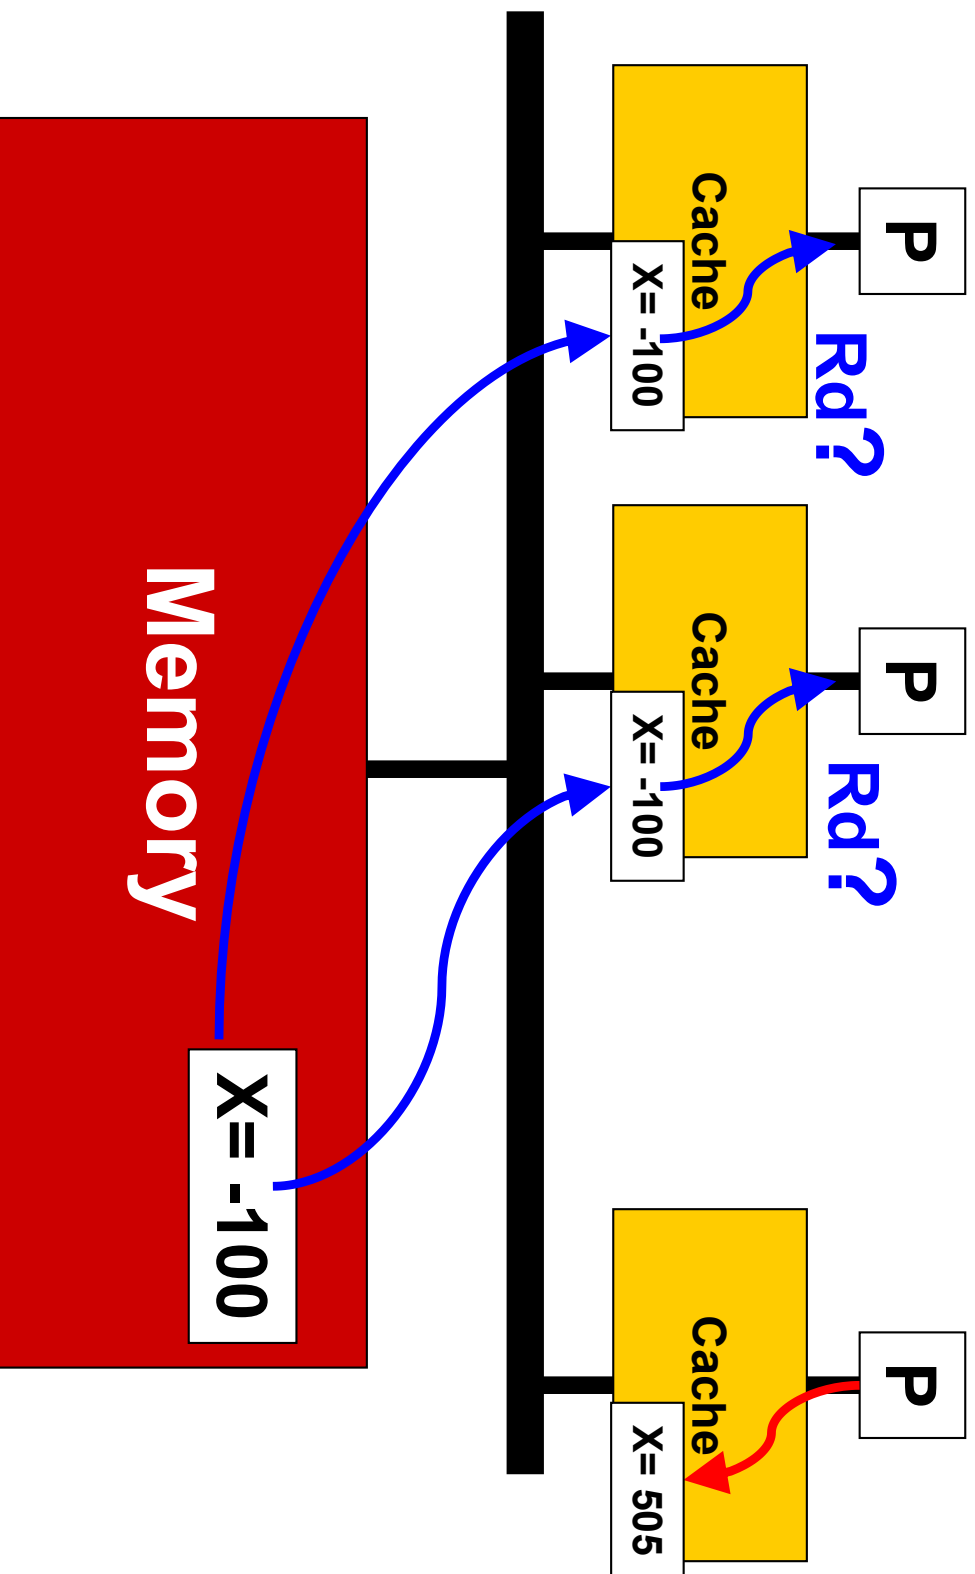
- Closest cache level is private
- Multiple copies of cache line can be present across different cores
- Local updates may lead to **incoherent state**

# Example (Writeback Cache)



- Closest cache level is private
- Multiple copies of cache line can be present across different cores
- Local updates may lead to **incoherent state**

# Example (Writeback Cache)




- Closest cache level is private
- Multiple copies of cache line can be present across different cores
- Local updates may lead to **incoherent state**

# Cache Coherence Protocol

- Write propagation
  - Writes are **visible** to other processes
- All the writes will be shown as a **transaction** on the shared bus to memory
- Two protocols
  - ➔ – **Update-based Protocol**
  - **Invalidation-based Protocol**


# Example (Writeback Cache)



- Closest cache level is private
- Multiple copies of cache line can be present across different cores
- Local updates may lead to **incoherent state**

# Bus Snooping

(Update-based Protocol on Write-Through cache)



- Each processor's cache controller constantly snoops on the bus
- **Update** local copies upon snoop hit
- What is the drawback of update-based protocol?
  - Heavy traffic on the bus to supply both address and data

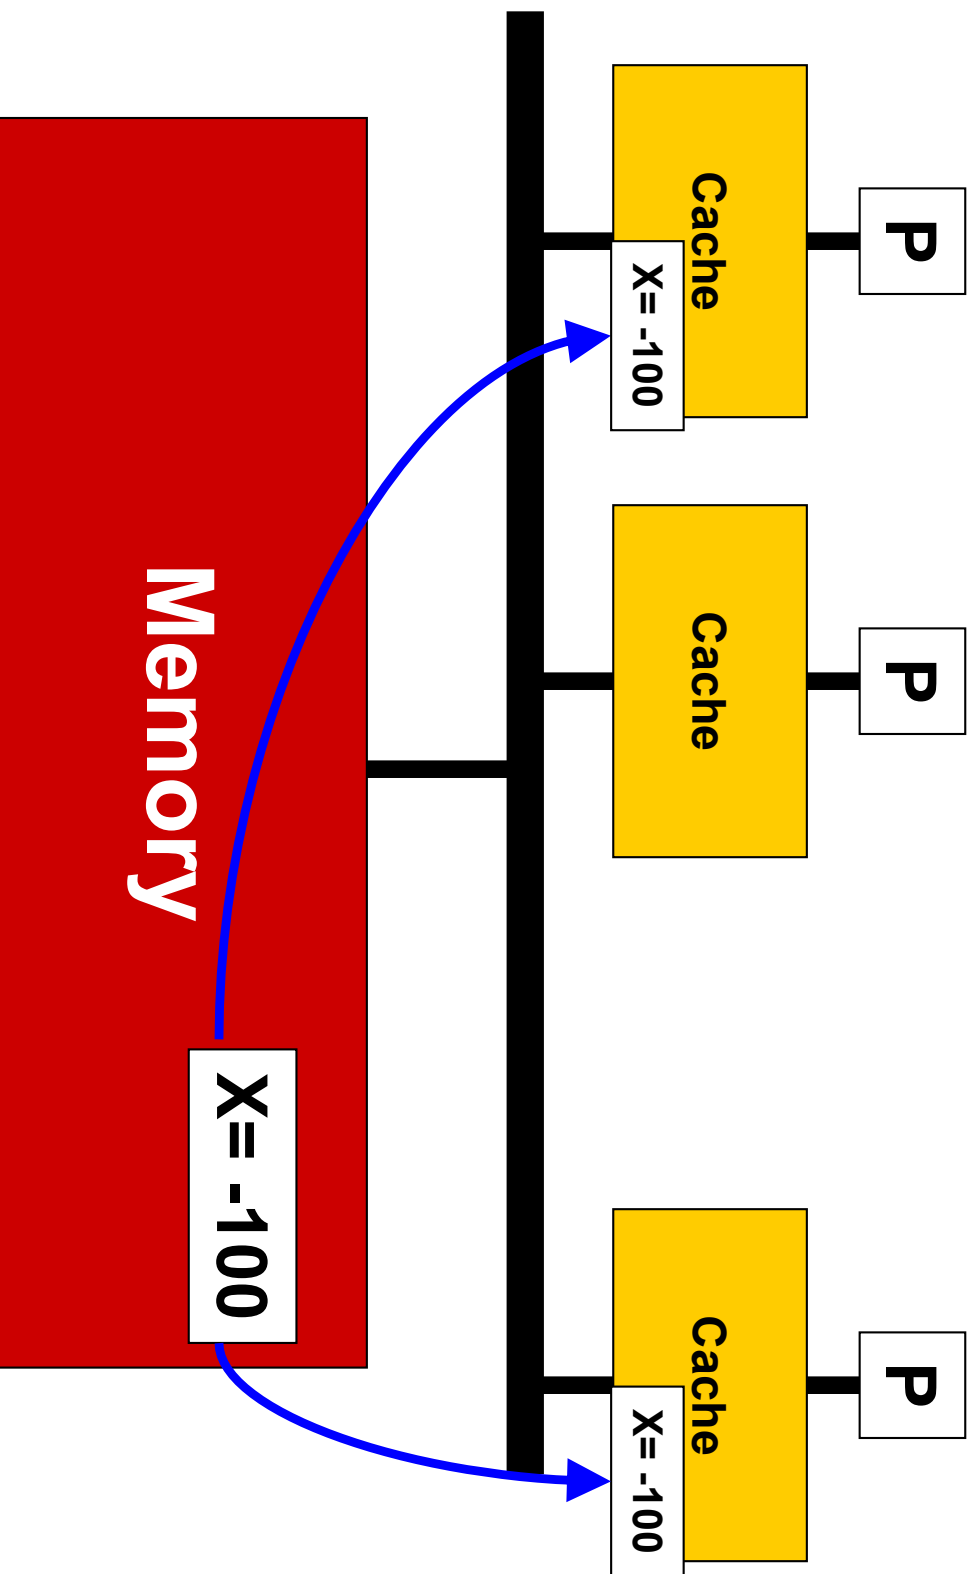


# Cache Coherence Protocol

- Write propagation
  - Writes are **visible** to other processes
- All the writes will be shown as a **transaction** on the shared bus to memory
- Two protocols
  - Update-based Protocol
  - Invalidation-based Protocol



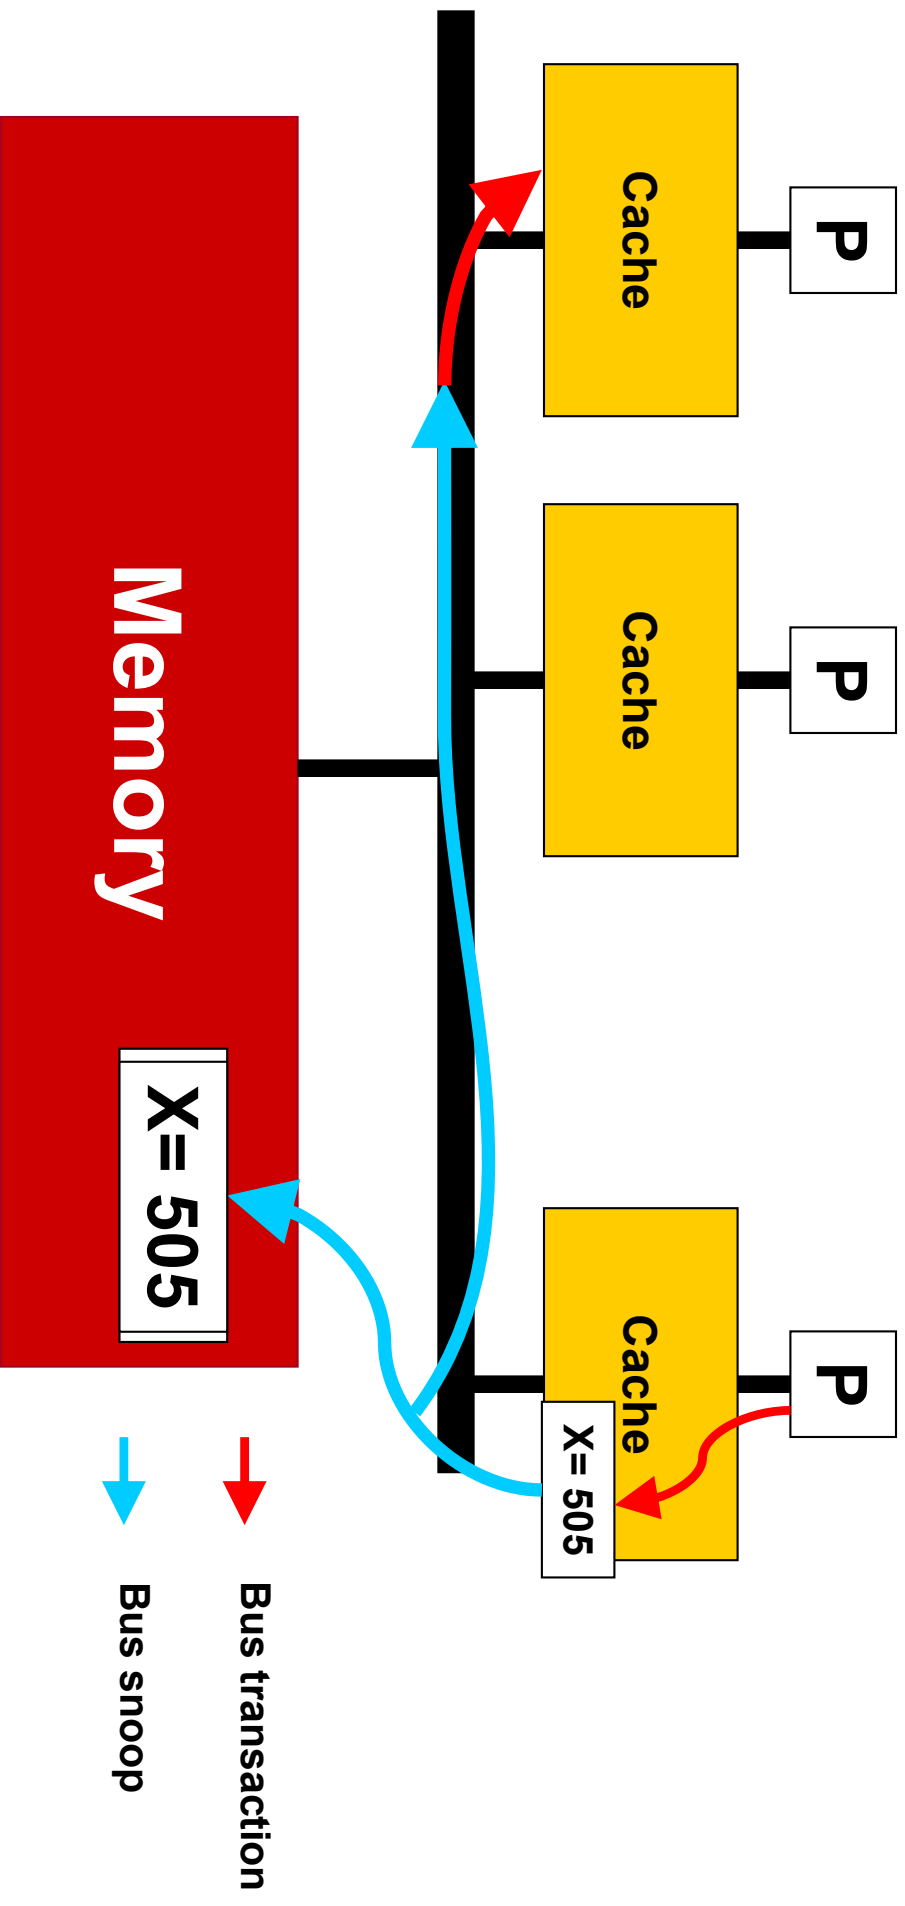
# Example (Writeback Cache)



- Closest cache level is private
- Multiple copies of cache line can be present across different cores
- Local updates may lead to **incoherent state**

# Bus Snooping


(Invalidation-based Protocol on Write-Through cache)



- Each processor's cache controller constantly snoops on the bus
- **Invalidate** local copies upon snoop hit
- What is the drawback of write-through cache based protocol?
  - Frequent access to memory

# Bus Snooping


(Invalidation-based Protocol on Write-Through cache)



- Each processor's cache controller constantly snoops on the bus
- **Invalidate** local copies upon snoop hit
- What is the drawback of write-through cache based protocol?
  - Frequent access to memory

# Cache Coherence Protocol

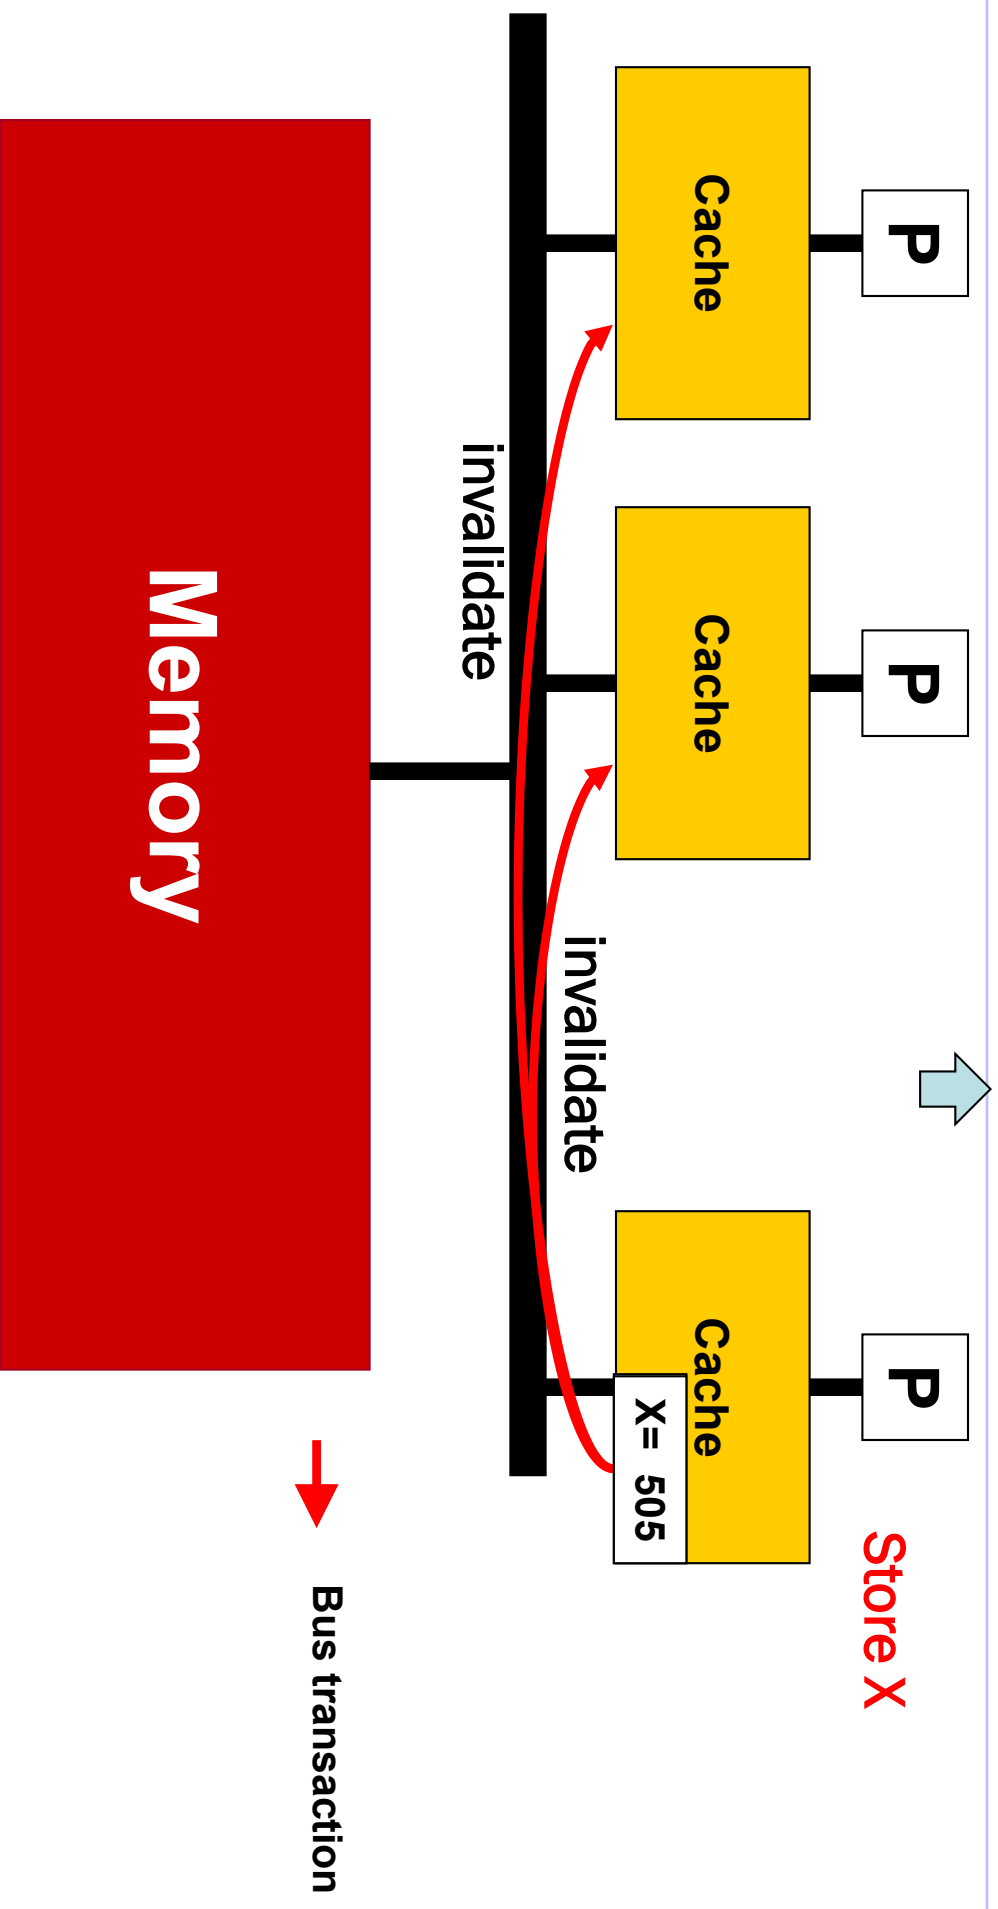
## (Invalidation-based Protocol on **Writeback** cache)



- Invalidate the data copies for the sharing processor nodes
- Reduced **traffic** when a processor node keeps updating the same memory location

# Cache Coherence Protocol

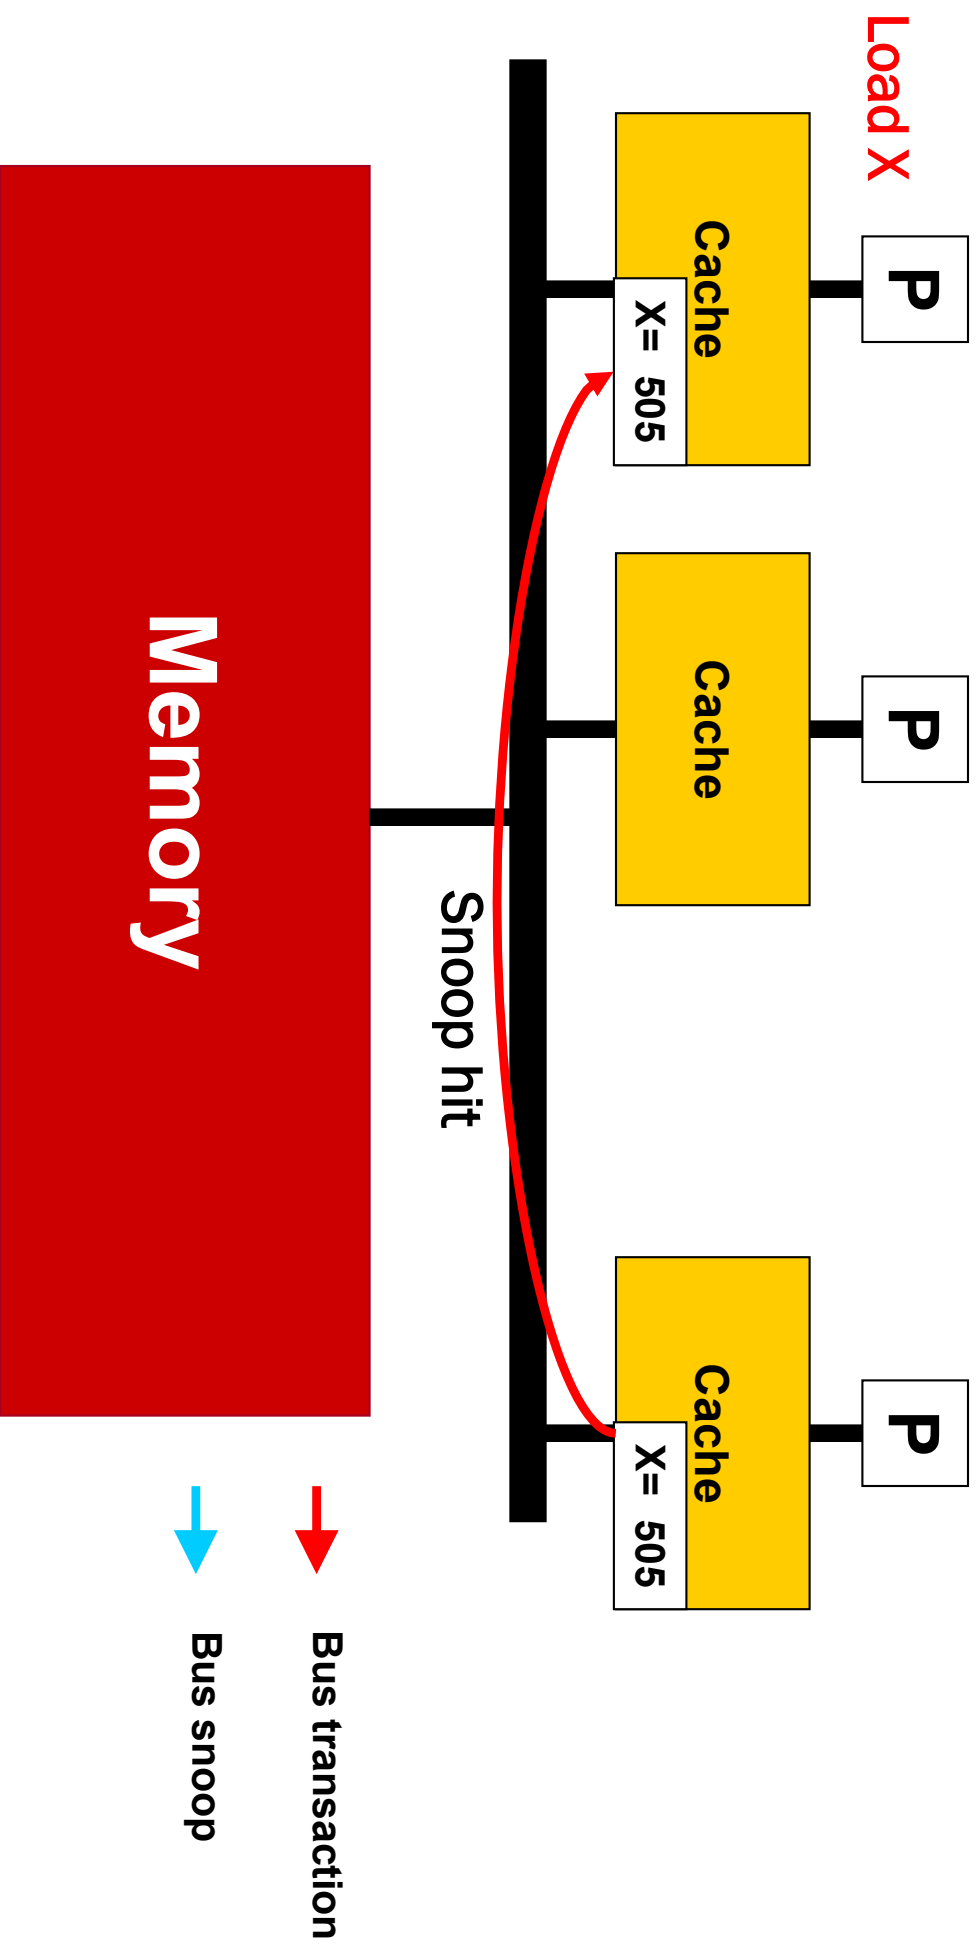
## (Invalidation-based Protocol on **Writeback** cache)



- Invalidate the data copies for the sharing processor nodes
- Reduced **traffic** when a processor node keeps updating the same memory location

# Cache Coherence Protocol

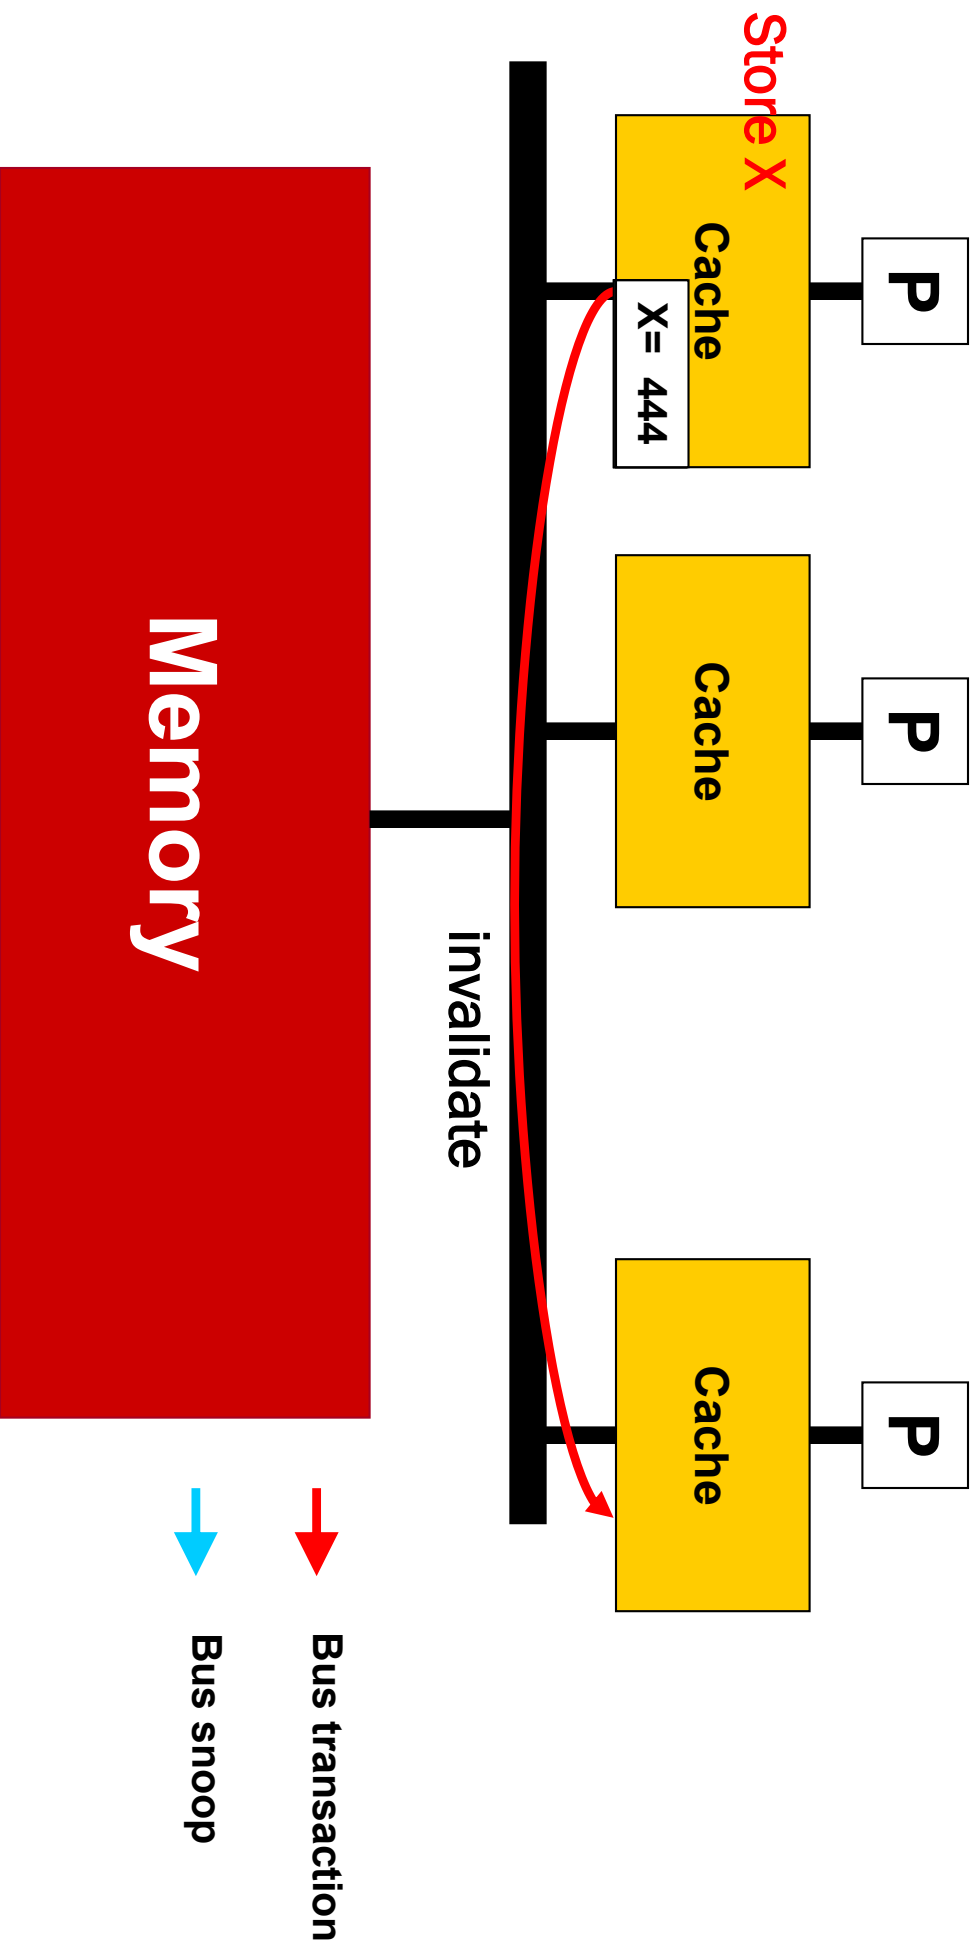
## (Invalidation-based Protocol on **Writeback** cache)



- Invalidate the data copies for the sharing processor nodes
- Reduced **traffic** when a processor node keeps updating the **same** memory location

# Cache Coherence Protocol

## (Invalidation-based Protocol on Writeback cache)

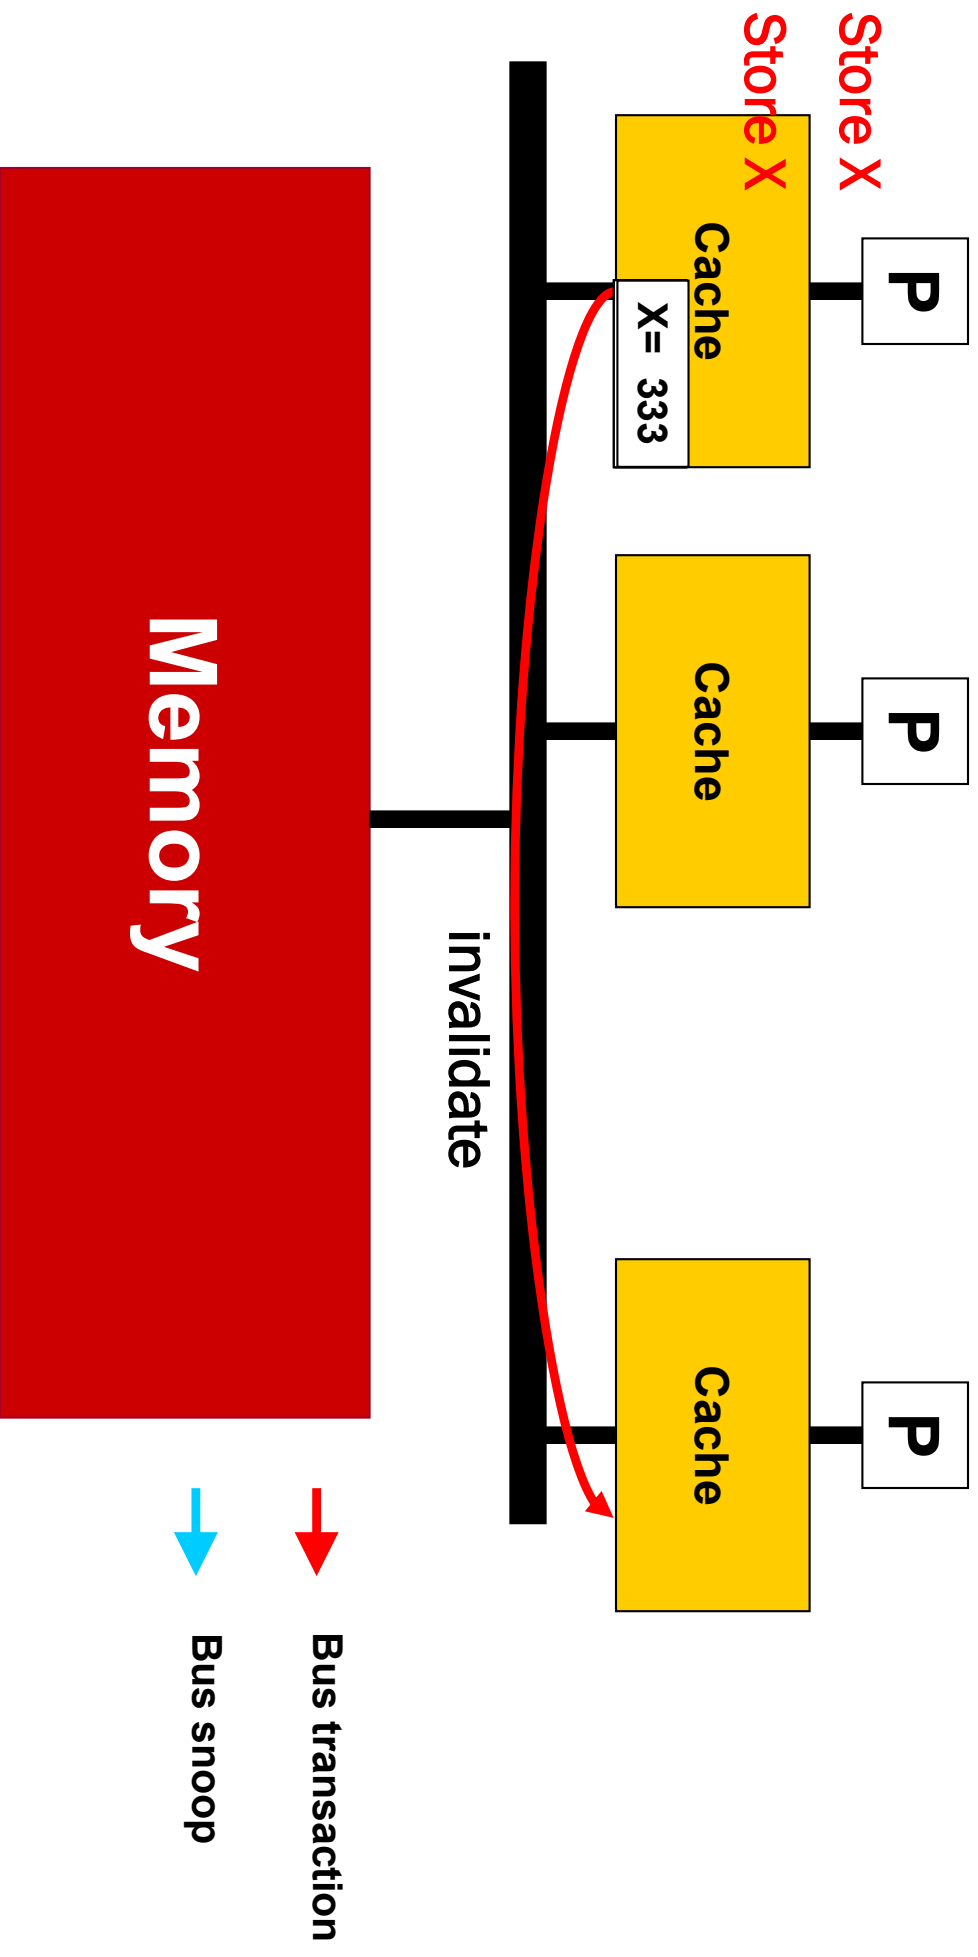


- Invalidate the data copies for the sharing processor nodes
- Reduced **traffic** when a processor node keeps updating the **same** memory location



# Cache Coherence Protocol

## (Invalidation-based Protocol on Writeback cache)



- Invalidate the data copies for the sharing processor nodes
- Reduced **traffic** when a processor node keeps updating the **same** memory location

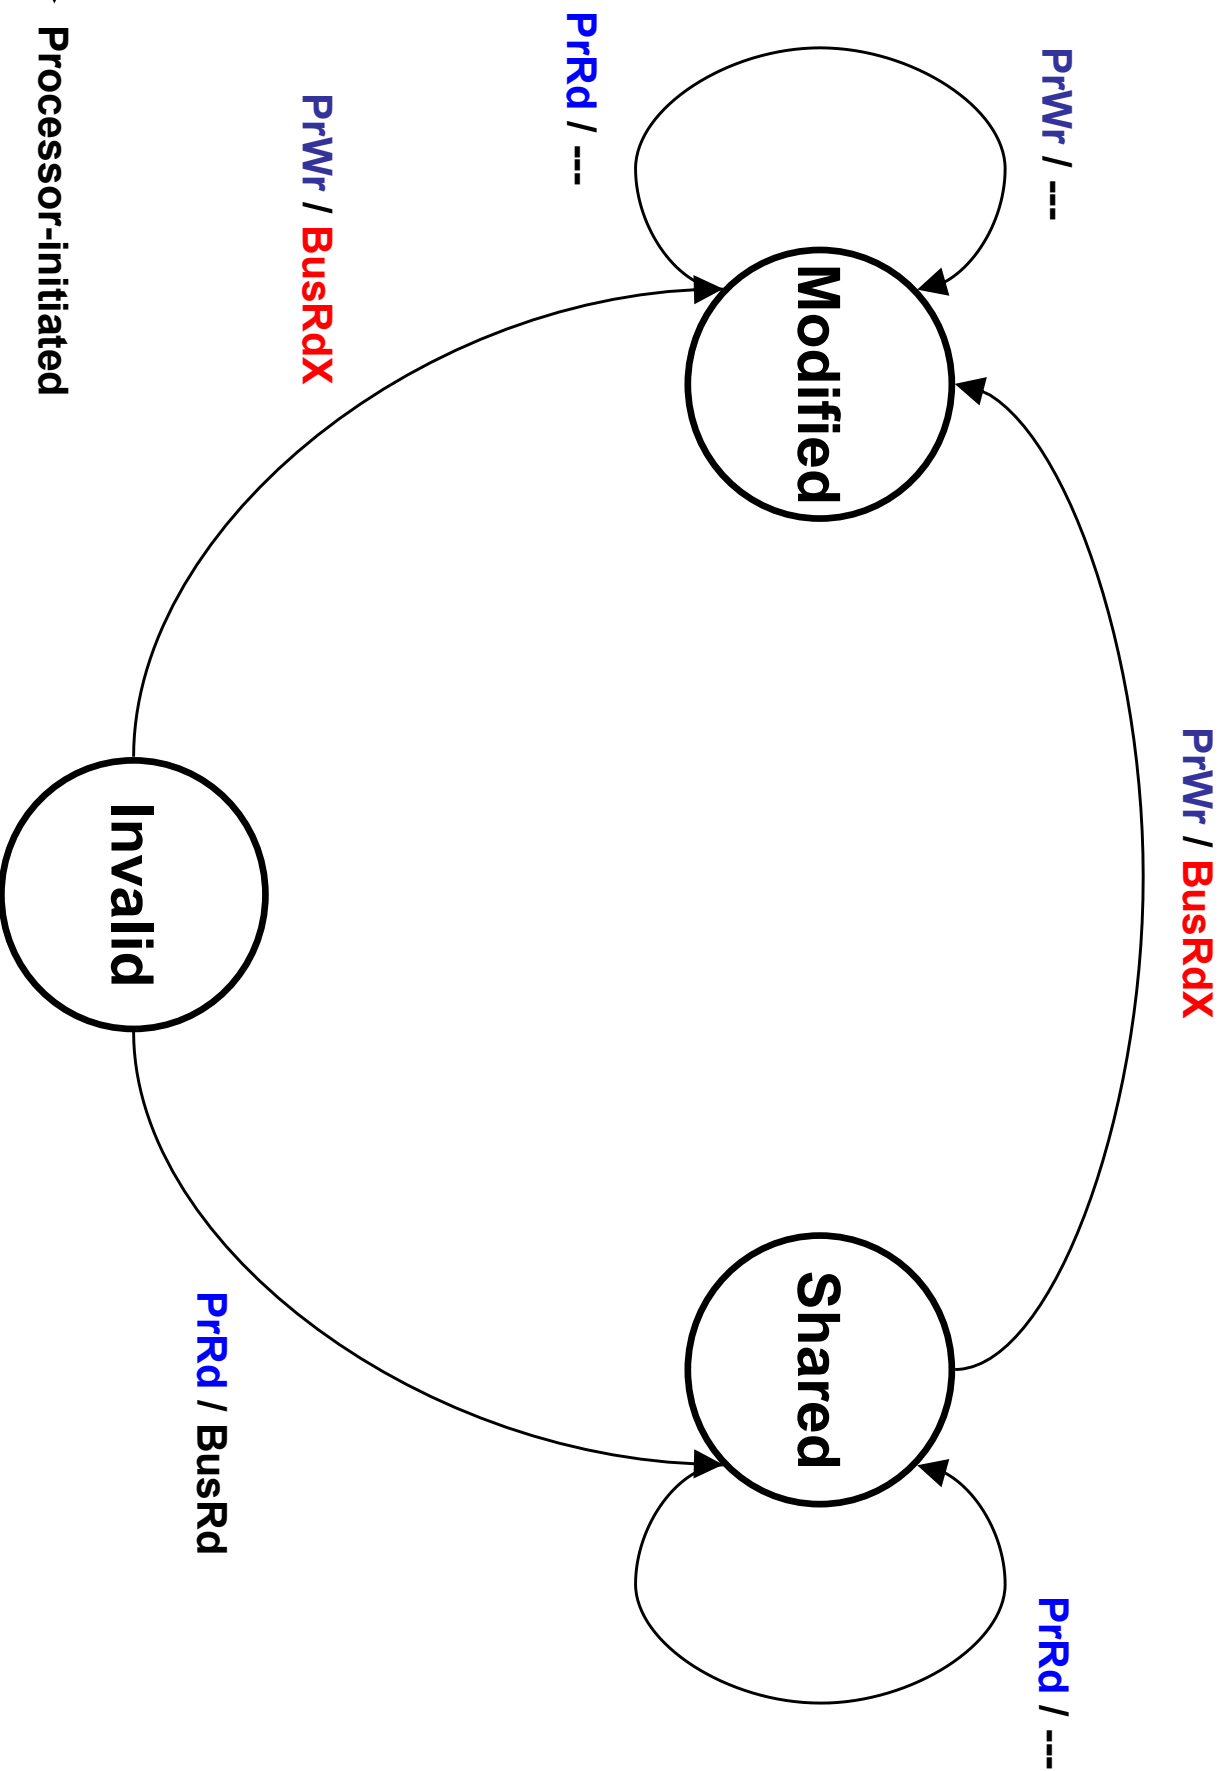
# MSI Writeback Invalidation Protocol

- **Modified**
  - Dirty
  - Only this cache has a valid copy
- **Shared**
  - Memory is consistent
  - One or more caches have a valid copy
- **Invalid**
- **Writeback protocol:** A cache line can be written multiple times before the memory is updated.

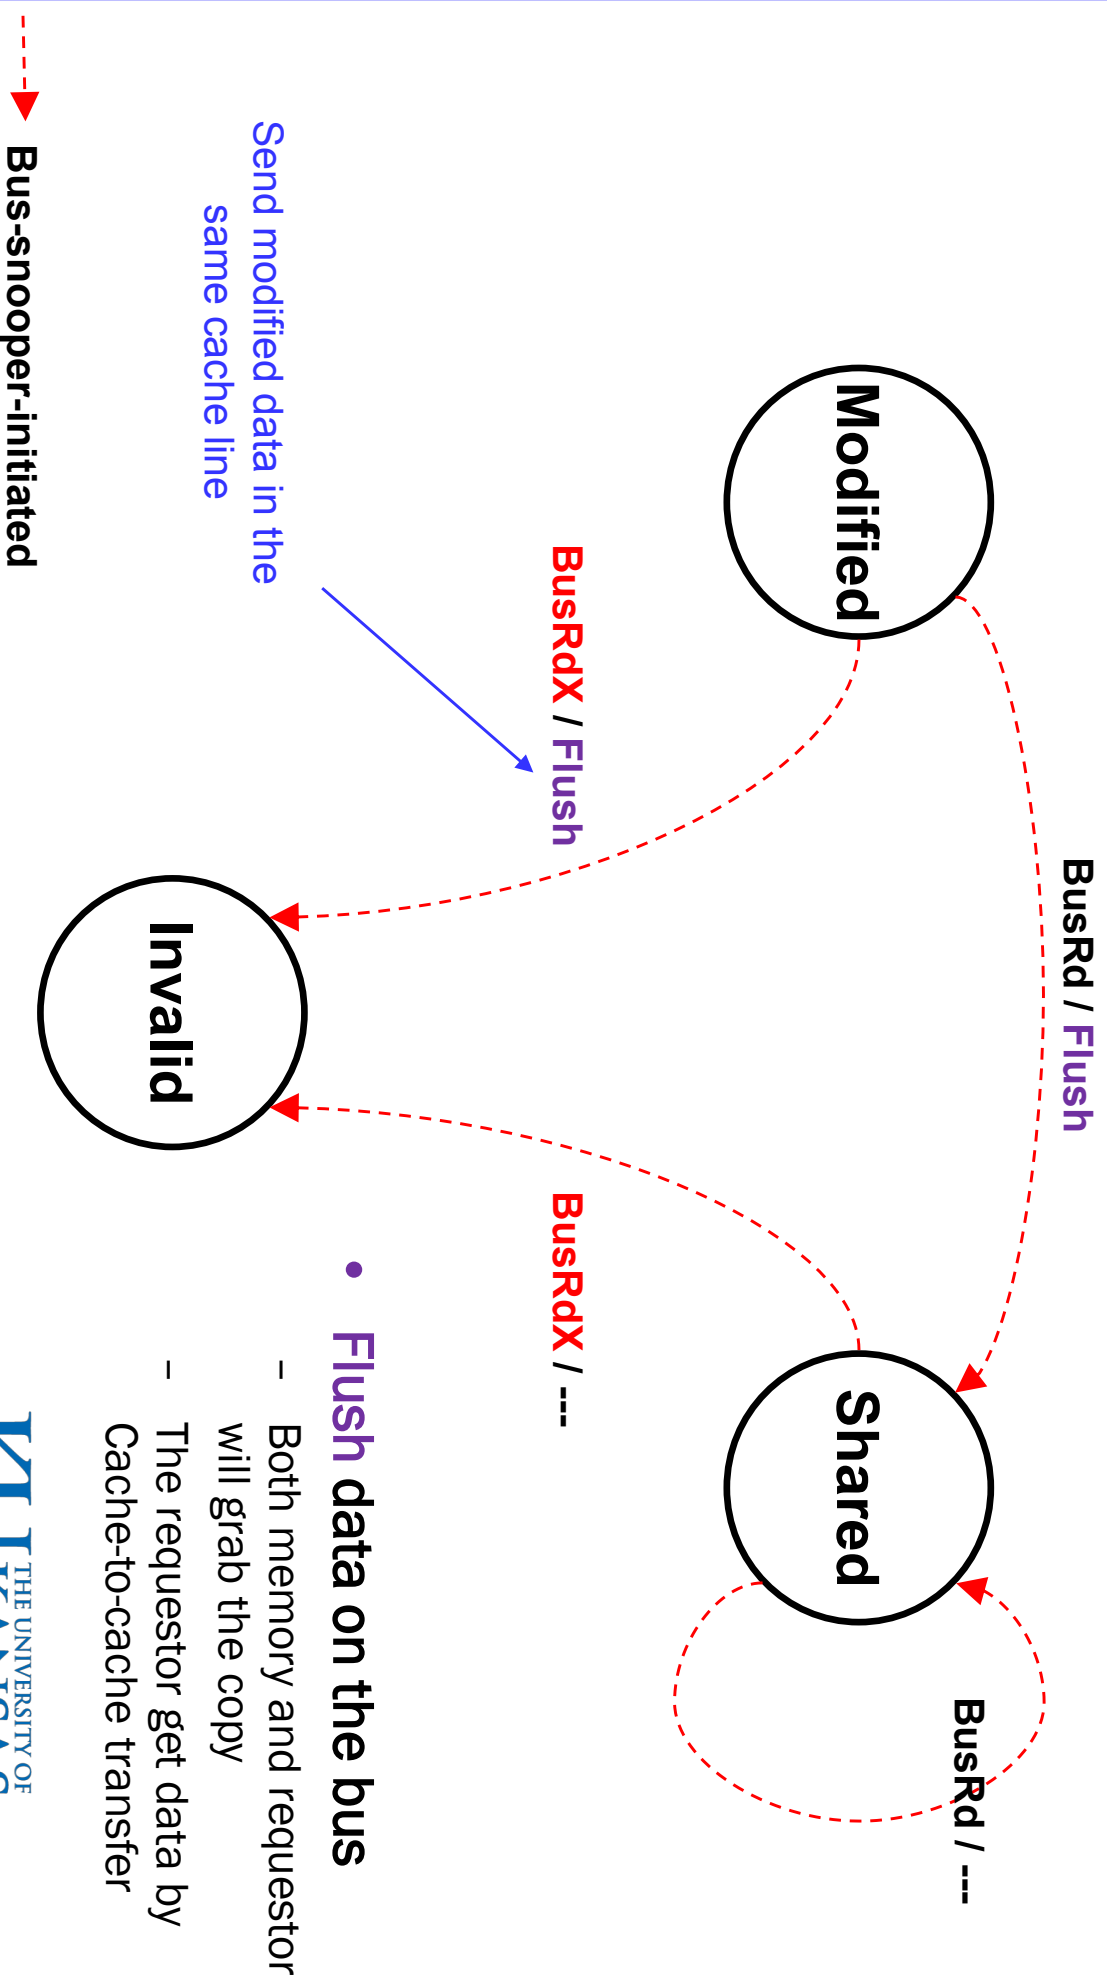
# MSI Writeback Invalidation Protocol

- Two types of local request from the **processor**
  - PrRd
  - PrWr
- Two types of **bus transactions** post by cache controller
  - **BusRd**
    - PrRd misses the cache
    - Memory or another cache supplies the line
  - **BusRd ex**clusive (Read-to-own)
    - PrWr is issued to a line which is **not** in the Modified state

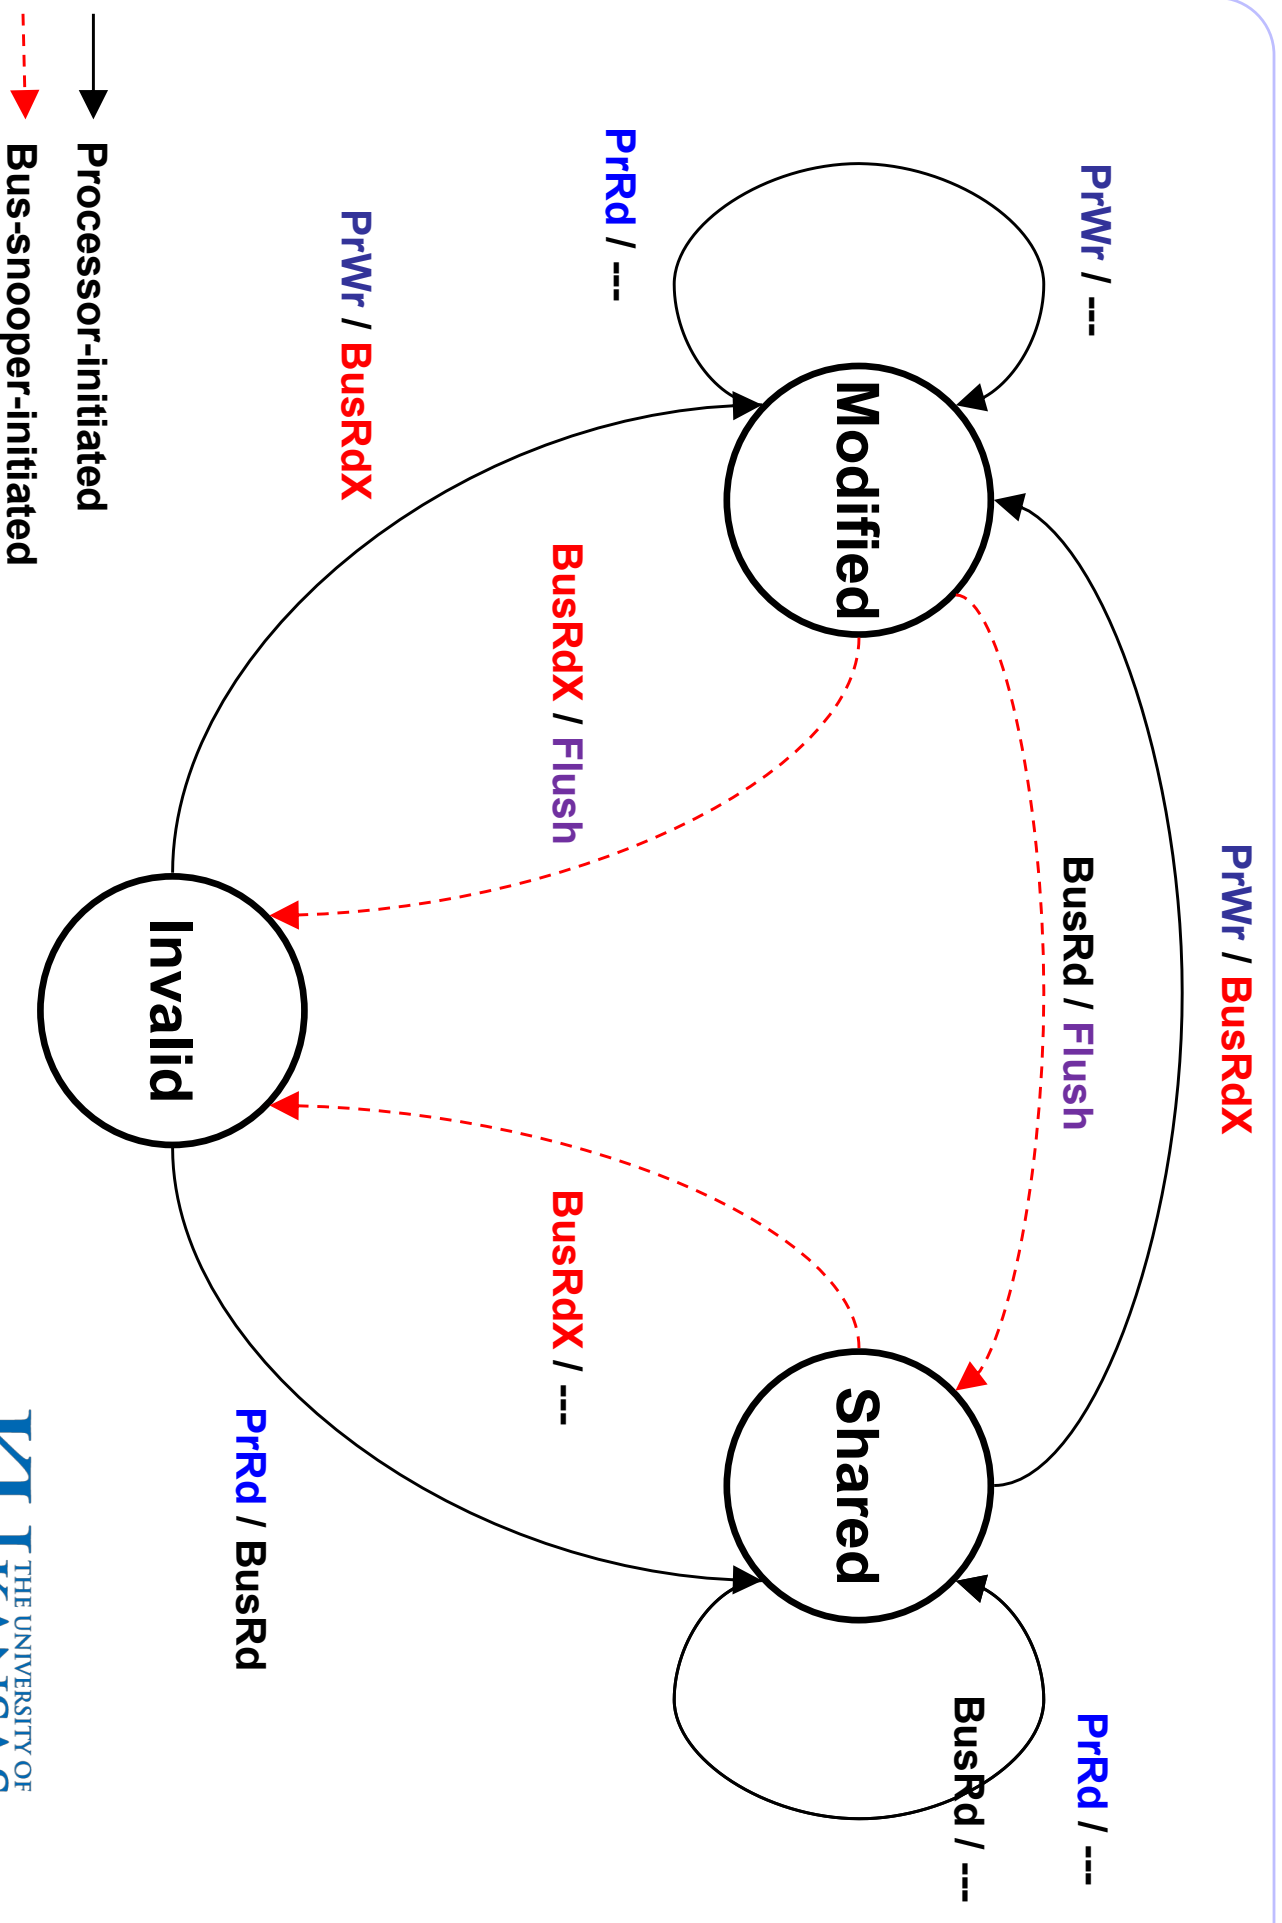
# MSI Writeback Invalidation Protocol (Initiated by Current Processor)




# MSI Writeback Invalidation Protocol (Bus Transaction Initiated by Other Processors)



# MSI Writeback Invalidation Protocol




# MSI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | S           | I           | I           | BusRd           | Memory        |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |

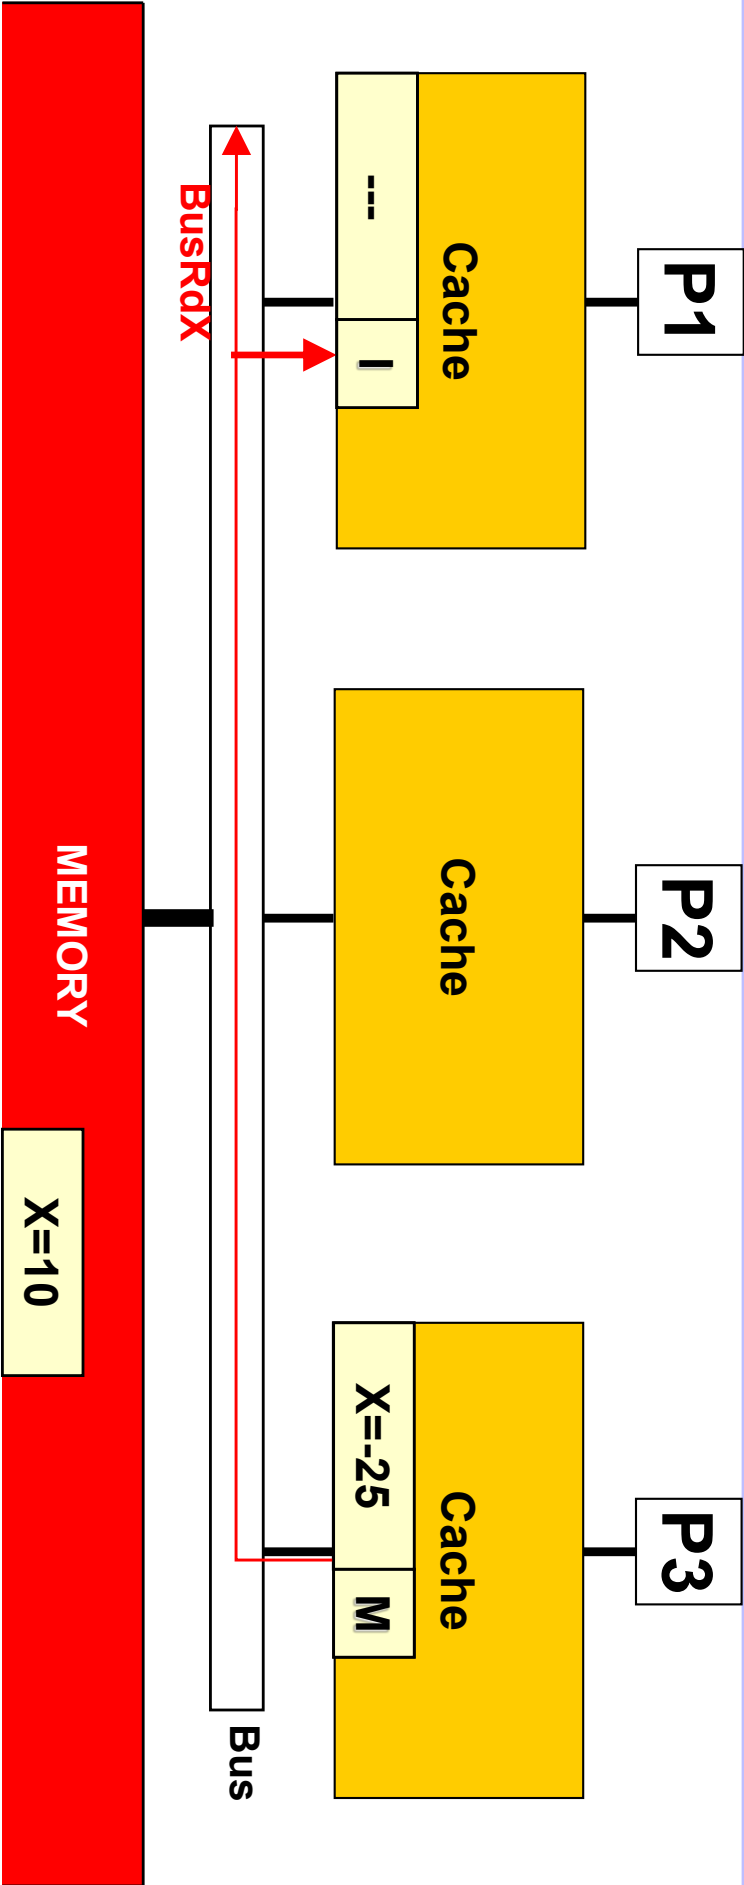
# MSI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | S           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | Memory        |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |



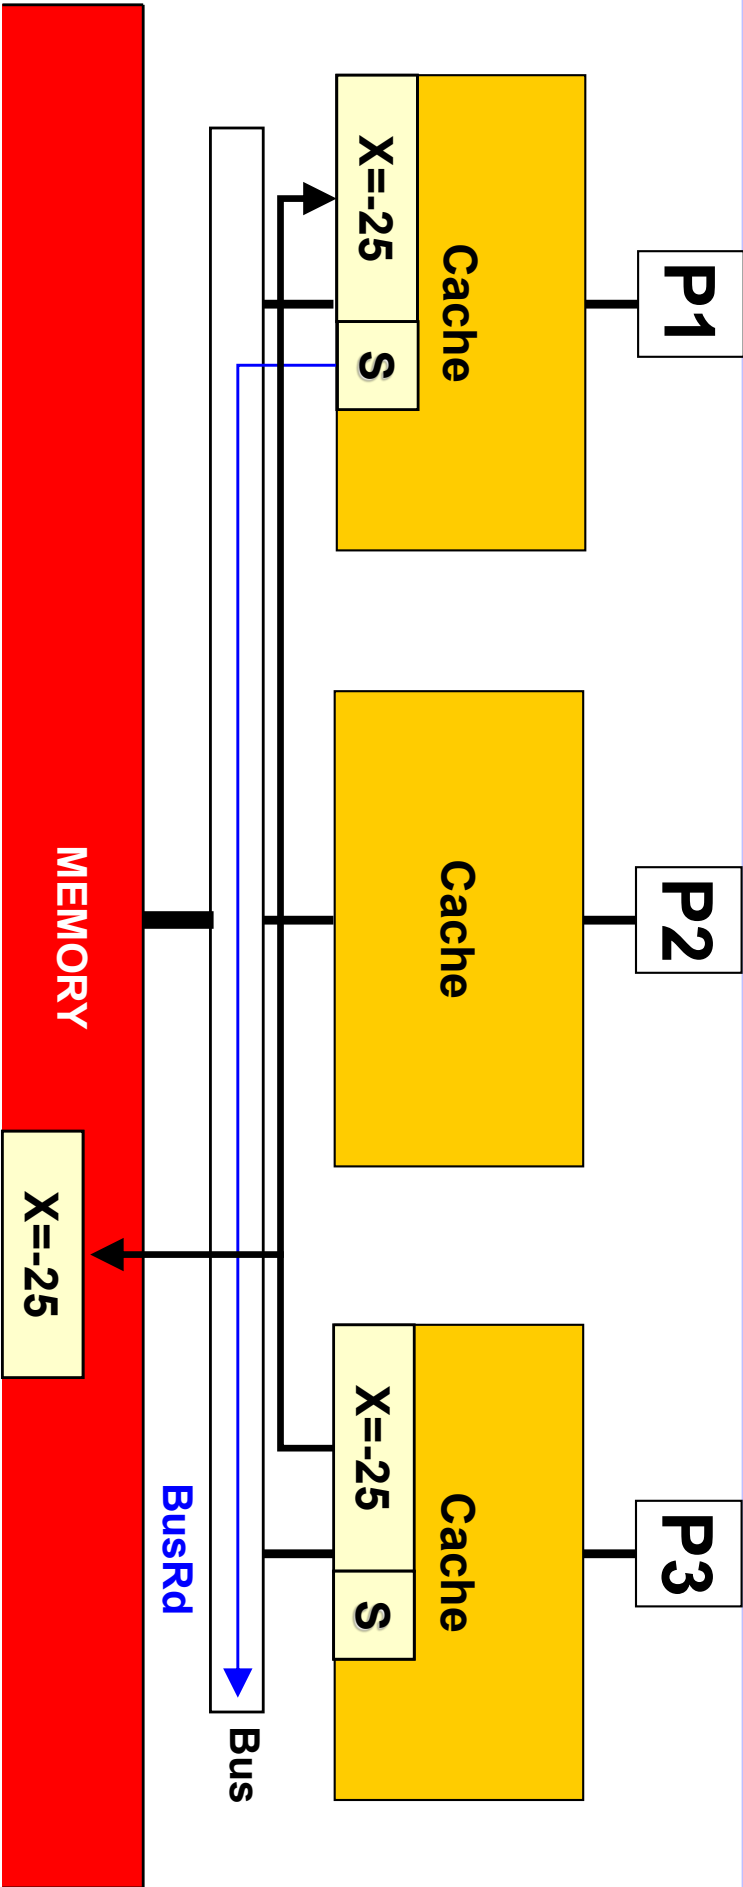
# MSI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | S           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | Memory        |
| P3 writes X      | I           | I           | M           | BusRdX          | ---           |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |

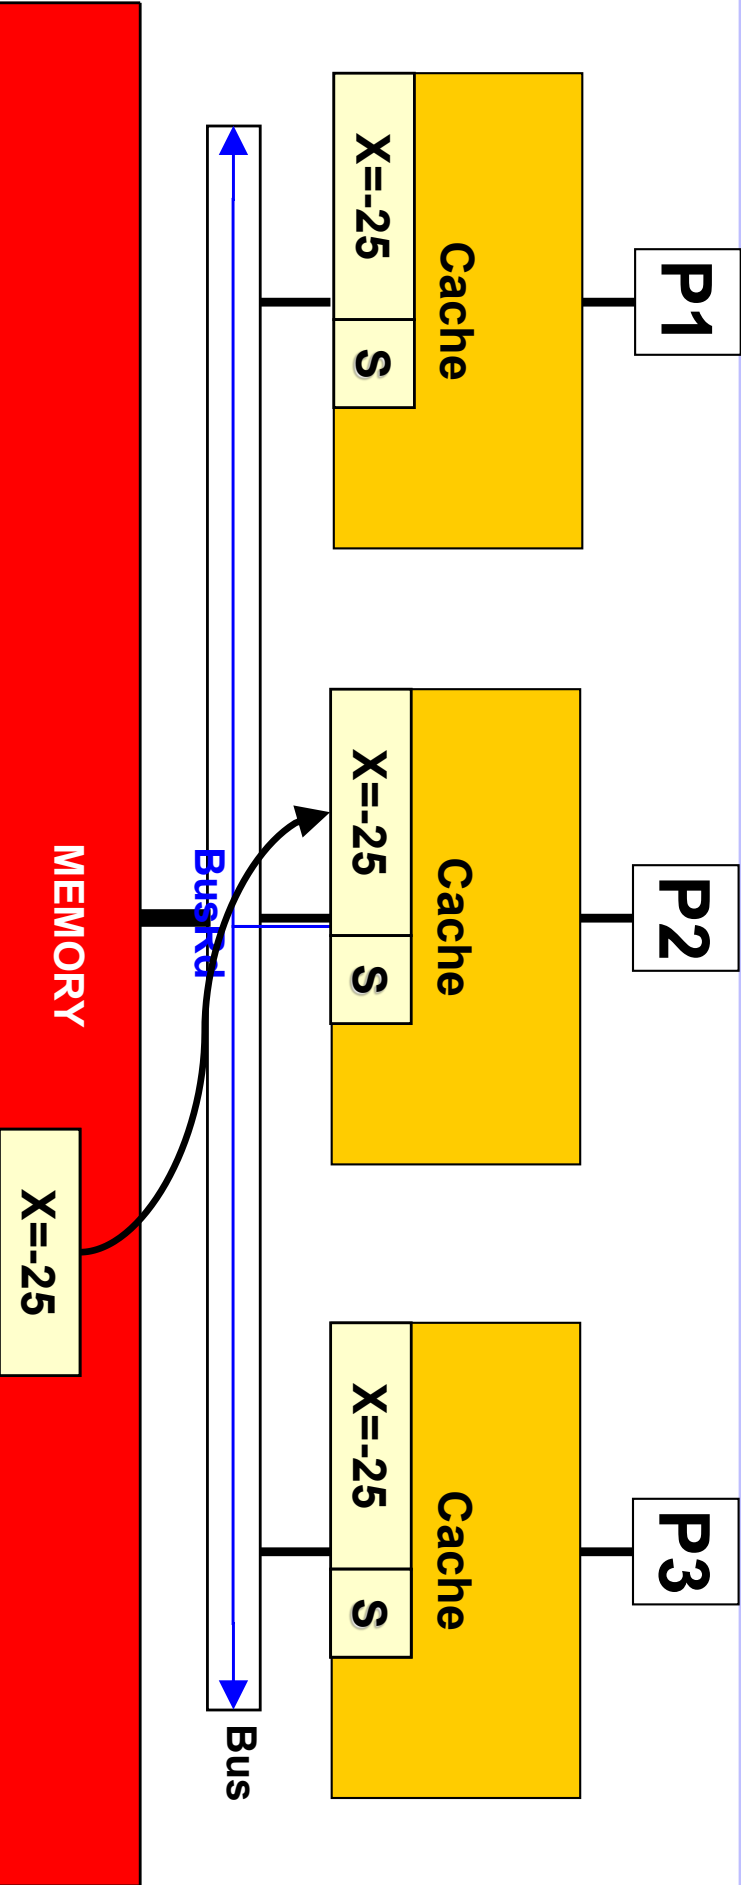
# MSI Example

Both P1 and memory get data during the flush



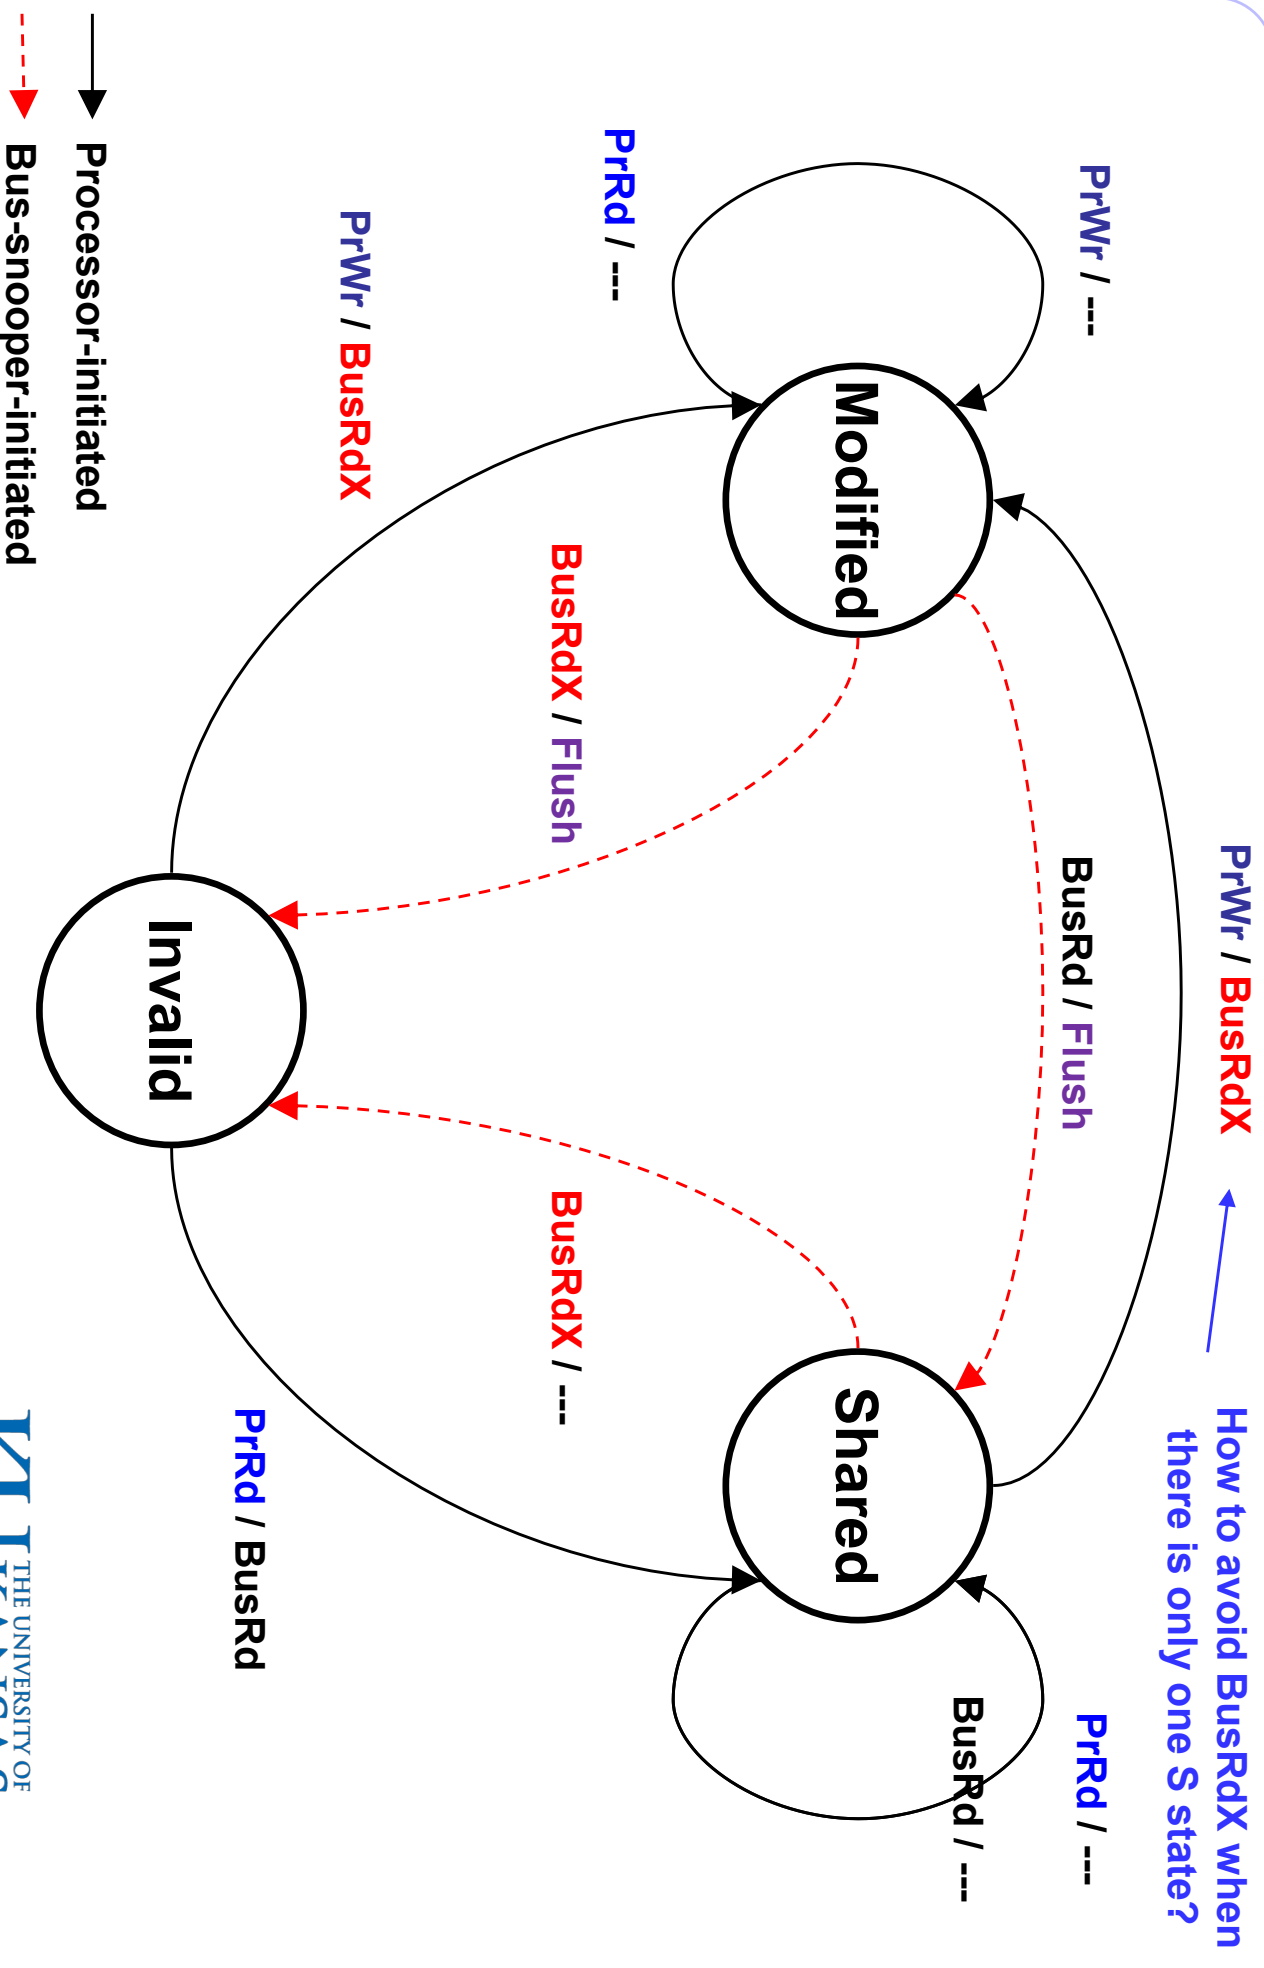
| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | S           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | Memory        |
| P3 writes X      | I           | I           | M           | BusRdX          | ---           |
| P1 reads X       | S           | I           | S           | BusRd           | P3 Cache      |
|                  |             |             |             |                 |               |

# MSI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | S           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | Memory        |
| P3 writes X      | I           | I           | M           | BusRdX          | ---           |
| P1 reads X       | S           | I           | S           | BusRd           | P3 Cache      |
| P2 reads X       | S           | S           | S           | BusRd           | Memory        |

# MSI Writeback Invalidation Protocol



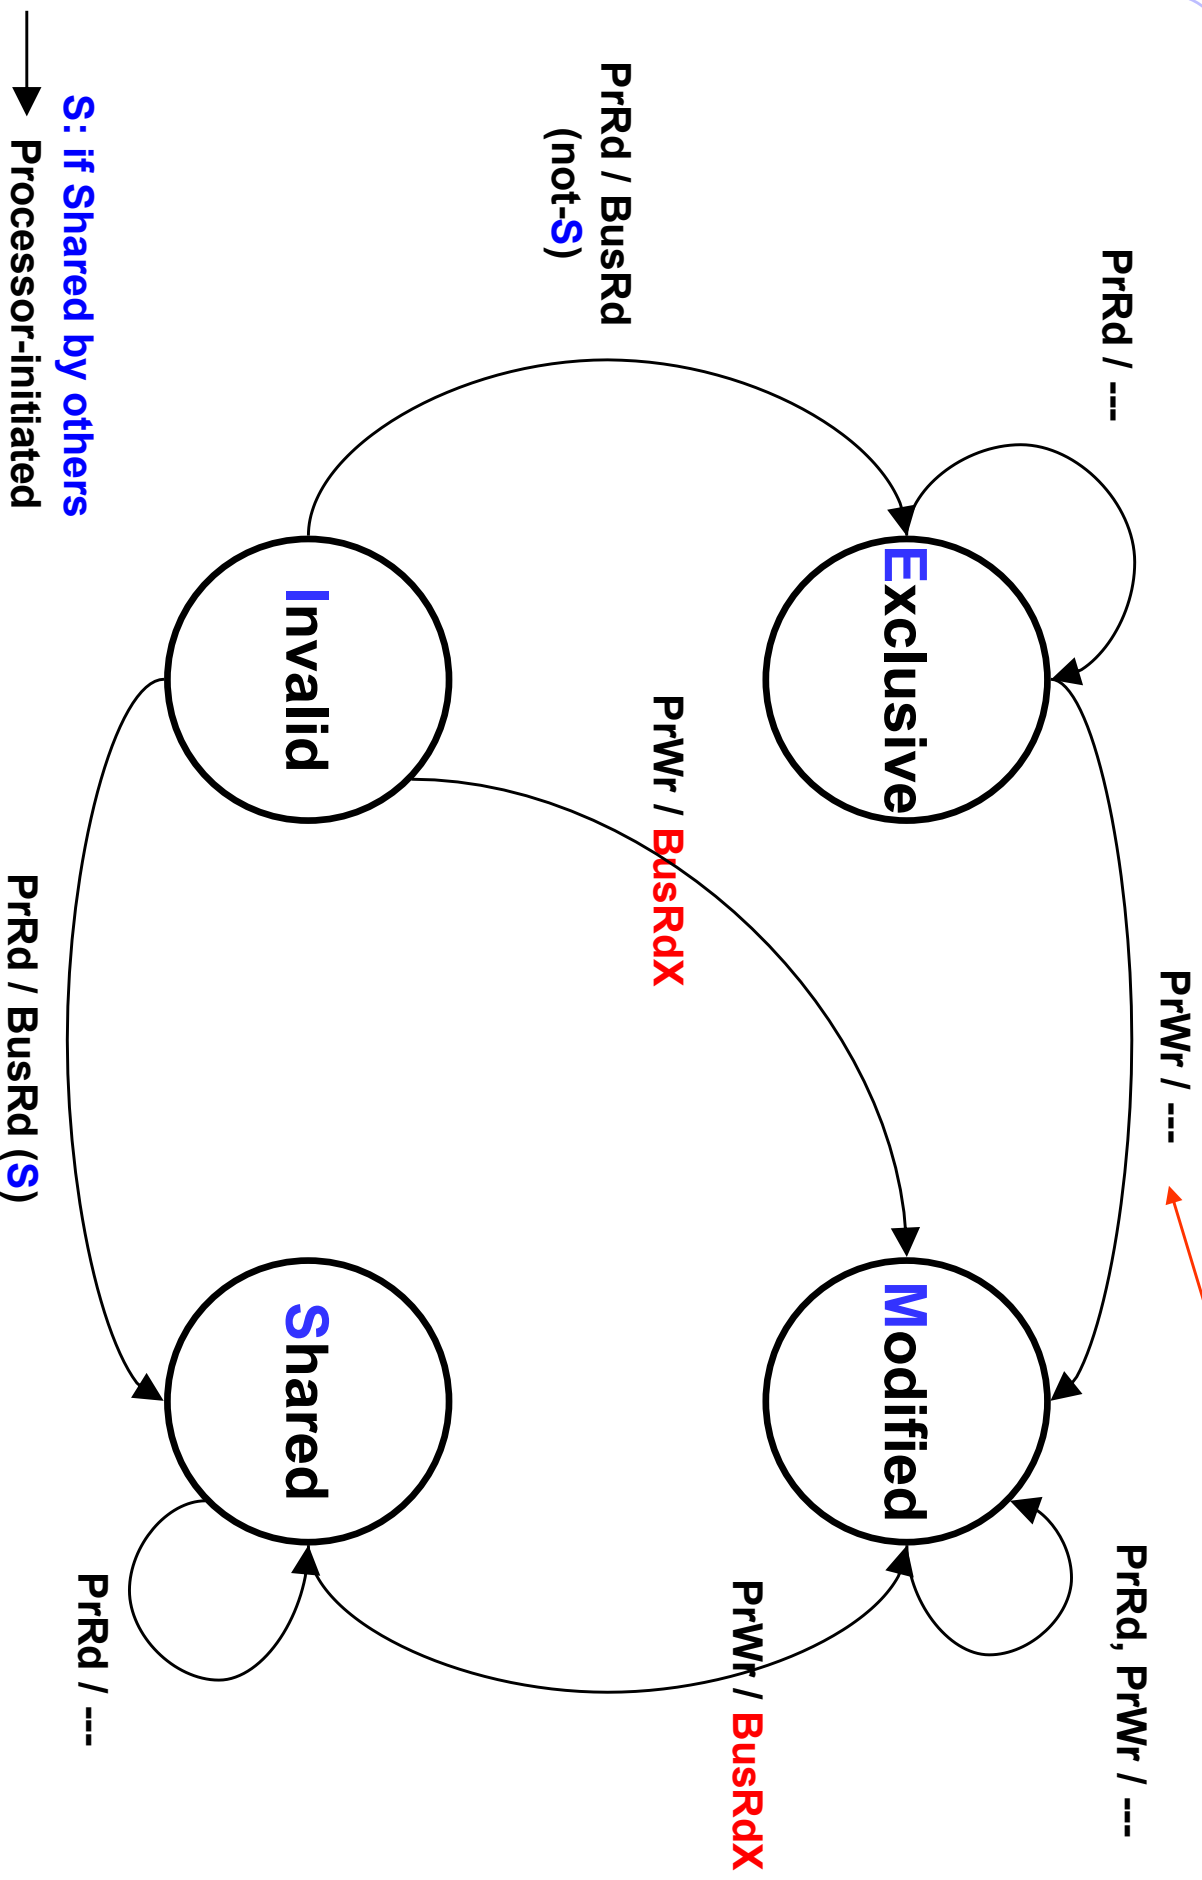
# MESI Writeback Invalidation Protocol

- Introduce the **Exclusive** state
  - One can write to the single copy **without** generating **BusRdX**
  - Very useful for running single-thread program
- Illinois Protocol: Proposed by Pamarcos and Patel in 1984
- Employed in Intel, PowerPC, MIPS

# MESI Writeback Invalidation Protocol

## Processor Request (Illinois Protocol)


Major Difference  
Save Snooping



# MESI Writeback Invalidation Protocol

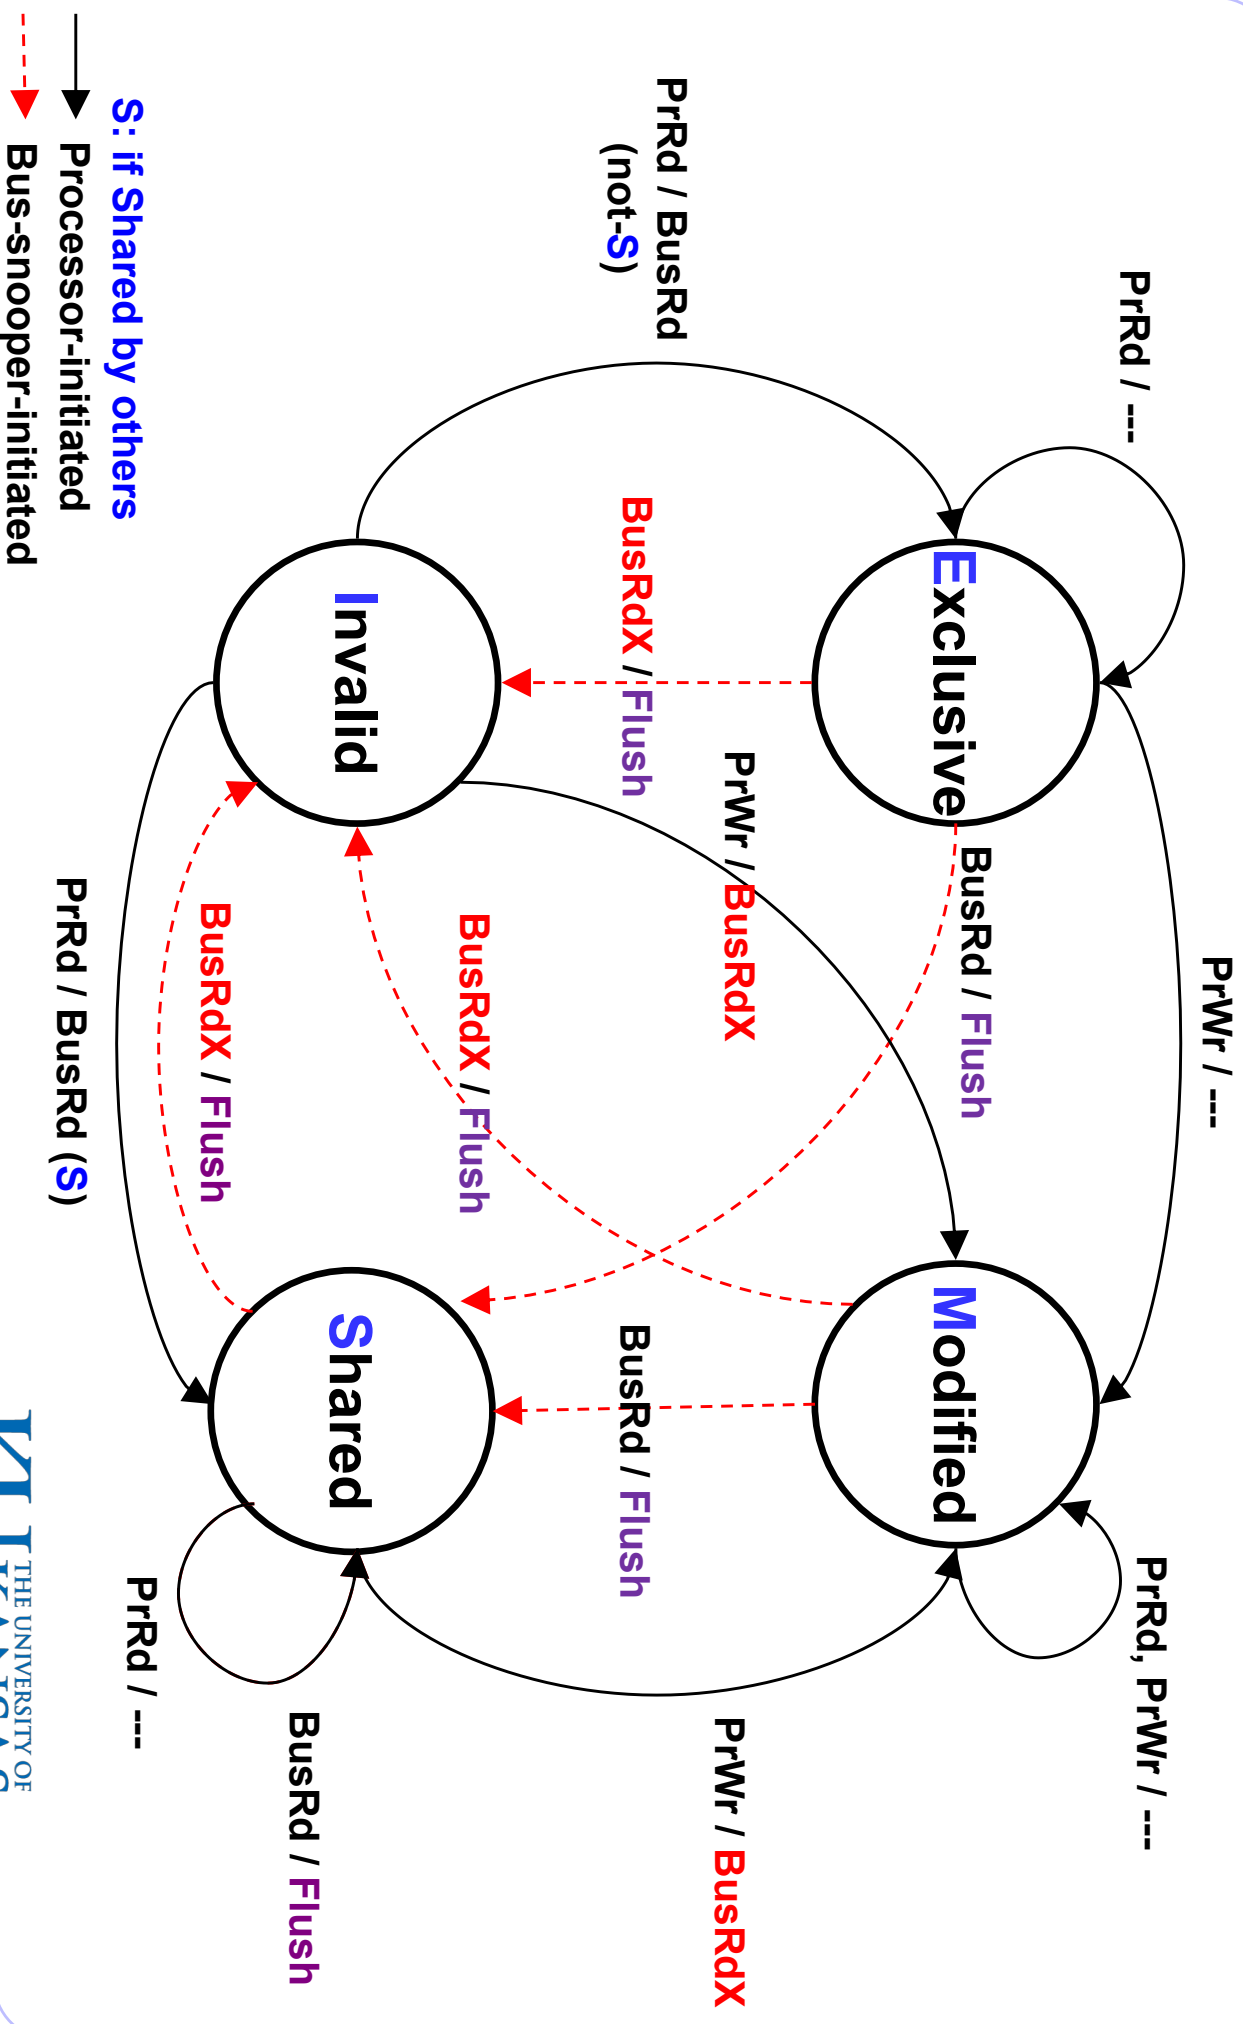
## Bus Transactions (Illinois Protocol)

**Whenever possible**, MESI protocol performs \$-to-\$ transfer rather than having memory to supply the data. Use a **Selection algorithm** if there are multiple suppliers



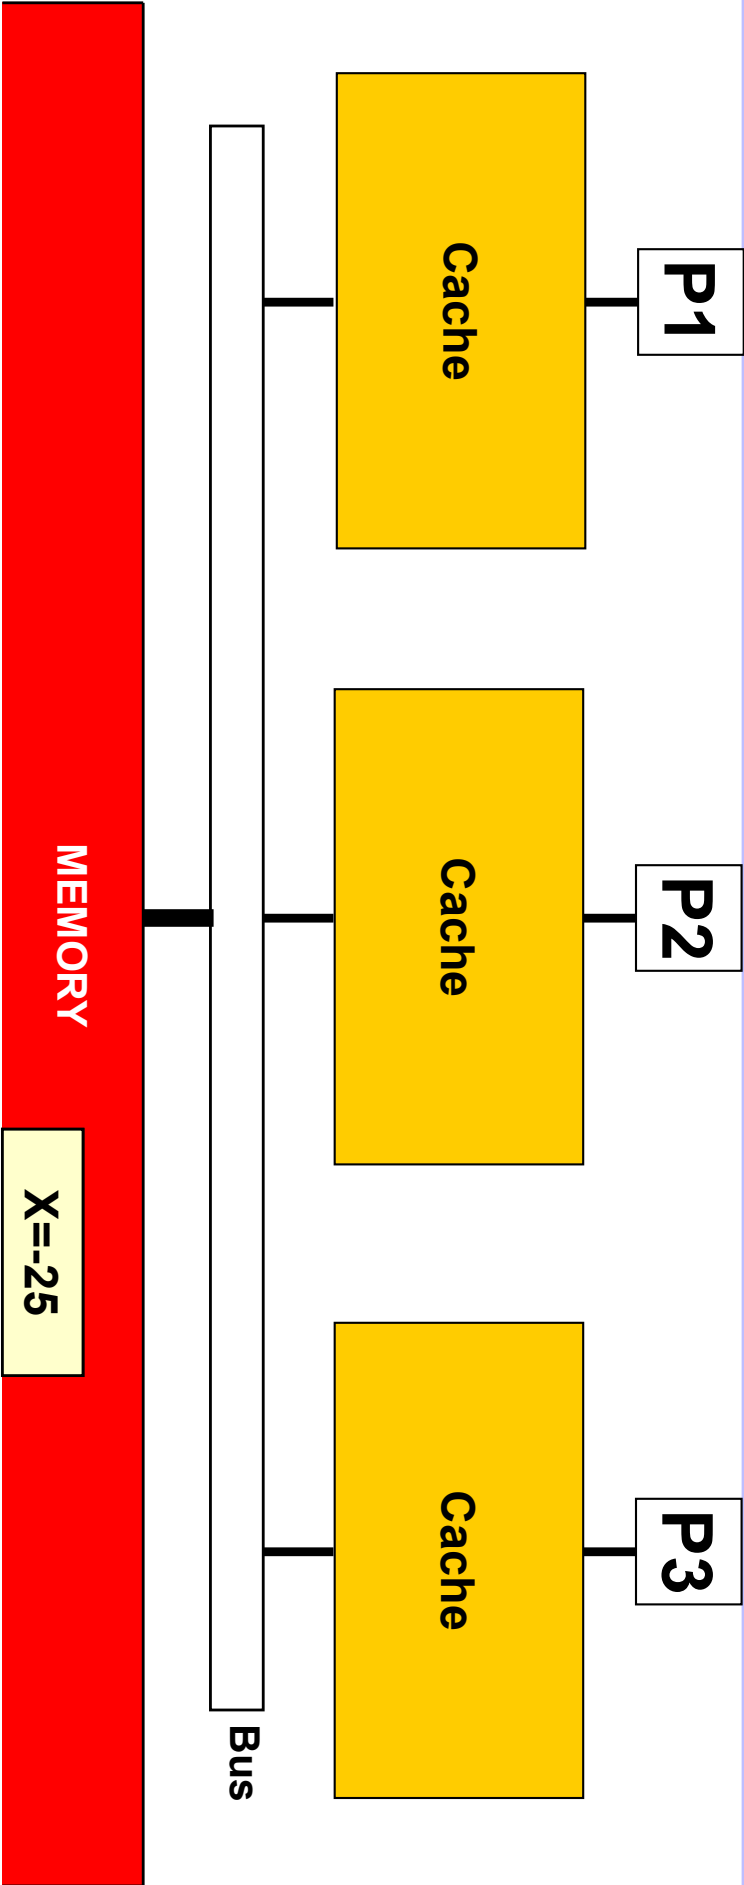
-----> Bus-snooper-initiated

# MESI Writeback Invalidation Protocol (Illinois Protocol)



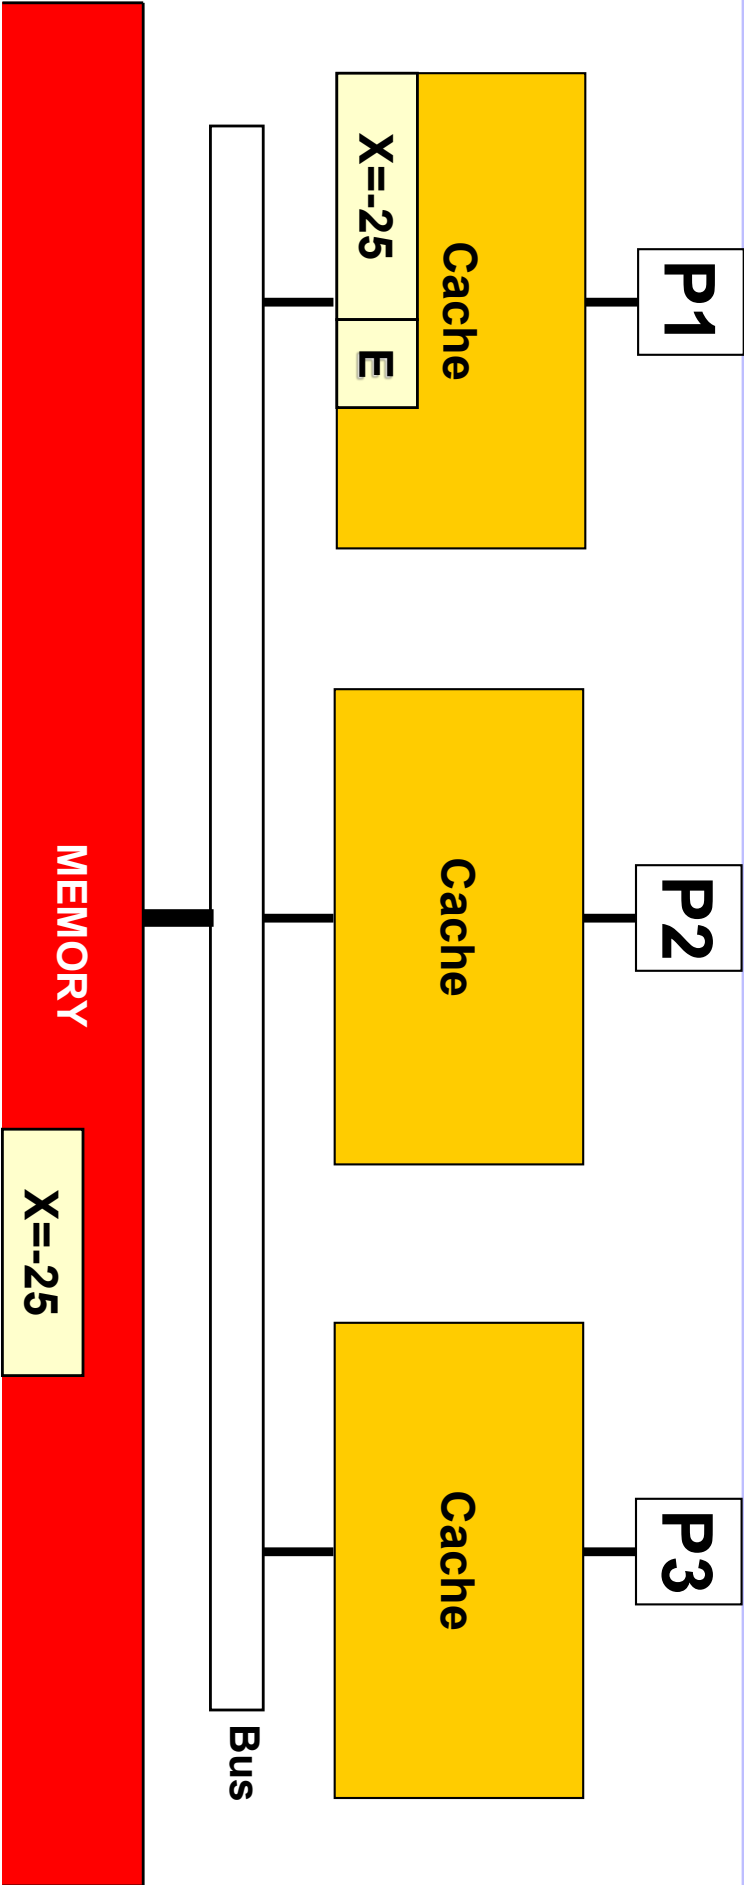


# MESI Example



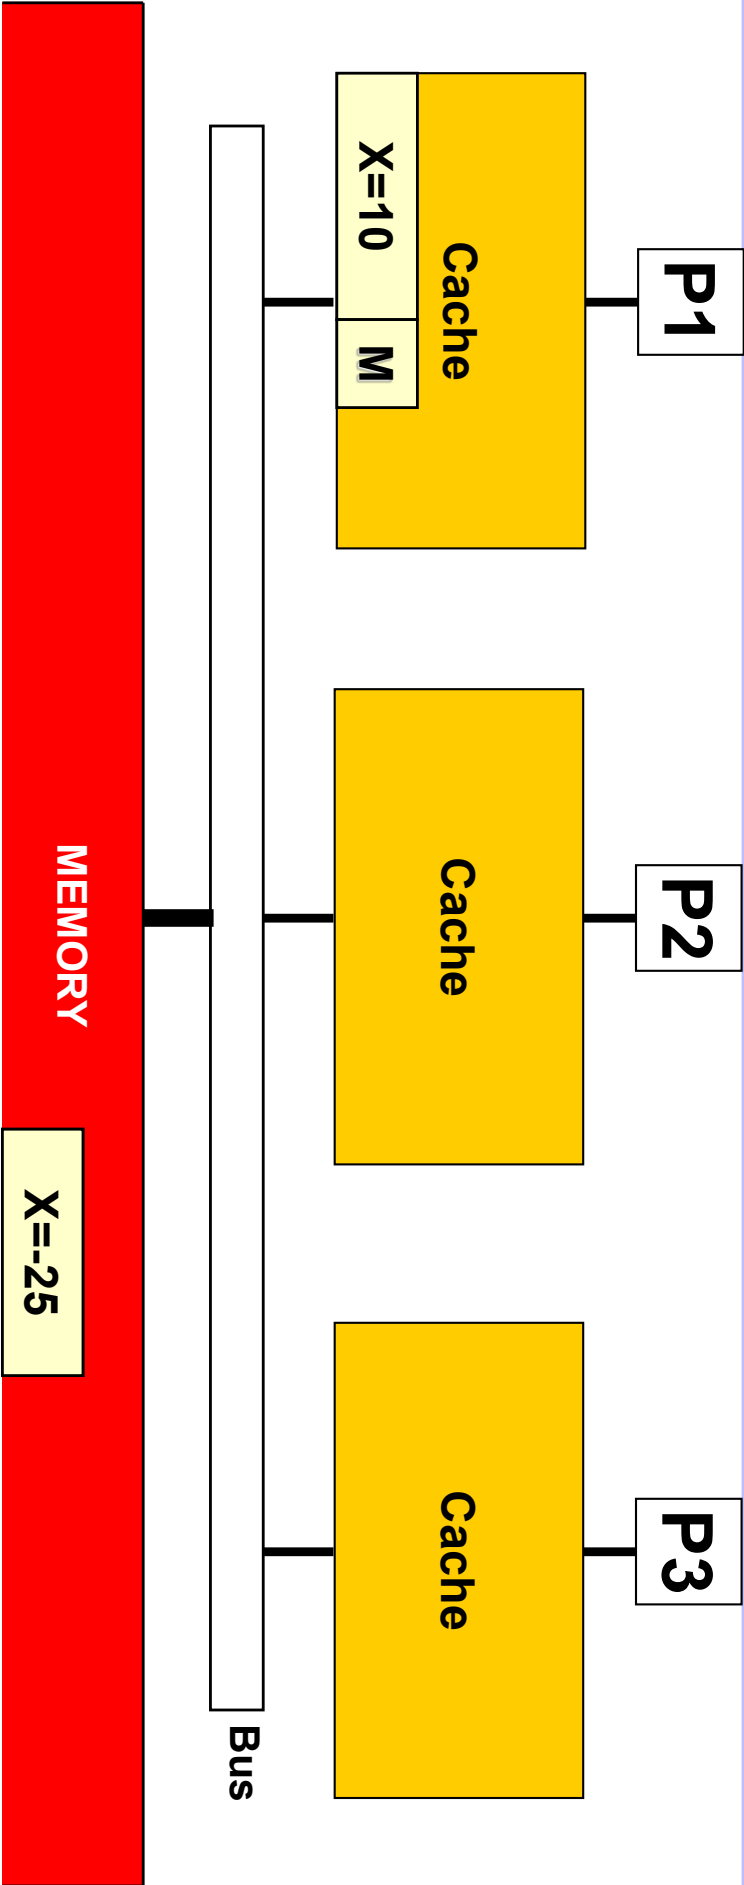
| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       |             |             |             |                 |               |
| P1 writes X      |             |             |             |                 |               |
| P3 writes X      |             |             |             |                 |               |
| P1 reads X       |             |             |             |                 |               |
| P2 reads X       |             |             |             |                 |               |

# MESI Example




| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P1 writes X      |             |             |             |                 |               |
| P3 writes X      |             |             |             |                 |               |
| P1 reads X       |             |             |             |                 |               |
| P2 reads X       |             |             |             |                 |               |

# MESI Example




| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P1 writes X      | M           | I           | I           | ---             | ---           |
| P3 writes X      |             |             |             |                 |               |
| P1 reads X       |             |             |             |                 |               |
| P2 reads X       |             |             |             |                 |               |

# MESI Example




| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P1 writes X      | M           | I           | I           | ---             | ---           |
| P3 writes X      | I           | I           | M           | BusRdX          | P1 Cache      |
| P1 reads X       |             |             |             |                 |               |
| P2 reads X       |             |             |             |                 |               |

# MESI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P1 writes X      | M           | I           | I           | ---             | ---           |
| P3 writes X      | I           | I           | M           | BusRdX          | P1 Cache      |
| P1 reads X       | S           | I           | S           | BusRd           | P3 Cache      |
| P2 reads X       |             |             |             |                 |               |

# MESI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P1 writes X      | M           | I           | I           | ---             | ---           |
| P3 writes X      | I           | I           | M           | BusRdX          | P1 Cache      |
| P1 reads X       | S           | I           | S           | BusRd           | P3 Cache      |
| P2 reads X       | S           | S           | S           | BusRd           | P1 or P3      |



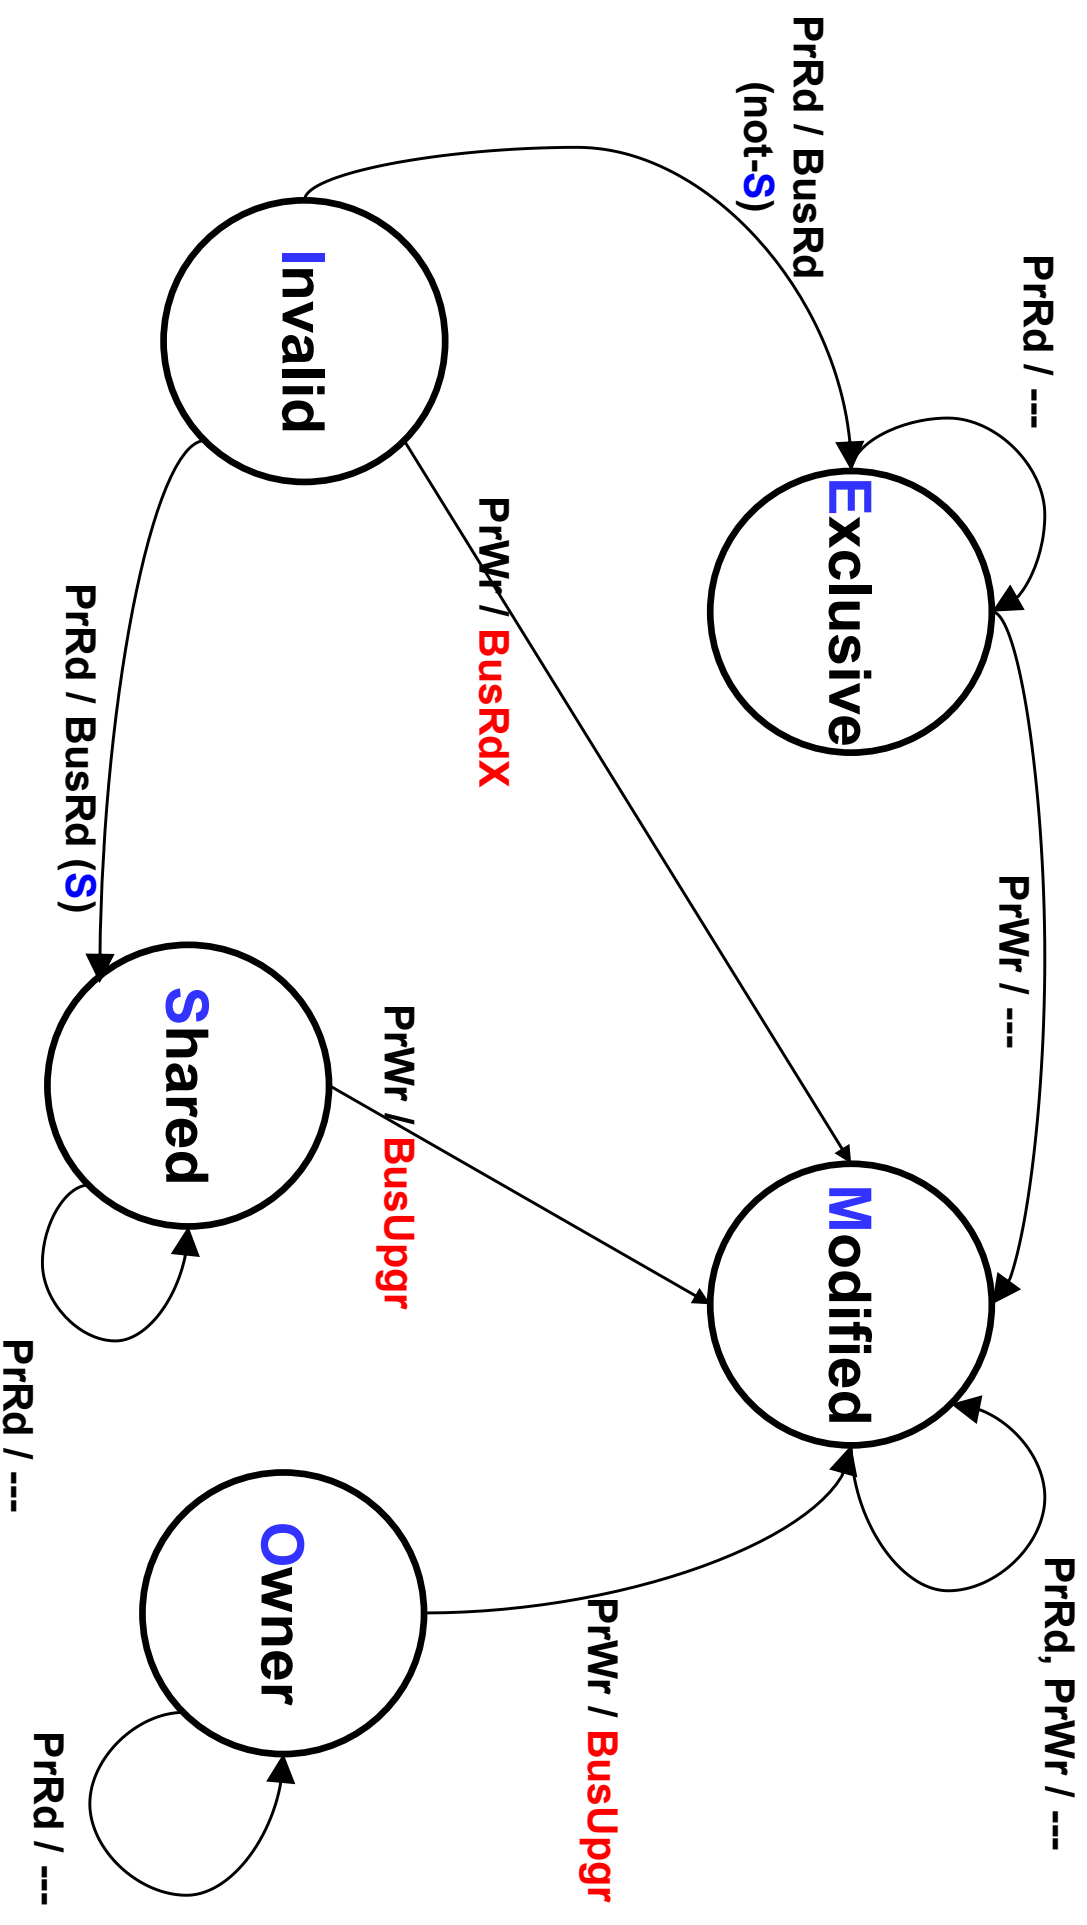
# MOESI Protocol

- Add one additional state — **Owner state** — to further save bus snooping signals
- Similar to **S**hared state, and **S** state in MOESI **may be** dirty, which is different from MESI
- The **O** state is **dirty** and will be responsible for supplying data and update memory
- Employed by
  - Sun UltraSparc
  - AMD Opteron

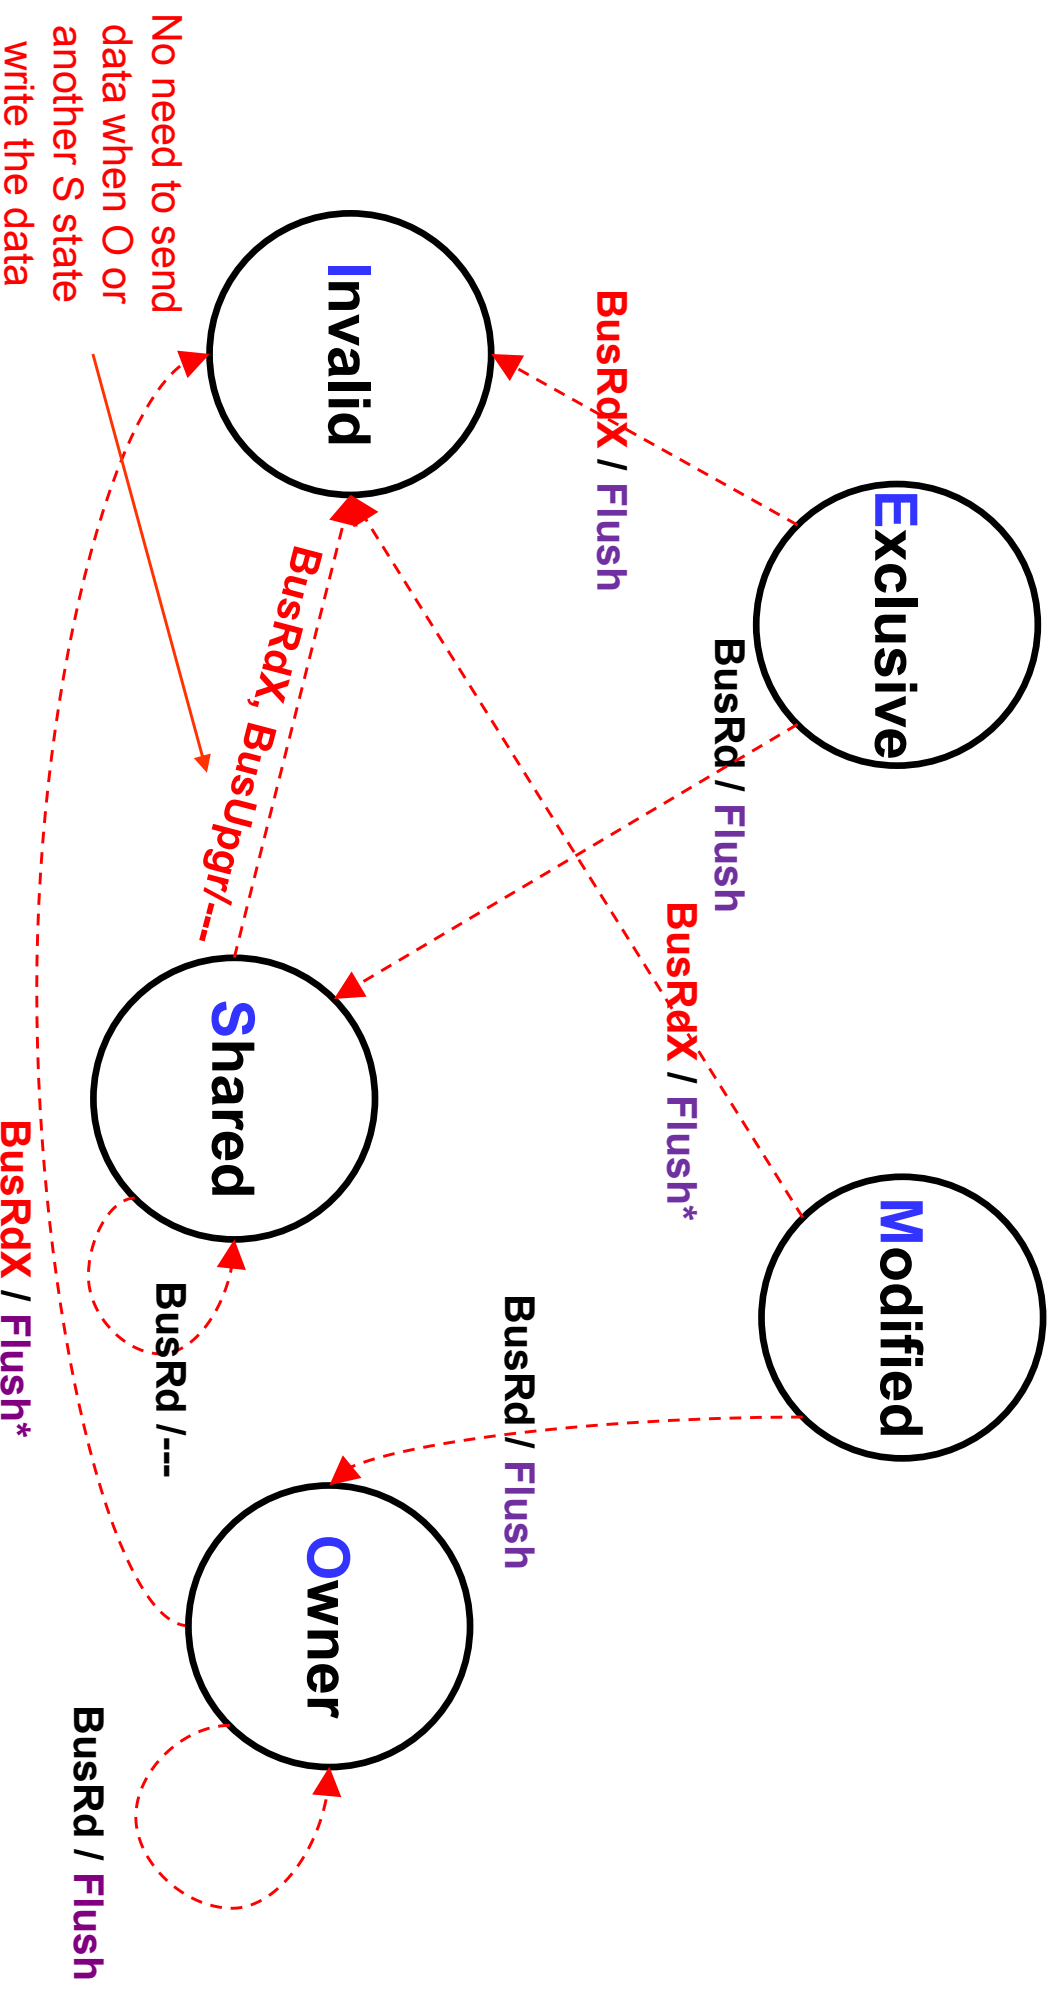




# MOESI Writeback Invalidation Protocol



# MOESI Writeback Invalidation Protocol



**BusUpgr**: invalidate other cache without incurring data transfer

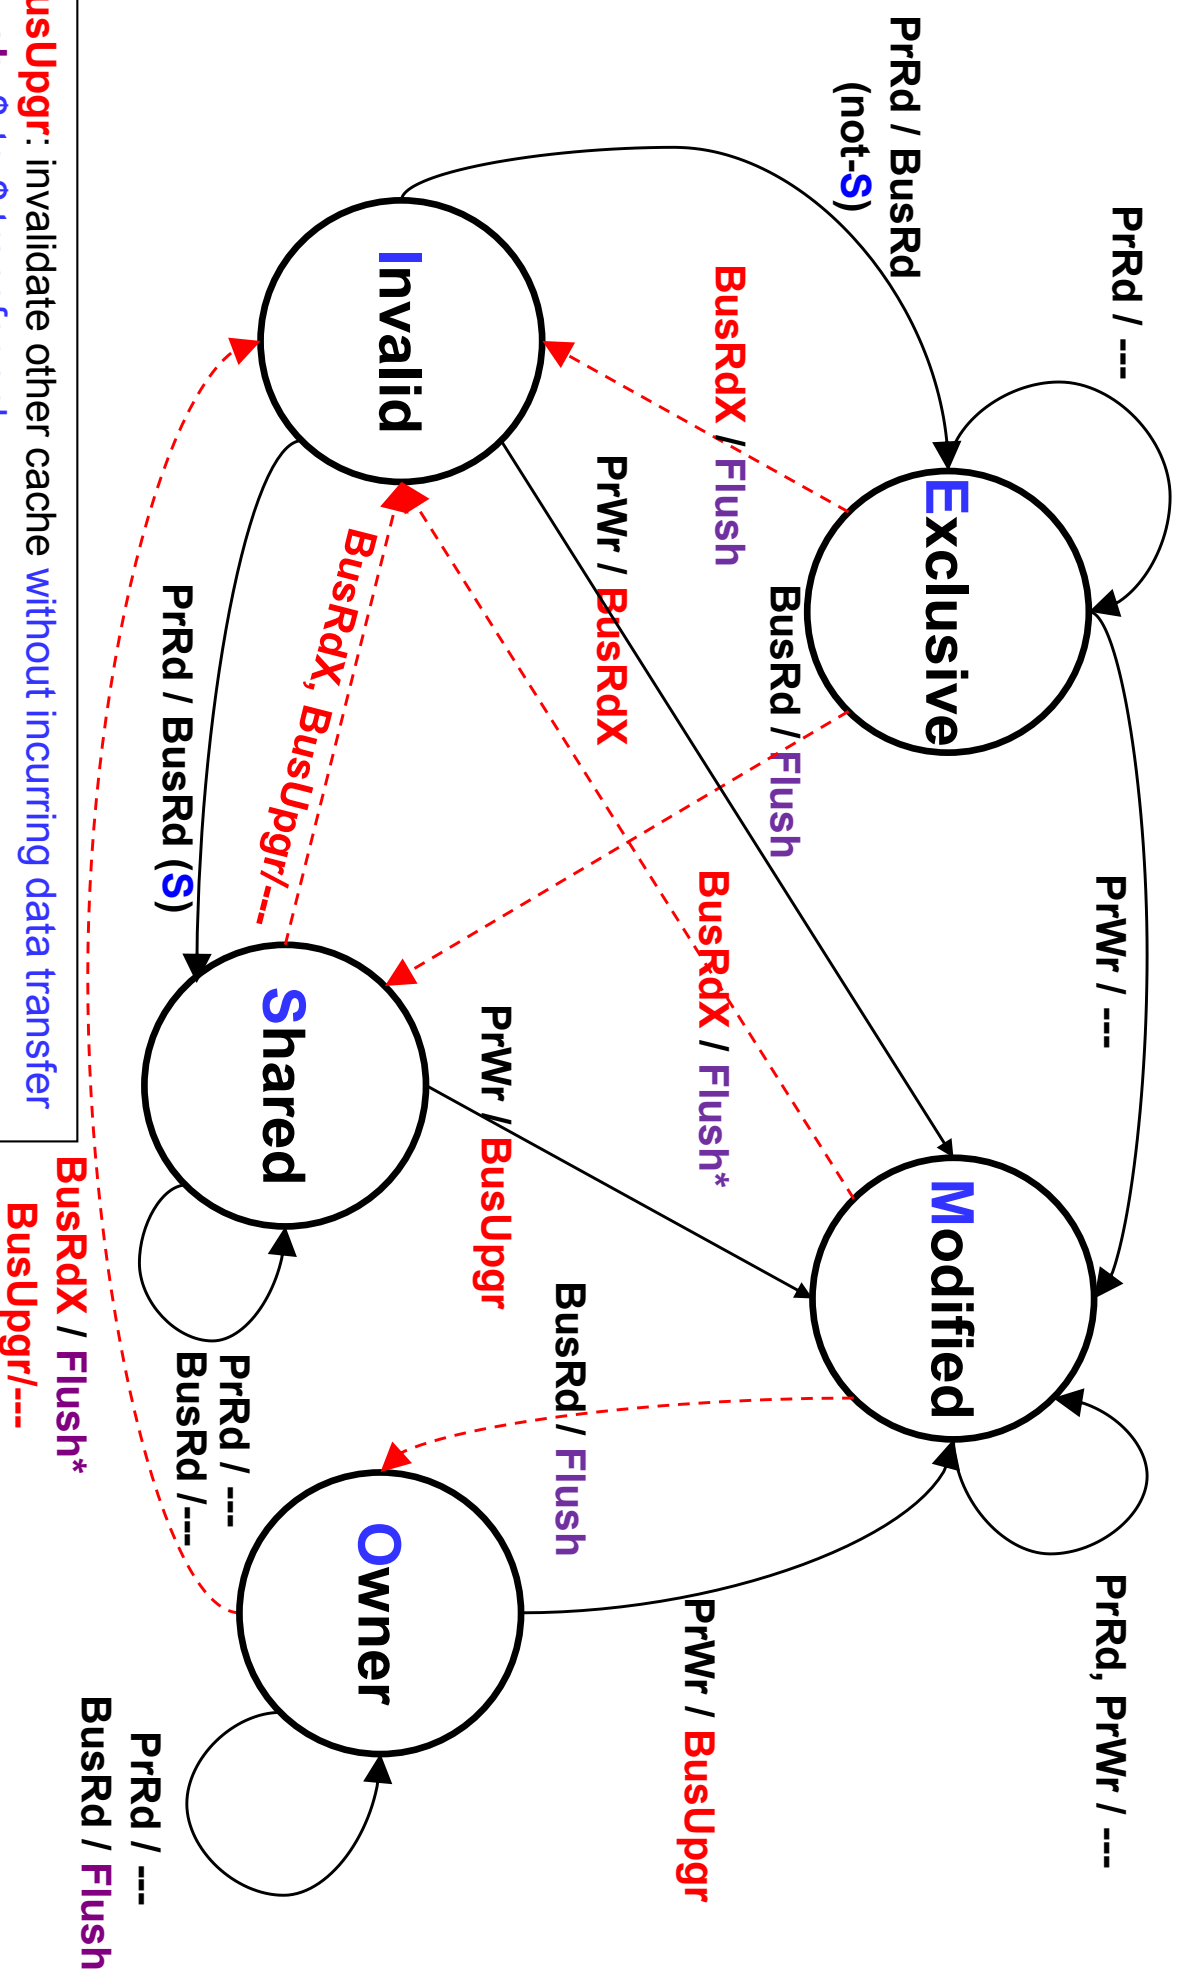
**Flush**: \$-to-\$ transfer only

**Flush\***: \$-to-\$ transfer and \$-to-mem transfer

**BusRdX / Flush\***  
**BusUpgr / ---**

No need to send data when another S state write the data


# MOESI Writeback Invalidation Protocol



**BusUpgr**: invalidate other cache without incurring data transfer  
**Flush**: \$-to-\$ transfer only  
**Flush\***: \$-to-\$ transfer and \$-to-mem transfer

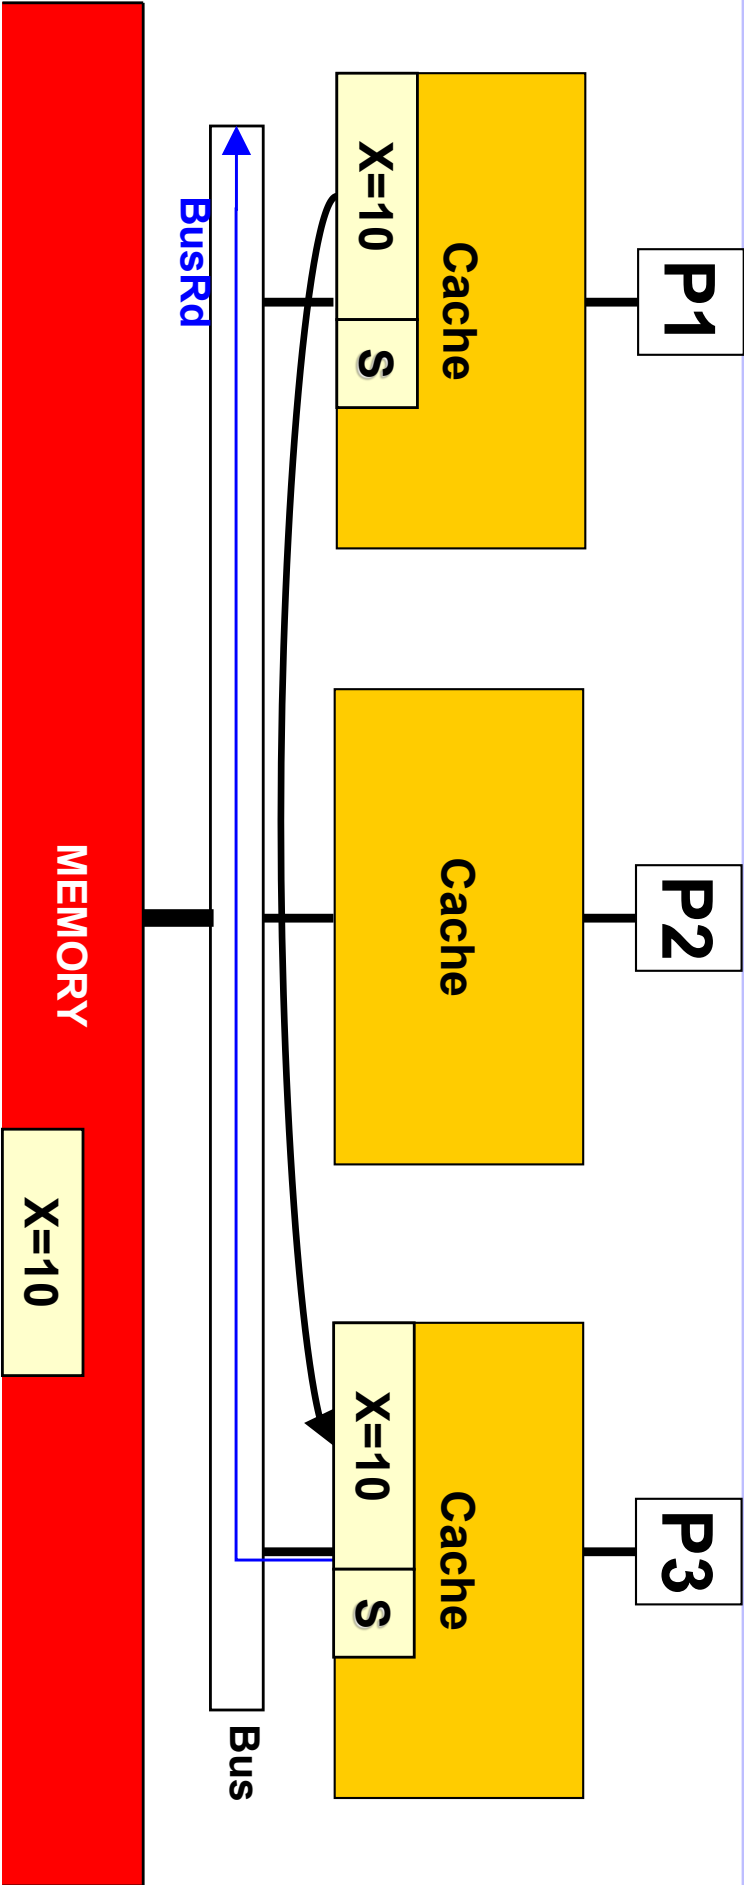
**BusRdX / Flush\***  
**BusUpgr/---**

# MOESI Example



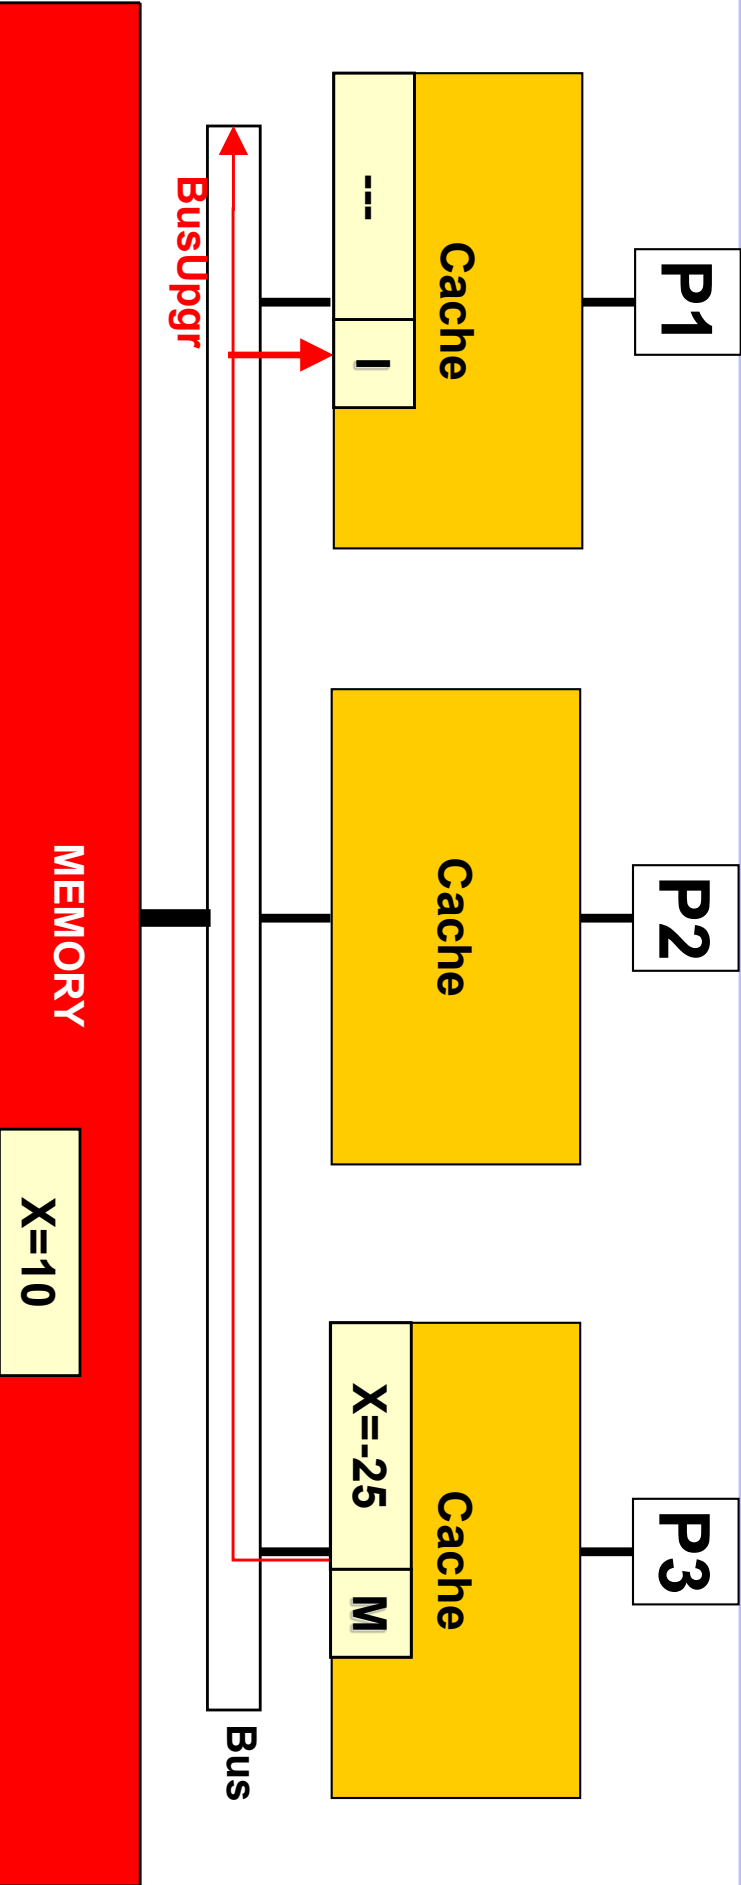
| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |

# MOESI Example



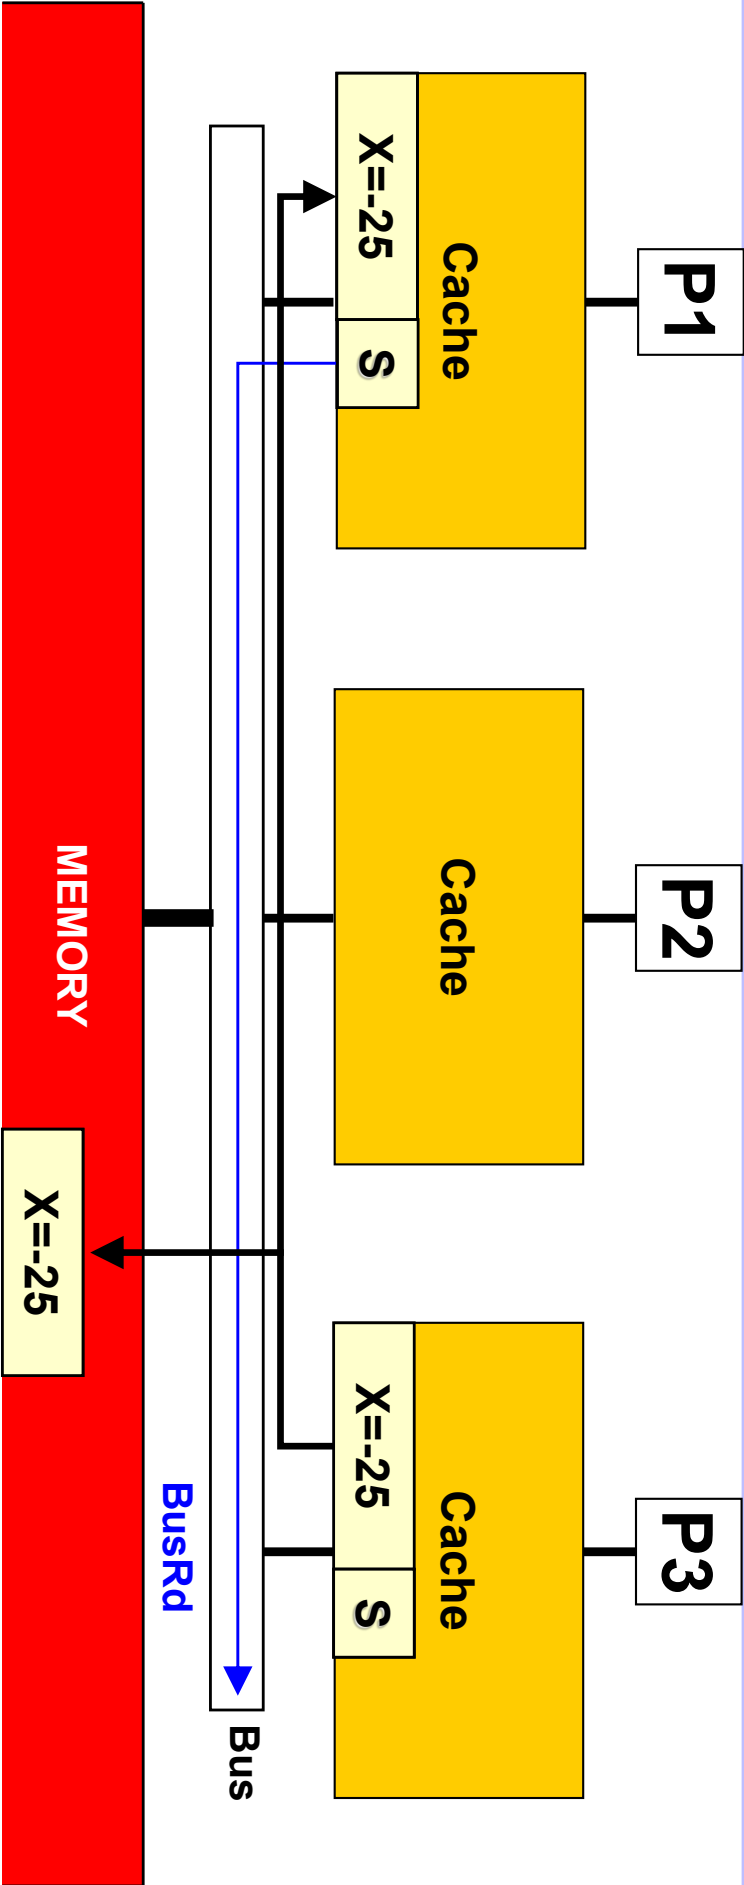
| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | P1 Cache      |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |

# MOESI Example



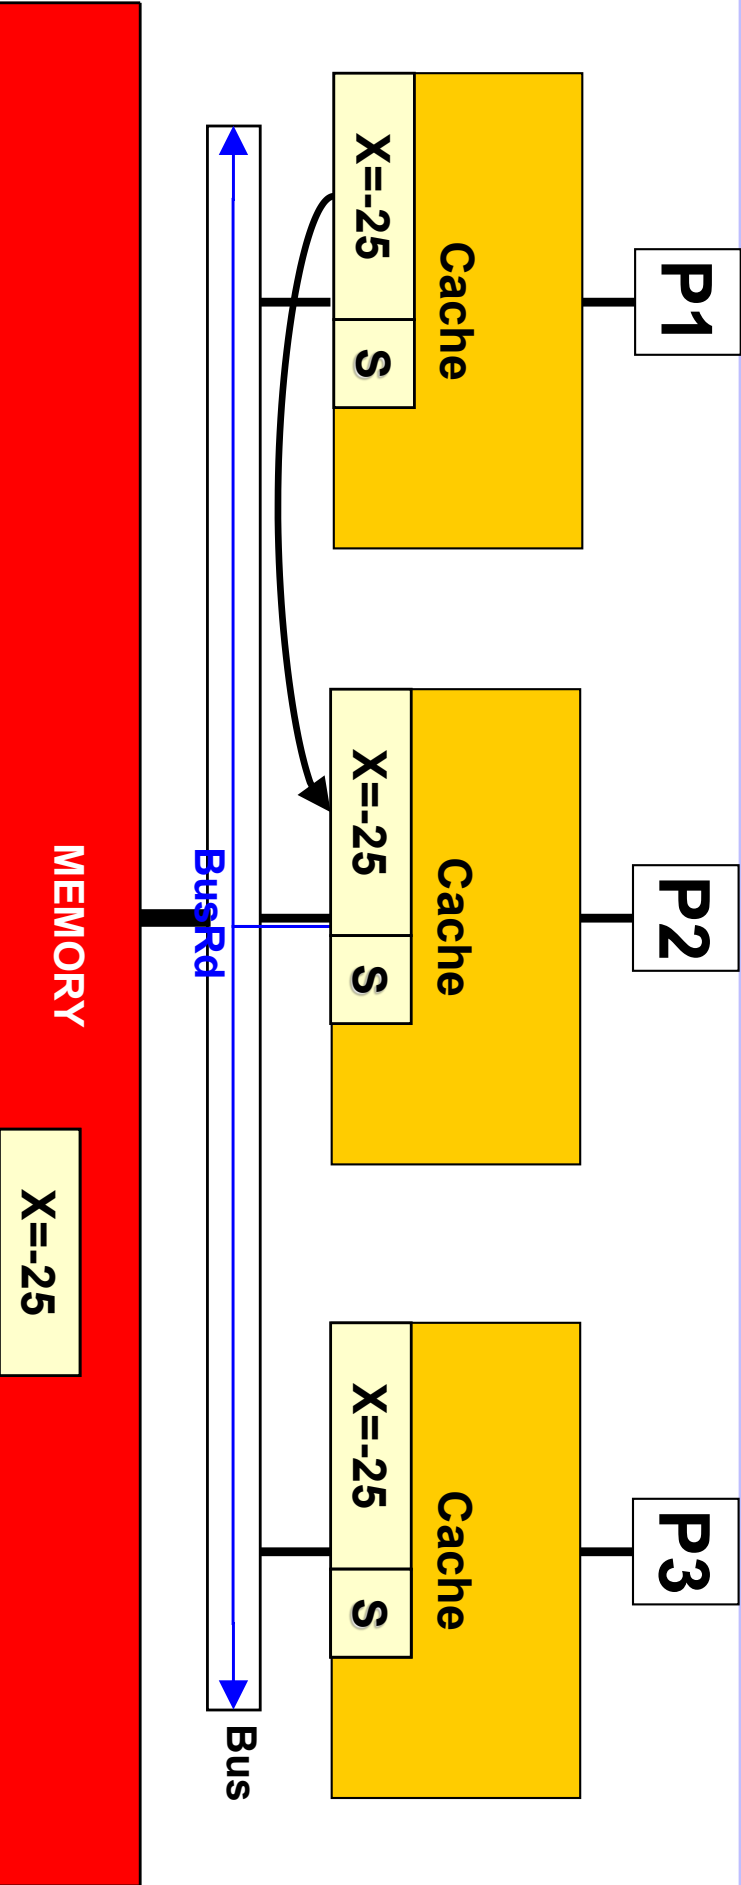
| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | P1 Cache      |
| P3 writes X      | I           | I           | M           | BusUpgr         | ---           |
|                  |             |             |             |                 |               |
|                  |             |             |             |                 |               |

# MOESI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | P1 Cache      |
| P3 writes X      | I           | I           | M           | BusUpgr         | ---           |
| P1 reads X       | S           | I           | O           | BusRd           | P3 Cache      |
|                  |             |             |             |                 |               |

# MOESI Example



| Processor Action | State in P1 | State in P2 | State in P3 | Bus Transaction | Data Supplier |
|------------------|-------------|-------------|-------------|-----------------|---------------|
| P1 reads X       | E           | I           | I           | BusRd           | Memory        |
| P3 reads X       | S           | I           | S           | BusRd           | P1 Cache      |
| P3 writes X      | I           | I           | M           | BusUpgr         | ---           |
| P1 reads X       | S           | I           | O           | BusRd           | P3 Cache      |
| P2 reads X       | S           | S           | O           | BusRd           | P3 Cache      |

How to avoid unnecessary bus snooping?