

# Midterm Review

# OS Structure

program always runs

- User mode/ kernel mode (Dual-Mode)
  - Memory protection, privileged instructions
    - Hardware support to distinguish app/kernel
    - Privileged instructions are only for kernel mode
    - Applications can enter into kernel mode only via pre-defined system calls

- System call
  - Definition, examples, how it works?

system calls

- Software interrupt. “int <num>” in Intel x86
- Programming interface to the services provided by the OS

- Other concepts to know

- Monolithic kernel vs. Micro kernel

Implements CPU scheduling, memory management, filesystems, and other OS modules all in a single big

chunk

Pros and Cons

+ Overhead is low

+ Data sharing among the modules is easy

– Too big.(devicedrivers!!!)

– A bug in one part of the kernel can crash the entire system

Moves as much from the kernel into user space

Communicate among kernels and user

via message passing

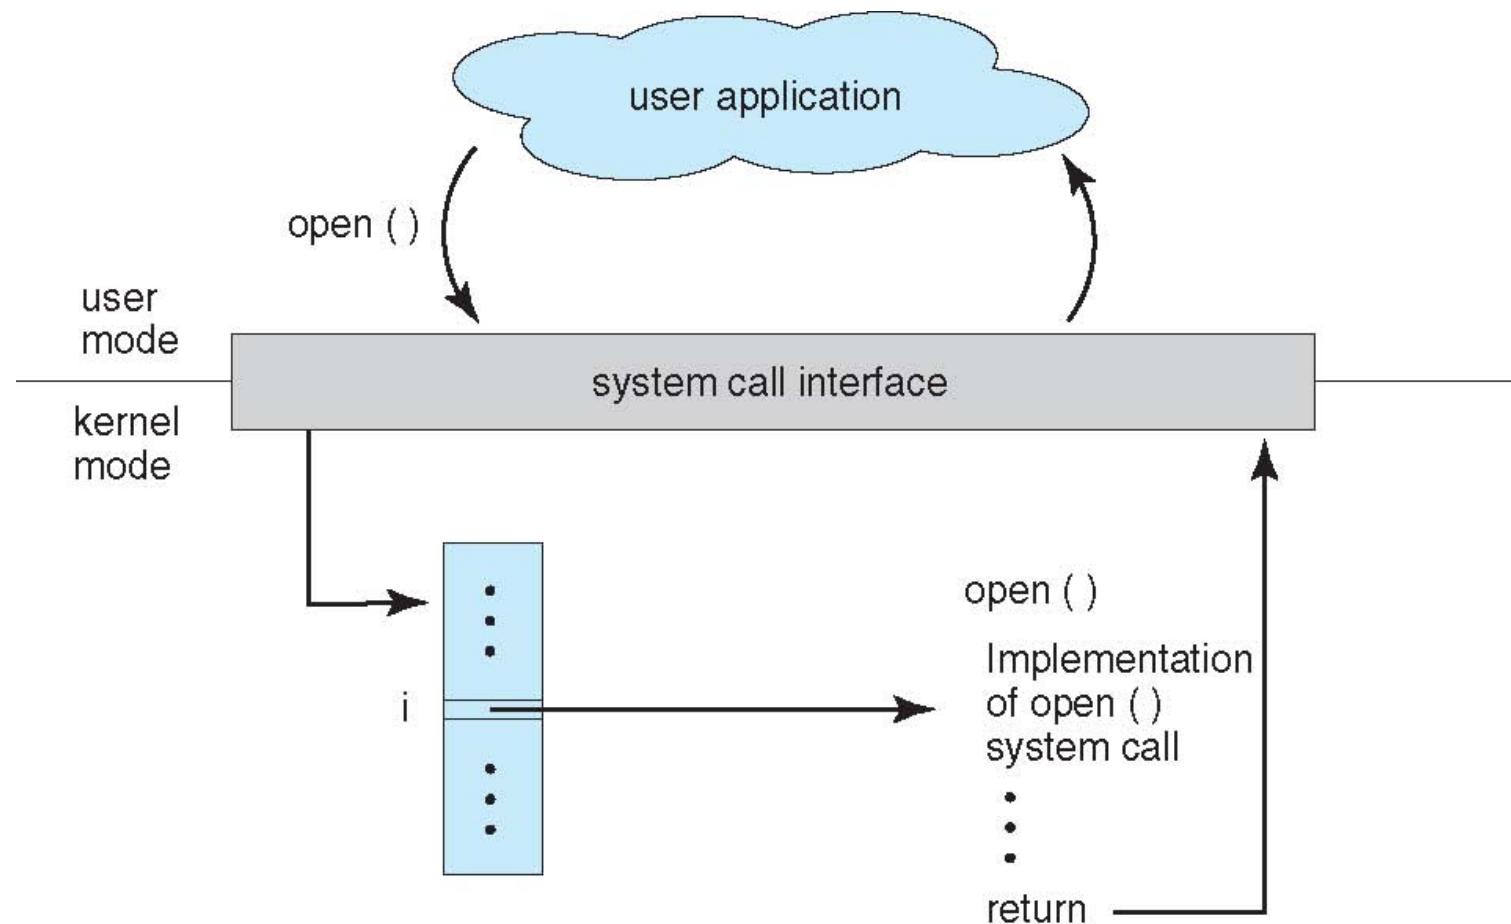
Pros and Cons

+ Easy to extend (user level driver)

+ More reliable (less code is running in kernel mode)

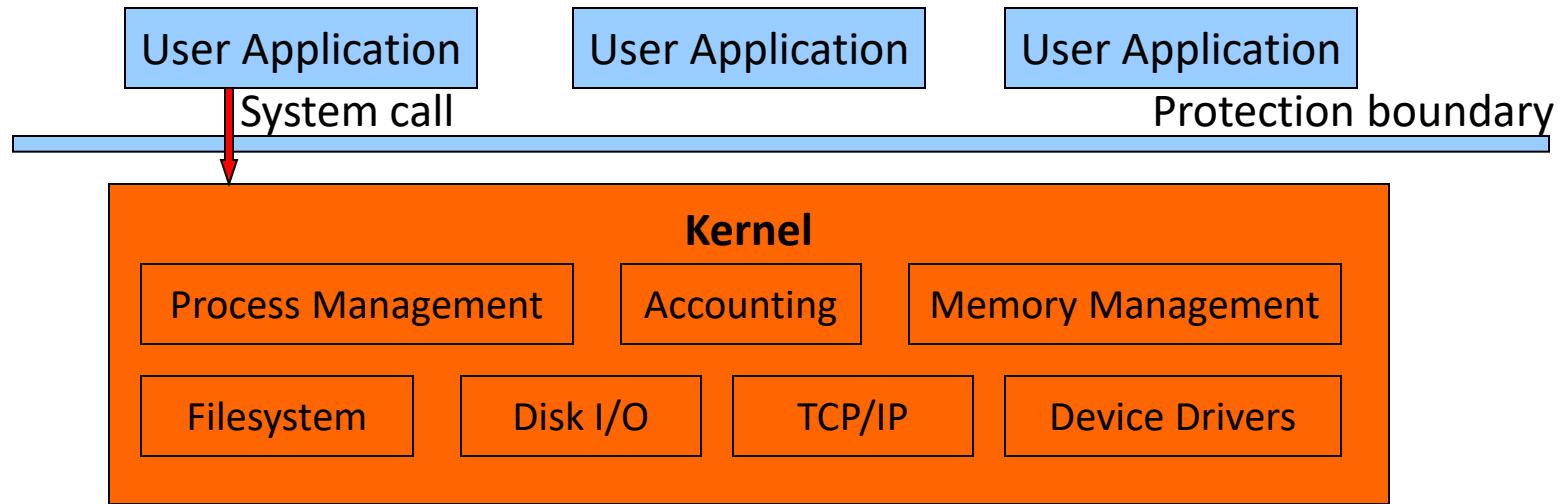
- Performance overhead of user space to kernel space communication

# API - System Call - OS



# UNIX: Monolithic Kernel

- Implements CPU scheduling, memory management, filesystems, and other OS modules all in a single big chunk



- Pros and Cons
  - + Overhead is low
  - + Data sharing among the modules is easy
  - Too big. (device drivers!!!)
  - A bug in one part of the kernel can crash the entire system

# Process

- Address space layout
  - Code, data, heap, stack
- Process states
  - new, ready, running, waiting, terminated

- Other concepts to know
  - Process Control Block
  - Context switch
  - Zombie, Orphan
  - Communication overheads of processes vs. threads

PCB: process state

PID

PC

register

CPU scheduling information

memory-management

information

Account information

OS resources

Suspend the current process and resume  
a next one from its last suspended state

Save&restoreCPUcontext

– Change address space and other info in  
the PCB

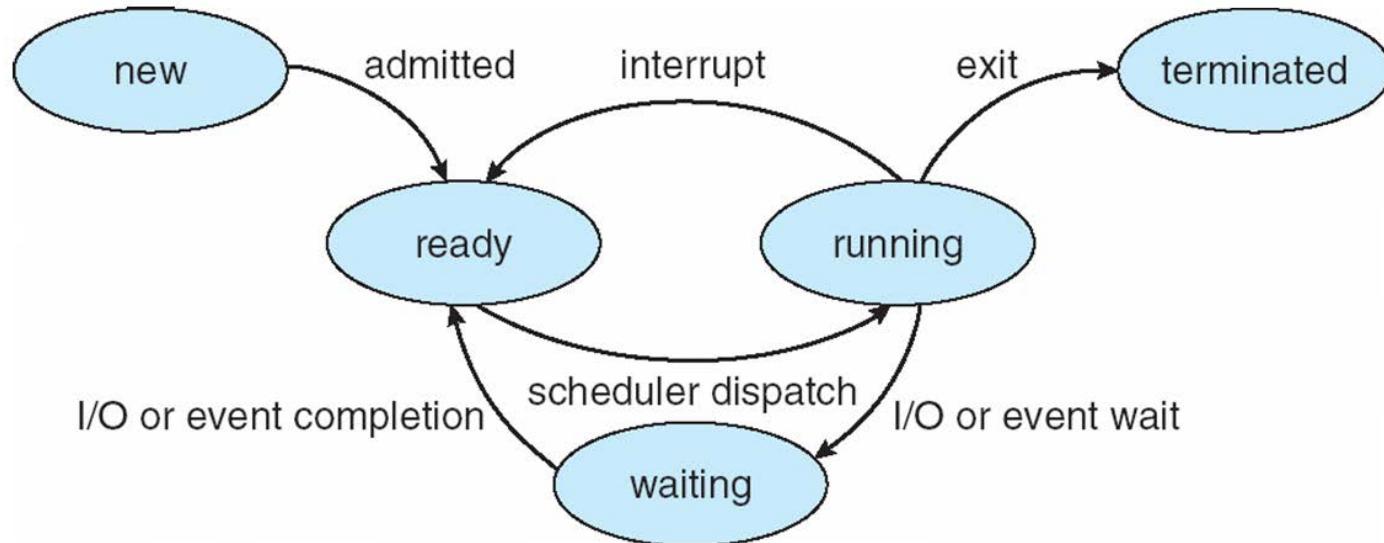
If no parent waiting (did not invoke wait())

If parent terminated without invoking wait

– Q: who will be the parent of a orphan process?

– A:Initprocess

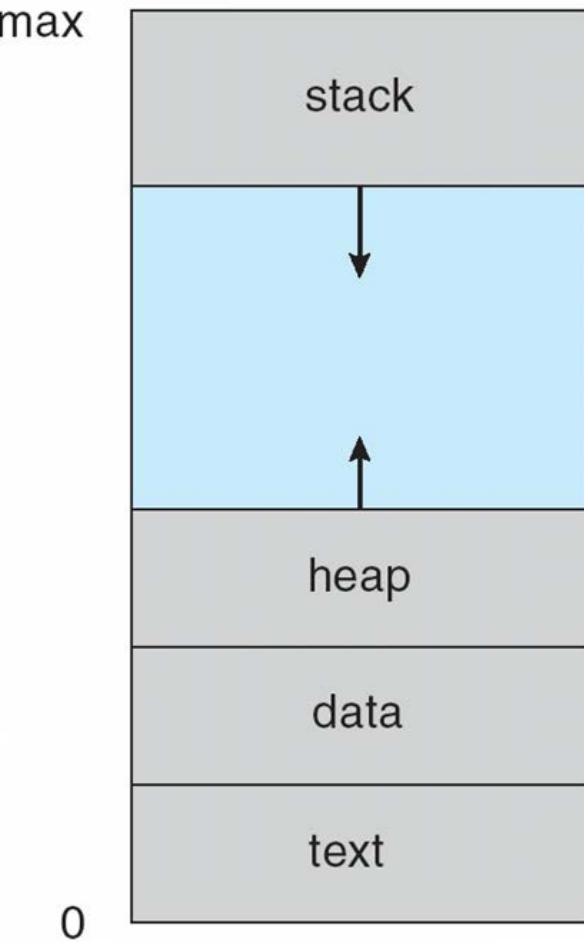
# Process State



- **running:** Instructions are being executed
- **waiting:** The process is waiting for some event to occur
- **ready:** The process is waiting to be assigned to a processor

# Process Address Space

- Text
  - Program code
- Data
  - Global variables
- Heap
  - Dynamically allocated memory
    - i.e., Malloc()
- Stack
  - Temporary data
  - Grow at each function call



# Quiz

```
int count = 0;  
int main()  
{  
    int pid = fork();  
    if (pid == 0){  
        count++;  
        printf("Child: %d\n", count);  
    } else{  
        wait(NULL);  
        count++;  
        printf("Parent: %d\n", count);  
    }  
    count++;  
    printf("Main: %d\n", count);  
    return 0;  
}
```

- Hints
  - Each process has its own private address space
  - Wait() blocks until the child finish
- Describe the output.
  - Child: 1
  - Main: 2
  - Parent: 1
  - Main: 2

# Inter-Process Communication

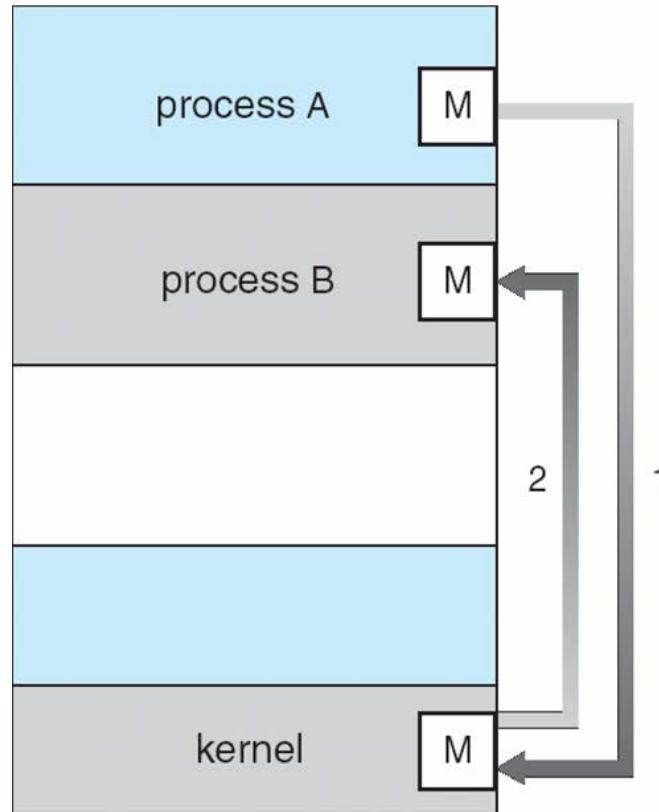
- Shared memory
- Message passing

## Shared memory

- share a region of memory between co-operating processes
  - read or write to the shared memory region –
    - ++ fast communication
    - -- synchronization is very difficult
      - Message passing
  - exchange messages (send and receive)
- typically involves data copies (to/from buffer) –
  - ++ synchronization is easier
  - -- slower communication

# Models of IPC

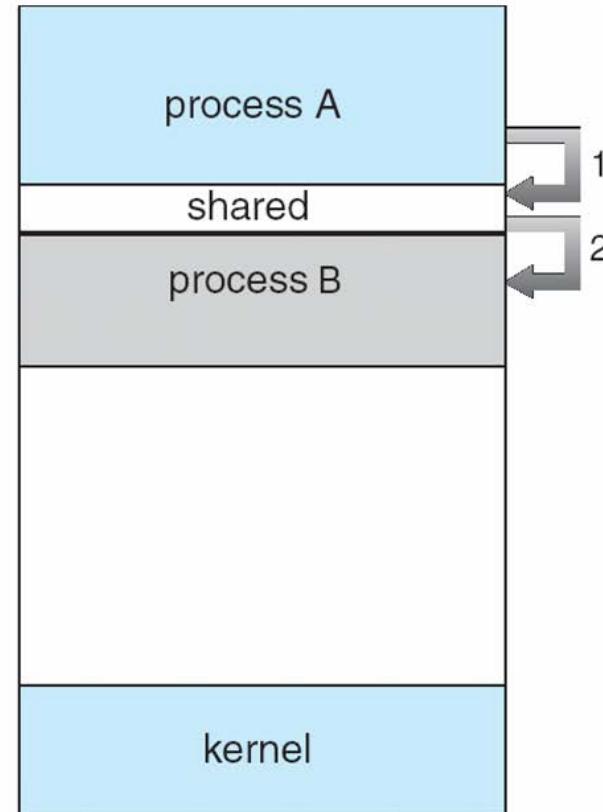
需要去内核取资源，交流缓慢



(a)

message passing

在user model 分享资源，但实现同步很困难，基本用来交流



(b)

shared memory

# Quiz

- A process produces 100MB data in memory. You want to share the data with two other processes so that each of which can access half the data (50MB each). What IPC mechanism will you use and why?
- IPC mechanism: POSIX Shared memory
- Reasons: (1) large data → need high performance, (2) no need for synchronization,

# Threads

- User threads vs. Kernel threads
  - user-threading: kernel don't understand threads
- Key benefits over processes?
  - + Performance
  - Protection



# Multi-threads vs. Multi-processes

- Multi-processes
  - (+) protection
  - (-) performance (?)
- Multi-threads
  - (+) performance
  - (-) protection

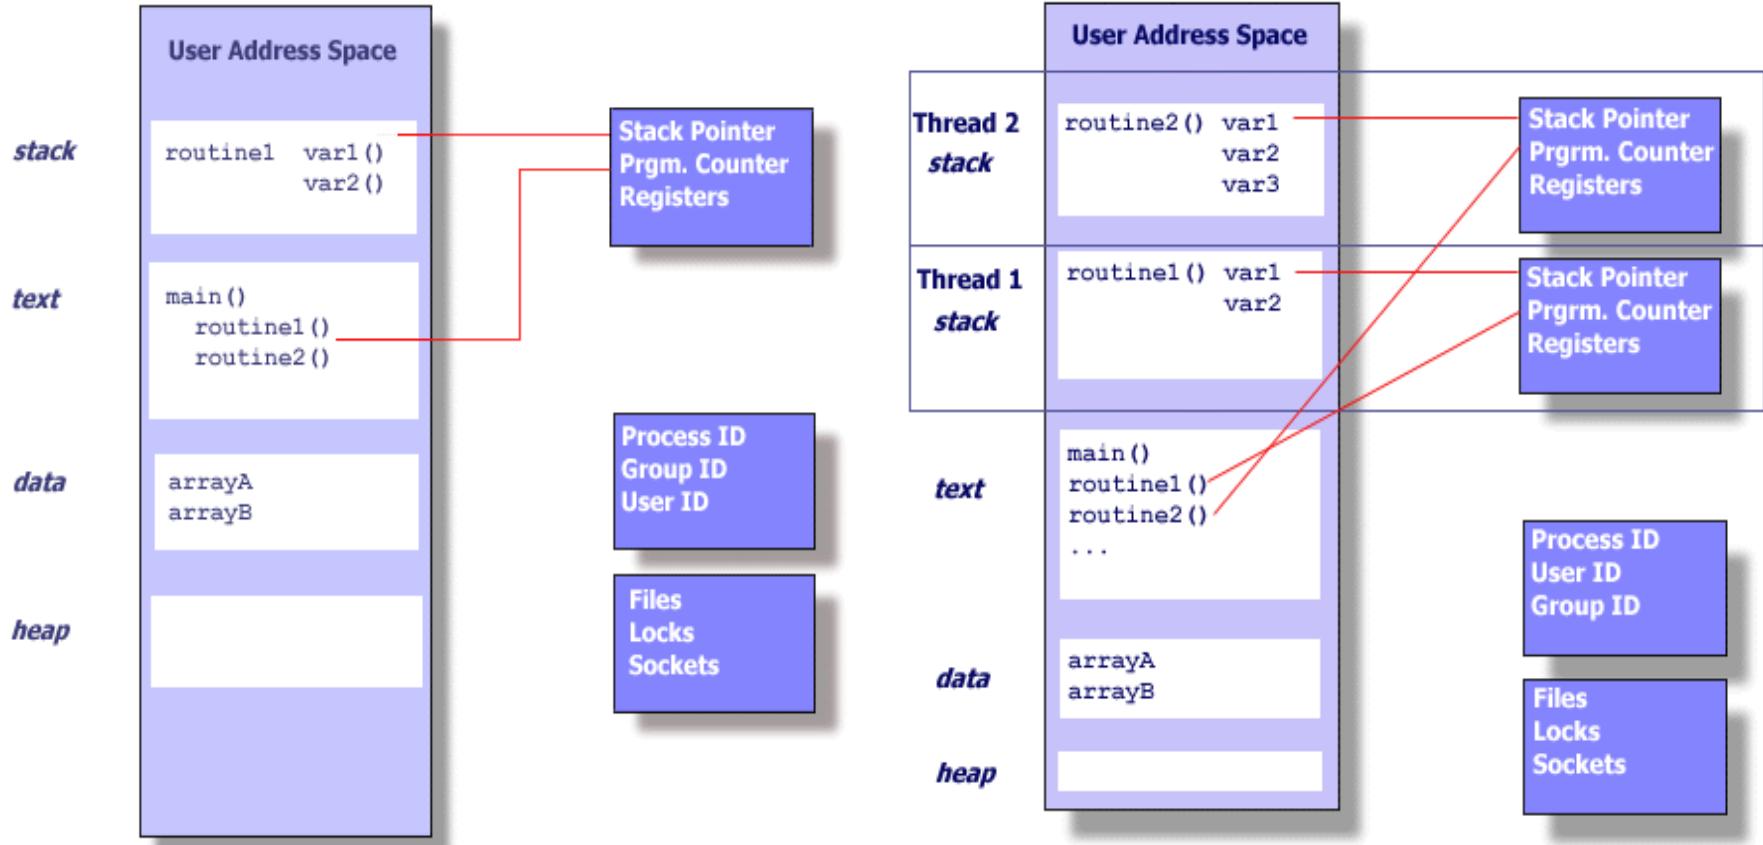


Process-per-tab



Single-process multi-threads

# Single and Multithreaded Process



source: <https://computing.llnl.gov/tutorials/pthreads/>

# Synchronization

- Race condition
  - A situation when two or more threads read and write shared data at the same time
- Synchronization instructions
  - *test&set, compare&swap*
- Spinlock
  - Spin on wait
  - Good for short critical section but can be wasteful
- Mutex
  - Block (sleep) on wait
  - Good for long critical section but bad for short one

# Race Condition

Initial condition: *counter* = 5

Thread 1

```
R1 = load (counter);  
R1 = R1 + 1;  
counter = store (R1);
```

Thread 2

```
R2 = load (counter);  
R2 = R2 - 1;  
counter = store (R2);
```

- What are the possible outcome?

5, 4, 6

# Quiz

- Write the pseudo-code definition of *TestAndSet* instruction
- Complete the following lock implementation using *TestAndSet* instruction

```
Int TestAndSet(int *lock)
{
    int ret;           ret = *lock
    _____           *lock = True
    _____
    return ret;
}
```

```
void init_lock(int *mutex)
{
    *mutex = 0;
}
void lock(int *mutex)
{
    while (_____);
}
void unlock(int *mutex)
{
    *metex = 0
    _____
}
```

# Quiz

- Pseudo code of test&set

```
int TestAndSet(int *lock) {  
    int ret = *lock;  
    *lock = 1;  
    return ret;  
}
```

```
int CAS(int *value, int oldval, int newval) {  
    int temp = *value;  
    if (*value == oldval)  
        *value = newval;  
    return temp;  
}
```

- Spinlock implementation using *TestAndSet* instruction

```
void init_lock(int *lock) {  
    *lock = 0;  
}
```

```
void lock(int *lock) {  
    while (TestAndSet(lock));  
}
```

```
void unlock(int *lock) {  
    *lock = 0;  
}
```

```
void mutex_init (mutex_t *lock)
{
    lock->value = 0;
    list_init(&lock->wait_list);
    spin_lock_init(&lock->wait_lock);
}

void mutex_lock (mutex_t *lock)
{
...
    while(TestAndSet(&lock->value)) {
        current->state = WAITING;
        list_add(&lock->wait_list, current);
...
        schedule();
...
    }
...
}

void mutex_unlock (mutex_t *lock)
{
...
    lock->value = 0;
    if (!list_empty(&lock->wait_list))
        wake_up_process(&lock->wait_list)
...
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

比spinlock更加有优势，睡眠或者等待，而不是无限循环

```
    list_init(&lock->wait_list); <--
```

Thread waiting list

```
    spin_lock_init(&lock->wait_lock); <--
```

To protect waiting list

```
}
```

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

对waiting list进行加锁

```
    while(TestAndSet(&lock->value)) {
```

Thread state change

```
        current->state = WAITING; <--
```

Add the current thread to the  
waiting list

```
        list_add(&lock->wait_list, current); <--
```

Sleep or schedule another thread

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule(); <--
```

```
        spin_lock(&lock->wait_lock);
```

```
}
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list)) <--
```

Someone is waiting for the lock

```
        wake_up_process(&lock->wait_list) <--
```

Wake-up a waiting thread

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

More reading: [mutex.c in Linux](#)

# Bounded Buffer Problem Revisit

## Monitor version

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    当queue为空，不应该dequeue，所以等待full  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```

## Semaphore version

```
Semaphore mutex = 1, full = 0,  
empty = N;  
  
produce (item)  
{  
    empty.P();  
    mutex.P();  
    queue.enqueue(item);  
    mutex.V();  
    full.V();  
}  
  
consume()  
{  
    full.P();  
    mutex.P();  
    item = queue.dequeue();  
    mutex.V();  
    empty.V();  
    return item;  
}
```

# Deadlock

DeadLock需要避免进程之间互相要  
锁然后都进入了等待睡眠的状态

- Deadlock conditions
- Resource allocation graph
- Banker's algorithm
- Dining philosopher example
- Other concepts to know
  - Starvation vs. deadlock

Starvation 饿死的场景很直接，避免资源的不公平利用，例如，不能某些任务  
总是获得资源，有些任务即使长期在等待下，却没有被分配到资源。



# Conditions for Deadlocks

- Mutual exclusion
  - only one process at a time can use a resource
- No preemption  
非抢占；自愿释放资源；非抢占（Non-Preemptive）内核中如果任务不主动放弃，就不会被其他的高优先级任务抢掉对于CPU的控制
  - resources cannot be preempted, release must be voluntary
- Hold and wait
  - a process must be holding at least one resource, and waiting to acquire additional resources held by other processes
- Circular wait
  - There must be a circular dependency. For example, A waits B, B waits C, and C waits A.
- **All four conditions must simultaneously hold**

# Quiz

- Determine whether this state is safe or unsafe.

Total resources: 12

Avail resources: 3

Process	Max	Alloc	
P <sub>0</sub>	10	4	$10 - 4 \leq 8$
P <sub>1</sub>	3	1	$3 - 1 \leq 7$
P <sub>2</sub>	6	4	$6 - 4 \leq 3$

Safe

# Quiz

- Can a request 3 by P0 be granted?

Total resources: 12

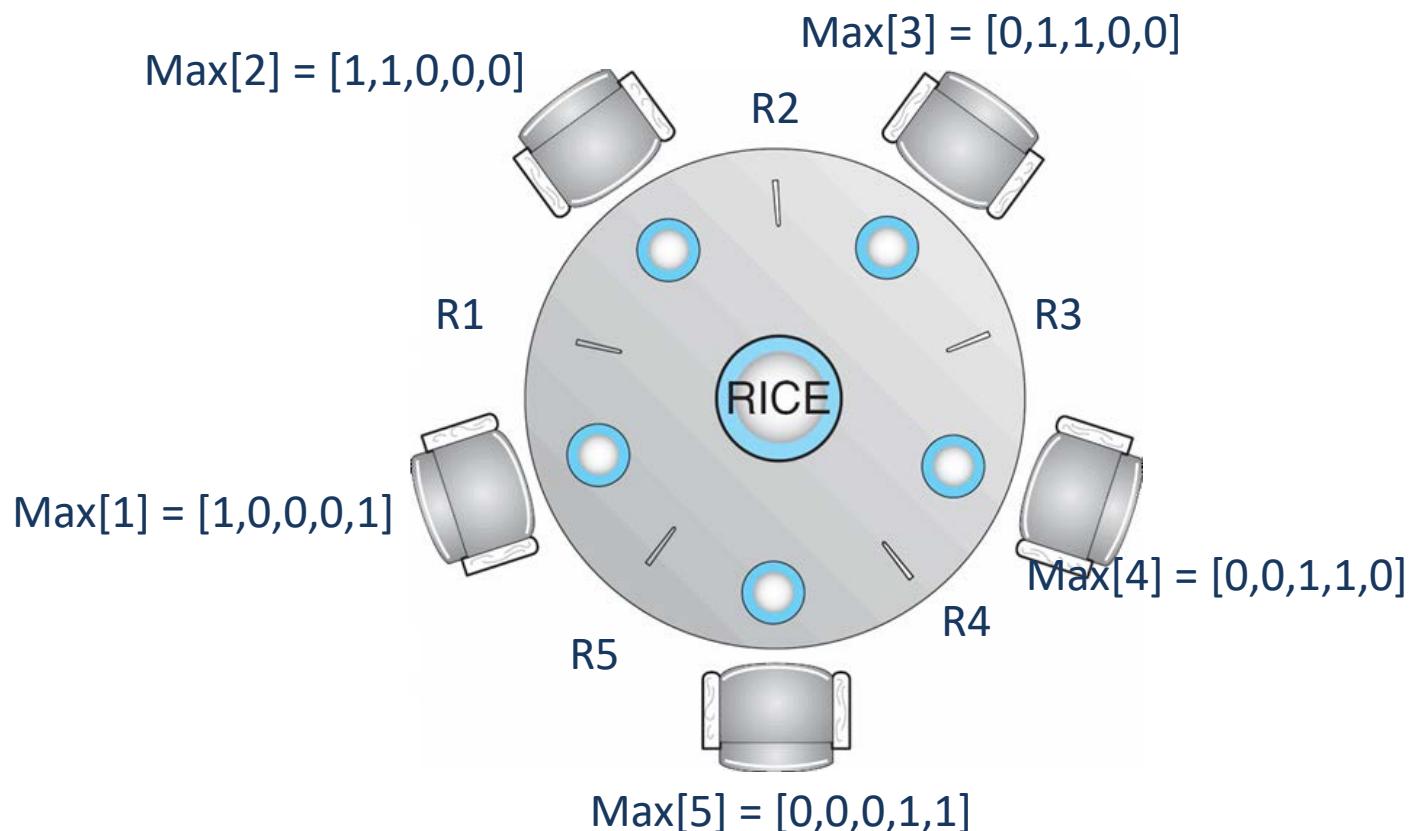
Avail resources: 0

Process	Max	Alloc	
P <sub>0</sub>	10	7	$10 - 7 \not\leq 0$
P <sub>1</sub>	3	1	$3 - 1 \not\leq 0$
P <sub>2</sub>	6	4	$6 - 4 \not\leq 0$

Unsafe

# Example

Free = [1,1,1,1,1]



# Example

Free = [0,0,0,0,1]

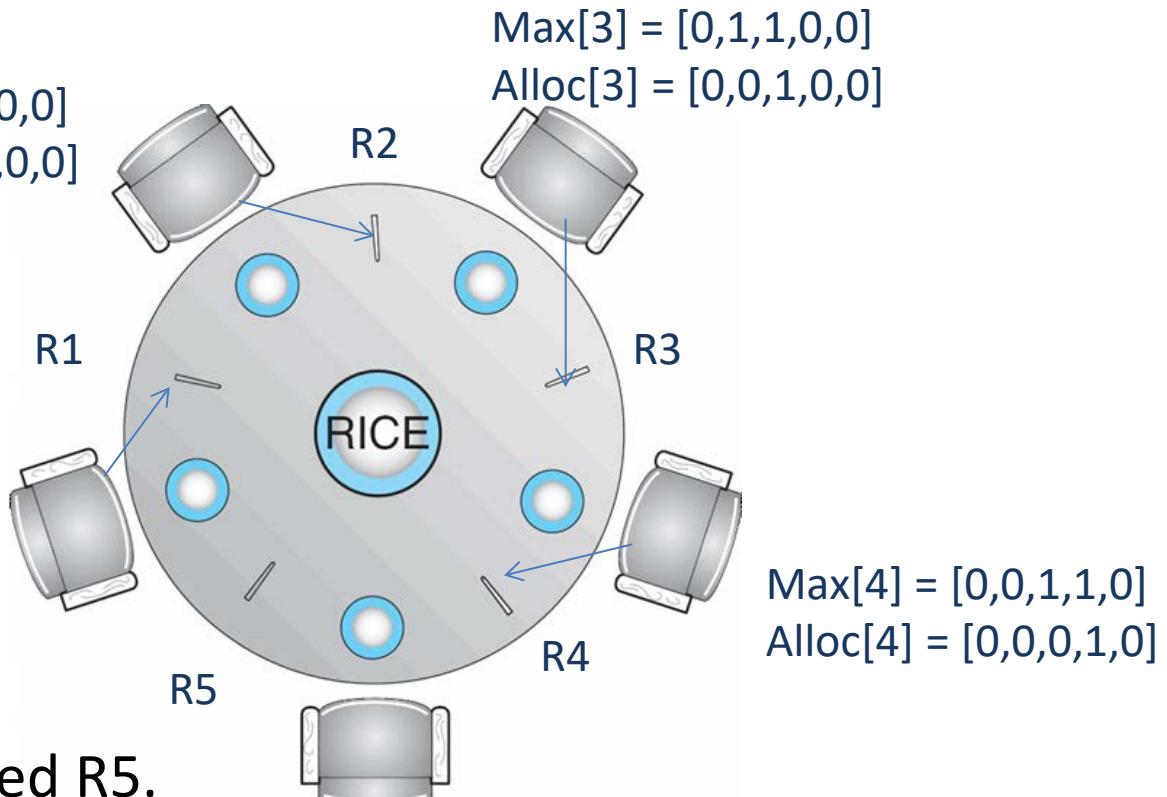
Avail = [0,0,0,0,1]

Max[2] = [1,1,0,0,0]

Alloc[2] = [0,1,0,0,0]

Max[1] = [1,0,0,0,1]

Alloc[1] = [1,0,0,0,0]



Max[5] = [0,0,0,1,1]

Alloc[5] = [0,0,0,0,0]

- Philosopher 5 requested R5.
- **Safe or Unsafe?**

# Example

Free = [0,0,0,0,0]

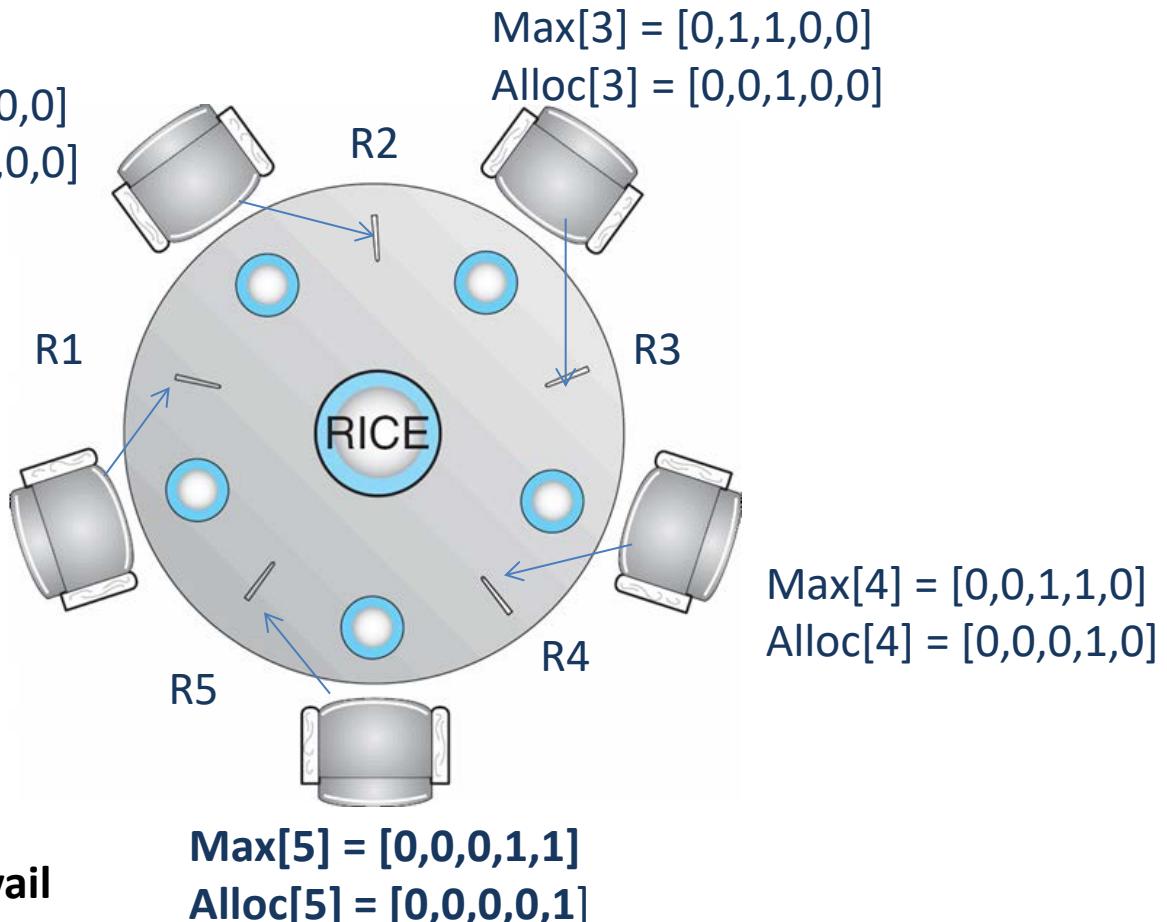
Avail = [0,0,0,0,0]

Max[2] = [1,1,0,0,0]

Alloc[2] = [0,1,0,0,0]

Max[1] = [1,0,0,0,1]

Alloc[1] = [1,0,0,0,0]



2. Find an index  $i$  such that

Finish[i] == false AND

**Max[i] – Alloc[i] ≤ Avail**

If no such  $i$  exists, go to step 4

# Scheduling

- Three main schedulers
  - FCFS, SJF/SRTF, RR
  - Gant chart examples
- Other concepts to know
  - Fair scheduling (CFS) not in text? 这块知识没学过！！大纲用以前的
  - Fixed priority scheduling
  - Multi-level queue scheduling
  - Load balancing and multicore scheduling

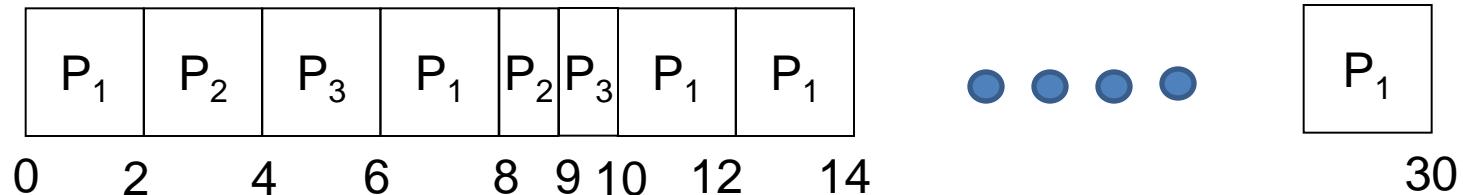
# Round-Robin (RR)

- Example

- Quantum size = 2

Process	Burst Times
P1	24
P2	3
P3	3

- Gantt chart



- Waiting time

- P1: 6, P2: 6, P3: 7. average waiting time =  $(6+6+7)/3 = 6.33$

- Sched. latency (between ready to first schedule)

- P1: 0, P2: 2, P3: 4. average sched. latency =  $(0+2+4)/3 = 2$