

# Inter-Process Communication

Disclaimer: some slides are adopted from the book authors' slides with permission

# Today

- Inter-Process Communication (IPC)
  - What is it?
  - What IPC mechanisms are available?

# Inter-Process Communication (IPC)


- What is it?
  - Communication among processes
- Why needed?
  - Information sharing
  - Modularity
  - Speedup

# Chrome Browser

- Multi-process architecture
- Each tab is a **separate** process
  - Why?
  - How to communicate among the processes?




# Models of IPC



(a)

message passing



(b)

shared memory

# Models of IPC

- Shared memory
  - share a region of memory between co-operating processes
  - read or write to the shared memory region
  - ++ fast communication
  - -- synchronization is very difficult
- Message passing
  - exchange messages (send and receive)
  - typically involves data copies (to/from buffer)
  - ++ synchronization is easier
  - -- slower communication

# Interprocess Communication in Unix (Linux)

- **Pipe**
- **FIFO**
- **Shared memory**
- **Socket**
- Message queue
- ...

# Pipes

- Most basic form of IPC on all Unix systems
  - Your shell uses this a lot (and your 1st project too)

ls | more

- Characteristics
  - Unidirectional communication
  - Processes must be in the same OS
  - Pipes exist only until the processes exist
  - Data can only be collected in FIFO order



# IPC Example Using Pipes


```
main()
{
    char *s, buf[1024];
    int fds[2];
    s = "Hello World\n";

    /* create a pipe */
    pipe(fds);

    /* create a new process using fork */
    if (fork() == 0) {

        /* child process. All file descriptors, including
           pipe are inherited, and copied.*/
        write(fds[1], s, strlen(s));
        exit(0);
    }

    /* parent process */
    read(fds[0], buf, strlen(s));
    write(1, buf, strlen(s));
}
```



(\*) Img. source: <http://beej.us/guide/bgipc/output/html/multipage/pipes.html>

# Pipes in Shells

- Example: `$ ls| more`
  - The output of 'ls' becomes the input of 'more'
- How does the shell realize this command?
  - Create a pipe
  - Create a process to run *ls*
  - Create a process to run *more*
  - The standard output of the process to run *ls* is redirected to a pipe streaming to the process to run *more*
  - The standard input of the process to run *more* is redirected to be the pipe from the process running *ls*

# FIFO (Named Pipe)

- Pipe with a name!
- More powerful than anonymous pipes
  - No parent-sibling relationship required
  - Allow bidirectional communication
  - FIFOs exists even after creating process is terminated
- Characteristics of FIFOs
  - Appear as typical files
  - Communicating process must reside on the same machine

# Example: Producer

```
main()
{
    char str[MAX_LENGTH];
    int num, fd;

    mkfifo(FIFO_NAME, 0666); // create FIFO file
    fd = open(FIFO_NAME, O_WRONLY); // open FIFO for writing

    printf("Enter text to write in the FIFO file: ");
    fgets(str, MAX_LENGTH, stdin);
    while(!(feof(stdin))){
        if ((num = write(fd, str, strlen(str))) == -1)
            perror("write");
        else
            printf("producer: wrote %d bytes\n", num);
        fgets(str, MAX_LENGTH, stdin);
    }
}
```


# Example: Consumer

```
main()
{
    char str[MAX_LENGTH];
    int num, fd;

    mkfifo(FIFO_NAME, 0666); // make fifo, if not already present
    fd = open(FIFO_NAME, O_RDONLY); // open fifo for reading

    do{
        if((num = read(fd, str, MAX_LENGTH)) == -1)
            perror("read");
        else{
            str[num] = '\0';
            printf("consumer: read %d bytes\n", num);
            printf("%s", str);
        }
    }while(num > 0);
}
```

# Shared Memory



# Shared Memory

- Kernel is not involved in data transfer
  - No need to copy data to/from the kernel
    - Very fast IPC
  - Pipes, in contrast, need to
    - Send: copy from user to kernel
    - Recv: copy from kernel to user
  - BUT, you have to *synchronize*
    - Will discuss in the next week

# POSIX Shared Memory

- Sharing between unrelated processes
- APIs
  - `shm_open()`
    - Open or create a shared memory object
  - `ftruncate()`
    - Set the size of a shared memory object
  - `mmap()`
    - Map the shared memory object into the caller's address space



# Example: Producer

```
$ ./writer /shm-name "Hello"
```

```
int main(int argc, char *argv[])
{
    char *addr;
    int fd;
    size_t len;

    fd = shm_open(argv[1], O_CREAT | O_RDWR, S_IRWXU | S_IRWXG);
    len = strlen(argv[2])+1;
    ftruncate(fd, len);
    addr = mmap(NULL, len, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    close(fd);

    memcpy(addr, argv[2], len);
    return 0;
}
```

# Example: Consumer

```
$ ./reader /shm-name
```

```
int main(int argc, char *argv[])
{
    char *addr;
    int fd;
    struct stat sb;

    fd = shm_open(argv[1], O_RDWR, 0);
    fstat(fd, &sb);
    addr = mmap(NULL, sb.st_size, PROT_READ, MAP_SHARED, fd, 0);
    close(fd);

    printf("%s\n", addr);
    return 0;
}
```

# Sockets

- Sockets
  - two-way communication pipe
  - Backbone of your internet services
- Unix Domain Sockets
  - communication between processes on the same Unix system
  - special file in the file system
- Client/Server
  - client sending requests for information, processing
  - server waiting for user requests
- Socket communication modes
  - connection-based, TCP
  - connection-less, UDP

# Example: Server

```
int main(int argc, char *argv[])
{
    int listenfd = 0, connfd = 0;
    struct sockaddr_in serv_addr;
    char sendBuff[1025];
    time_t ticks;

    listenfd = socket(AF_INET, SOCK_STREAM, 0);
    memset(&serv_addr, '0', sizeof(serv_addr));
    memset(sendBuff, '0', sizeof(sendBuff));

    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);
    serv_addr.sin_port = htons(5000);

    bind(listenfd, (struct sockaddr*)&serv_addr, sizeof(serv_addr));
    listen(listenfd, 10);

    while(1)
    {
        connfd = accept(listenfd, (struct sockaddr*)NULL, NULL);
        snprintf(sendBuff, "Hello. I'm your server.");
        write(connfd, sendBuff, strlen(sendBuff));
        close(connfd);
    }
}
```

# Example: Client

```
int main(int argc, char *argv[])
{
    int sockfd = 0, n = 0;
    char recvBuff[1024];
    struct sockaddr_in serv_addr;

    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    memset(&serv_addr, '0', sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_port = htons(5000);

    inet_pton(AF_INET, argv[1], &serv_addr.sin_addr);
    connect(sockfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr));

    while ( (n = read(sockfd, recvBuff, sizeof(recvBuff)-1)) > 0)
    {
        recvBuff[n] = 0;
        printf("%s\n" recvBuff);
    }
    return 0;
}
```

\$ ./client 127.0.0.1  
Hello. I'm your server.

# Quiz

- A process produces 100MB data in memory. You want to share the data with two other processes so that each of which can access half the data (50MB each). What IPC mechanism will you use and why?

# Project 1: Quash

- Goals
  - Learn to use UNIX system calls
  - Learn the concept of processes
- Write your own shell (like *bash* in Linux)
  - Run external programs
    - Use `fork()/execve()` system calls
  - Support built-in commands
  - Support pipe and redirections

# Thread

Disclaimer: some slides are adopted from the book authors' slides with permission



# Recap

- IPC
  - Shared memory
    - share a memory region between processes
    - read or write to the shared memory region
    - ++ fast communication
    - synchronization is very difficult
  - Message passing
    - exchange messages (send and receive)
    - typically involves data copies (to/from buffer)
    - ++ synchronization is easier
    - slower communication

# Recap

- **Process**
  - **Address space**
    - The process's view of memory
    - Includes program code, global variables, dynamic memory, stack
  - **Processor state**
    - Program counter (PC), stack pointer, and other CPU registers
  - **OS resources**
    - Various OS resources that the process uses
    - E.g.) open files, sockets, accounting information

# Recap: Pipes


```
main()
{
    char *s, buf[1024];
    int fds[2];
    s = "Hello World\n";

    /* create a pipe */
    pipe(fds);

    /* create a new process using fork */
    if (fork() == 0) {

        /* child process. All file descriptors, including
           pipe are inherited, and copied.*/
        write(fds[1], s, strlen(s));
        exit(0);
    }

    /* parent process */
    read(fds[0], buf, strlen(s));
    write(1, buf, strlen(s));
}
```



(\*) Img. source: <http://beej.us/guide/bgipc/output/html/multipage/pipes.html>

# Concurrent Programs



- Objects (tanks, planes, ...) are moving simultaneously
- Now, imagine you implement each object as a process. Any problems?

# Why Processes Are Not Always Ideal?

- Not memory efficient
  - Own address space (page tables)
  - OS resources: open files, sockets, pipes, ...
- Sharing data between processes is not easy
  - No direct access to others' address space
  - Need to use IPC mechanisms

# Better Solutions?


- We want to run things concurrently
  - i.e., multiple independent flows of control
- We want to share memory easily
  - Protection is not really big concern
  - Share code, data, files, sockets, ...
- We want do these things efficiently
  - Don't want to waste memory
  - Performance is very important

# Thread



# Thread in OS

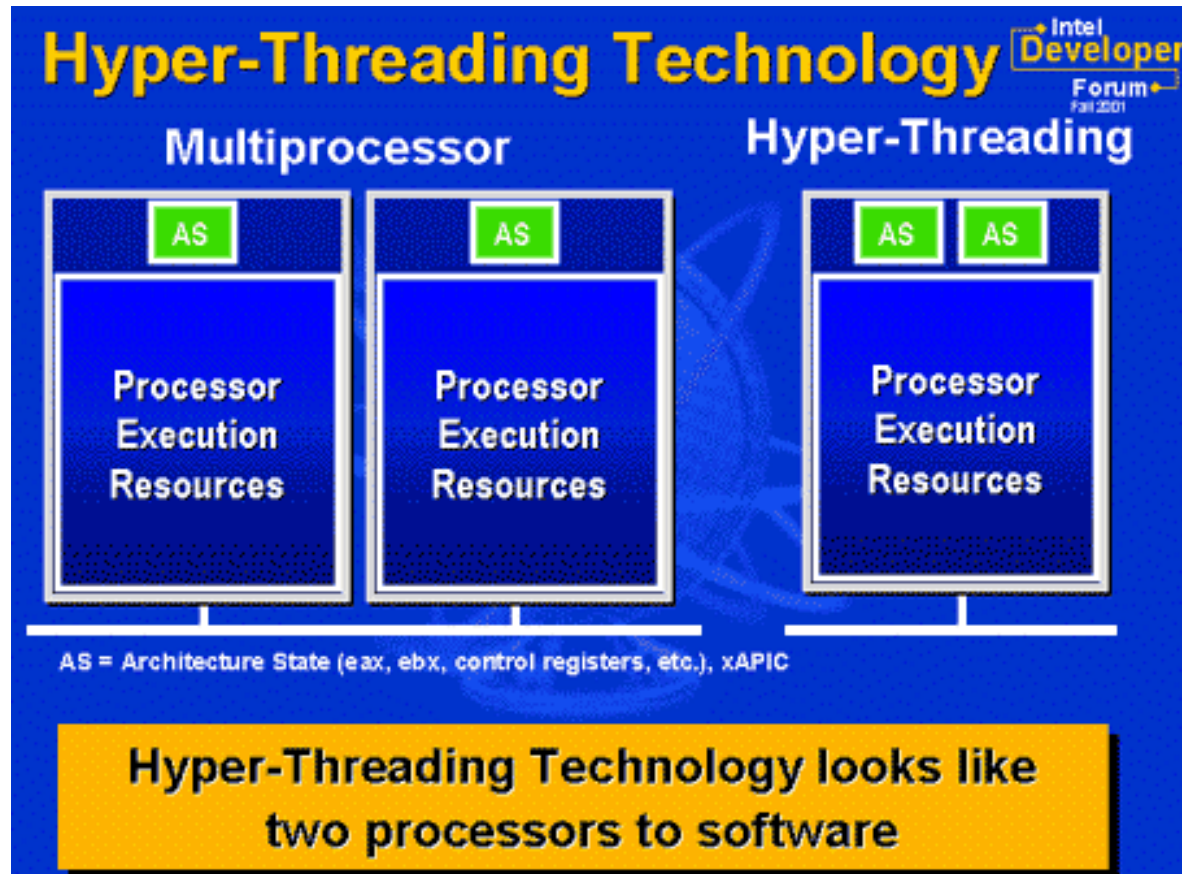
- Lightweight process
- Process
  - Address space
  - CPU context: PC, registers, stack, ...
  - OS resources
- Thread
  - ~~Address space~~
  - CPU context: PC, registers, stack, ...
  - ~~OS resources~~





# Thread in Architecture

- Logical processor



# Thread

- Lightweight process
  - Own independent flow of control (execution)
  - Stack, thread specific data (tid, ...)
  - Everything else (address space, open files, ...) is shared

## Shared

- Program code
- (Most) data
- Open files, sockets, pipes
- Environment (e.g., HOME)

## Private

- Registers
- Stack
- Thread specific data
- Return value

# Process vs. Thread

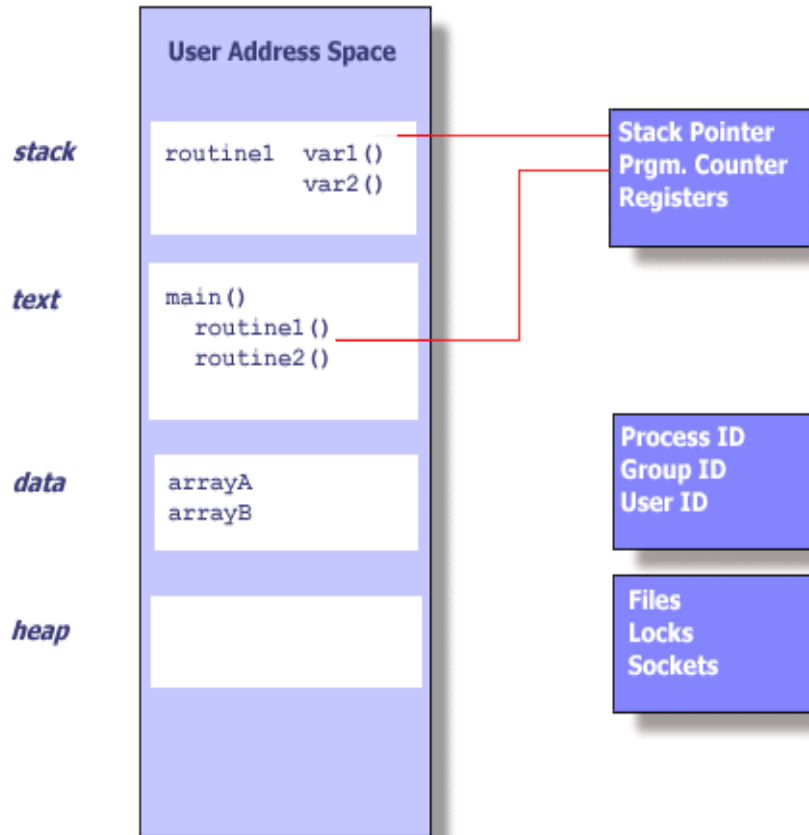


Figure source: <https://computing.llnl.gov/tutorials/pthreads/>

# Process vs. Thread

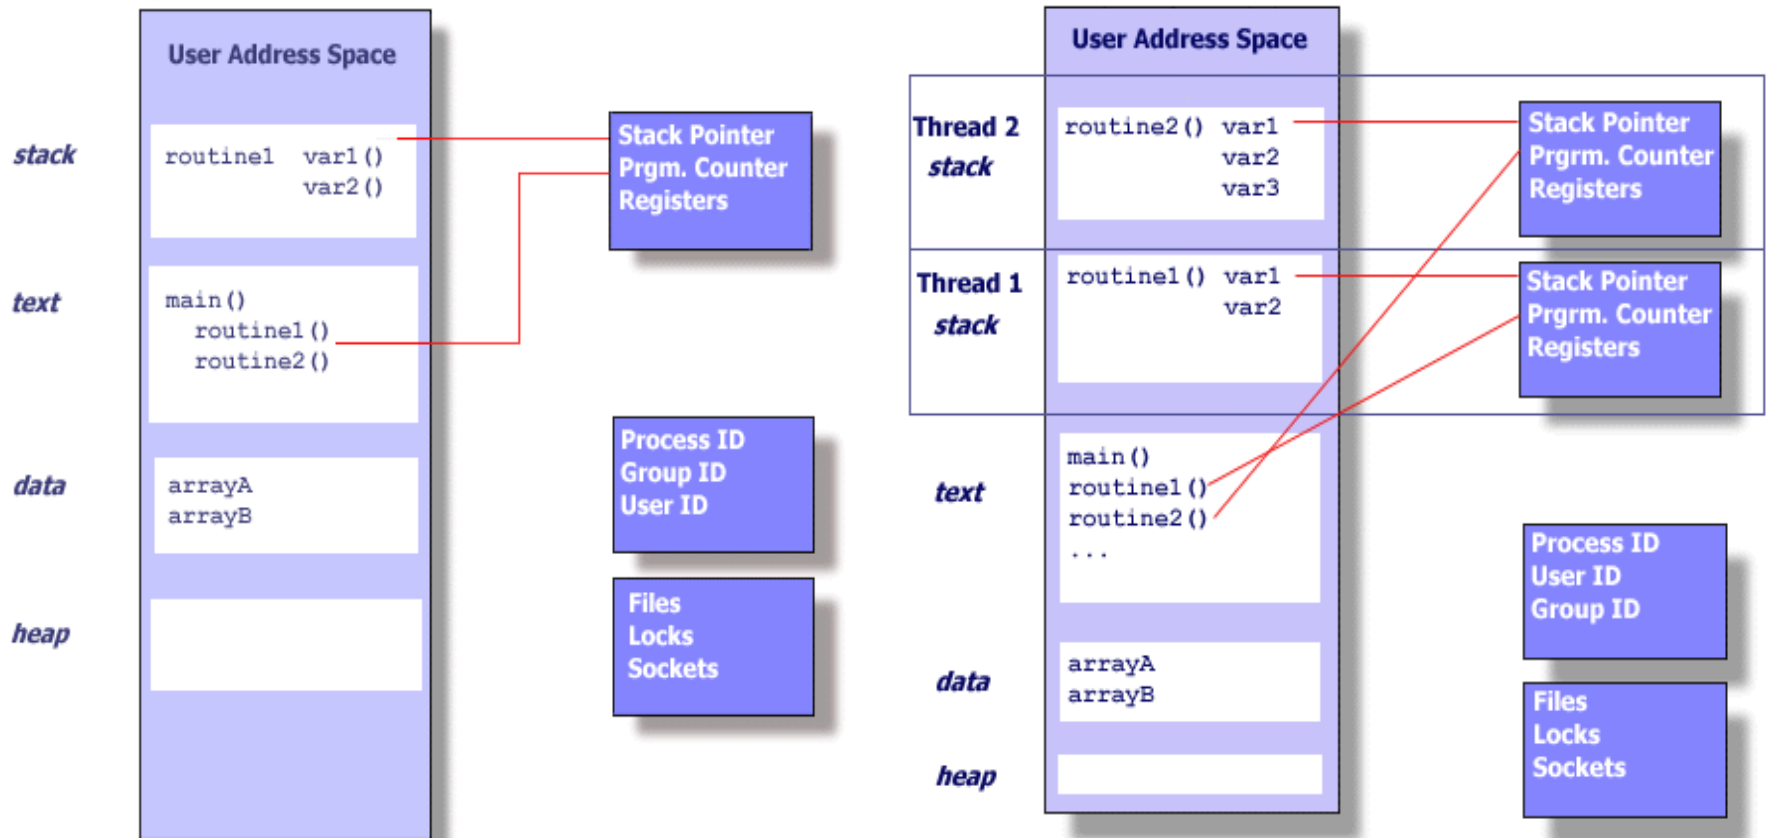


Figure source: <https://computing.llnl.gov/tutorials/pthreads/>

# Thread Benefits

- Responsiveness
  - Simple model for concurrent activities.
  - No need to block on I/O
- Resource Sharing
  - Easier and faster memory sharing (but be aware of synchronization issues)
- Economy
  - Reduces context-switching and space overhead <sup>降低开销</sup> → better performance
- Scalability
  - Exploit multicore CPU

# Thread Programming in UNIX

- Pthread
  - IEEE POSIX standard threading API
- Pthread API
  - Thread management
    - create, destroy, detach, join, set/query thread attributes
  - Synchronization
    - Mutexes –lock, unlock
    - Condition variables – signal/wait

# Pthread API

- `pthread_attr_init` – initialize the thread attributes object
  - `int pthread_attr_init(pthread_attr_t *attr);`
  - defines the attributes of the thread created
- `pthread_create` – create a new thread
  - `int pthread_create(pthread_t *restrict thread, const pthread_attr_t *restrict attr, void *(*start_routine)(void*), void *restrict arg);`
  - upon success, a new thread id is returned in thread
- `pthread_join` – wait for thread to exit
  - `int pthread_join(pthread_t thread, void **value_ptr);`
  - calling process blocks until thread exits
- `pthread_exit` – terminate the calling thread
  - `void pthread_exit(void *value_ptr);`
  - make return value available to the joining thread

# Pthread Example 1

```
#include <pthread.h>
#include <stdio.h>

int sum; /* data shared by all threads */
void *runner (void *param)
{
    int i, upper = atoi(param);
    sum = 0;
    for(i=1 ; i<=upper ; i++)
        sum += i;
    pthread_exit(0);
}

int main (int argc, char *argv[])
{
    pthread_t tid; /* thread identifier */
    pthread_attr_t attr;
    pthread_attr_init(&attr);

    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);
    fprintf(stdout, "sum = %d\n", sum);
}
```

Quiz: Final output?

\$/a.out 10

sum = 55



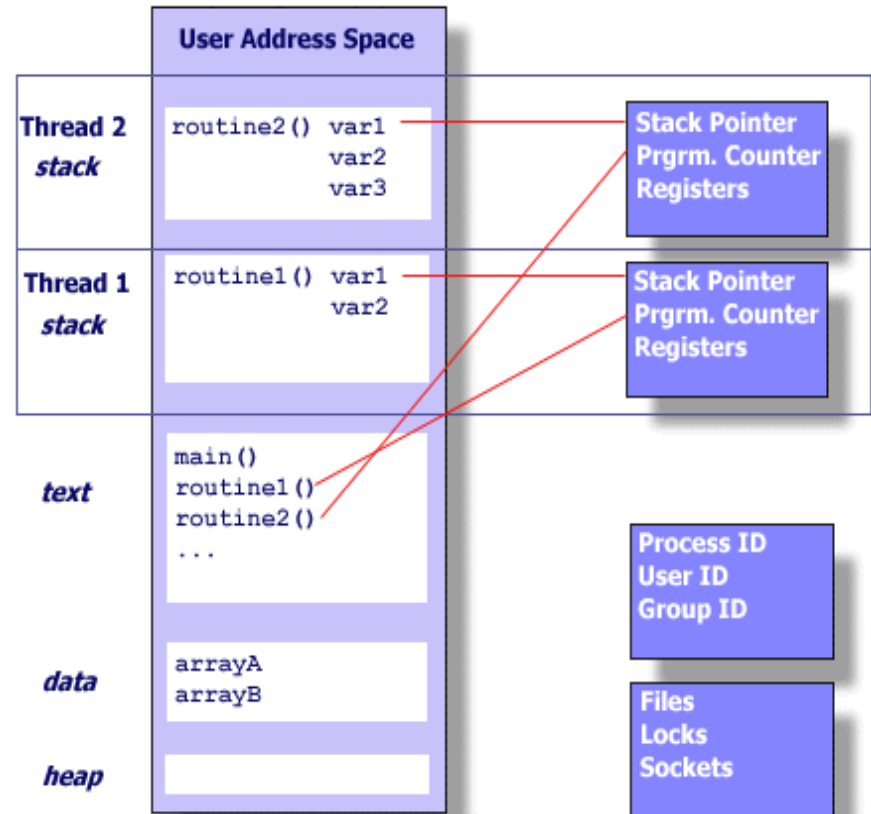
# Pthread Example 2

```
#include <pthread.h>
#include <stdio.h>

int arrayA[10], arrayB[10];


void *routine1(void *param)
{
    int var1, var2
    ...
}
void *routine2(void *param)
{
    int var1, var2, var3
    ...
}

int main (int argc, char *argv[])
{
    /* create the thread */
    pthread_create(&tid[0], &attr, routine1, NULL);
    pthread_create(&tid[1], &attr, routine2, NULL);
    pthread_join(tid[0]); pthread_join(tid[1]);
}
```




# User-level Threads

- Kernel is unaware of threads
  - Early UNIX and Linux did not support threads
- Threading runtime
  - Handle context switching
    - Setjmp/longjmp, ...
- Advantage
  - No kernel support
  - Fast (no kernel crossing)
- Disadvantage
  - Blocking system call. What happens?




# Kernel-level Threads

- Native kernel support for threads
  - Most modern OS (Linux, Windows NT)
- Advantage
  - No threading runtime
  - Native system call handling
- Disadvantage
  - Overhead



# Hybrid Threads

- Many kernel threads to many user threads
  - Best of both worlds?



# Threads: Advanced Topics

- Semantics of Fork/exec()
- Signal handling
- Thread pool
- Multicore

# Semantics of fork()/exec()

- Remember fork(), exec() system calls?
  - Fork: create a child process (a copy of the parent)
  - Exec: replace the address space with a new pgm.
- Duplicate *all* threads or *the caller* only?
  - Linux: the calling thread only
  - Complicated. [Don't do it!](#)
    - Why? Mutex states, library, ...
    - Exec() immediately after Fork() may be okay.

# Signal Handling


- What is *Singal*?
  - \$ man 7 signal
  - OS to process notification
    - “hey, wake-up, you’ve got a packet on your socket,”
    - “hey, wake-up, your timer is just expired.”
- Which *thread* to deliver a signal?
  - Any thread
    - e.g., kill(pid)
  - Specific thread
    - E.g., pthread\_kill(tid)

# Thread Pool


- Managing threads yourself can be cumbersome and costly
  - Repeat: create/destroy threads as needed.
- Let's create a set of threads ahead of time, and just ask them to execute my functions
  - #of thread  $\sim$  #of cores
  - No need to create/destroy many times
  - Many high-level parallel libraries use this.
    - e.g., Intel TBB (threading building block), ...



# Single Core Vs. Multicore Execution



*Single core execution*



*Multiple core execution*

# Challenges for Multithreaded Programming in Multicore

- How to divide activities?
- How to divide data?
- **How to synchronize accesses to the shared data? → next class**
- **How to test and debug?**

**EECS750**

# Summary

- Thread
  - What is it?
    - Independent flow of control.
  - What for?
    - Lightweight programming construct for concurrent activities
  - How to implement?
    - Kernel thread vs. user thread
- Next class
  - How to synchronize?