

Synchronization

Disclaimer: some slides are adopted from the book authors' slides with permission

Recap: Thread

- What is it?
 - Independent flow of control
- What does it need (thread private)?
 - Stack
- What for?
 - Lightweight programming construct for concurrent activities
- How to implement?
 - Kernel thread vs. user thread

Recap: Process vs. Thread

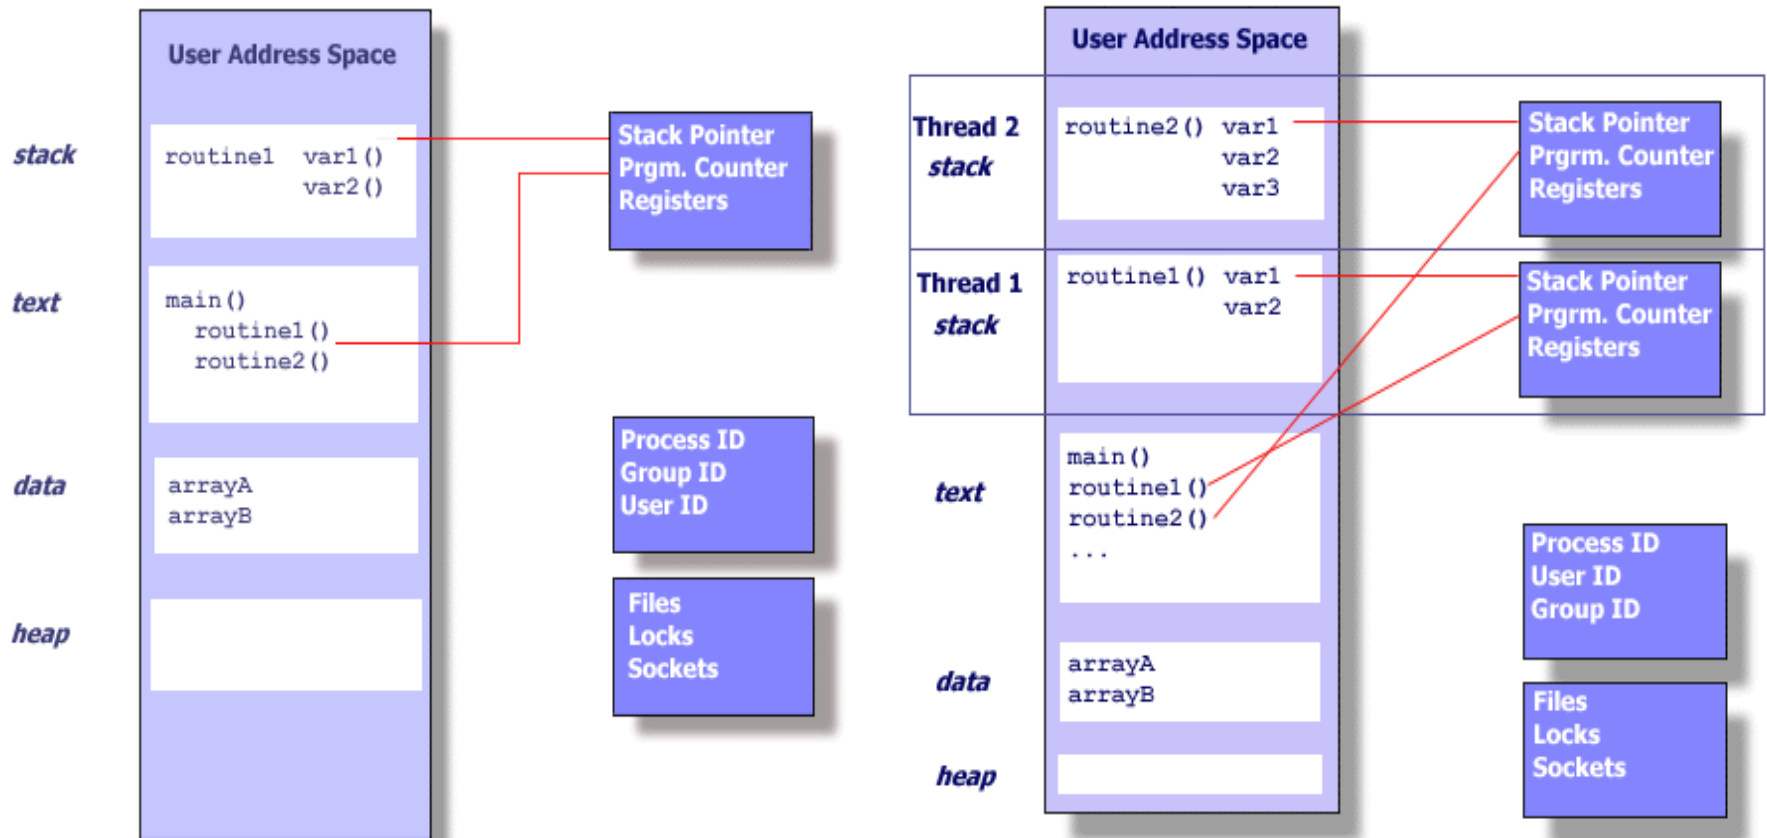


Figure source: <https://computing.llnl.gov/tutorials/pthreads/>

Recap: Multi-threads vs. Multi-processes

- Multi-processes
 - (+) protection
 - (-) performance (?)
- Multi-threads
 - (+) performance
 - (-) protection



Process-per-tab



Single-process multi-threads

Agenda

- **Mutual exclusion**
 - **Peterson's algorithm (Software)**
 - Synchronization instructions (Hardware)
 - Spinlock 自旋锁
- High-level synchronization mechanisms
 - Mutex
 - Semaphore
 - Monitor

Producer/Consumer

线程不安全

计算机在执行程序时，每条指令都是在CPU中执行的，而执行指令过程中，势必涉及到数据的读取和写入



Producer/Consumer

Producer

```
while (true){  
  
    /* wait if buffer full */  
    while (counter == 10);  
  
    /* produce data */  
    buffer[in] = sdata;  
    in = (in + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    counter++;  
}
```

Problem: 当同步工作的时候，有时候counter的值还没有完全被确定就被另一边更改了 (Race condition)

Consumer

```
while (true){  
  
    /* wait if buffer empty */  
    while (counter == 0);  
  
    /* consume data */  
    sdata = buffer[out];  
    out = (out + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    counter--;  
}
```

Producer/Consumer

Producer

```
while (true){  
  
    /* wait if buffer full */  
    while (counter == 10);  
  
    /* produce data */  
    buffer[in] = sdata;  
    in = (in + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    R1 = load (counter);  
    R1 = R1 + 1;  
    counter = store (R1);  
}
```

Consumer

```
while (true){  
  
    /* wait if buffer empty */  
    while (counter == 0);  
  
    /* consume data */  
    sdata = buffer[out];  
    out = (out + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    R2 = load (counter);  
    R2 = R2 - 1;  
    counter = store (R2);  
}
```


Check Yourself

```
int count = 0;
int main()
{
    count = count + 1;
    return count;
}
```

```
$ gcc -O2 -S sync.c
```

```
...
movl    count(%rip), %eax
addl    $1, %eax
movl    %eax, count(%rip)
...
```

Race Condition

Initial condition: *counter* = 5

Thread 1

```
R1 = load (counter);  
R1 = R1 + 1;  
counter = store (R1);
```


Thread 2

```
R2 = load (counter);  
R2 = R2 - 1;  
counter = store (R2);
```

- What are the possible outcome?

Race Condition

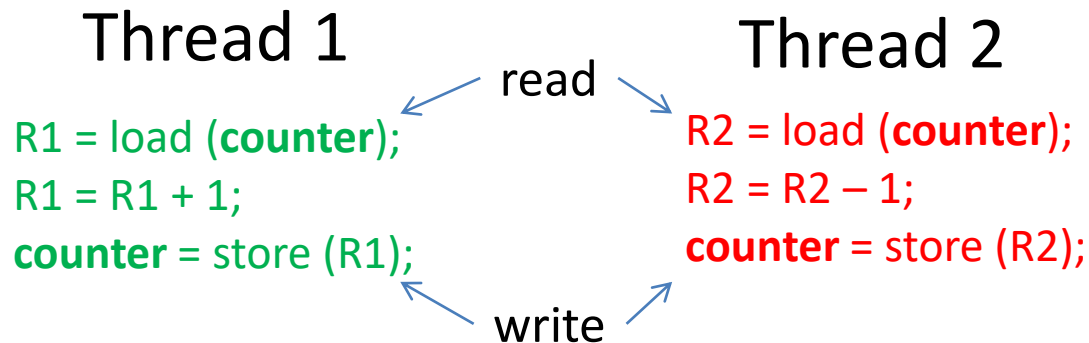
Initial condition: *counter* = 5



- Why this happens?

Race Condition

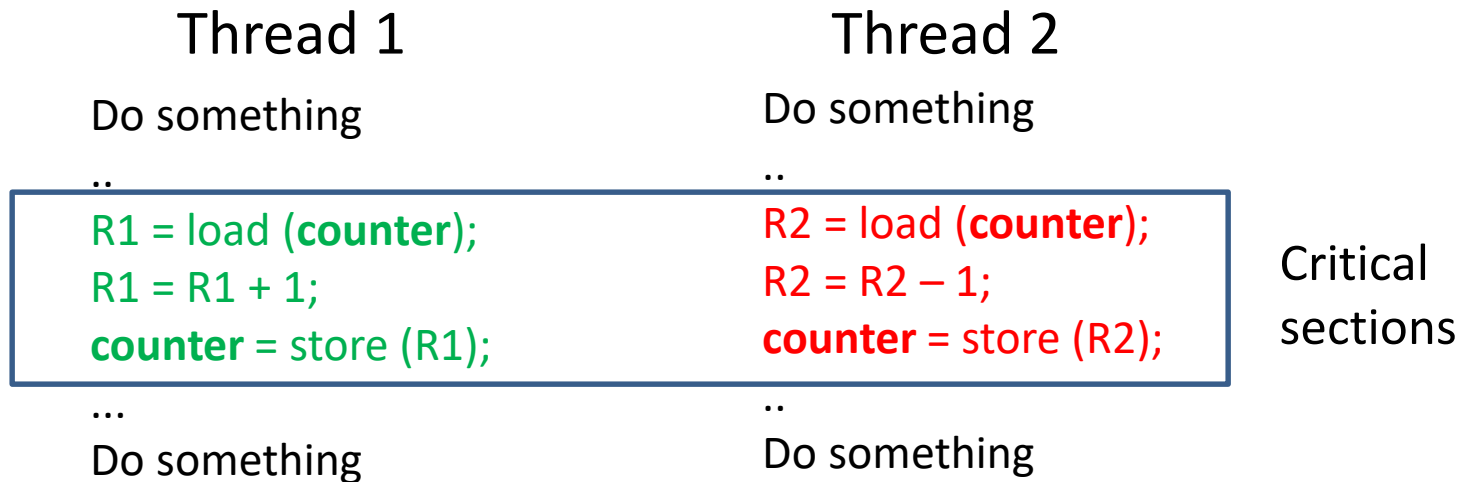
- A situation when two or more threads **read and write** shared data at the same time
- Correctness depends on the execution order



- How to prevent race conditions?

Critical Section

- Code sections of potential race conditions



Solution Requirements

- Mutual Exclusion
 - If a thread executes its critical section, *no other threads* can enter their critical sections
- Progress
 - If no one executes a critical section, someone can enter its critical section
- Bounded waiting
 - Waiting (time/number) must be bounded

Simple Solution (?): Use a Flag

```
// wait
while (in_cs)
    ;
// enter critical section
in_cs = true;

Do something

// exit critical section
in_cs = false;
```

T1
while(in_cs){};

in_cs = true;
//enter

T2

while(in_cs){};
in_cs = true;

//enter

- Mutual exclusion is not guaranteed

Peterson's Solution

- Software solution (no h/w support)
- Two process solution
 - Multi-process extension exists
- The two processes share two variables:
 - int turn;
 - The variable turn indicates whose turn it is to enter the critical section
 - Boolean flag[2]
 - The flag array is used to indicate if a process is ready to enter the critical section.

Peterson's Solution

Thread 1

```
do {  
    flag[0] = TRUE;  
    turn = 1;  
    while (flag[1] && turn==1)  
    {};  
    // critical section  
  
    flag[0] = FALSE;  
  
    // remainder section  
} while (TRUE)
```

Thread 2

```
do {  
    flag[1] = TRUE;  
    turn = 0;  
    while (flag[0] && turn==0)  
    {};  
    // critical section  
  
    flag[1] = FALSE;  
  
    // remainder section  
} while (TRUE)
```

- Solution meets all **three** requirements
 - Mutual exclusion: P0 and P1 cannot be in the critical section at the same time
 - Progress: if P0 does not want to enter critical region, P1 does no waiting
 - Bounded waiting: process waits for at most one turn

Peterson's Solution

- Limitations
 - Only supports two processes
 - generalizing for more than two processes has been achieved, but not very efficient
 - Assumes LOAD and STORE instructions are atomic
 - In reality, no guarantees
 - Assumes memory accesses are not reordered
 - compiler re-orders instructions (gcc -O2, -O3, ...)
 - Out-of-order processor re-orders instructions

并不重新排序

Atomic: Reading and writing this data type is guaranteed to happen in a single instruction, so there's no way for a handler to run "in the middle" of an access.

程序的原子性指：整个程序中的所有操作，要么全部完成，要么全部不完成，不可能停滞在中间某个环节

Reordering by the CPU

Initially $X = Y = 0$

注意读x和写回x是原子操作，两个线程
不能同时执行，无法在同一行

Thread 0 Thread 1

$X = 1$ $Y = 1$
 $R1 = Y$ $R2 = X$

Thread 0 Thread 1

$R1 = Y$
 $X = 1$
 $R2 = X$
 $Y = 1$

- Possible values of R1 and R2?
 - 0,1
 - 1,0
 - 1,1
 - **0,0** ← possible on PC ?

Summary

- Peterson's algorithm
 - Software-only solution
 - Turn based
 - Pros
 - No hardware support
 - Satisfy all requirements
 - Mutual exclusion, progress, bounded waiting
 - Cons
 - Complicated
 - Assume program order
 - May not work on out-of-order processors

Recap

- Race condition
 - A situation when two or more threads **read and write** shared data at the same time
 - Correctness depends on the execution order
- Critical section
 - Code sections of potential race conditions
- Mutual exclusion
 - If a thread executes its critical section, *no other threads* can enter their critical sections

Recap

- Peterson's algorithm
 - Software-only solution
 - Turn based
 - Pros
 - No hardware support
 - Satisfy all requirements
 - Mutual exclusion, progress, bounded waiting
 - Cons
 - Complicated
 - Assume program order
 - May not work on out-of-order processors

Today

- Hardware support
 - Synchronization instructions
- Lock
 - Spinlock
 - Mutex

Lock

spin lock是一种死等的锁机制。当发生访问资源冲突的时候，可以有两个选择：一个是死等，一个是挂起当前进程，调度其他进程执行。spin lock是一种死等的机制，当前的执行thread会不断的重新尝试直到获取锁进入临界区。
保护共享资源

- General solution
 - Protect critical section via a lock
 - Acquire on enter, release on exit

```
do {  
    acquire lock;  
  
    critical section  
  
    release lock;  
  
    remainder section  
} while(TRUE);
```


How to Implement a Lock?

- Unicore processor
 - No true concurrency
one thread at a time
 - Threads are *interrupted* by the OS
 - scheduling events: timer interrupt, device interrupts

- Disabling interrupt

- Threads can't be interrupted

do {

disable interrupts;
critical section
enable interrupts;

remainder section

} while(TRUE);

How to Implement a Lock?

- Multicore processor
 - True concurrency
 - More than one active threads sharing memory
 - Disabling interrupts don't solve the problem
 - More than one threads are executing at a time
- Hardware support
 - Synchronization instructions
 - **Atomic** *test&set* instruction 原子性操作: 不可被中断的一个或一系列操作
 - **Atomic** *compare&swap* instruction
- What do we mean by **atomic**?
 - All or nothing

TestAndSet Instruction

- Pseudo code

之所以叫硬同步，是因为TestAndSet()函数中的三条语句是由硬件保证同步的，即硬件保证这三条语句必须原子运行，中间不发生任何中断，如同一条语句

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

<https://zhidao.baidu.com/question/2117056499260142027.html>

如果criticalsection没有被锁住，即lock值为false，那么传入TestAndSet()的参数为false，
【在函数体内：rv先被赋为false，然后*target被赋为true，然后返回rv（值为false）。运行的结果为：返回值为false，使线程跳出while循环，得以进入criticalsection；同时参数（即lock）被改为true，表示criticalsection被锁住】，其他线程在当前线程没有出criticalsection之前（即运行lock=false这句之前），不能进入criticalsection。为什么呢，下面说

Spinlock using *TestAndSet*

spinlock 在 Linux/底下是以 spinlock_t 来表示的。使用spinlock
必须包含#include <linux/spinlock>

```
int mutex;  
init_lock (&mutex);
```

```
do {
```

```
    lock (&mutex);  
    critical section  
    unlock (&mutex);
```

```
        remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex)  
{  
    *mutex = 0;  
}  
  
void lock (int *mutex)  
{  
    while(TestAndSet(mutex))  
        ; while的死循环代表  
        另一个线程被锁住了  
}  
  
void unlock (int *mutex)  
{  
    *mutex = 0;  
}
```

CAS (Compare & Swap) Instruction

CAS 操作包含三个操作数 —— 内存位置 (V)、预期原值 (A) 和新值(B)。

- Pseudo code

如果内存位置的值与预期原值相匹配，那么处理器会自动将该位置值更新为新值。

否则，处理器不做任何操作。无论哪种情况，它都会在 CAS 指令之前返回该位置的值。

<https://www.cnblogs.com/shangdawei/p/3917117.html>

```
int CAS(int *value, int oldval, int newval)
{
    int temp = *value;
    if (*value == oldval)
        *value = newval;
    return temp;
}
```

Spinlock using CAS

```
int mutex;  
init_lock (&mutex);  
  
do {  
    lock (&mutex);  
    critical section  
    unlock (&mutex);  
    remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex) {  
    *mutex = 0;  
}  
  
void lock (int *mutex) {  
    while(CAS(&mutex, 0, 1) != 0);  
}  
  
void unlock (int *mutex) {  
    *mutex = 0;  
}
```

What's Wrong With **Spinlocks**?

- **Very wasteful**
 - Waiting thread continues to use CPU cycles
 - While doing absolutely nothing but wait
 - 100% CPU utilization, but no useful work done
 - Power consumption, fan noise, ...
- Useful when
 - You hold the lock only *briefly*
- Otherwise
 - A better solution is needed