

Synchronization

Disclaimer: some slides are adopted from the book authors' slides

Recap: *TestAndSet* Instruction

- Pseudo code

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

Recap: **Spinlock** using *TestAndSet*

```
int mutex;  
init_lock (&mutex);  
  
do {  
    lock (&mutex);  
        critical section  
    unlock (&mutex);  
        remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex)  
{  
    *mutex = 0;  
}  
  
void lock (int *mutex)  
{  
    while(TestAndSet(mutex))  
        ;  
}  
  
void unlock (int *mutex)  
{  
    *mutex = 0;  
}
```

Recap

- Synchronization instructions
 - *test&set, compare&swap*
 - All or nothing
- Spinlock
 - Spin on wait
 - Good for short critical section but can be wasteful

```
void mutex_init (mutex_t *lock)
{
    lock->value = 0;
    ...
    ...
}

void mutex_lock (mutex_t *lock)
{
    ...
    while(TestAndSet(&lock->value)) {
        ...
        ...
        ...
        ...
    }
    ...
}

void mutex_unlock (mutex_t *lock)
{
    ...
    lock->value = 0;
    ...
    ...
    ...
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list))
```

```
        wake_up_process(&lock->wait_list)
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Someone is waiting for the lock

← Wake-up a waiting thread

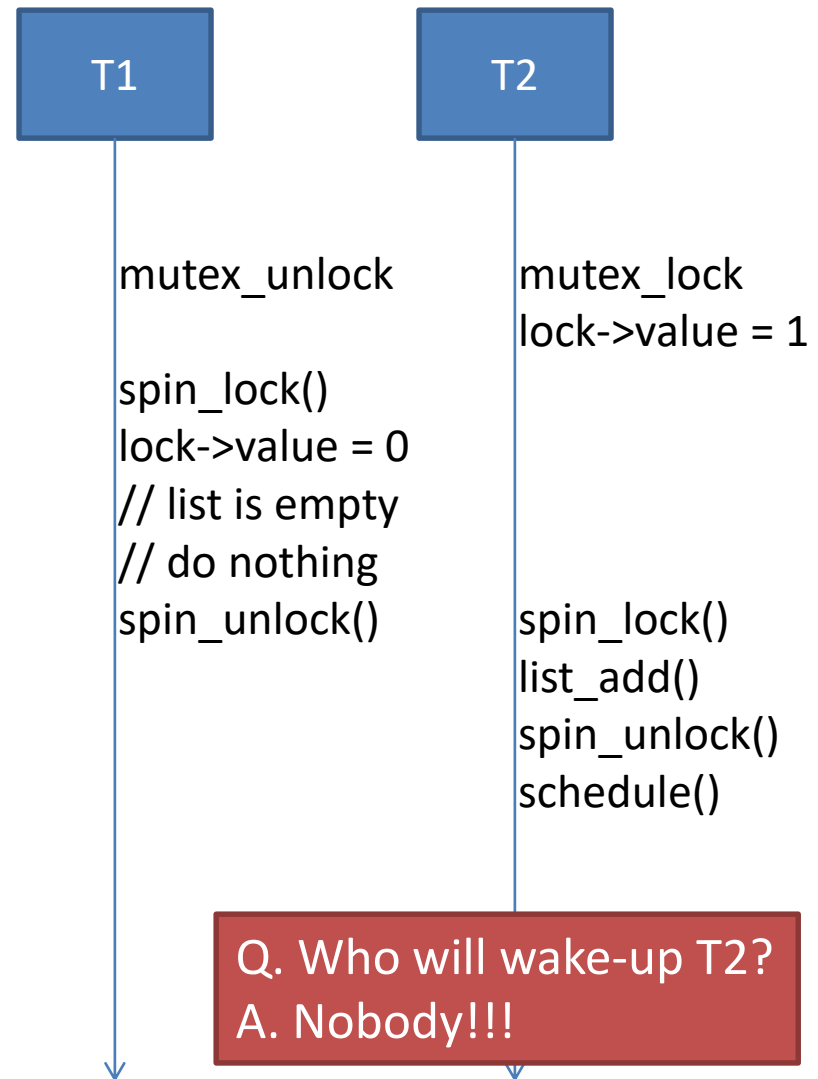
```
void mutex_init (mutex_t *lock)
{
    lock->value = 0;
    list_init(&lock->wait_list);
    spin_lock_init(&lock->wait_lock);
}
```

```
void mutex_lock (mutex_t *lock)
{
    while(TestAndSet(&lock->value)) {
        current->state = WAITING;
        spin_lock(&lock->wait_lock);
        list_add(&lock->wait_list, current);
        spin_unlock(&lock->wait_lock);
        schedule();
    }
}
```

Correct?

```
void mutex_unlock (mutex_t *lock)
{
    spin_lock(&lock->wait_lock);
    lock->value = 0;
    if (!list_empty(&lock->wait_list))
        wake_up_process(&lock->wait_list)
    spin_unlock(&lock->wait_lock);
}
```

More reading: [mutex.c in Linux](#)



Agenda

- High-level synchronization mechanisms
 - Mutex 互斥锁
 - **Semaphore** 信号量
 - Monitor

Semaphore

- High-level synchronization primitive
 - Designed by Dijkstra in 1960'
- Definition
 - Semaphore is an integer variable
 - Only two operations are possible:
 - P() or wait() or down()
 - V() or signal() or up()

P(sv): 如果sv的值大于零, 就给它减1; 如果它的值为零, 就挂起该进程的执行; 运行 P (wait()), 信号量s的值将被减少。企图进入临界区块的进程, 需要先运行 P (wait())。当信号量s减为负值时, 进程会被挡住, 不能继续; 当信号量s不为负值时, 进程可以获准进入临界区块。

V(sv): 如果有其他进程因等待sv而被挂起, 就让它恢复运行, 如果没有进程因等待sv而挂起, 就给它加1。

Simple Semaphore Implementation

- P() operation

```
P(semaphore *S) {  
    S->value—;  
    If (S->value < 0) {  
        S->list->addQ(P);  
        schedule( );  
    }  
}
```

执行，当<0 被挂起，等待前面一个完成

*schedule () – schedule
another thread*

举个例子，就是 两个进程共享信号量sv，一旦其中一个进程执行了P(sv)操作，它将得到信号量，并可以进入临界区，使sv减1。而第二个进程将被阻止进入临界区，因为 当它试图执行P(sv)时，sv为0，它会被挂起以等待第一个进程离开临界区域并执行V(sv)释放信号量，这时第二个进程就可以恢复执行

<https://www.cnblogs.com/fangshenghui/p/4039946.html>

- V() operation

```
V(semaphore *S) {  
    S->value++;  
    if(S->value <= 0) {  
        P = delQ(&S->list);  
        wakeup(P);  
    }  
}
```

wakeup() – wake up a thread

What's **wrong** with the code?

Simple Semaphore Implementation

- P() operation

```
P(semaphore *S) {  
    S->lock->Acquire();  
    S->value— ;  
    if(S->value < 0) {  
        addQ(&S->list, P);  
        S->lock->Release();  
        schedule( );  
        S->lock->Acquire();  
    }  
    S->lock->Release();  
}
```

- V() operation

```
V(semaphore *S) {  
    S->lock->Acquire();  
    S->value++ ;  
    if(S->value <= 0) {  
        P = delQ(&S->list);  
        wakeup(P);  
    }  
    S->lock->Release();  
}
```

Train Control Problem

- No more than two trains can enter the rail

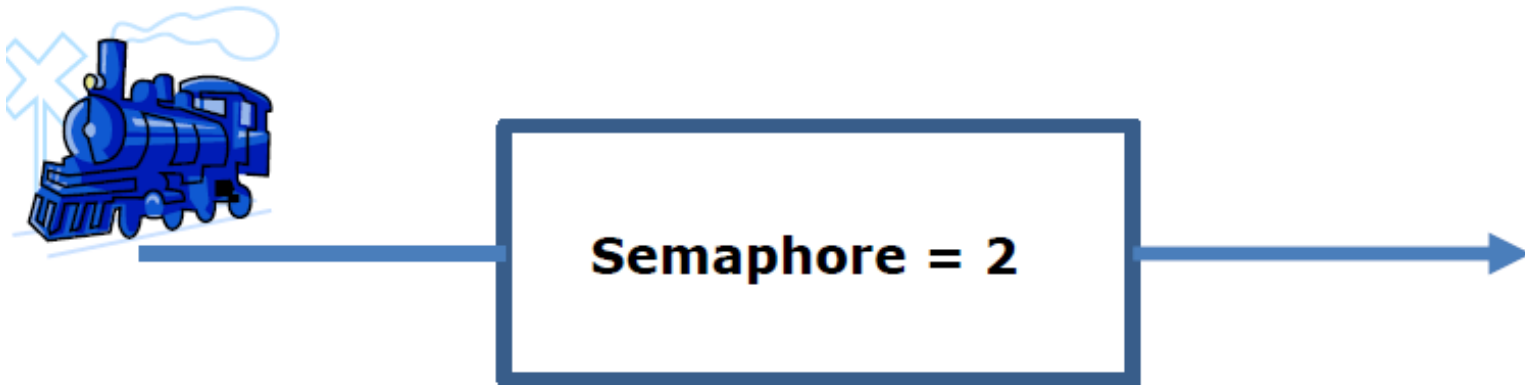
```
Enter() {  
  Leave()  
}
```



Train Control Problem


- No more than two trains can enter the rail

```
Enter() { P(); }  
Leave() { V(); }
```



Train Control Problem

- No more than two trains can enter the rail

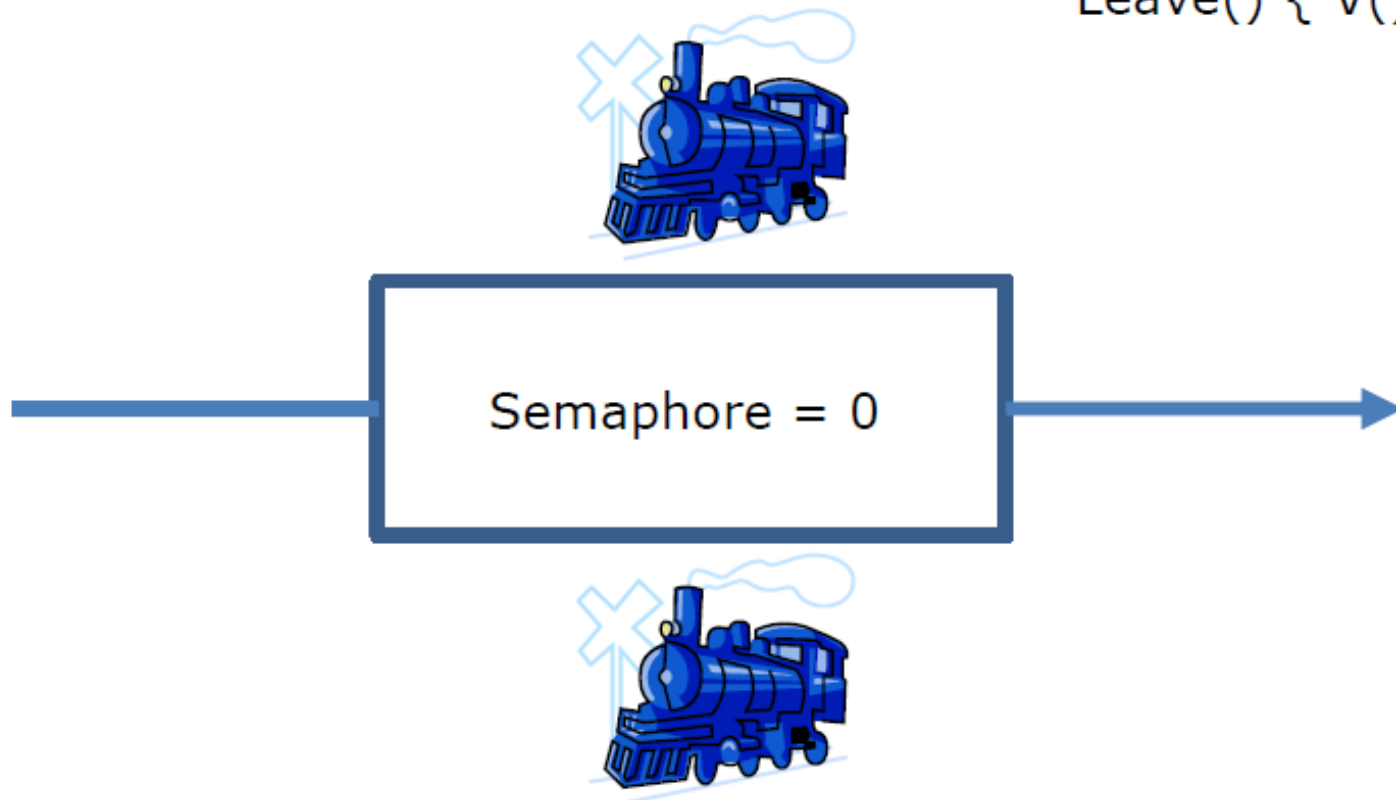


```
Enter() { P(); }  
Leave() { V(); }
```

Train Control Problem


- No more than two trains can enter the rail

```
Enter() { P(); }  
Leave() { V(); }
```



Bounded Buffer Problem

- Problem synopsis
 - A buffer is shared between a producer and a consumer
 - The size of the buffer is N .
 - Producer inserts resources into the buffer
 - If the buffer is **not full**; (otherwise wait)
 - Consumer removes resources from the buffer
 - If the buffer is **not empty**; (otherwise wait)
 - producer and consumer execute **asynchronously**
 - CPU scheduler determines what runs when



Bounded Buffer Problem

- We use *two integer semaphores* and *one binary semaphore*
- *mutex* 互斥操作
 - Mutual exclusion for accessing the buffer
- *empty* ==0 代表空的
 - Available space, initialized as N
 - If zero, the producer has to wait
- *full* 是否有东西, ==0 代表里面没有数据
 - #of data in the buffer; initialized as 0
 - If zero, the consumer has to wait

对于生产者:

判断缓冲区是否为满, 如果为满, 则等待; 否则, 允许一个生产者写入。

对于消费者:

判断缓冲区是否为空, 如果为空, 则等待; 否则, 允许一个消费者读取。

Bounded Buffer Problem

- Semaphore bool mutex;
- Semaphore int **full** = 0;
- Semaphore int **empty** = N;

Producer

do {

Produce new resource

信号量的加锁，每次只能放入指定数量的东西，否则等待

P (empty);

P (mutex);

Add resource to next buffer

V (mutex);

V (full);

} while (TRUE);

加锁
和解锁，每一次要么进行
producer要么进行
consumer

Consumer

do {

P (full);

P (mutex);

Remove resource from buffer

V (mutex);

V (empty);

Consume resource

} while (TRUE);

Reader-Writer Problem

- Problem synopsis
 - An object is shared among several threads
 - Some threads only read the object (Readers)
 - Some threads only write the object (Writers)
- Correctness constraints
 - **Multiple readers** can access the shared resource simultaneously
 - But only **one writer** can update the object, when there is no other reader or writer
 - Readers can't access the object when a writer updates it

Reader-Writer Problem

- Problem synopsis
 - An object is shared among several threads
 - Some threads only read the object (Readers)
 - Some threads only write the object (Writers)
- Correctness constraints
 - **Multiple readers** can access the shared resource simultaneously
 - But only **one writer** can update the object, when there is no other reader or writer
 - Readers can't access the object when a writer updates it

Recap: Semaphore

- High-level synchronization primitive
 - Designed by Dijkstra in 1960'
- Definition
 - Semaphore is an integer variable
 - Only two operations are possible:
 - P() or wait() or down()
 - V() or signal() or up()

Reader-Writer Problem

- A solution using two binary semaphores
 - mutex
 - ensure mutual exclusion for the **readcount** variable
 - mutex semaphore, initialized to 1
 - wrt
 - ensure **mutual exclusion for writers**
 - ensure **mutual exclusion between readers and writer**
 - mutex semaphore, initialized to 1

Reader-Writer Problem

```
semaphore mutex = 1, wrt = 1;  
int readcount = 0;
```

Writer

```
do {  
    P(wrt);  
    write object resource  
    V(wrt);  
} while (TRUE);
```

Reader-Writer Problem

```
semaphore mutex = 1, wrt = 1;  
int readcount = 0;
```

Writer

```
do {  
    P(wrt);  
    write object resource  
    V(wrt);  
} while (TRUE);
```

What's wrong with
this code?

Reader

```
do {  
    readcount是共享资源  
    readcount++;  
    if (readcount == 1)  
        P(wrt);  
    read from object resource  
  
    readcount--;  
    if (readcount == 0)  
        V(wrt);  
} while (TRUE);
```


Reader-Writer Problem

```
semaphore mutex = 1, wrt = 1;  
int readcount = 0;
```

Writer

```
do {  
    P(wrt);  
    write object resource  
    V(wrt);  
} while (TRUE);
```

Reader

```
do {  
    P(mutex);  
    readcount++;  
    if (readcount == 1)  
        P(wrt);  
    V(mutex)  
    read from object resource  
    P(mutex);  
    readcount--;  
    if (readcount == 0)  
        V(wrt);  
    V(mutex)  
} while (TRUE);
```

Semaphore Review

- Semaphores can solve many synchronization problems
 - A “huge” step up from locks
- Drawbacks of semaphores
 - Semaphores are still low-level primitives
 - Used for both mutual exclusion and scheduling
 - Very easy to screw up

Agenda

- High-level synchronization mechanisms
 - Mutex
 - Semaphore
 - **Monitor**


Monitor

管程 或者 监视器cheng

- Monitor
 - A lock (mutual exclusion) + condition variables (scheduling)
 - Some languages like Java natively support this, but you can use monitors in other languages like C/C++
- Lock: mutual exclusion
 - Protects the shared data structures inside the monitor
 - Always acquire it to enter the monitor
 - Always release it to leave the monitor
- Condition Variable: scheduling
 - Allow thread to wait on certain events inside the monitor
 - Key idea: to wait (sleep) inside the monitor, it first releases the lock and go to sleep atomically

Monitor

- Lock: mutual exclusion
 - Only one thread can execute any monitor procedure at a time.
 - Other threads invoking a monitor procedure when one is already executing some monitor procedure must wait.
 - When the active thread exits the monitor procedure, one other waiting thread can enter.



A Simple Monitor Example

C++

```
Mutex lock;  
Queue queue;  
  
produce (item)  
{  
    lock.acquire();  
    queue.enqueue(item);  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    item = queue.dequeue(item);  
    lock.release();  
    return item;  
}
```

A Simple Monitor Example

C++

```
Mutex lock;
Queue queue;

produce (item)
{
    lock.acquire();
    queue.enqueue(item);
    lock.release();
}

consume()
{
    lock.acquire();
    item = queue.dequeue(item);
    lock.release();
    return item;
}
```

Java

```
Queue queue;

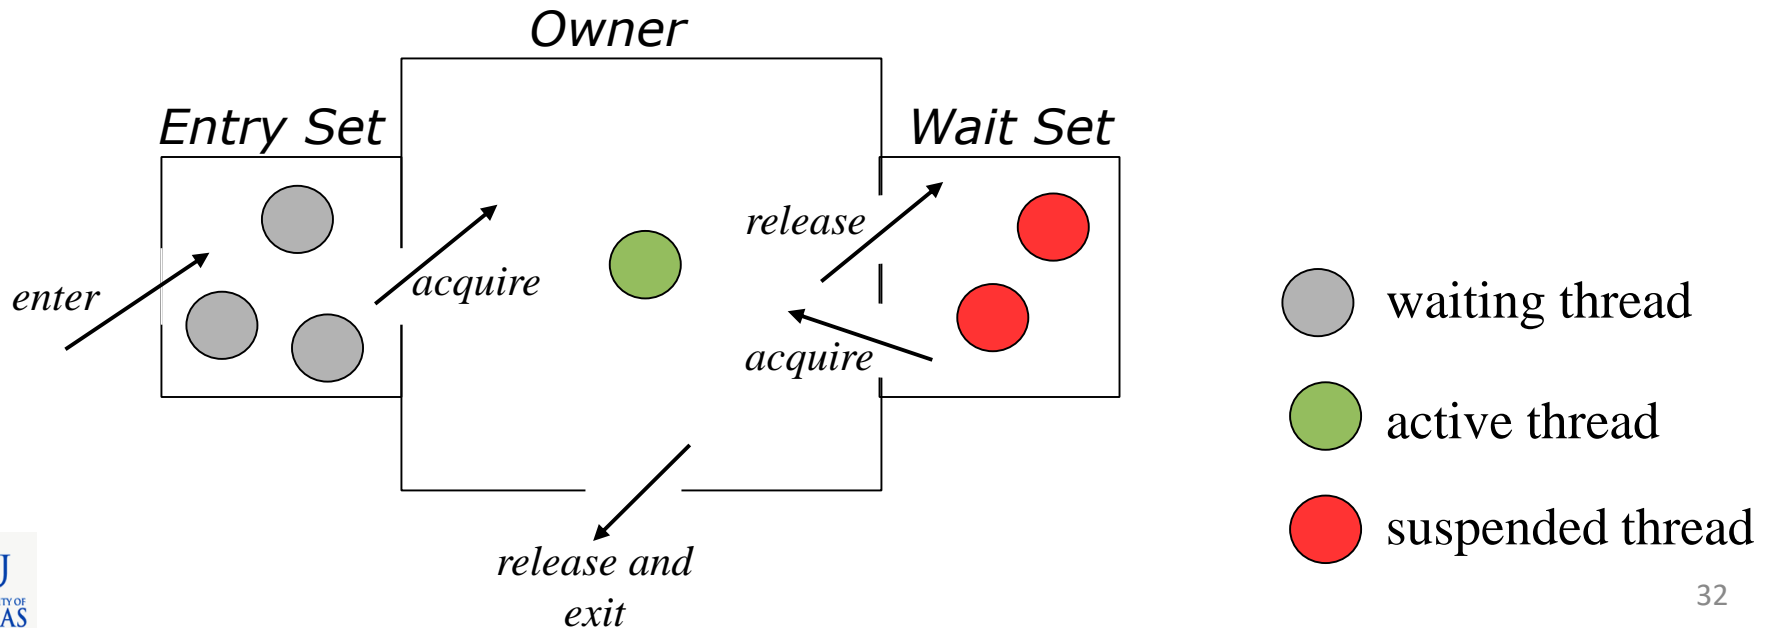
Synchronized produce (item)
{
    queue.enqueue(item);
}

Synchronized consume()
{
    item = queue.dequeue(item);

    return item;
}
```

Monitor

- What if a thread needs to wait inside a monitor?
- Condition variable: Scheduling
 - Wait(&lock): atomically release the lock and sleep; Re-acquire the lock on returning.
 - Signal(): wake-up one waiter, if exists
 - Broadcast(): wake-up all waiters



Monitor with Condition Variable

```
Mutex lock;  
Condition full;  
Queue queue;  
  
produce (item)  
{  
    lock.acquire();  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    lock.release();  
    return item;  
}
```

Why not 'if'?

idk, 可能就是想一直check lock的状态

Semantics

- Hoare monitors (original)
 - Signaler immediately switches to the waiting thread
 - Waiter's condition is guaranteed to hold when it resumes
- Mesa monitors
 - Waiter is simply placed on ready queue, signaler continues
 - Waiter's condition may no longer be true when it resumes
- Almost always Mesa style in practice

Bounded Buffer Problem Revisit

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    ...  
    ...  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    ...  
    lock.release();  
    return item;  
}
```

Bounded Buffer Problem Revisit

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```

Bounded Buffer Problem Revisit

Monitor version

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```

Semaphore version

```
Semaphore mutex = 1, full = 0,  
empty = N;  
  
produce (item)  
{  
    P(&empty);  
    P(&mutex);  
    queue.enqueue(item);  
    V(&mutex);  
    V(&full);  
}  
  
consume()  
{  
    P(&full);  
    P(&mutex);  
    item = queue.dequeue();  
    V(&mutex);  
    V(&empty);  
    return item;  
}
```

More Synchronization Primitives

- RCU (Read-Copy-Update)
 - Optimized for frequent read, infrequent write
 - Read is zero cost
 - Write can be costly
 - Heavily used in Linux kernel
- Transactional memory (Intel TSX instruction)
 - Opportunistic synchronization
 - Declare a set of instructions as a transaction
 - If no other CPUs update the shared data while executing the transaction, it is committed
 - If not (i.e., someone tries to modify in the middle of the transaction), the transaction will be aborted
 - If successful, there's no synchronization overhead

Summary

- Synchronization
 - Spinlock
 - Implement using h/w instructions (e.g., test-and-set)
 - Mutex
 - Sleep instead of spinning.
 - Semaphore
 - powerful tool, but often difficult to use
 - **Monitor**
 - Powerful and (relatively) easy to use

Agenda

- Famous Synchronization Bugs
 - THERAC-25
 - Mars Pathfinder

Therac 25



Image source: http://idg.bg/test/cwd/2008/7/14/21367-radiation_therapy.JPG

- Computer controlled medical X-ray treatments
- Six people died/injured due to massive overdoses (1985-1987)

Accident History

Date	What happened
June 1985	First overdose
July-Dec 1985	2nd and 3rd overdoses. Lawsuit against the manufacturer and hospital
Jan-Feb 1986	Manufacturer denied the possibility of overdoses
Mar-Apr 1986	Two more overdoses
May-Dec 1986	FDA orders corrective action plans to the manufacturer
Jan 1987	Sixth overdose
Nov 1988	Final safety analysis report

The Problem




Image source: <http://radonc.wdfiles.com/local--files/radiation-accident-therac25/Therac25.png>


- X-ray must be dosed with the filter in place
- But sometimes, X-ray was dosed w/o the filter

The Bug

```
unsigned char in_progress = 1;
```

```
Thread 1 : // tray movement thread (periodic)
  if (system_is_ready())
    in_progress = 0;
  else
    in_progress++;
```

```
Thread 2 : // X-ray control thread.
  if (in_progress == 0)
    start_radiation();
```



tray including the target, a flattening filter, the collimator jaws and an ion chamber was moved OUT for "electron" mode, and IN for "photon" mode.


- Can you spot the bug?

Fixed Code

```
unsigned char in_progress;

Thread 1 : // tray movement thread (periodic)
    if (system_is_ready())
        in_progress = 0;
    else
        in_progress = 1;

Thread 2 : // X-ray control thread.
    if (in_progress == 0)
        start_radiation();
```



- Can you do better using a monitor?

Monitor Version

```
Mutex lock;  
Condition ready;  
unsigned char in_progress;
```

```
Thread 1 : // on finishing tray movement  
    lock.acquire();  
    in_progress = 0;  
    ready.signal();  
    lock.release();
```

```
Thread 2 : // X-ray control thread.  
    lock.acquire();  
    while (in_progress)  
        ready.wait(&lock);  
    start_radiation();  
    lock.release();
```

- No periodic check is needed.

Mars Pathfinder

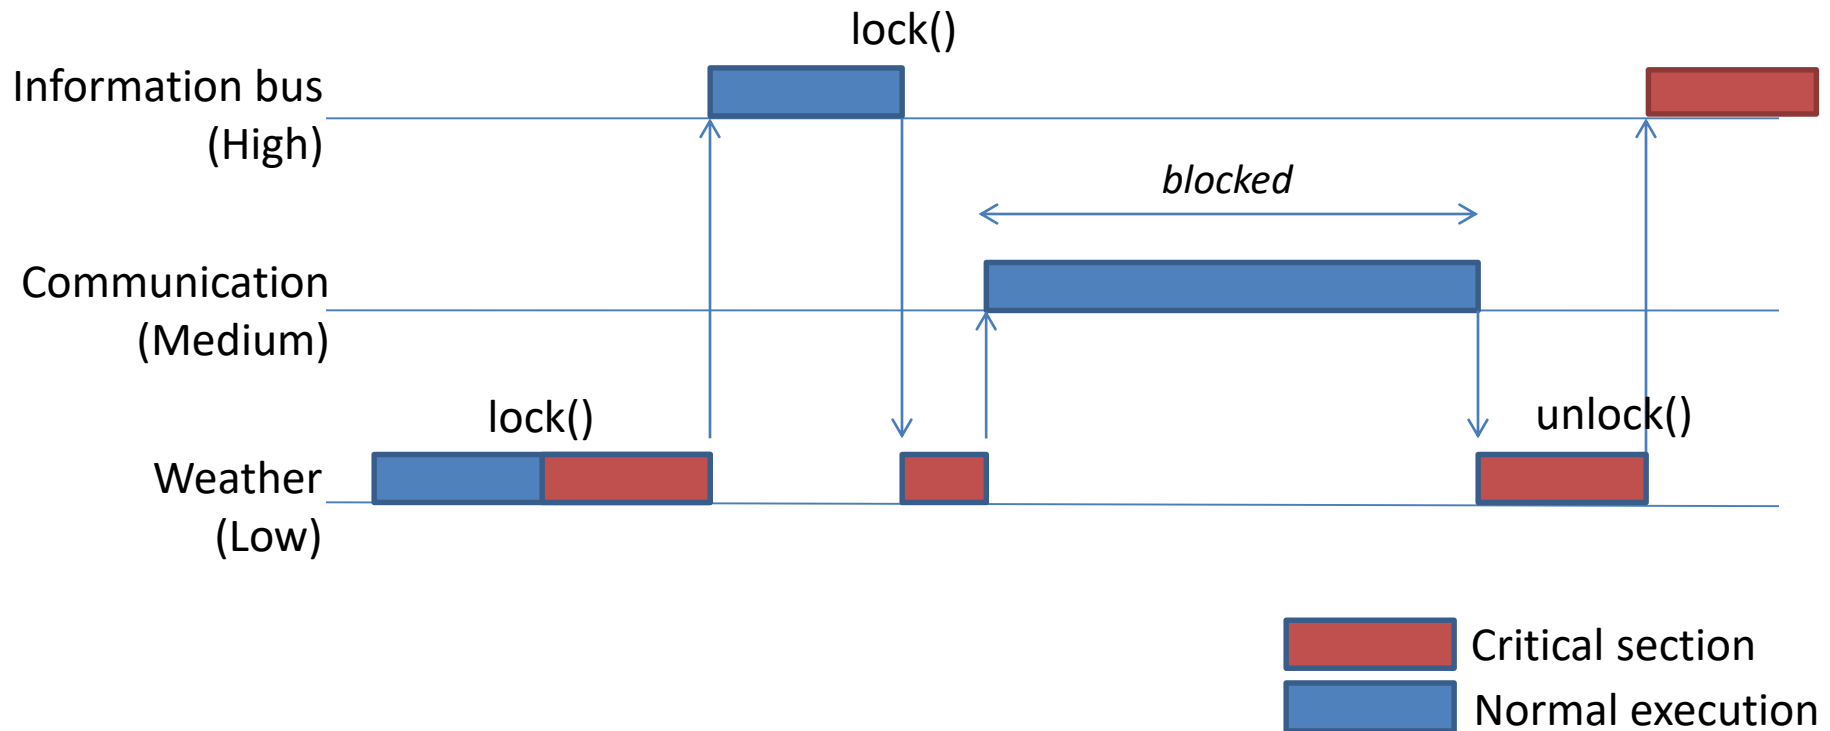


- Landed on Mars, July 4, 1997
- After operating for a while, it rebooted itself

The Bug

- Three threads with priorities
 - Weather data thread (low priority)
 - Communication thread (medium priority)
 - Information bus thread (high priority)
- Each thread obtains a lock to write data on the shared memory
- High priority thread can't acquire the lock for a very long time → something must be wrong. Let's reboot!

Priority Inversion



- High priority thread is delayed by the medium priority thread (potentially) indefinitely!!!

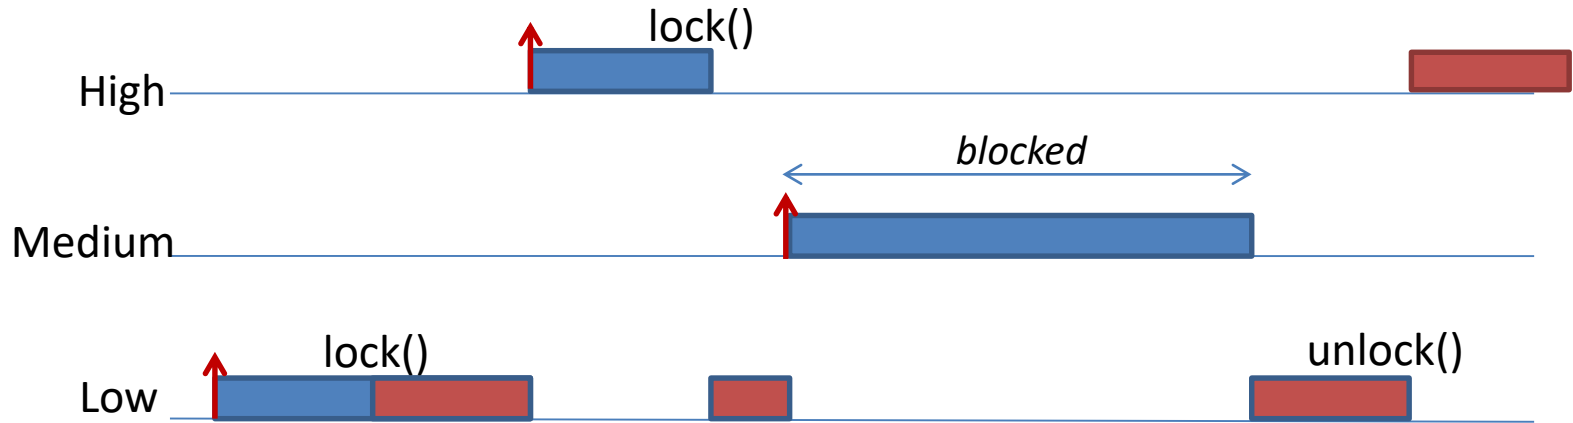
Solution

- Priority inheritance protocol [Sha'90]
 - If a high priority thread is waiting on a lock, boost the priority of the lock owner thread (low priority) to that of the high priority thread.
- Remotely patched the code
 - To use the priority inheritance protocol in the lock
 - First-ever(?) interplanetary remote debugging


L. Sha, R. Rajkumar, and J. P. Lehoczky. Priority Inheritance Protocols: An Approach to Real-Time Synchronization. In IEEE Transactions on Computers, vol. 39, pp. 1175-1185, Sep. 1990.

Priority Inheritance

Old



New



Summary

- Race condition
 - A situation when two or more threads **read and write** shared data at the same time
- Critical section
 - Code sections of potential race conditions
- Mutual exclusion
 - If a thread executes its critical section, *no other threads* can enter their critical sections
- Peterson's solution
 - Software only solution providing mutual exclusion

Summary

- Spinlock
 - Spin on waiting
 - Use synchronization instructions (test&set)
- Mutex
 - Sleep on waiting
- Semaphore
 - Powerful tool, but often difficult to use
- Monitor
 - Powerful and (relatively) easy to use

Quiz

- Semaphore mutex = 1;
- Semaphore **full** = 0;
- Semaphore **empty** = N;

Producer

```
do {  
    Produce new resource  
    _____;  
    mutex.P();  
    Add resource to next buffer  
    mutex.V();  
    _____;  
} while (TRUE);
```

Consumer

```
do {  
    _____;  
    mutex.P();  
    Remove resource from buffer  
    mutex.V();  
    _____;  
    Consume resource  
} while (TRUE);
```

Quiz

- Semaphore mutex = 1;
- Semaphore **full** = 0;
- Semaphore **empty** = N;

Producer

```
do {  
    Produce new resource  
    empty.P();  
    mutex.P();  
    Add resource to next buffer  
    mutex.V();  
    full.V();  
} while (TRUE);
```

Consumer

```
do {  
    full.P();  
    mutex.P();  
    Remove resource from buffer  
    mutex.V();  
    empty.V();  
    Consume resource  
} while (TRUE);
```

Quiz

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    _____  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    _____  
}  
  
consume()  
{  
    _____  
    while (queue.isEmpty())  
        _____  
    item = queue.dequeue(item);  
    _____  
    return item;  
}
```


Quiz

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```