

Deadlock

Disclaimer: some slides are adopted from Dr. Kulkarni's and book authors' slides with permission

Recap: Synchronization

- Race condition
 - A situation when two or more threads **read and write** shared data at the same time
- Critical section
 - Code sections of potential race conditions
- Mutual exclusion
 - If a thread executes its critical section, *no other threads* can enter their critical sections
- Peterson's solution
 - Software only solution providing mutual exclusion

Recap: Synchronization

- Spinlock
 - Spin on waiting
 - Use synchronization instructions (test&set)
- Mutex
 - Sleep on waiting
- Semaphore
 - Powerful tool, but often difficult to use
- Monitor
 - Powerful and (relatively) easy to use

Quiz

```
Mutex lock;
Condition full, empty;

produce (item)
{
    _____
    while (queue.isFull())
        empty.wait(&lock);
    queue.enqueue(item);
    full.signal();
    _____
}

consume()
{
    _____
    while (queue.isEmpty())
        _____
    item = queue.dequeue(item);
    _____
    return item;
}
```

```
Semaphore mutex = 1, full = 0,
empty = N;

produce (item)
{
    P (&empty) _____;
    P(&mutex);
    queue.enqueue(item);
    V(&mutex);
    _____ V (&full);
}

consume()
{
    _____ P (&full);
    P(&mutex);
    item = queue.dequeue();
    V(&mutex);
    _____ V (&empty);
    return item;
}
```

Quiz

Mutex lock;
Condition full, empty;

```
produce (item)
{
    lock.acquire();
    while (queue.isFull())
        empty.wait(&lock);
    queue.enqueue(item);
    full.signal();
    lock.release();
}

consume()
{
    lock.acquire();
    while (queue.isEmpty())
        full.wait(&lock);
    item = queue.dequeue(item);
    empty.signal();
    lock.release();
    return item;
}
```

Semaphore mutex = 1, full = 0,
empty = N;


```
produce (item)
{
    P(&empty);
    P(&mutex);
    queue.enqueue(item);
    V(&mutex);
    V(&full);
}

consume()
{
    P(&full);
    P(&mutex);
    item = queue.dequeue();
    V(&mutex);
    V(&empty);
    return item;
}
```

Agenda

- Deadlock
 - Starvation vs. deadlock
 - Deadlock conditions
 - General solutions: detection and prevention
 - Detection algorithm
 - Banker's algorithm

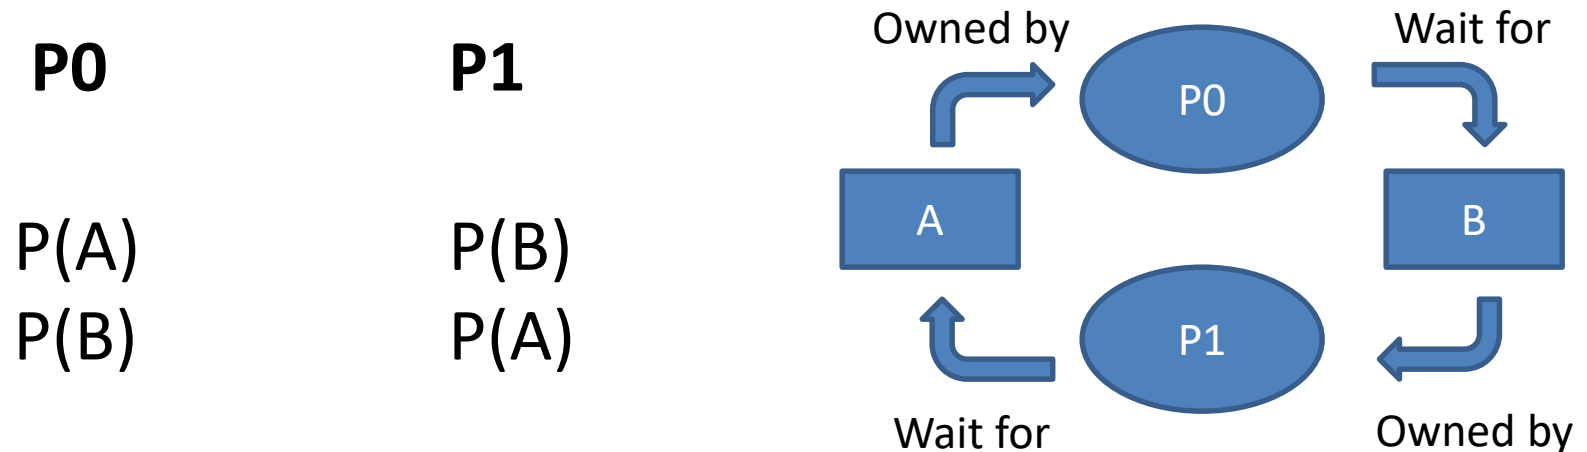
Starvation



- Starvation
 - Wait potentially **indefinitely**, but it **can end**

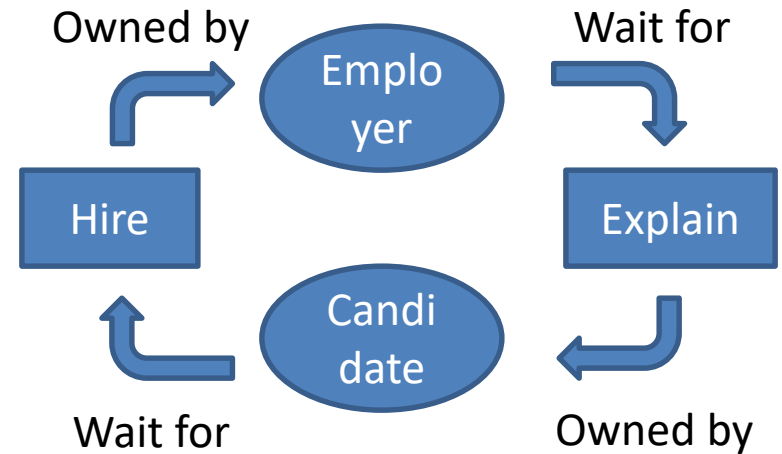

Starvation vs. Deadlock

- Deadlock: circular waiting for resources
 - Example: semaphore A = B = 1




- Deadlock → Starvation
 - But reverse is not true
 - Deadlock can't end but starvation can

Deadlock




Bridge Crossing



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- If a deadlock occurs, how to fix it?
 - Make one car backs up
 - Several cars may have to be backed up if a deadlock occurs

Dining Philosophers

- Problem synopsis
 - Need two chopsticks to eat
 - Grab one chopsticks at a time
- What happens if all grab left chopstick at the same time??
 - Deadlock!!!
- How to fix it?
- How to avoid it?



Conditions for Deadlocks

- Mutual exclusion
 - only one process at a time can use a resource
- No preemption
 - resources cannot be preempted, release must be voluntary
- Hold and wait
 - a process must be holding at least one resource, and waiting to acquire additional resources held by other processes
- Circular wait
 - There must be a circular dependency. For example, A waits B, B waits C, and C waits A.
- **All four conditions must simultaneously hold**

Resource-Allocation Graph

- To illustrate deadlock conditions.
- Graph consists of a set of vertices V and a set of edges E
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- request edge – directed edge $P_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph


- Process




- Resource Type with 4 instances




- P_i requests instance of R_j




- P_i is holding an instance of R_j




Resource Allocation Graph



Simple example



Deadlock example



With cycle, but
no deadlock

- request edge – directed edge $P_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow P_i$

Methods for Handling Deadlocks

- Detection and recovery
 - Allow a system to enter a deadlock and then recover
 - Need a *detection algorithm*
 - Somehow “preempt” resources
- Prevention and avoidance
 - Ensure a system never enter a deadlock
 - Possible solutions
 - have “Infinite resources”
 - prevent “hold and wait”
 - prevent “circular wait”

Recall four deadlock conditions:

(1) Mutual exclusion, (2) no preemption, (3) hold and wait, (4) circular wait


Deadlock Detection

- Deadlock detection algorithms
 - Single instance for each resource type
 - Multiple instances for each resource type


Single Instance Per Resource

- Each resource is unique
 - E.g., one printer, one audio card, ...
- Wait-for-graph
 - Variant of the simplified resource allocation graph
 - Remove resource nodes, collapse corresponding edges
- Detection algorithm
 - Searches for a cycle in the wait-for graph
 - Presence of a cycle points to the existence of a deadlock

Wait-for Graph




Resource-Allocation Graph



Corresponding wait-for graph

Multiple Instances Per Resource

- **n** processes, **m** resources
- **FreeResources**: resource vector (of size m)
 - indicates the number of available resources of each type
 - $[R_1, R_2] = [0, 0]$
- **Alloc[i]**: process i's allocated resource vector
 - defines the number of resources of each type currently allocated to each process
 - $\text{Alloc}[1] = [0, 1]$,
 - $\text{Alloc}[2] = [1, 0], \dots$
- **Request[i]**: process i's requesting resource vector
 - indicates the resources each process requests
 - $\text{Request}[1] = [1, 0]$,
 - $\text{Request}[2] = [0, 0], \dots$



Detection Algorithm

1. Initialize **Avail** and **Finish** vectors

Avail = FreeResources;

For $i = 1, 2, \dots, n$, Finish[i] = false

2. Find an index i such that

Finish[i] == false AND Request[i] \leq Avail

If no such i exists, go to step 4

3. Avail = Avail + Alloc[i], Finish[i] = true

Go to step 2


4. If Finish[i] == false, for some i , $1 \leq i \leq n$,

(a) then the system is in deadlock state

■ **FreeResources:** resource vector
 $[R_1, R_2] = [0, 0]$

■ **Alloc[i]:** process i 's allocated resource vector:
 $\text{Alloc}[1] = [0, 1]$, $\text{Alloc}[2] = [1, 0]$

■ **Request[i]:** process i 's requesting vector:
 $\text{Request}[1] = [1, 0]$
 $\text{Request}[2] = [0, 0]$

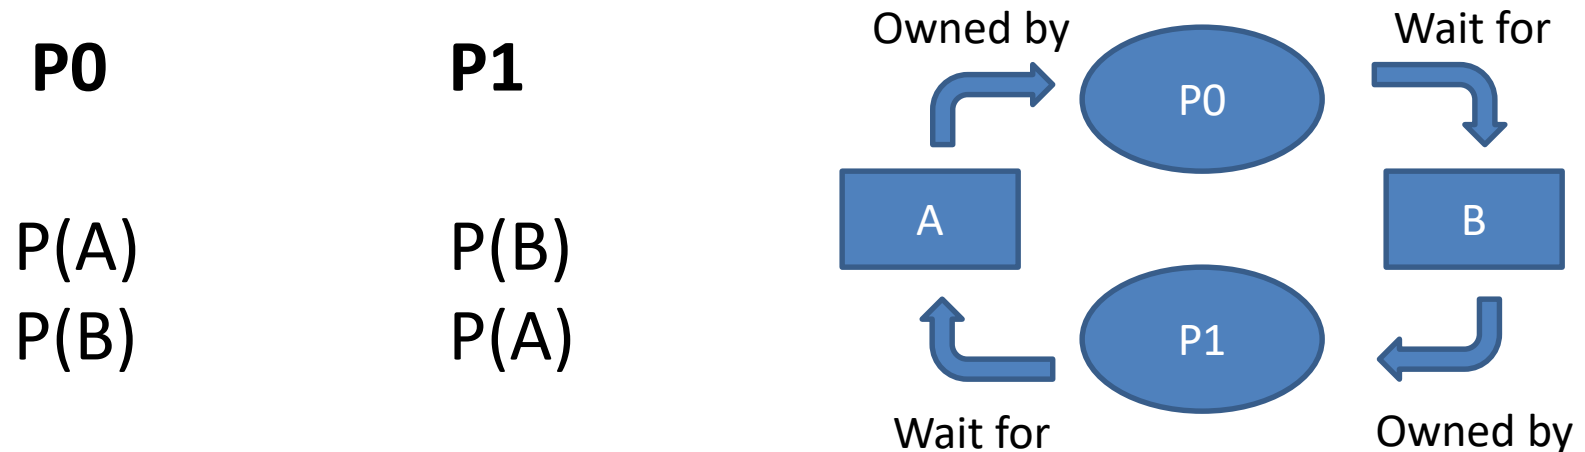


Recovery from Deadlock

- Terminate
 - Preempt the resources
 - Bridge example: throw the car to the river
 - Kill the deadlocked threads and return the resources
- Rollback
 - Return to a known safe state
 - Bridge example: move one car backward
 - Dining philosopher: make one philosopher give up a chopstick
- Not always possible!

Recap: Starvation vs. Deadlock

- Deadlock: circular waiting for resources
 - Example: semaphore $A = B = 1$



- Deadlock → Starvation
 - But reverse is not true
 - Deadlock can't end but starvation can

Recap: Conditions for Deadlocks

- Mutual exclusion
 - only one process at a time can use a resource
- No preemption
 - resources cannot be preempted, release must be voluntary
- Hold and wait
 - a process must be holding at least one resource, and waiting to acquire additional resources held by other processes
- Circular wait
 - There must be a circular dependency. For example, A waits B, B waits C, and C waits A.
- **All four conditions must simultaneously hold**

Recap: Detection Algorithm

1. Initialize **Avail** and **Finish** vectors

Avail = FreeResources;

For $i = 1, 2, \dots, n$, Finish[i] = false

2. Find an index i such that

Finish[i] == false AND Request[i] \leq Avail

If no such i exists, go to step 4

3. Avail = Avail + Alloc[i], Finish[i] = true

Go to step 2


4. If Finish[i] == false, for some i , $1 \leq i \leq n$,

(a) then the system is in deadlock state

■ **FreeResources:** resource vector
 $[R_1, R_2] = [0, 0]$

■ **Alloc[i]:** process i 's allocated resource vector:
 $\text{Alloc}[1] = [0, 1]$, $\text{Alloc}[2] = [1, 0]$

■ **Request[i]:** process i 's requesting vector:
 $\text{Request}[1] = [1, 0]$
 $\text{Request}[2] = [0, 0]$



Recap: Recovery from Deadlock

- Terminate
 - Preempt the resources
 - Bridge example: throw the car to the river
 - Kill the deadlocked threads and return the resources
- Rollback
 - Return to a known safe state
 - Bridge example: move one car backward
 - Dining philosopher: make one philosopher give up a chopstick
- Not always possible!

Deadlock Prevention

- Break any of the four deadlock conditions
 - Mutual exclusion
 - No preemption
 - Hold and wait
 - Circular wait

Deadlock Prevention

- Break any of the four deadlock conditions
 - **Mutual exclusion → allow sharing**
 - Well, not all resources are sharable
 - No preemption
 - Hold and wait
 - Circular wait

Deadlock Prevention

- Break any of the four deadlock conditions
 - Mutual exclusion → allow sharing
 - Well, not all resources are sharable
 - **No preemption → allow preemption**
 - This is also quite hard (kill the threads)
 - Hold and wait
 - Circular wait

Deadlock Prevention

- Break any of the four deadlock conditions
 - Mutual exclusion → allow sharing
 - Well, not all resources are sharable
 - No preemption → allow preemption
 - This is also quite hard (kill the threads)
 - **Hold and wait → get all resources at once**
 - Dining philosopher: get *both* chopsticks or *none*
 - Circular wait

Deadlock Prevention

- Break any of the four deadlock conditions
 - Mutual exclusion → allow sharing
 - Well, not all resources are sharable
 - No preemption → allow preemption
 - This is also quite hard (kill the threads)
 - Hold and wait → get all resources at once
 - Dining philosopher: get *both* chopsticks or *none*
 - **Circular wait → prevent cycle**
 - Dining philosopher: change the chopstick picking order;
if grabbing a chopstick will form a cycle, prevent it.

Banker's Algorithm

- General idea
 - Assume that each process's maximum resource demand is known in advance
 - $\text{Max}[i]$: process i 's maximum resource demand vector
 - **Pretend** each request is granted, then run the deadlock detection algorithm
 - If a deadlock is detected, the do not grant the request to keep the system in a **safe** state

Banker's Algorithm

1. Initialize **Avail** and **Finish** vectors

Avail = FreeResources;

For $i = 1, 2, \dots, n$, Finish[i] = false

2. Find an index i such that

Finish[i] == false AND

$$\text{Max}[i] - \text{Alloc}[i] \leq \text{Avail}$$

If no such i exists, go to step 4

3. Avail = Avail + Alloc[i] , Finish[i] = true

Go to step 2

4. If Finish[i] == false, for some i , $1 \leq i \leq n$,

(a) then the system is in deadlock state

(b) if Finish[i] == false, then P_i is deadlocked

■ FreeResources: resource vector
[R1, R2] = [0,0]


■ Alloc[i]: process i 's allocated resource vector:
Alloc[1] = [0,1], Alloc[2] = [1, 0]

■ Request[i]: process i 's requesting vector:
Request[1] = [1,0]
Request[2] = [0,0]

■ **Max[i]**: process i 's maximum resource demand vector

Example

Free = [1,1,1,1,1]



Example

Free = [0,0,0,0,1]

Avail = [0,0,0,0,1]


Max[2] = [1,1,0,0,0]
Alloc[2] = [0,1,0,0,0]

Max[3] = [0,1,1,0,0]
Alloc[3] = [0,0,1,0,0]

Max[1] = [1,0,0,0,1]
Alloc[1] = [1,0,0,0,0]

Max[4] = [0,0,1,1,0]
Alloc[4] = [0,0,0,1,0]

Max[5] = [0,0,0,1,1]
Alloc[5] = [0,0,0,0,0]



- Philosopher 5 requested R5.
- **Safe or Unsafe?**

Example

Free = [0,0,0,0,0]

Avail = [0,0,0,0,0]


Max[2] = [1,1,0,0,0]
Alloc[2] = [0,1,0,0,0]

Max[3] = [0,1,1,0,0]
Alloc[3] = [0,0,1,0,0]

Max[1] = [1,0,0,0,1]
Alloc[1] = [1,0,0,0,0]

Max[4] = [0,0,1,1,0]
Alloc[4] = [0,0,0,1,0]

Max[5] = [0,0,0,1,1]
Alloc[5] = [0,0,0,0,1]



2. Find an index i such that
Finish[i] == false AND
Max[i] - Alloc[i] ≤ Avail
If no such i exists, go to step 4

Quiz

- Determine whether this state is safe or unsafe.

Total resources: 12

Avail resources: 3

Process	Max	Alloc
P_0	10	4
P_1	3	1
P_2	6	4

$$10 - 4 \leq 8$$

$$3 - 1 \leq 7$$

$$6 - 4 \leq 3$$

Safe

Quiz

- Suppose P0 requested 3 additional resources. Should this request be granted?

Total resources: 12

Avail resources: 0

Process	Max	Alloc		
P ₀	10	7	$10 - 7 \leq 0$	Unsafe
P ₁	3	1	$3 - 1 \leq 0$	
P ₂	6	4	$6 - 4 \leq 0$	

Quiz

- Suppose there are three resource types, which are needed by four processes. Is it safe now?

R1	R2	R3
2	1	2

← Free Resources

	Current Allocation			Max		
Process	R1	R2	R3	R1	R2	R3
P1	0	0	1	0	0	3
P2	2	0	0	2	4	5
P3	0	0	3	6	3	5
P4	2	3	5	4	3	5

← Current and Maximum Allocations

Quiz

- If P2 requests (1, 1, 2), should it be granted?

R1	R2	R3
2	1	2

← Free Resources

	Current Allocation			Max		
Process	R1	R2	R3	R1	R2	R3
P1	0	0	1	0	0	3
P2	2	0	0	3	4	5
P3	0	0	3	6	3	5
P4	2	3	5	4	3	5

← Current and
Maximum Allocations

Quiz

- If P2 requests (1, 1, 2), should it be granted?

R1	R2	R3
1	0	0

← Free Resources

No

	Current Allocation			Max		
Process	R1	R2	R3	R1	R2	R3
P1	0	0	1	0	0	3
P2	3	1	2	3	4	5
P3	0	0	3	6	3	5
P4	2	3	5	4	3	5

← Current and Maximum Allocations

Summary

- Four deadlock conditions:
 - Mutual exclusion
 - No preemption
 - Hold and wait
 - Circular wait
- Detection
- Avoidance
 - Banker's algorithm