

CPU Scheduling

Heechul Yun



Agenda

- Introduction to CPU scheduling
- Classical CPU scheduling algorithms

CPU Scheduling

- CPU scheduling is a **policy** to decide
 - **Which** thread to run next?
 - **When** to schedule the next thread?
 - **How long?**
- Context switching is a **mechanism**
 - To change the running thread

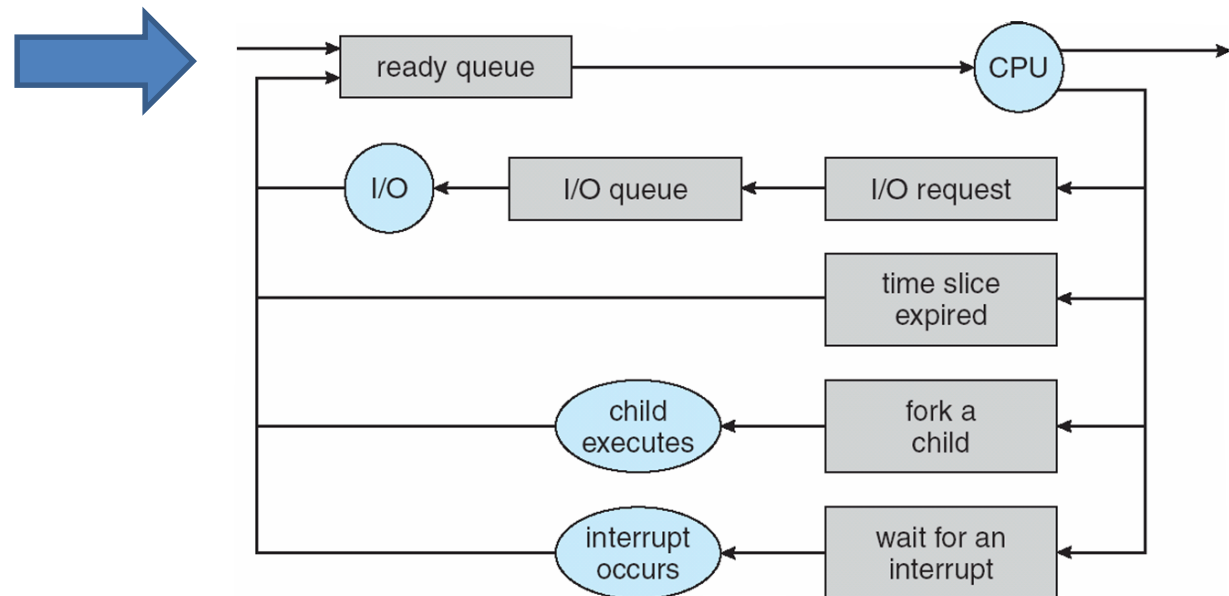
Assumption: CPU Bursts



- Execution model
 - Program uses the CPU for a while and then does some I/O, back to use CPU, ..., keep alternating


CPU Scheduler

- An OS component that determines which thread to run, at what time, and how long
 - Among threads in the **ready queue**



CPU Scheduler

- When the scheduler runs?




- The running thread finishes
- The running thread voluntarily gives up the CPU
 - yield, block on I/O, ...
- The OS **preempts** the current running thread
 - quantum expire (timer interrupt)

Performance Metrics for CPU Scheduling

- CPU utilization
 - % of time the CPU is busy doing something
- Throughput
 - #of jobs done / unit time
- **Response time (Turn-around time)**
 - Time to complete a task (ready -> complete)
- **Waiting time**
 - Time spent on waiting in the ready queue
- **Scheduling latency**
 - Time to schedule a task (ready -> first scheduled)

Example


- Assumption: A, B, C are released at time 0



- The times of Process **A**
 - Response time: 9
 - Wait time: 5
 - Sched. latency: 0

Example


- Assumption: A, B, C are released at time 0



- The times of Process **B**
 - Response time: 5
 - Wait time: 3
 - Latency: 1

Example

- Assumption: A, B, C are released at time 0




- The times of Process **C**
 - Response time: 10
 - Wait time: 6
 - Latency: 2

Recap: CPU Scheduling

- CPU scheduling is a **policy** to decide
 - **Which** thread to run next?
 - **When** to schedule the next thread?
 - **How long?**
- Context switching is a **mechanism**
 - To change the running thread

Recap: Example Metrics

- Assumption: A, B, C are released at time 0



- The times of Process **C**
 - Response time: 10
 - Wait time: 6
 - Latency: 2


Workload Model and Gantt Chart

- Workload model

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	1	10
P4	6	2

- Gantt chart

– bar chart to illustrate a particular schedule



Agenda

- Basic scheduling policies
 - First-come-first-serve (FCFS)
 - Shortest-job-first (SJF)
 - Shortest-remaining-time-first (SRTF)
 - Round-robin (RR)

Scheduling Policy Goals

- Maximize throughput (minimize avg. waiting time)
 - High throughput (#of jobs done / time) is good
- Minimize scheduling latency
 - Important to interactive applications (games, editor, ...)
- Fairness
 - Make all threads progress equally
- Goals often conflicts
 - Frequent context switching may be good for reducing response time, but not so much for maximizing throughput

First-Come, First-Served (FCFS)

- FCFS
 - Assigns the CPU based on the order of the requests.
 - Implemented using a FIFO queue.




© bnpdesignstudio * www.ClipartOf.com/226258

FCFS

- Example

Process	Arrival Time	Burst Time
P1	0	24
P2	0	3
P3	0	3

– Suppose that the processes arrive in the order: P_1, P_2, P_3



– Waiting time?

- $P1 = 0$; $P2 = 24$; $P3 = 27$

– Average waiting time


- $(0 + 24 + 27)/3 = 17$

FCFS

- Example 2

Process	Arrival Time	Burst Time
P1	0	24
P2	0	3
P3	0	3

– Suppose that the processes arrive in the order: P_2, P_3, P_1



– Waiting time?

- $P_1 = 6$; $P_2 = 0$; $P_3 = 3$

– Average waiting time

- $(6 + 0 + 3)/3 = 3$

– Much better than previous case → performance varies greatly depending on the scheduling order

Shortest Job First (SJF)


- Can we always do better than FIFO?
 - Yes: **if** you know the tasks' CPU burst times
- Shortest Job First (SJF)
 - Order jobs based on their burst lengths
 - Executes the job with the shortest CPU burst first
 - SJF is optimal
 - Achieves minimum average waiting time

Shortest Job First (SJF)

- Example

Process	Arrival Time	Burst Time
P1	0	6
P2	0	8
P3	0	7
P4	0	3

- Gantt chart



- Average waiting time?


- $(3 + 16 + 9 + 0) / 4 = 7$

- How to know the CPU burst time **in advance**?

Determining CPU Burst Length

- Can only estimate the length
 - Next CPU burst similar to previous CPU bursts ?
 - Predict based on the past history
- Exponential weighted moving average (EWMA)
 - of past CPU bursts

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.




CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

Shortest Job First (SJF)

- What if jobs don't arrive at the same time?

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5




– Average waiting time

- $(0+7+15+9)/4 = 7.5$

Shortest Remaining Time First (SRTF)

- Preemptive version of SJF
- New shorter job preempt longer running job

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5

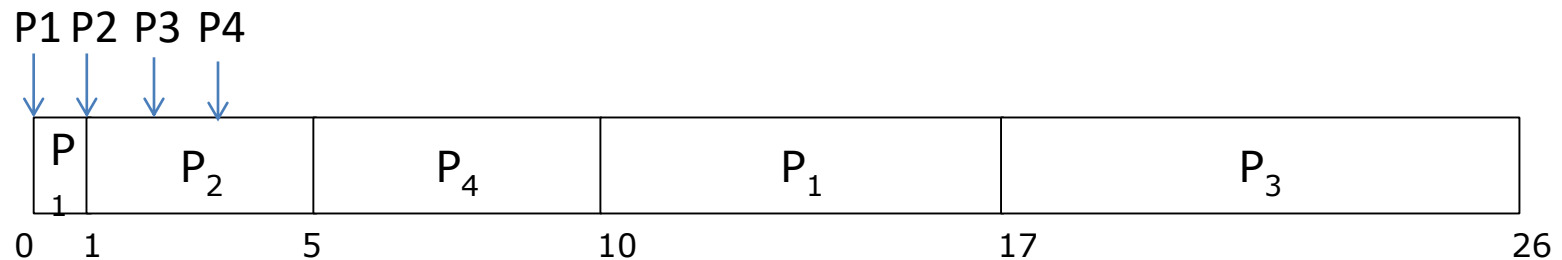


- Average waiting time
– $(9 + 0 + 15 + 2) / 4 = 6.5$

Quiz: SRTF

- Average waiting time?

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5



- $(9 + 0 + 15 + 2) / 4 = 6.5$

So Far...

- FIFO
 - In the order of arrival
 - Non-preemptive
- SJF
 - Shortest job first.
 - Non preemptive
- SRTF
 - Preemptive version of SJF

Issues

- FIFO
 - Bad average response time
- SJF/SRTF
 - Good average waiting time
 - IF you know or can predict the future
- Time-sharing systems
 - Multiple users share a machine
 - Need high interactivity → low **scheduling latency**

Round-Robin (RR)

- FIFO with preemption
- Simple, fair, and easy to implement
- Algorithm
 - Each job executes for a fixed time slice: **quantum**
 - When quantum expires, the scheduler preempts the task
 - Schedule the next job and continue...


Round-Robin (RR)

- Example

- Quantum size = 4

Process	Burst Times
P1	24
P2	3
P3	3

- Gantt chart



- Sched. Latency (between ready to first schedule)


- P1: 0, P2: 4, P3: 7. average response time = $(0+4+7)/3 = 3.67$

- Waiting time

- P1: 6, P2: 4, P3: 7. average waiting time = $(6+4+7)/3 = 5.67$

How To Choose Quantum Size?

- Quantum length
 - Too short → high overhead (why?)
 - Too long → bad scheduling latency
 - Very long quantum → FIFO




Round-Robin (RR)

- Example

- Quantum size = 2

Process	Burst Times
P1	24
P2	3
P3	3

- Gantt chart



- Scheduling latency

- P1: 0, P2: 2, P3: 4. average response time = $(0+2+4)/3 = 2$

- Waiting time

- P1: 6, P2: 6, P3: 7. average waiting time = $(6+6+7)/3 = 6.33$

Discussion

- Comparison between FCFS, SRTF(SJF), and RR
 - What to choose for smallest average waiting time?
 - SRTF (SFJ) is the optimal
 - What to choose for better interactivity?
 - RR with small time quantum (or SRTF)
 - What to choose to minimize scheduling overhead?
 - FCFS

Example




- Task A and B
 - CPU bound, run an hour
- Task C
 - I/O bound, repeat(1ms CPU, 9ms disk I/O)
- FCFS?
 - If A or B is scheduled first, C can begin an hour later
- RR and SRTF?

Example Timeline




RR with 100ms time quantum

C A B ... A B C A B A B ... C A B ...



RR with 1ms time quantum

C A C A C A ...



SRTF

Summary

- First-Come, First-Served (FCFS)
 - Run to completion in order of arrival
 - Pros: simple, low overhead, good for batch jobs
 - Cons: short jobs can stuck behind the long ones
- Round-Robin (RR)
 - FCFS with preemption. Cycle after a fixed time quantum
 - Pros: better interactivity (low average scheduling latency)
 - Cons: performance is dependent on the quantum size
- Shortest Job First (SJF)/ Shorted Remaining Time First (SRTF)
 - Shorted job (or shortest remaining job) first
 - Pros: optimal average waiting time
 - Cons: you need to know the future, long jobs can be starved by short jobs