

Memory Management

Disclaimer: some slides are adopted from book authors' slides with permission

Roadmap

- CPU management
 - Process, thread, synchronization, scheduling
- **Memory management**
 - Virtual memory
- Disk management
- Other topics

Memory Management

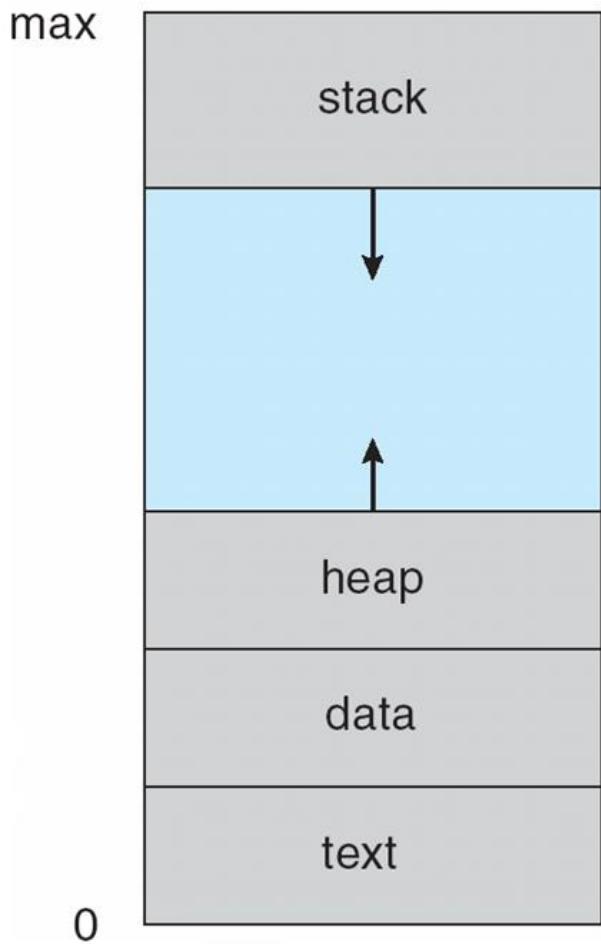
- Goals
 - Easy to use abstraction
 - Same virtual memory space for all processes
 - Isolation among processes
 - Don't corrupt each other
 - Efficient use of capacity limited physical memory
 - Don't waste memory

Concepts to Learn

- Virtual address translation
- Paging and TLB
- Page table management
- Swap

Virtual Memory (VM)

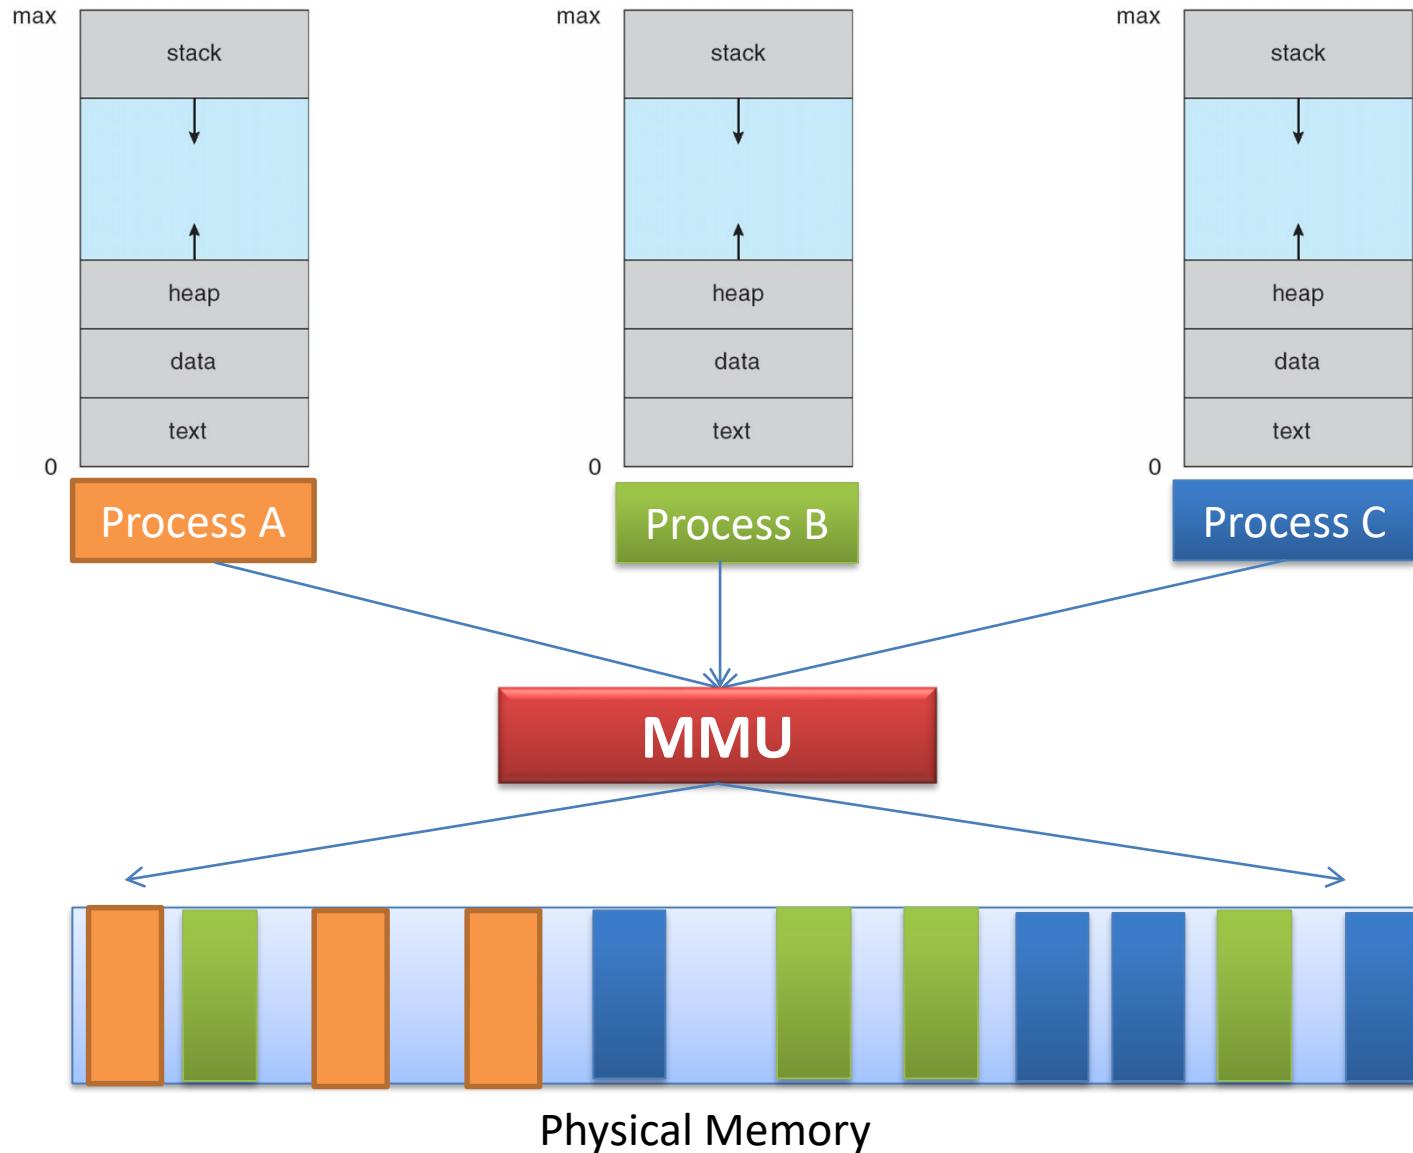
- Abstraction
 - 4GB linear address space for each process
- Reality
 - 1GB of actual physical memory shared with 20 other processes
- How?



Virtual Memory

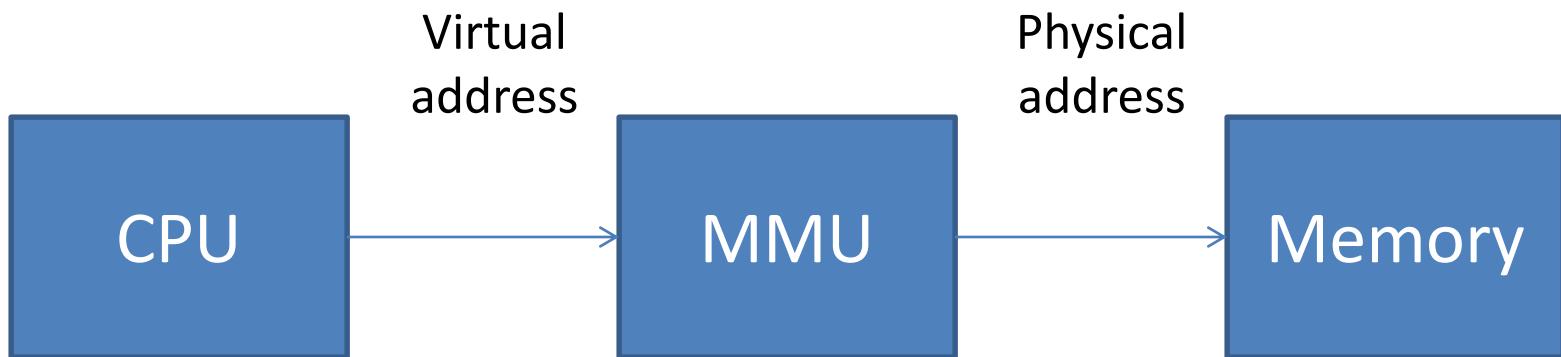
- **Hardware support**
 - MMU (memory management unit)
 - TLB (translation lookaside buffer)
- OS support
 - Manage MMU (sometimes TLB)
 - Determine address mapping
- Alternatives
 - No VM: many real-time OS (RTOS) don't have VM

Virtual Address



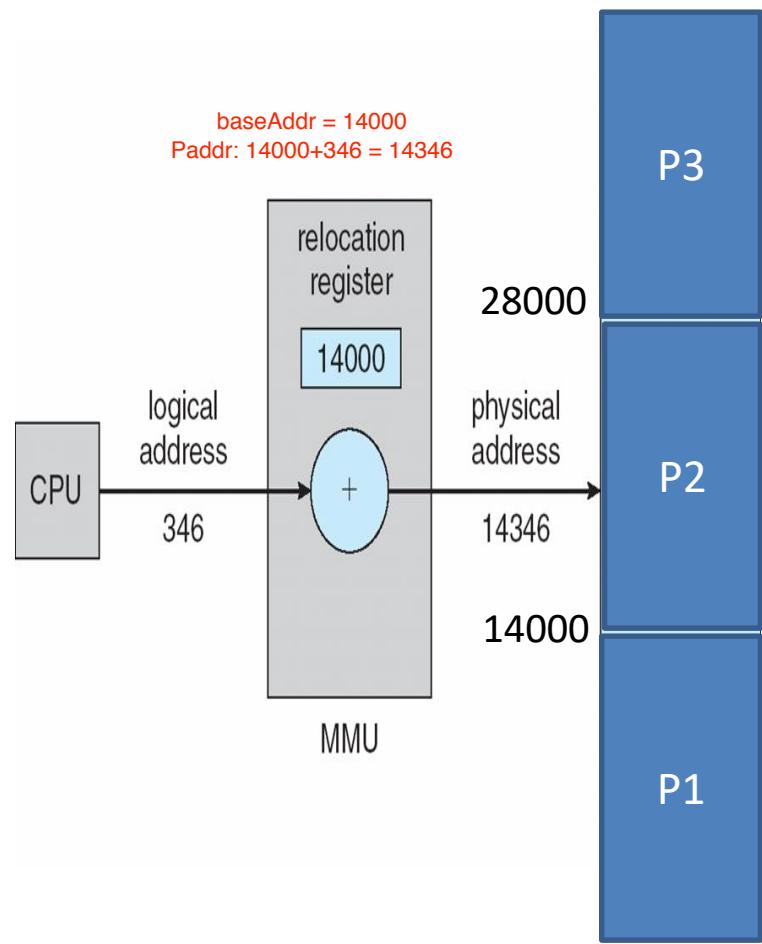
MMU

- Hardware unit that translates *virtual address* to *physical address*



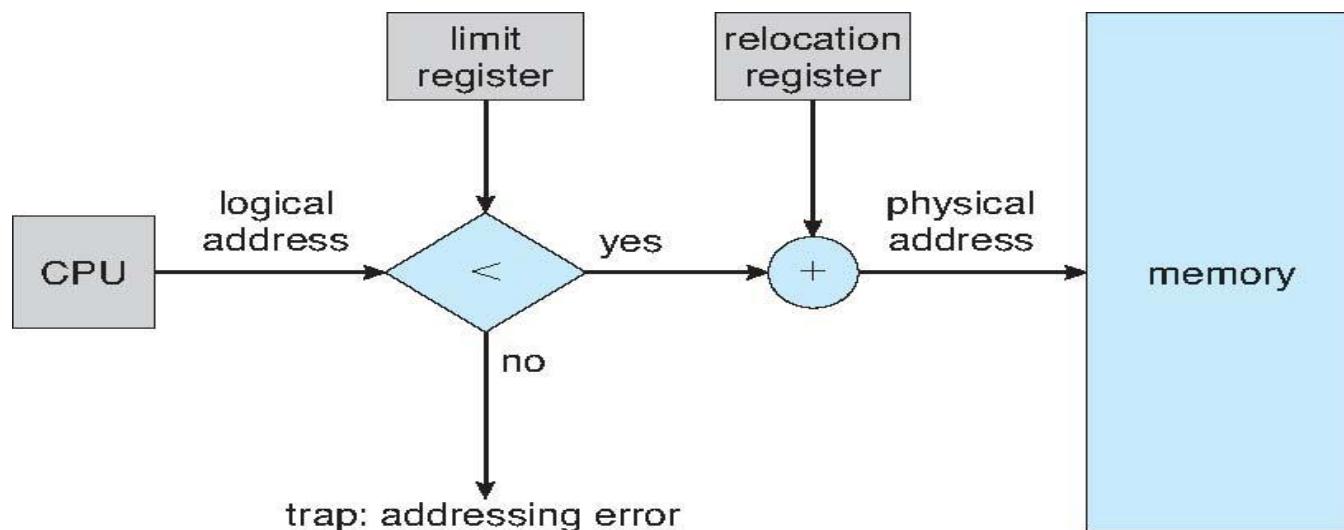
A Simple MMU

- BaseAddr: base register
- $P_{addr} = V_{addr} + BaseAddr$
- Advantages
 - Fast
- Disadvantages
 - No protection
 - Wasteful



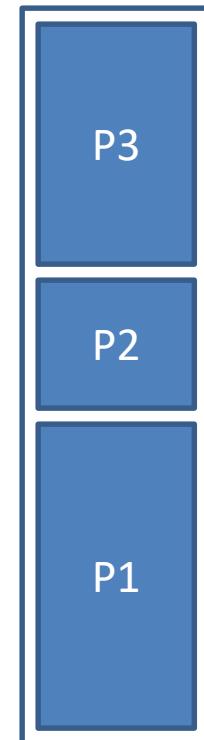
A Better MMU

- Base + Limit approach
 - If $Vaddr > limit$, then trap to report error
 - Else $Paddr = Vaddr + BaseAddr$



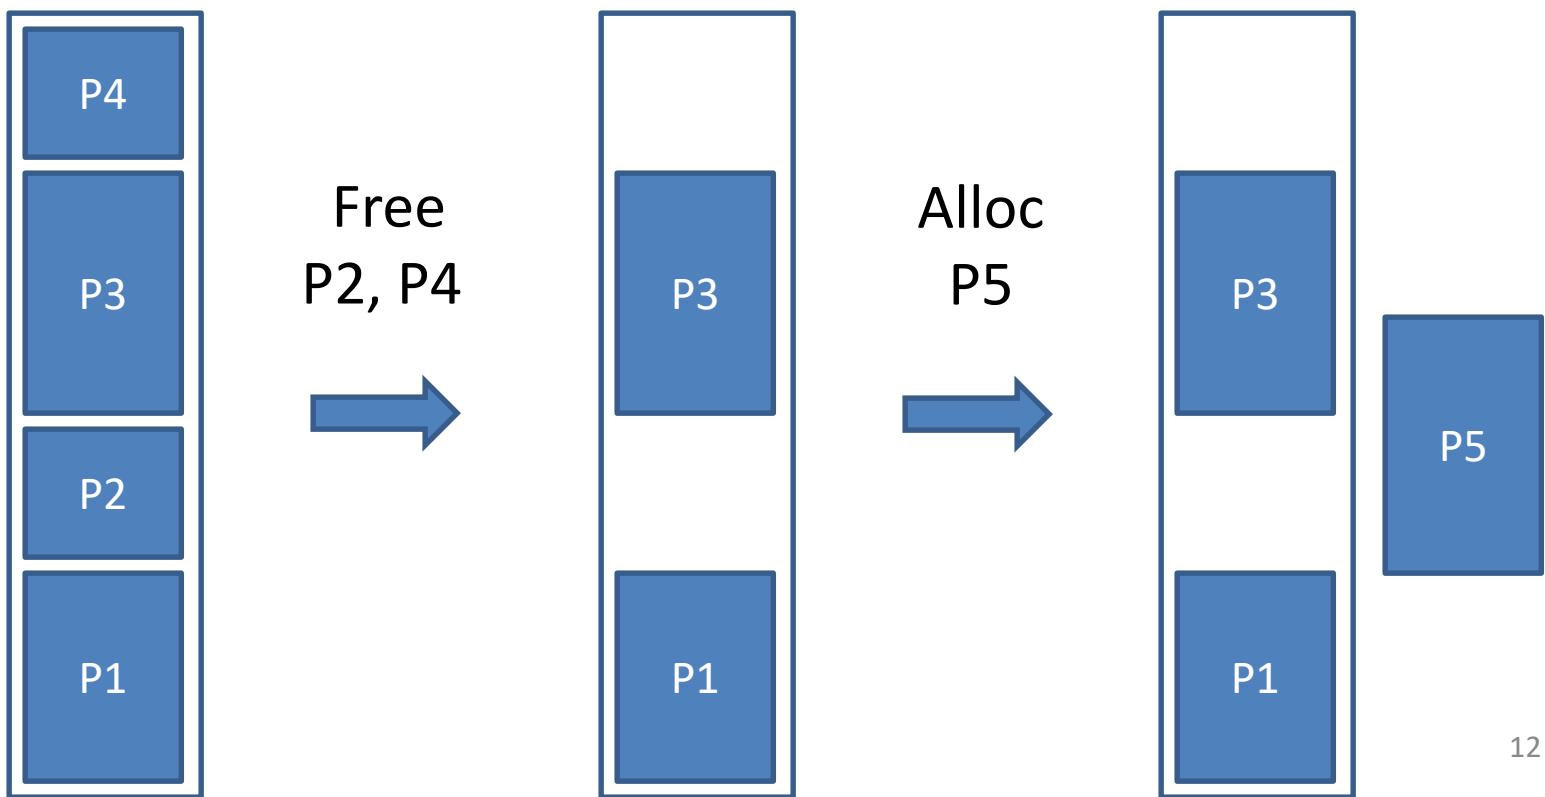
A Better MMU

- Base + Limit approach
 - If $Vaddr > limit$, then trap to report error
 - Else $Paddr = Vaddr + BaseAddr$
- Advantages
 - Support protection
 - Support variable size partitions
- Disadvantages
 - Fragmentation



Fragmentation

- External fragmentation
 - total available memory space exists to satisfy a request, but it is **not contiguous** 连续的



Modern MMU

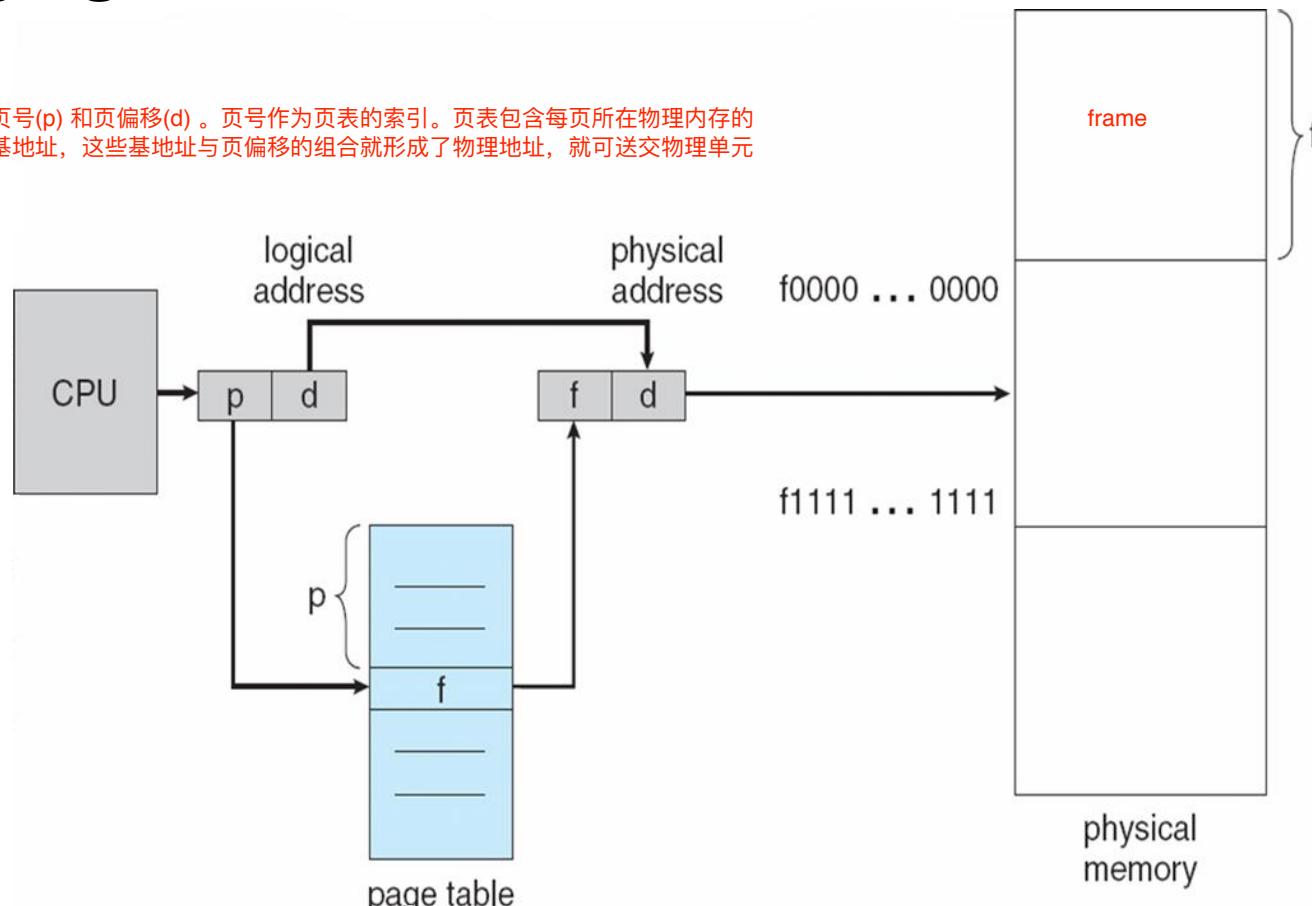
- **Paging approach**
 - Divide physical memory into fixed-sized blocks called frames (e.g., 4KB each)
2^2 * 2^10 = 4KB
1*10^3 = 1K
 - Divide logical memory into blocks of the same size called **pages** (page size = frame size)
 - Pages are mapped onto frames via a table → page table

<https://blog.csdn.net/qqqqq1993qqqq/article/details/74199784>

Modern MMU

- Paging hardware

页号(p) 和页偏移(d)。页号作为页表的索引。页表包含每页所在物理内存的基地址，这些基地址与页偏移的组合就形成了物理地址，就可送交物理单元

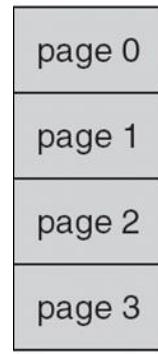


Modern MMU

$n=2$ $m=4$ page size = 4 bytes Memory = 32 bytes; so memory was divided into 8 part; Actually we have 4 pages

- Memory view

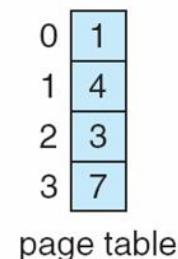
logical memory = $n+m$
 n 决定 page 数量
 m 决定页偏移量



logical
memory

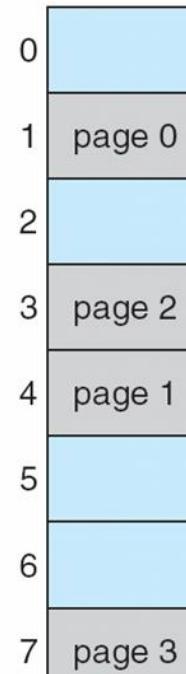
divide logical memory into blocks = pages

each page size == each frame size



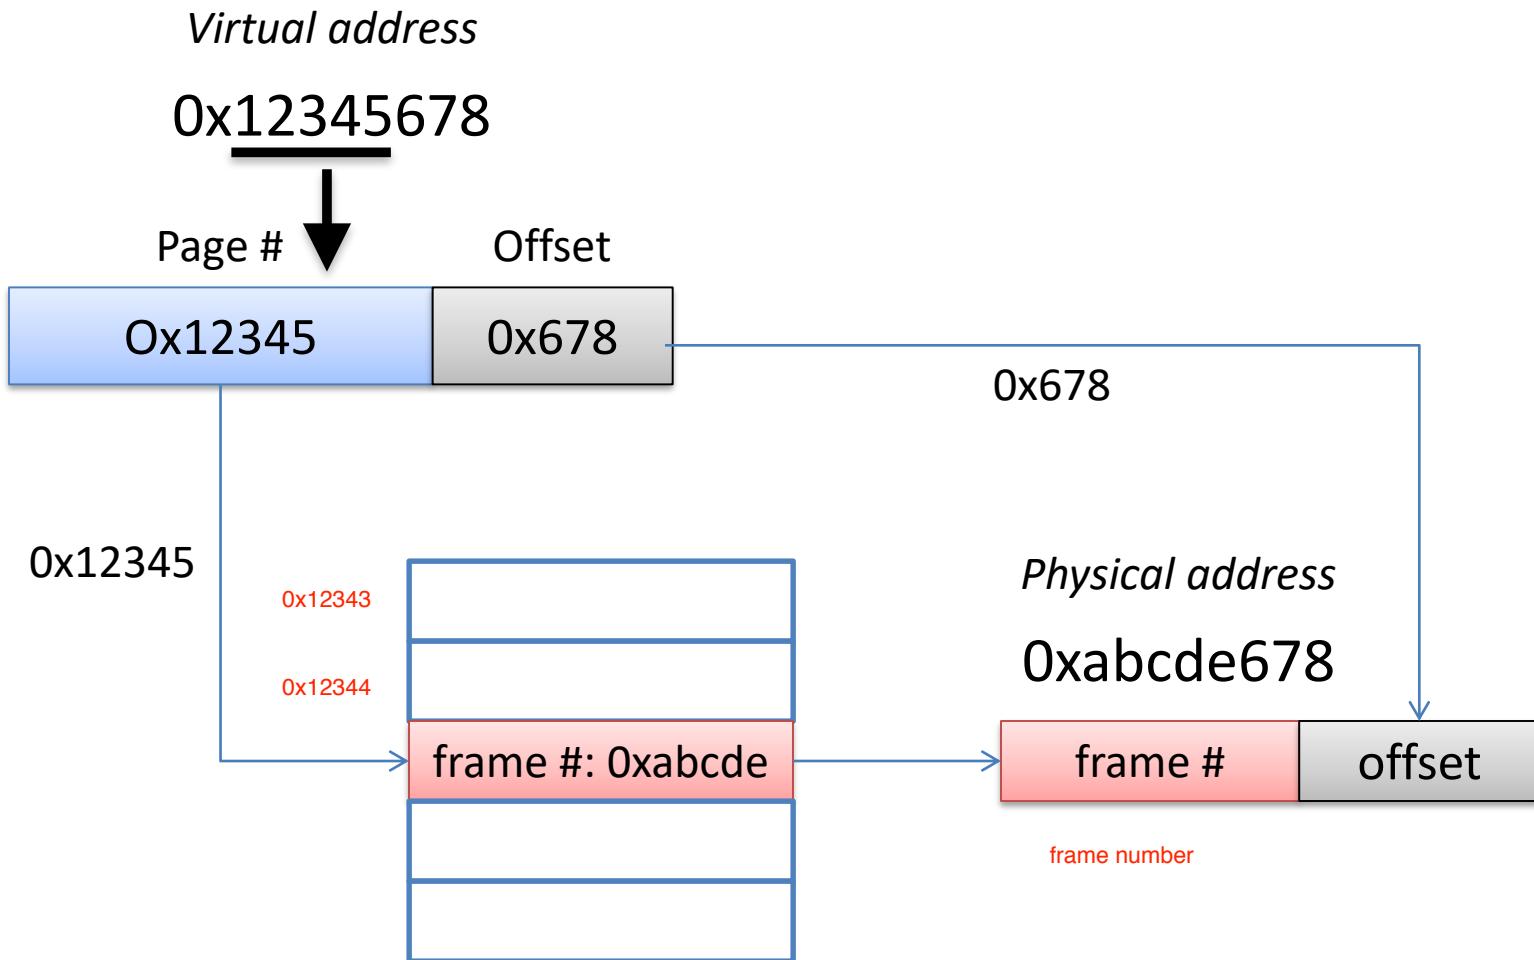
page table

frame
number 32byte physical memory / 4 byte PTE = 8 pages



physical
memory

Virtual Address Translation



Advantages of Paging

- No external fragmentation
 - Efficient use of memory
 - Internal fragmentation (waste within a page) still exists

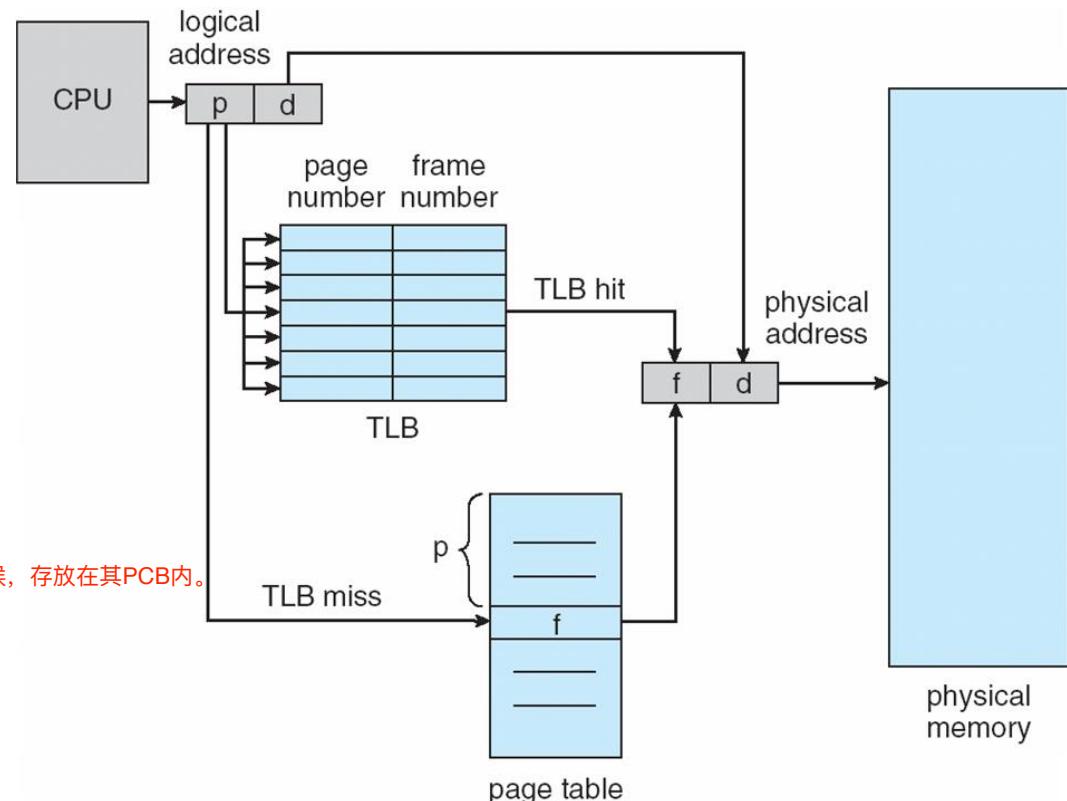
Issues of Paging

- Translation speed
 - Each load/store instruction requires a translation
 - Table is stored in **memory**
 - Memory is **slow** to access
 - ~100 CPU cycles to access DRAM

Translation Lookaside Buffer (TLB)

Important 后面会学到很多替换方法当TLB满了
现查找TLB的，再找Page table

- Cache frequent address translations
 - So that CPU don't need to access the page table all the time
 - Much faster



<https://blog.csdn.net/u014338577/article/details/82750771>

所以页表被放在物理内存中，由操作系统维护。
程都有页表，页表起始地址和页表长度的信息在进程不被CPU执行的时候，存放在其PCB内。

Issues of Paging

Memory 内存基本单元: Byte

Since we have a virtual address space of 2^{32} and each page size is 2^{12} , we can store $(2^{32}/2^{12}) = 2^{20}$ pages. Since each entry into this page table has an address of size 4 bytes, then we have $2^{20} * 4 = 4MB$. So the page table takes up 4MB in memory.

- **Page size**

- Small: minimize space waste, requires a large table
- Big: can waste lots of space, the table size is small
- Typical size: **4KB**
 - 32位操作系统最低12位的页偏移量—>所以PTE = $2^{12} * 1B = 4KB$
- How many pages are needed for 4GB (32bit)?
 - $4GB/4KB = 1M$ pages
 - B —> KB —> MB —> GB(GigaBytes) => $2^{32}/2^{12} = 2^{20}$
 - $4GB/4KB == 1(M) = 2^{20}$
- What is the required page table size?
 - assume 1 page table entry (PTE) is 4bytes
 - $1M * 4bytes = 4MB$
- Btw, this is for each process. What if you have 100 processes? Or what if you have a 64bit address?

4MB page table size can represent 4GB spaces

$100 * 4MB = 400$ MB too large

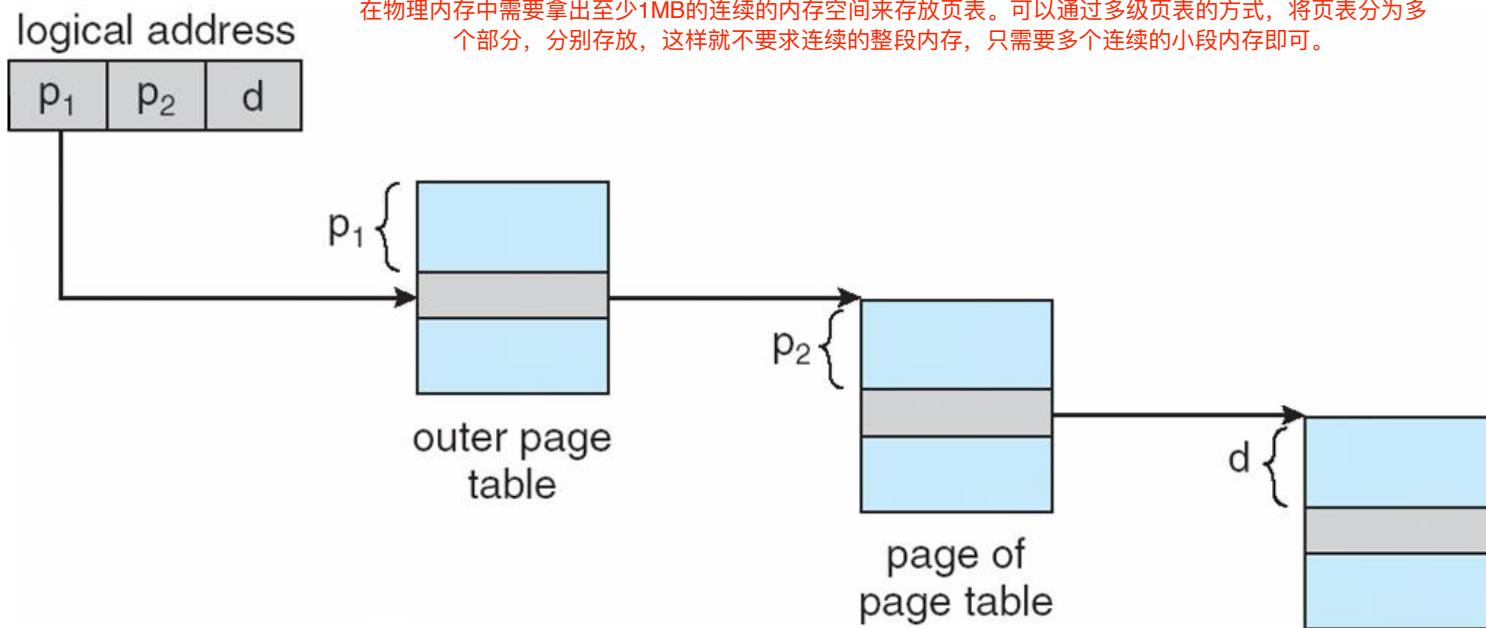
$2^{64}/2^{12} = 2^{52}$ too large

Paging

- Advantages
 - No external fragmentation
- Two main Issues
 - Translation speed can be slow
 - TLB
 - Table size is big

Multi-level Paging

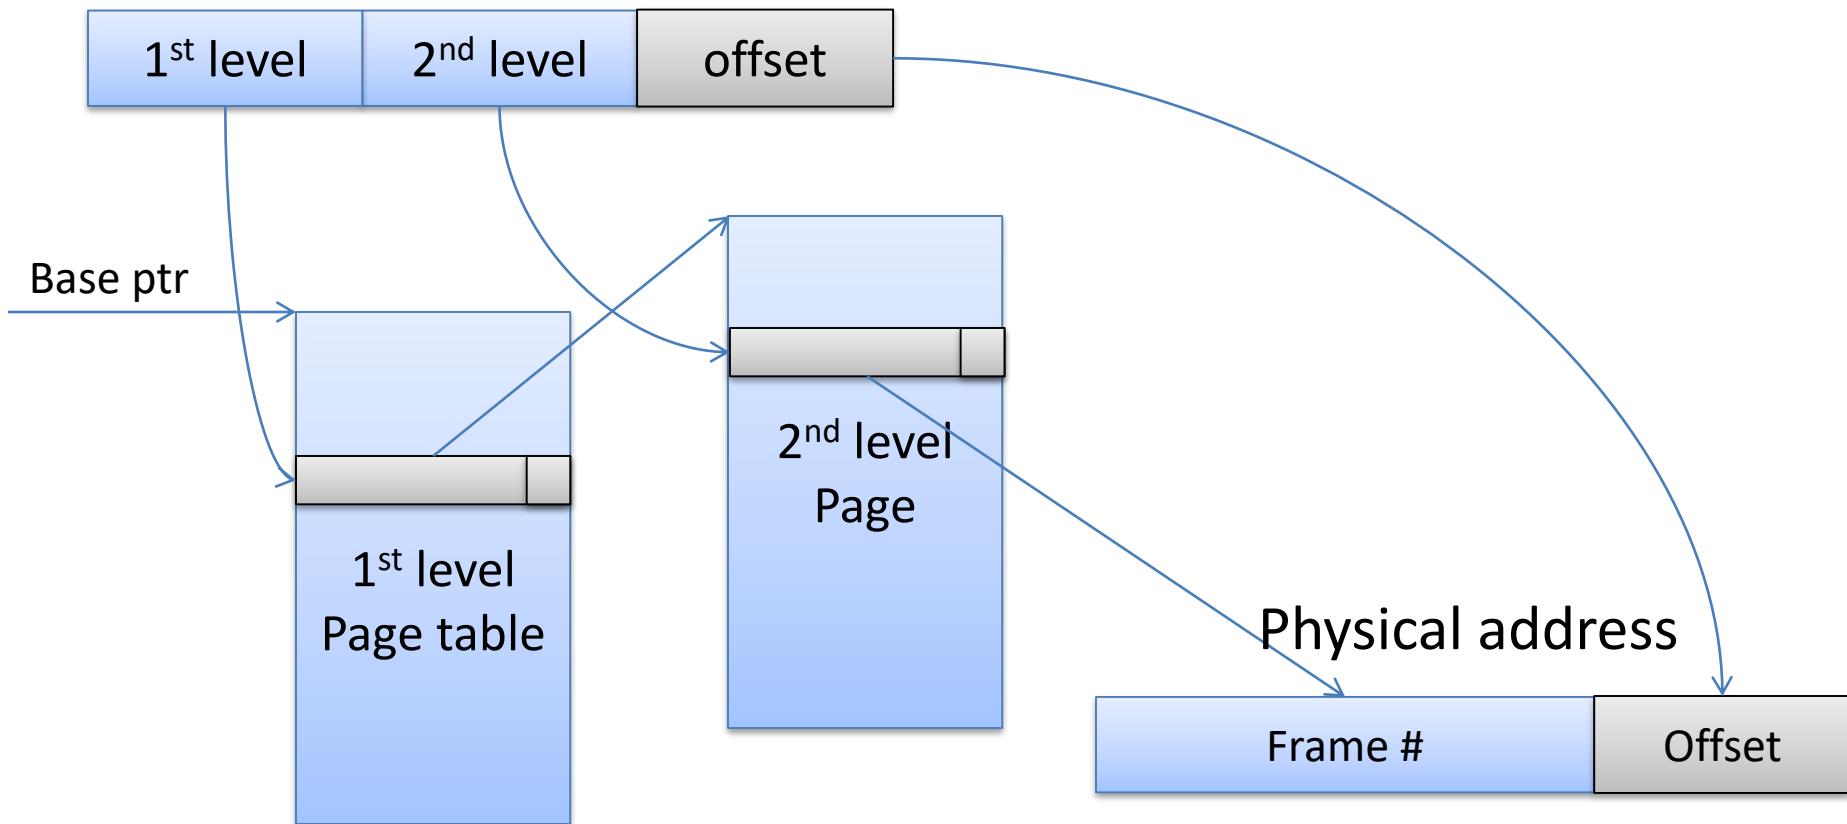
- Two-level paging



Two Level Address Translation

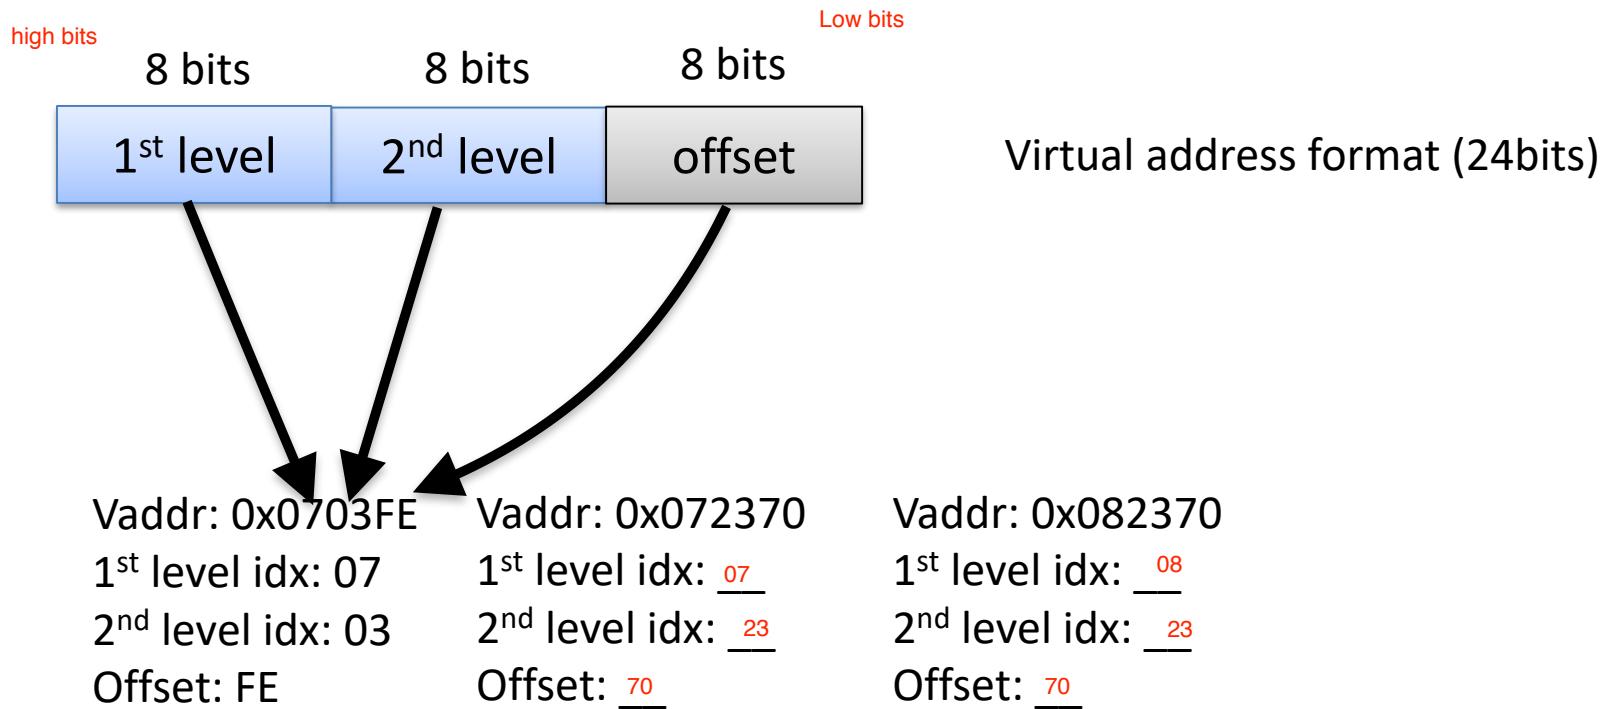
第一层所有指针不是全部被使用，根据内存的不同，使用不同数量的level 2

Virtual address



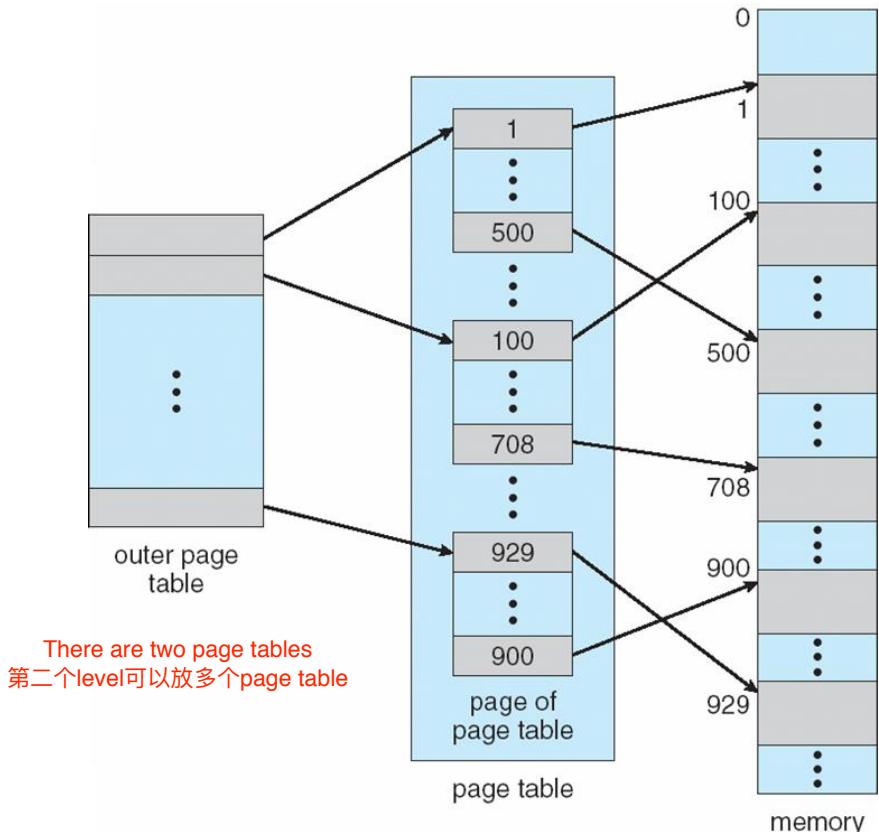
1st level 里面存的是level 2 中第几个page

Example



Multi-level Paging

- Can save table space
- How, why?
- Don't need to create all mappings in the outer page table



Summary

- MMU
 - Virtual address → physical address
 - Various designs are possible, but
- Paged MMU
 - Memory is divided into fixed-sized pages
 - Use page table to store the translation table
 - No external fragmentation

Summary

- Paged MMU: Two main Issues
 - Translation speed can be slow
 - TLB
 - Table size is big
 - Multi-level page table

Quiz

- What is the minimum page table size of a process that uses only 4MB memory space?
 - assume a PTE size is 4B



$$4 * 2^{20} = 4\text{MB}$$



$$4 * 2^{10} + 4 * 2^{10} = 8\text{KB}$$

Quiz

- What is the page table size for a process that only uses 8MB memory?
 - Common: 32bit address space, 4KB page size
 - Case 1) 1-level page table
 - Assume each page table entry is 4 bytes
 - Answer: $2^{20} \times 4 \text{ byte} = 4\text{MB}$
1M
 - Case 2) two-level page table
 - Assume first 10 bits are used as the index of the first-level page table, next 10 bits are used as the index of the second-level page table. In both-levels, single page table entry size is 4 bytes
 - Answer: $2^{10} \times 4 + 2 \times (2^{10} \times 4) = 4\text{KB} + 8\text{KB} = 12\text{KB}$

Quiz

- What is the page table size for a process that only uses 16MB memory?
 - Common: 32bit address space, 4KB page size
 - Case 1) 1-level page table
 - Assume each page table entry is 4 bytes
 - Answer: $2^{20} \times 4 \text{ byte} = 4\text{MB}$
 - Case 2) two-level page table
 - Assume first 10 bits are used as the index of the first-level page table, next 10 bits are used as the index of the second-level page table. In both-levels, single page table entry size is 4 bytes
 - Answer: $2^{10} \times 4 + 4 \times (2^{10} \times 4) = 4\text{KB} + 16\text{KB} = 20\text{KB}$