

EECS 678: Introduction to Operating Systems

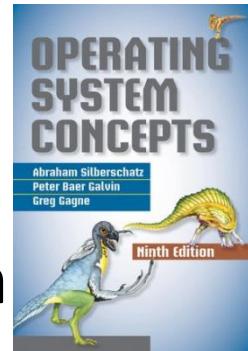
Heechul Yun

About Me

- Heechul Yun, Assistant Prof., Dept. of EECS
 - Office: 3040 Eaton, 236 Nichols
 - Email: heechul.yun@ku.edu
- Research Areas
 - Operating systems and architecture support for embedded/real-time systems
 - To improve time predictability, energy efficiency, and throughput
 - Multicore, memory systems
- Previously
 - Worked as a systems software engineer at Samsung Electronics
 - mainly worked on Linux kernel
- More Information
 - <http://ittc.ku.edu/~heechul>

About This Class

- Textbook: Operating System Concepts
- Objectives: Learn OS basics and practical system programming skills
 - Understand *how it works!*
- Audience: Senior and Junior undergraduate (grad students)
- Course website:
<http://ittc.ku.edu/~heechul/courses/eecs678/>



Course Structure

- Lectures
 - Discuss OS concepts and the design of major OS components
- Quiz
 - Weekly online quizzes to check your understanding
- Lab
 - Hands-on system programming experiences.
 - Each lab includes lab discussion and an assignment
- Programming projects
 - Design and implement some parts of OS.
 - 3 projects: 1) Shell, 2) CPU scheduler, 3) Memory allocator
 - To do in groups of two persons. Solo project is also allowed.

Grading

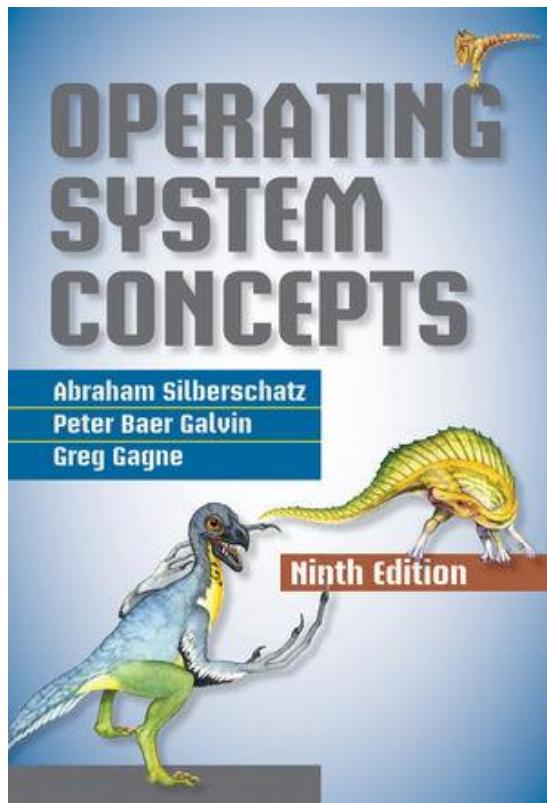
- Attendance: 5%
- Exam: 50% (Mid:20%, Final:30%)
- Quiz: 5%
- Lab: 15%
- Projects: 25%

Grading

- 90+ : A
- 80-89: B
- 70-79: C
- 50-69: D
- 0-49: F

Policy

- Late submissions
 - Lab assignments: **not allowed.**
 - Projects: 20% off each additional 24 hours delay (~24h = 80%, ~48h = 60%, ~72h=40%, ~96h=20%, >96h = 0%)
- Cheating
 - You can discuss about code and help find bugs of your peers. However, copying another's code (e.g., from github) or writing code for someone else is cheating and, if identified, **the involved students will be notified to the department chair**



Operating Systems Are Everywhere

- Computers
- Smart phones
- Cars
- Airplanes
- ...
- Almost everything

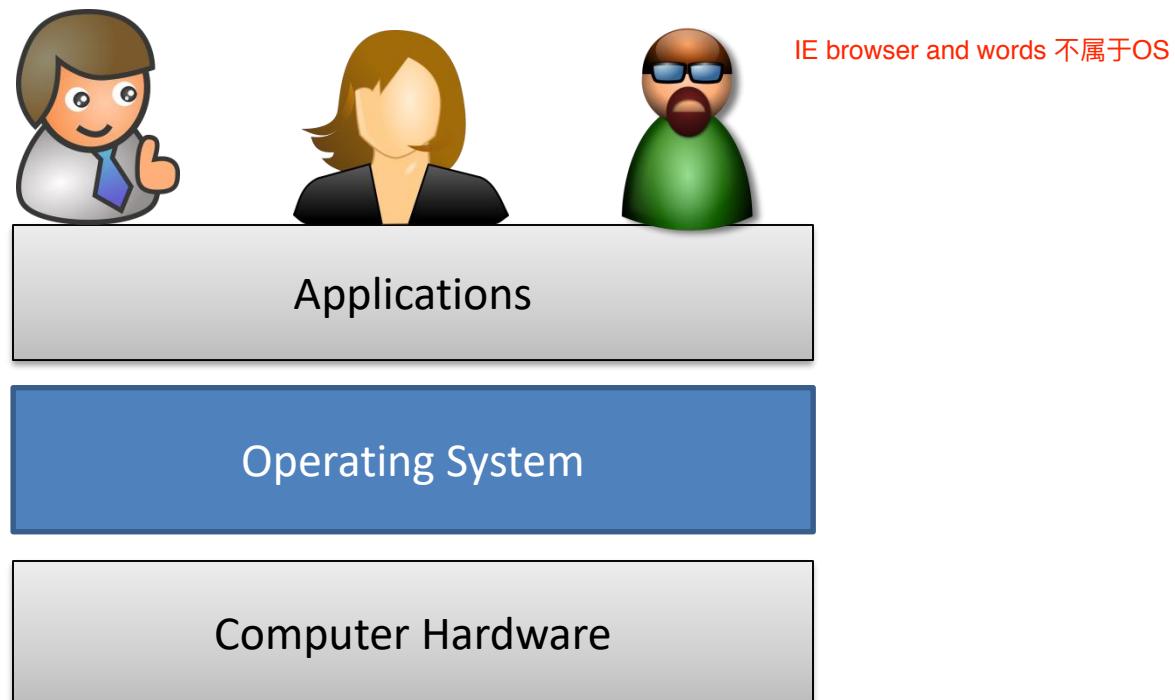


What is an Operating System?



What is an Operating System?

- A program that acts as an intermediary between users and the computer hardware

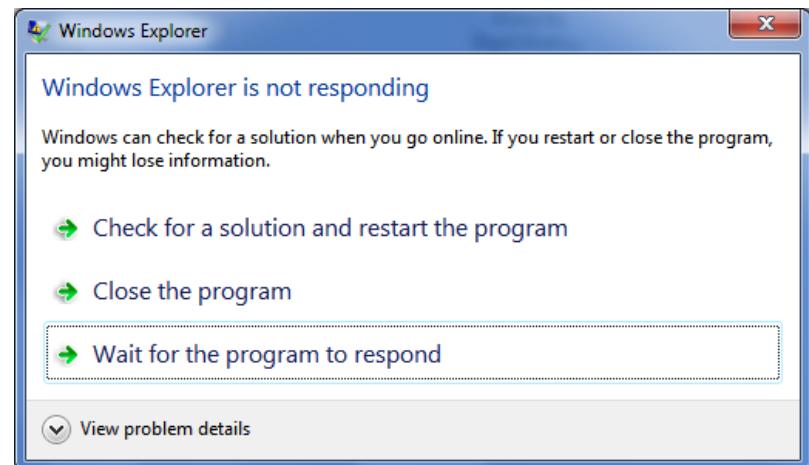


What is an Operating System?

- An easy to use virtual machine – **User's view**
 - Hide complex details for you.
 - What CPU am I using? Intel or AMD?
 - How much memory do I have?
 - Where and how to store my data on the disk?
 - Provide APIs and services
 - `read(...)`, `write(..)`
 - Virtual memory, filesystems, ...

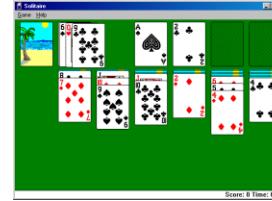
What is an Operating System?

- A resource manager – **System's view**
 - Make everybody get a fair share of resources
 - Time and space multiplexing hardware resources
 - Monitor/prevent error or improper use



What is an Operating System?

- Is an internet browser part of an OS?
 - Everything that shipped by the OS vendor?
 - What about ‘solitaire’?
- The program that always runs
 - Typically in kernel mode (we will learn it later)



Why Needed?

- Programmability
 - You don't need to know hardware details to do stuffs
- Portability
 - You can run the same program on different hardware configurations
- Efficiency
 - Multiple programs/users can share the same hardware efficiently
- Safety & Security
 - The OS protects your program from faults and attacks

Memory Isolation

- Data stored in memory can only be accessed by the authorized entities
- Foundation for most systems security
- Enforced by hardware (e.g., MMU) with runtime (OS) support
- Successful attacks breaking memory isolation by exploiting side-channels exist
 - e.g., RowHammer, Meltdown, Spectre



Meltdown Attack

<https://meltdownattack.com/>

CPU漏洞或者缺陷

- What is it?
 - An attack that exploits Intel CPU's flaw that allows any user-level process to read the content of the kernel-only accessible memory, **breaking memory isolation**
- What's the impact?
 - An attacker can dump the entire memory, including password and other confidential information
- This course will help you better understand how these attacks work and how they are mitigated

What to Study?

- Not “how to use”
 - I’m sure you know better than me about how to use the iOS in your iPhone.
- But “how it works!”
 - We will study the underlying concepts, standard OS components and their designs

OS Design Issues

- Structure
 - How to organize the OS?
- Communication
 - How to exchange data among different programs?
- Performance
 - How to maximize/guarantee performance and fairness?
- Naming
 - How to name/access resources ?
- Protection
 - How to protect with each other?
- Security
 - How to prevent unauthorized access?
- Reliability
 - How to prevent system crash?

Why Study?

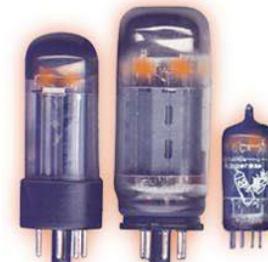
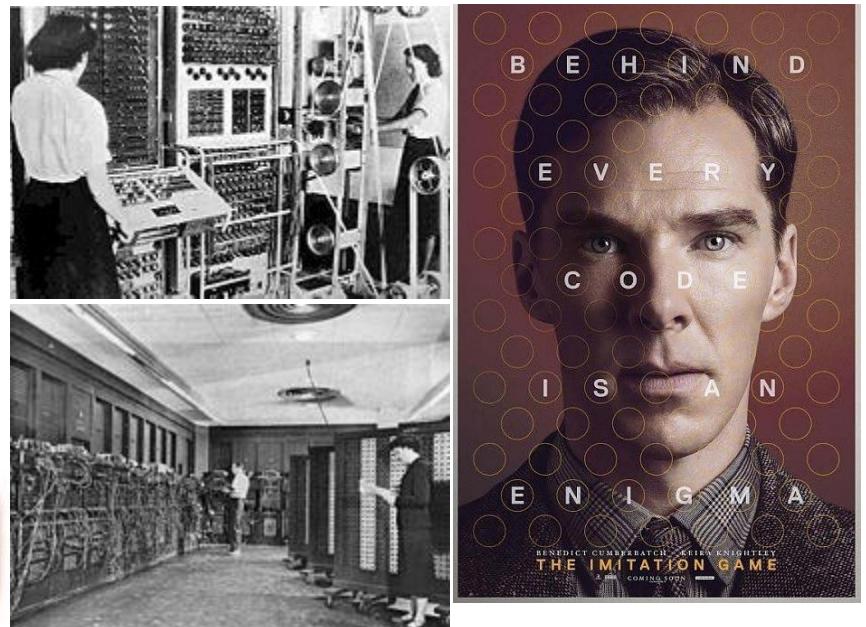
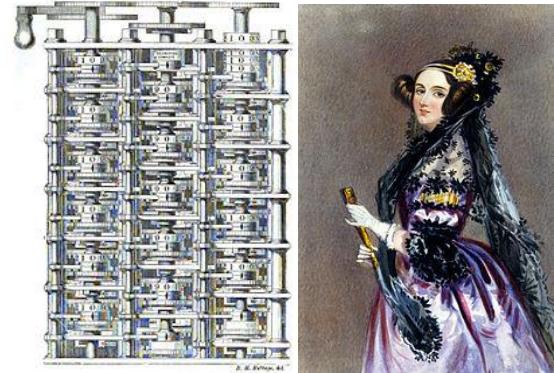
- I'm a user
 - Have you ever wondered how it works?
 - You can better tune the OS to improve performance (or save energy)
- I'm a system programmer
 - You can write more efficient programs by knowing how the OS works.
- I'm a hacker
 - You need to know the enemy (the OS) to beat it

Recap

- What is an OS?
 - An intermediary between users and hardware
 - A program that is always running
 - A resource manager
 - Manage resources efficiently and fairly
 - A easy to use virtual machine
 - providing APIs and services

Brief History of Computers

- Early computing machines
 - Babbage's analytical engine
 - First programmer: Ada Lovelace
- Vacuum tube machines
 - 1940s ~ 1950s
 - Used to break code in WWII
 - No OS, No PL



Brief History of Computers

- Vacuum tubes → Transistors → IC → VLSI
 - Smaller, faster, and more reliable
 - Enable smaller computers
- 1960s Mainframes
- 1970s Minicomputers
- 1980s Microprocessor, Apple, IBM PC
- 1990s PC, Internet
- 2000s Cloud computing
- 2010s Mobile, Internet-of-things (IoT)



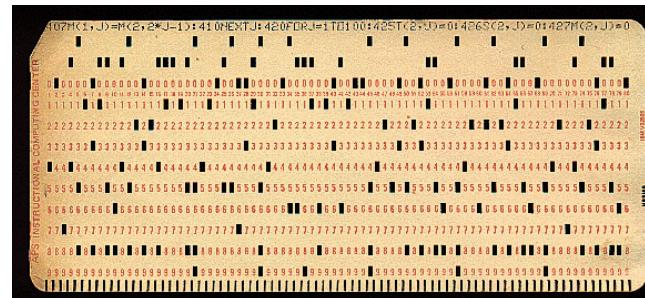
Evolution of Operating Systems

- Batch systems
 - Each user submits her job on punch cards
 - Collect a batch of jobs, read the batch before start processing
 - The ‘OS’ processes each job at a time
 - Problems
 - CPU is underutilize to wait I/O operations
 - No interactivity



IBM 029 card punch

<http://www.catb.org/esr/writings/taoou/html/ch02s01.html>



Evolution of Operating Systems

- Multiprogramming
 - Multiple runnable jobs at a time
 - I/O and compute can overlap
 - OS goal: maximize system throughput
 - IBM OS/360



Evolution of Operating Systems

- Timesharing
 - Multiple interactive users sharing a machine
 - Each user accesses the machine via a terminal
 - Provide each user an illusion of using the entire machine
 - OS goal: optimize response time
 - UNIX

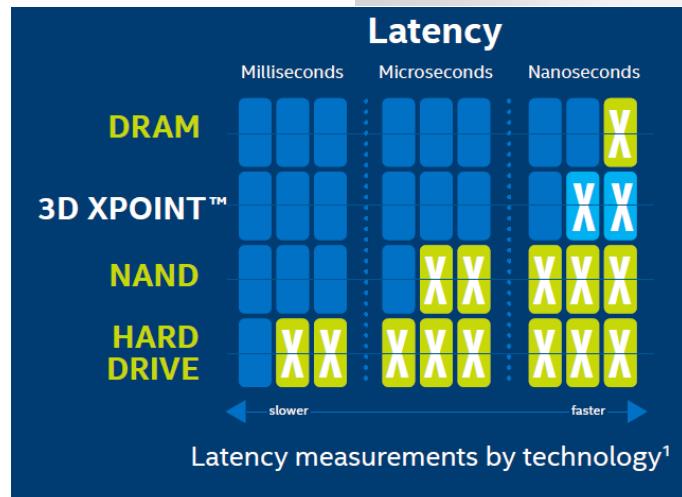
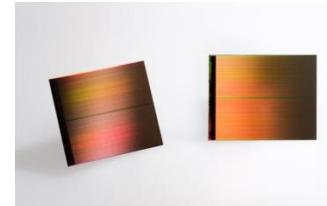


Evolution of Operating Systems

- Parallel computing
 - Use multiple CPUs/cores to speed up performance
 - OS goal: fast synchronization, max utilization
- Distributed computing
 - Physically separate networked computers
- Virtualization
 - Multiple OSes on a single machine

Challenges for Future OS

- New kinds of hardware are keep coming
 - Heterogeneous multicore processors (e.g., ARM big.LITTLE)
 - Storage Class Memory (SCM): non-volatile DRAM-like memories
- New computing paradigms
 - Cloud computing
 - Internet-of-Things (IoT)



Summary

- In this class, you will learn
 - Major OS components
 - Their structure, interface, mechanisms, policies, and algorithms
- This class will (hopefully) help you
 - Understand the foundation of computing systems
 - Understand various engineering trade-offs in designing complex systems you would build in future

Computer Architecture and OS

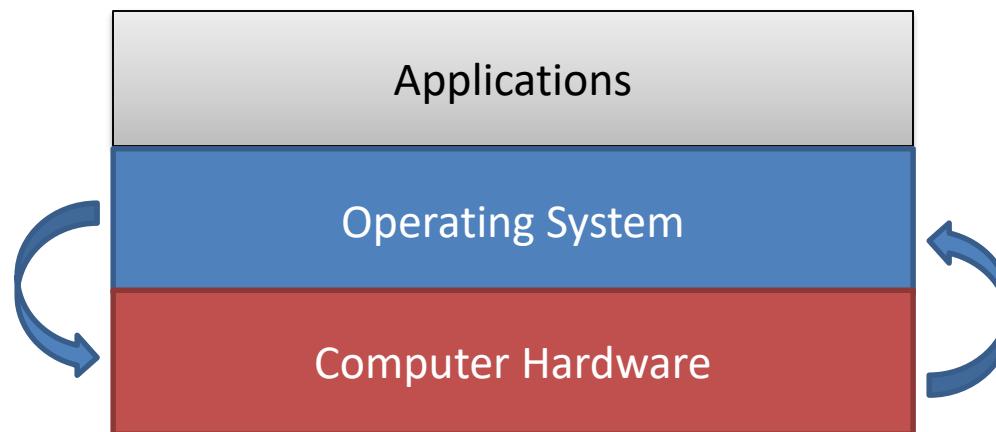
EECS678

Agenda

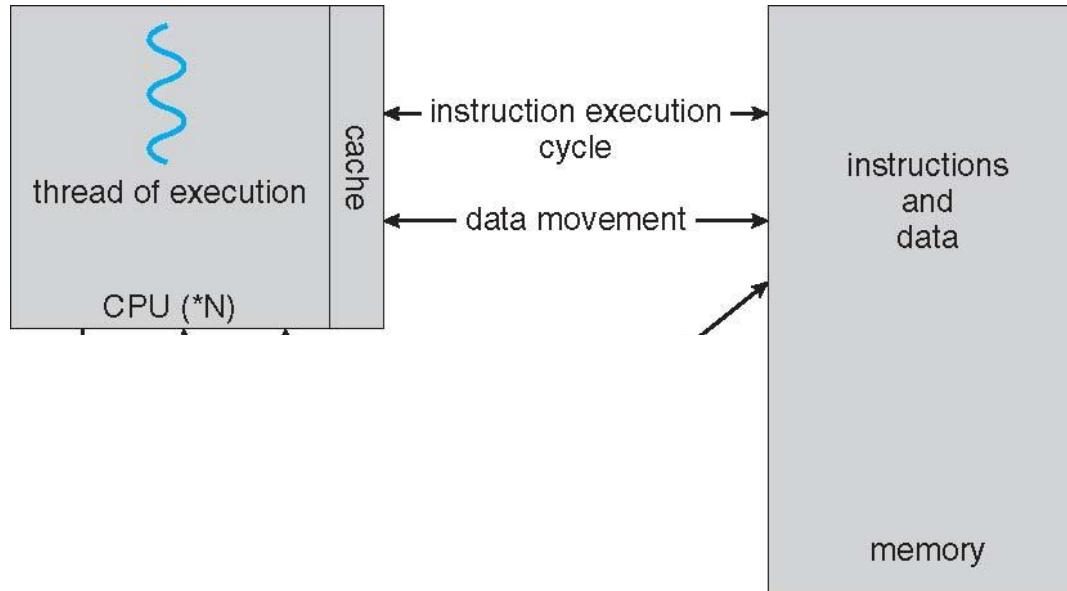
- Computer architecture and OS
 - CPU, memory, disk
 - Architecture trends and their impact to OS
 - Architectural support for OS

Computer Architecture and OS

- OS talks to hardware
 - OS needs to know the hardware features
 - OS drives new hardware features



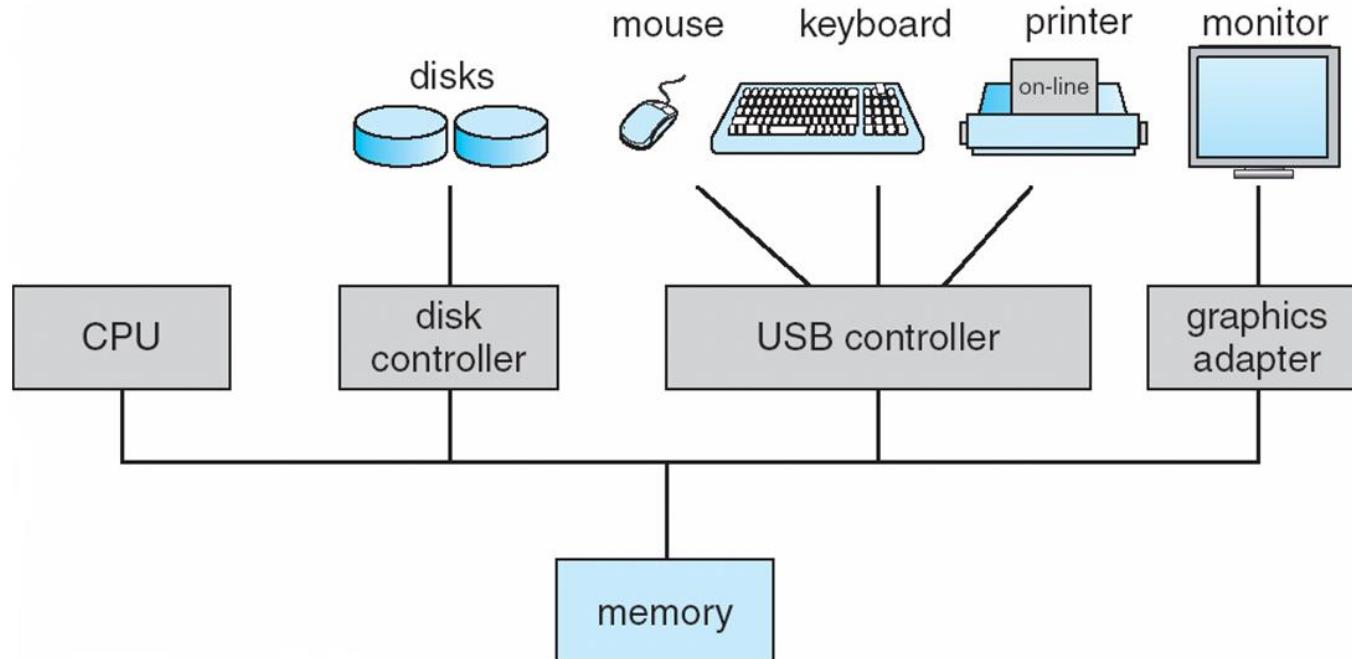
Simplified Computer Architecture



A von Neumann architecture

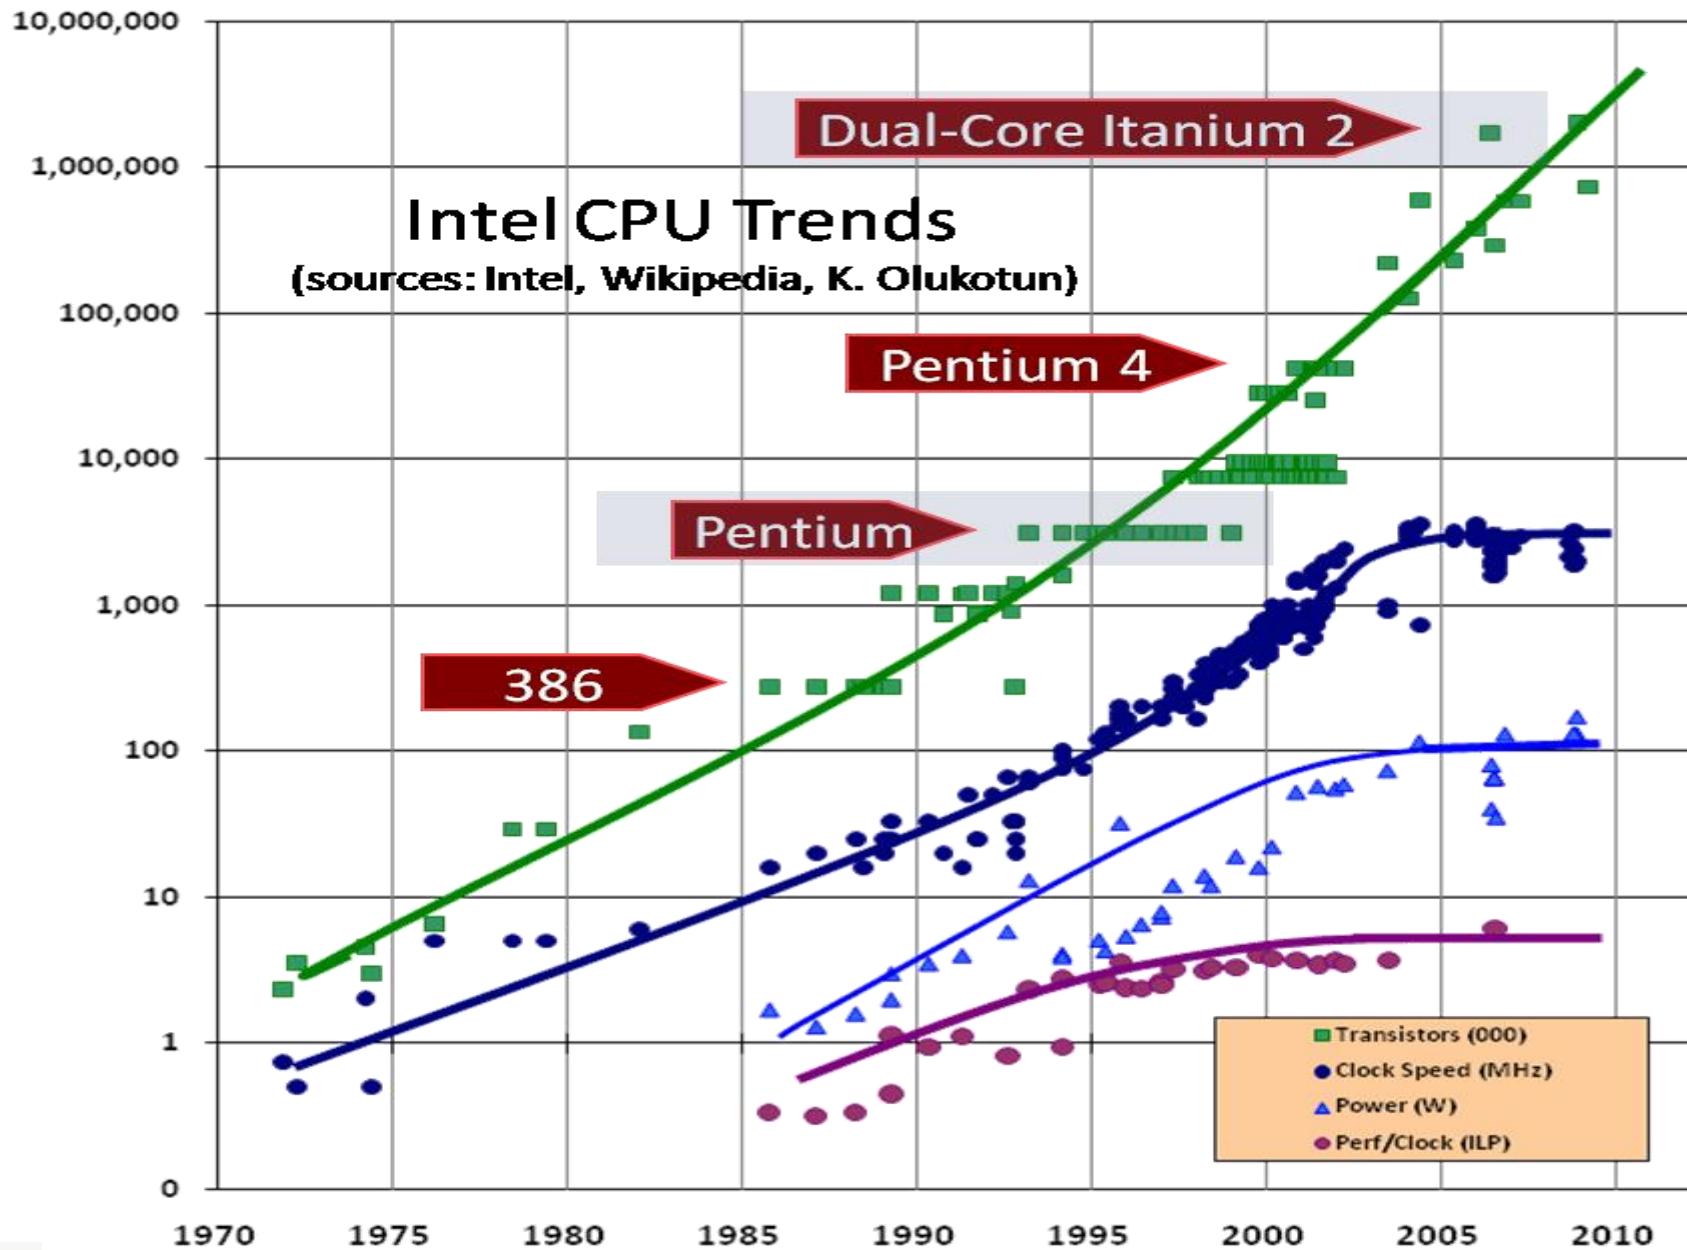
A Computer System

- Essentials: CPU, Memory, Disk
- Others: graphic, USB, keyboard, mouse, ...



Central Processing Unit (CPU)

- The brain of a computer
 - Fetch instruction from memory
 - Decode and execute
 - Store results on memory/registers
- Moore's law
 - Transistors double every 1~2yr
 - 5.56 billion in a 18-core Intel Xeon Haswell-E5

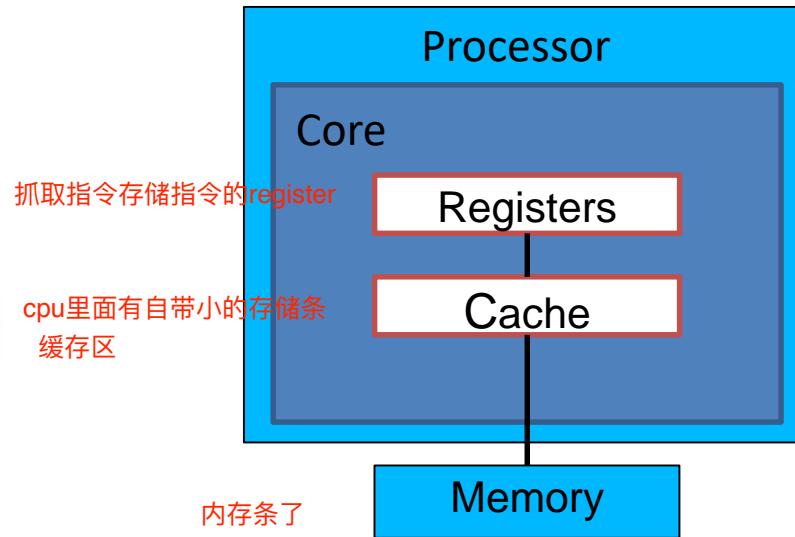


H Sutter, "The Free Lunch Is Over", Dr. Dobb's Journal, 2009

Review

- Batch computing
- Multiprogramming
- Time sharing
- Parallel & distributed computing
- Mobile computing

Single-core CPU

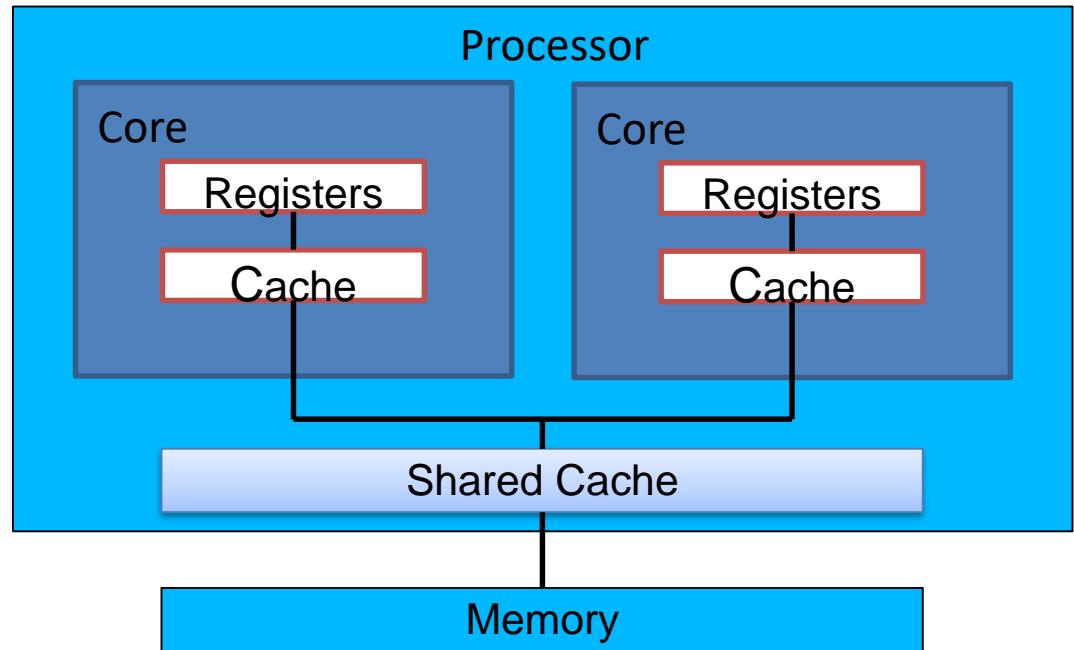


- Time sharing
 - When to schedule which task?

Multicore CPU



多核

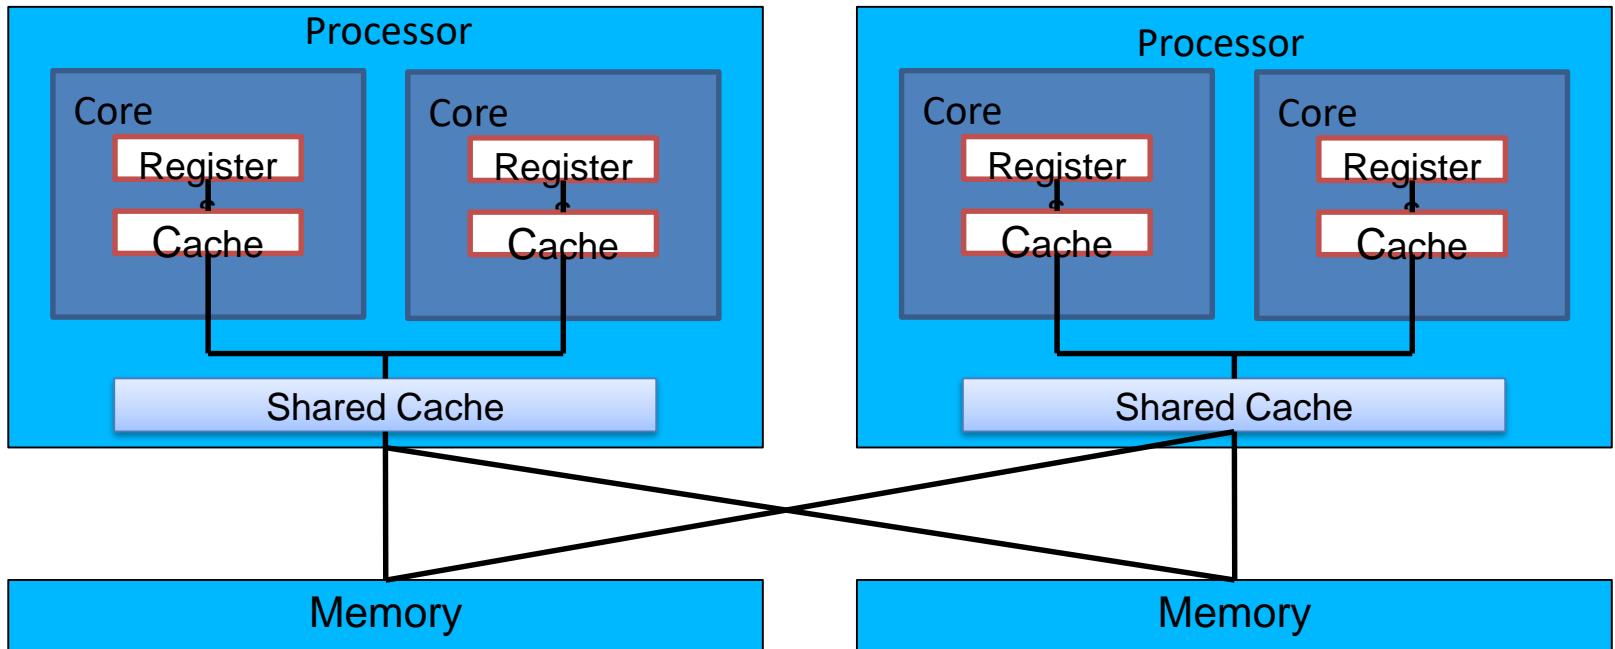


- Parallel processing
 - Which tasks to which cores?
 - May have performance implication due to cache contention → contention-aware scheduling

Multiprocessors



多核多cpu



- Non-uniform memory access (NUMA) architecture
 - Memory access cost varies significantly: local vs. remote
 - Which tasks to which processors?

Memory Hierarchy

- Main memory
 - DRAM

内存条不易存储数据，和cpu交换频繁
 - Fast, volatile, expensive

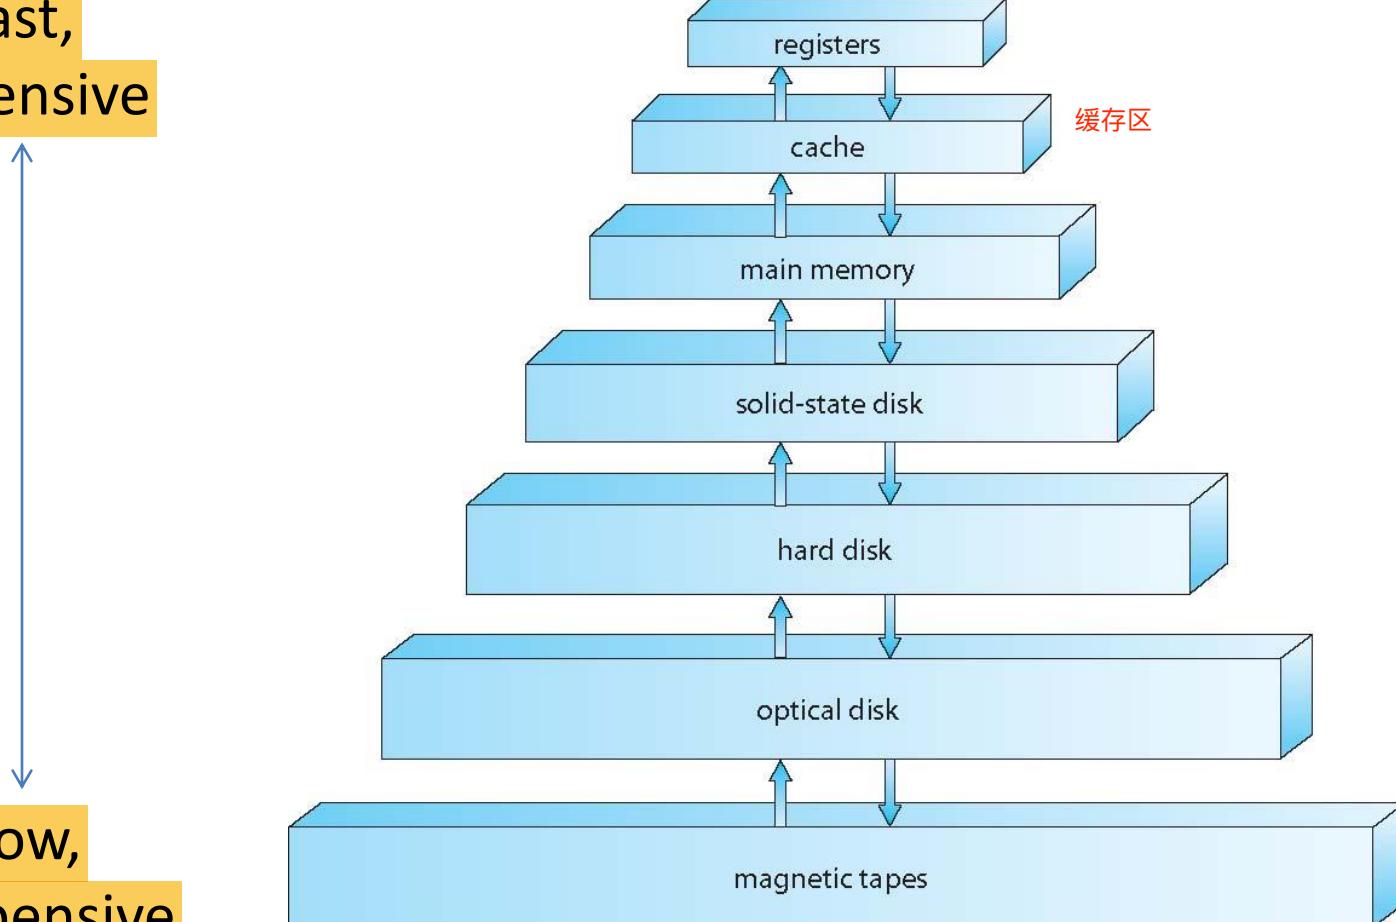
数据易丢失
 - CPU has direct access
- Disk
 - Hard disks, solid-state disks
 - Slow, non-volatile, inexpensive
 - CPU doesn't have direct access.



Memory Hierarchy

Fast,
Expensive

Slow,
Inexpensive



缓存区

Storage Performance

- Performance of various levels of storage depends on
 - distance from the CPU, size, and process technology used
- Movement between levels of storage hierarchy can be explicit or implicit

内存是与CPU之间的桥梁或与仓库。显然，内存的容量决定“仓库”的大小，而内存的带宽决定“桥梁”的宽窄，两者缺一不可

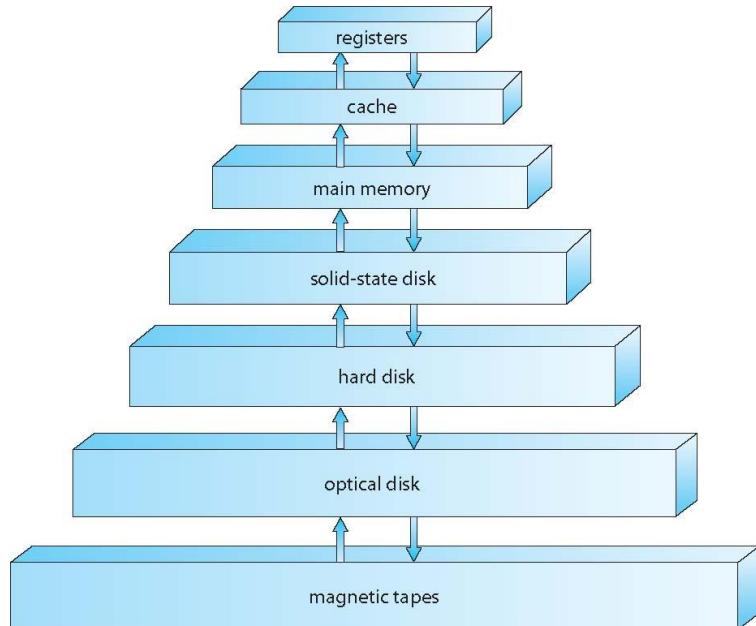
Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Caching

- A very important principle applied in all layers of hardware, OS, and software
 - Put frequently accessed data in a small amount of faster memory
 - Fast, most of the time (hit)
 - Copy from slower memory to the cache (miss)

命中：可以直接通过缓存获取到需要的数据。
读写命中！！

不命中：无法直接通过缓存获取到想要的数据，需要再次查询数据库或者执行其它的操作。原因可能是由于缓存中根本不存在，或者缓存已经过期。



个人理解就是你磁盘或者内存上的存储区域之前有没有写过数据，如果有，这次再写到相同的区域叫写命中，如果写到其他区域，叫写未命中，在数据恢复方面，如果写命中了，那之前的数据被覆盖，就很难再恢复回来，如果写未命中，那么之前的数据就容易被找回

Architectural Support for OS

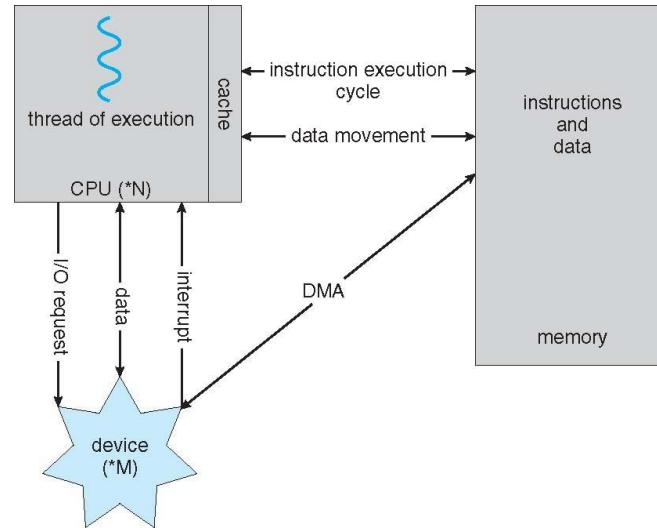
- Interrupts and exceptions
- Protected modes (kernel/user modes)
- Memory protection and virtual memory
- Synchronization instructions

Interrupt

- What is an interrupt?
 - A signal to the processor telling “do something now!”
- Hardware interrupts
 - Devices (timer, disk, keyboard, ...) to CPU
- Software interrupts (exceptions)
 - Divide by zero, special instructions (e.g., int 0x80)

主动的：中断，打断
是被动的不好理解

Interrupt Handling



- save CPU states (registers)
- execute the associated interrupt service routine (ISR)
- restore the CPU states
- return to the interrupted program

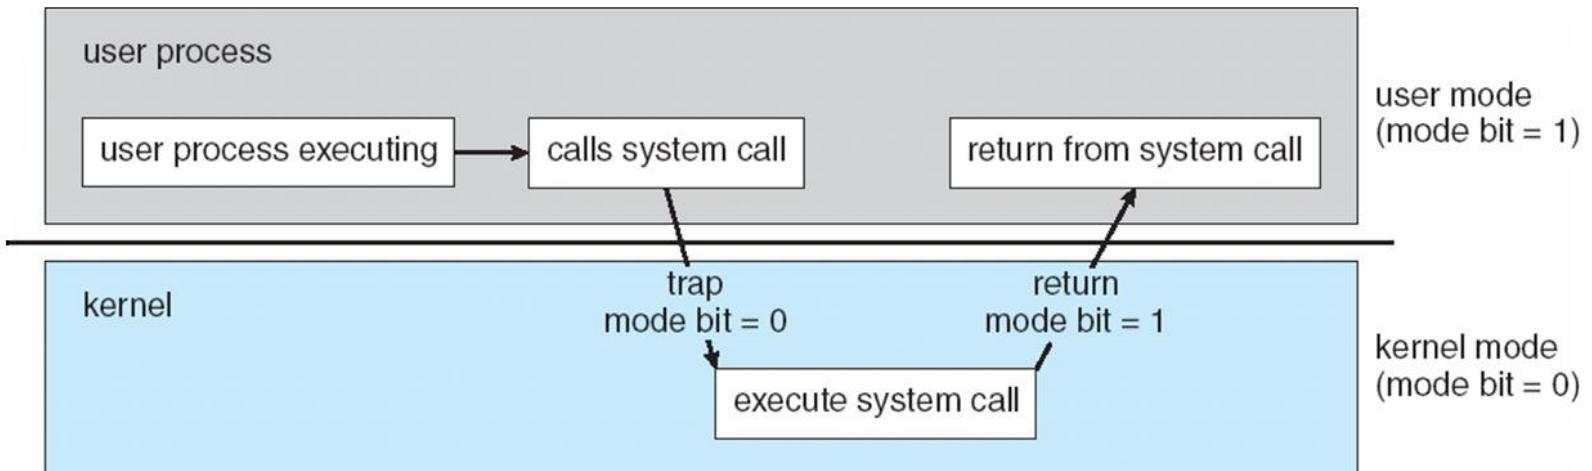
Timesharing

- Multiple tasks share the CPU at the same time
 - But there is only one CPU (assume single-core)
 - Want to schedule different task at a regular interval of 10 ms, for example.
- Timer and OS scheduler tick
 - The OS programs a timer to generate an interrupt at every 10 ms.

Dual (User/Kernel) Mode

- Some operations must be restricted to the OS
 - accessing registers in the disk controller
 - updating memory management unit states
MMU: memory management unit
- User/Kernel mode
 - Hardware support to distinguish app/kernel
 - Privileged instructions are only for kernel mode
 - Applications can enter into kernel mode only via pre-defined **system calls**

User/Kernel Mode Transition



- System calls
 - Programs ask OS services (privileged) via system calls
 - Software interrupt. “int <num>” in Intel x86

Memory Protection

- How to protect memory among apps/kernel?
 - Applications shouldn't be allowed to access kernel's memory
 - An app shouldn't be able to access another app's memory

Virtual Memory

内存条太小了，需要虚拟地址来帮忙

所谓直接运行在物理内存上，是指程序在运行时所访问的地址都是物理地址

- How to overcome memory space limitation?
 - Multiple apps must share limited memory space
 - But they want to use memory as if each has dedicated and big memory space
 - E.g.,) 1GB physical memory and 10 programs, each of which wants to have a linear 4GB address space

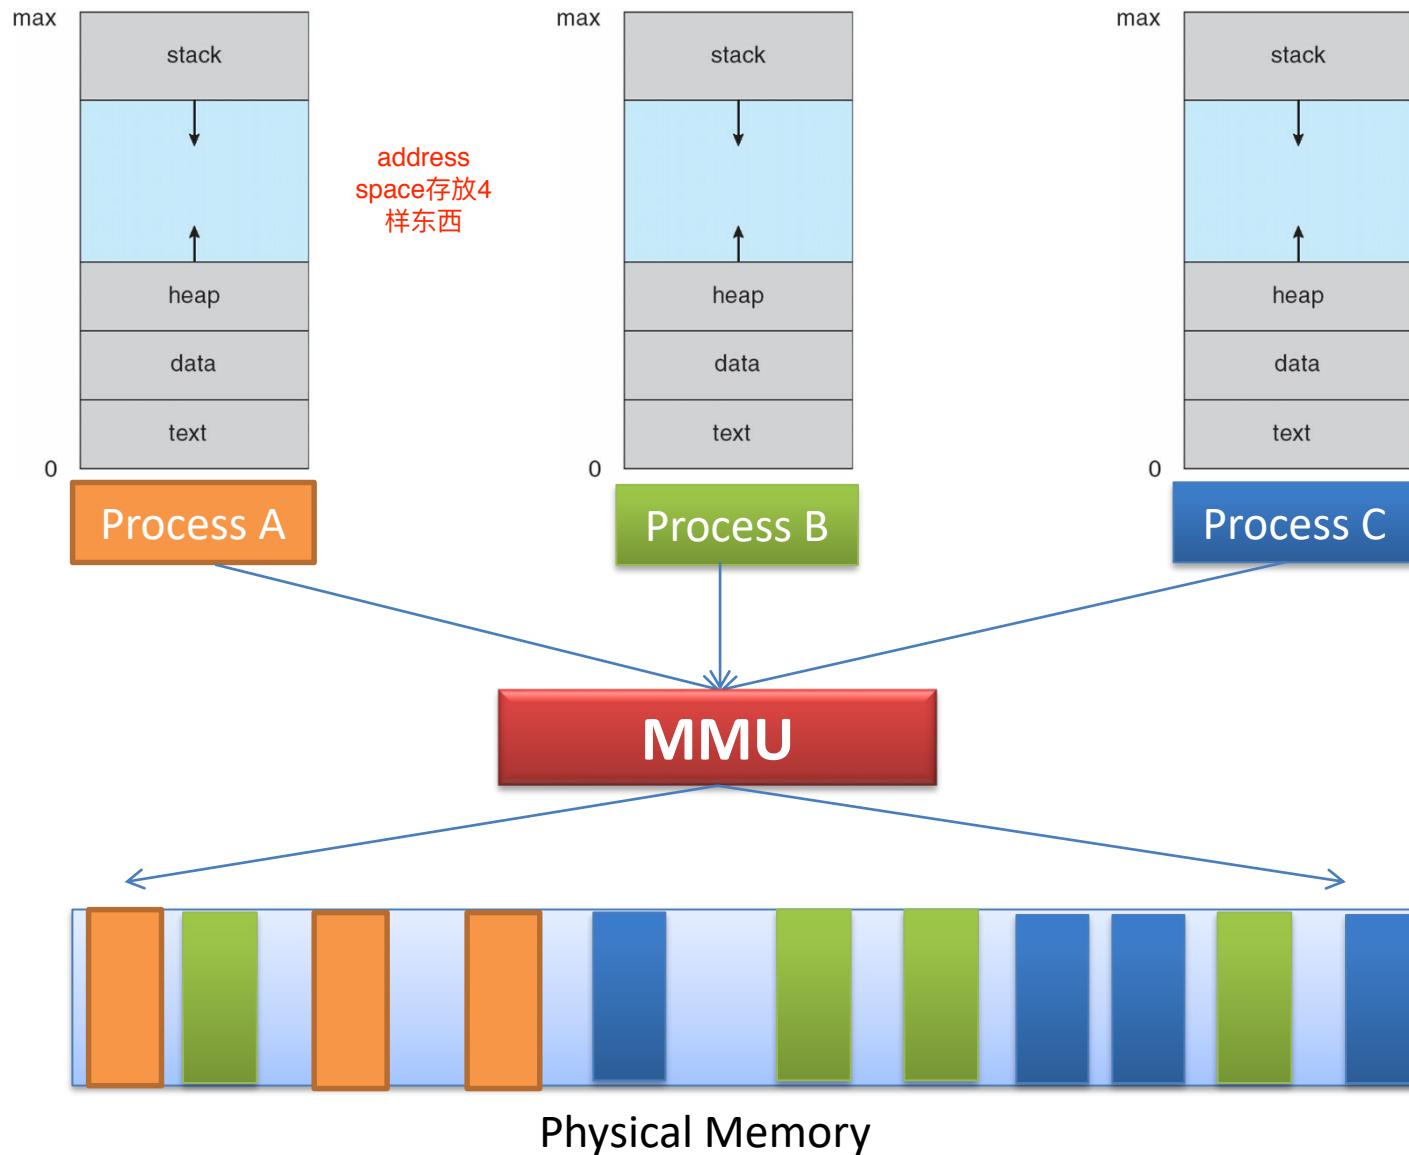
物理内存不足。例如，某个程序运行需要64K的内存，而机器上只有32K的物理内存。

程序运行的地址不确定。同一个程序，每次被装载到内存的地址可能不一样。

内存使用率低。需要运行某个程序，就需要将整个程序装入内存才能够运行。

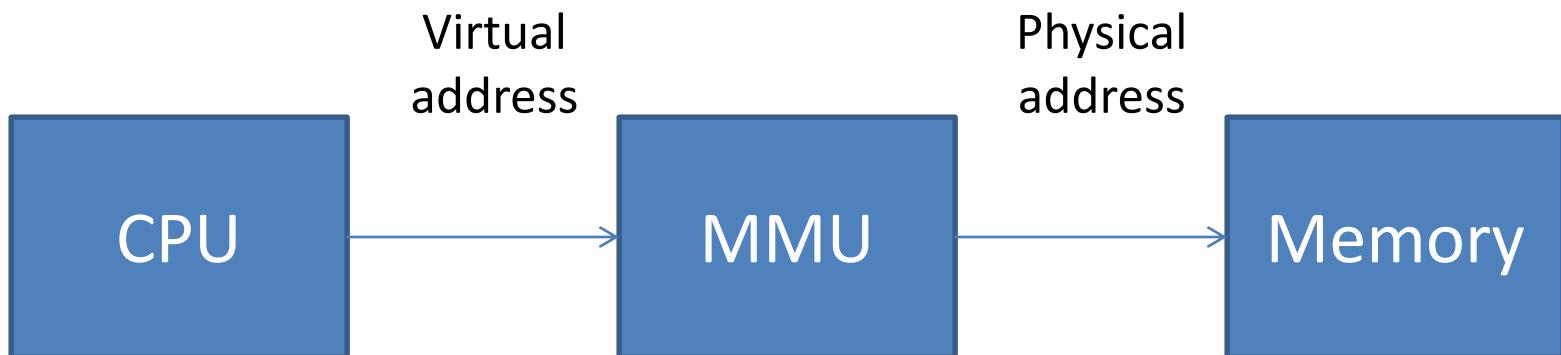
对于多任务OS，存在进程间地址空间不隔离的问题。这样一个任务失败了，可能会导致整个系统宕机
<https://blog.csdn.net/u014085791/article/details/39298639>

Virtual Memory



MMU

- Hardware unit that translates *virtual address* to *physical address*
 - Defines the boundaries of kernel/apps
 - Enable efficient use of physical memory



Synchronization

- Synchronization problem with threads

```
Deposit(account, amount) {  
{  
    account->balance += amount;  
}
```

Thread 1: Deposit(acc, 10)

LOAD R1, account->balance

ADD R1, amount

STORE R1, account->balance

Thread 2: Deposit(acc, 10)

LOAD R1, account->balance

ADD R1, amount

STORE R1, account->balance



Synchronization Instructions

- **Hardware support** for synchronization
 - *TestAndSet, CompareAndSwap* instructions
 - Atomic load and store
 - Used to implement **lock** primitives
 - New TSX instruction → hardware transaction
- Another methods to implement locks in single-core systems
 - Disabling interrupts

Summary

- OS needs to understand architecture
 - Hardware (CPU, memory, disk) trends and their implications in OS designs
- Architecture needs to support OS
 - Interrupts and timer
 - User/kernel mode and privileged instructions
 - MMU
 - Synchronization instructions

OS Abstractions

<i>Reality</i>	<i>Abstraction</i>
A single computer	Multiple computers
Limited RAM capacity	Infinite capacity
Mechanical disk	File system
Insecure and unreliable networks	Reliable and secure